

ATTORNEY WALKTHROUGH

What We're Building for You, and How the Drafting Works

AI-drafted first-pass appeal letters and medical-review rebuttals, written in the firm's own voice from the firm's own past work. You review and edit. You never redraft from scratch, and nothing goes out without an attorney.

 Attorneys · Tue Jun 9, 11:30 AM

 Eddie · Kevin · Bryn

 First AI appeals · mid-July

The point of today: see what the appeals build does for you as the attorneys, and how the drafting actually works, because that is where you will spend your time. This is not a decision meeting. The seed appeals and sample files come from Eddie and Bryn directly, not the room.

1 · Why we're here

Two things. First, to show you what the appeals build actually does for you: AI-drafted first-pass appeal letters and medical-review rebuttals, in the firm's own voice, from the firm's own past work, that you review and edit rather than draft from scratch. Second, to do this rollout the right way.

The HubSpot build got thrown at the team last time. This time it's phased, you control the substance, and nobody learns a new tool overnight. Yesterday we walked your appeals team through the full eight-stage workflow, and they gave us a lot of useful feedback. Today is about the two pieces that take the most time off *your* plate: the appeal letter and the rebuttal.

2 • What you get first

Two headline deliverables land in Wave 1A, around mid-July. These are the pieces you spend the most time on.



First-pass appeal letters (A-16)

A drafted appeal waiting in the matter folder when the file is ready to appeal.



Medical-review rebuttals (A-23)

A drafted rebuttal when a peer review, IME, or FCE comes back.

Everything else (the request letters, the homework tracking, the calendaring, the resolution letters) automates a few weeks later in Wave 1B. None of it changes the law you write. The next two sections walk exactly how each draft gets built.

3 • How the appeal letter gets drafted (A-16)

When a file is ready to appeal, the draft assembles from four inputs and produces a structured letter you edit.

What it reads

The claim-file summary

Key dates, the policy provisions in play, the denial rationale, treating-physician statements, and any IME or peer-review conclusions.

The denial letter itself

Exactly what the carrier said and the basis they gave, so the appeal answers it directly.

The policy and plan provisions

Definition of disability, own-occ versus any-occ, exclusions, and limitations (for example the 24-month mental-health cap).

The firm's past appeals

The firm's own prior appeals on the same carrier and condition, from a growing index of the firm's work. It keeps learning from every attorney's edits, so over time it reflects the firm's voice.

What it produces

A full first draft in the firm's own appeal structure. This is confirmed against the firm's own appeal letters (verified across Unum, Aetna, MetLife, Sun Life, Hartford, Lincoln, and The Standard, and conditions from Long COVID to migraine to TBI):

1

Header block and opening

Claimant, claim and policy numbers, policyholder, and an opening that names the carrier's denial or termination letter by date.

2

Background

The claim's denial or termination history: what the carrier approved, what it denied or terminated, and (on a termination) what supposedly changed.

3

Medical evidence of the client's disability

The core of the letter, organized by source: treating-provider support, objective testing (CPET, neuropsych, imaging), the IME or FCE, a rebuttal of the carrier's peer review, the vocational assessment, the claimant's personal statement and symptom diary, and third-party declarations from family and friends.

4

The carrier's opportunity to reverse

The firm's recurring "[Carrier]'s Opportunity to Reverse / Reinstate / Approve" section: why the denial fails and what it takes to pay the claim.

5

Conclusion and demand

6

The close

Physicians available for questions, the enclosed Request for Information and Documents, and the evidence list (Attachment "A").



What makes it the firm's draft, not generic AI

Carrier- and condition-aware. It knows a Hartford fight reads differently than a Unum fight, and a Long COVID case argues differently than a mental-health case.

Grounded in the firm's own work. Because it pulls from the firm's past appeals, it argues the way the firm argues, in the firm's own structure and language, not boilerplate. And because it learns from each attorney's edits, it reflects the firm's voice over time.

4 • How the rebuttal gets drafted (A-23)

When a medical review lands (a peer review, an IME, or an FCE), the same engine drafts the firm's response, which the firm files as a supplement to the appeal.

What it reads: the adverse report, the claim file, and the firm's past supplement and response letters on similar reports and conditions.

What it produces

A response built the way the firm rebuts these reports (drawn from the firm's own supplement letters), typically working point by point:

Identifies the reviewer's conclusions

Pulls out exactly what the peer reviewer, IME, or FCE examiner concluded, so the response answers each point.

Attacks the basis

A records-only review with no exam, cherry-picked records, ignored treating-physician opinions, methodology gaps.

Counters with your evidence

The treating evidence and the functional record.

In the firm's voice

Drafted from the firm's own past rebuttals.

It is the same retrieval-grounded engine as the appeal letter, pointed at the report instead of the denial. This is the second piece that eats attorney time, which is why A-23 ships in the same wave as A-16.

5 • What every draft has in common



Lands where you expect

A `.docx` in the matter folder (AI output subfolder), with a Teams alert and a mobile link.



Shows its sources

Which past appeals, rebuttals, and templates it pulled from, so you verify provenance before you trust a line.



Learns from your edits

Rate it A through F; it learns from the diff between its draft and your final version. It improves on your corpus, not in the abstract.

6 • The quality bar

We're being honest about what to expect. The first drafts will be rough, and that's fine. The work is in the tuning: we run the engine on the matters you pick and dial it in against your edits until a draft is a real starting point, not a blank page.



We dial it in together

We'll agree on what "good enough" looks like for your appeals, section by section, as we go. And if the drafts aren't earning their place by the end of that window, we deal with that before you commit further. Everything built in this phase is yours to keep regardless.

7 • You stay in control



It's always a first draft

Nothing drafted by AI reaches a carrier or a client without an attorney editing and sending it. A firm rule, and the system is built around it.



Human review on everything

The AI drafts and assembles. An attorney reads, edits, and sends. Always, for anything client- or carrier-facing.



Calendaring stays human

The system may read a denial and propose the appeal deadline. A human always confirms it against the letter. We never call calendaring automated.



You own the substance

A faster starting point with its sources shown. The legal judgment, the argument, the final letter, those are yours.

8 • Doing the rollout right

The honest part. The first HubSpot rollout under-communicated the why, and it landed on the team as extra work. The fix is built into how we run every phase.

Phased, one department at a time

Appeals first. Nothing drops on the whole firm at once.

Training before anything goes live

A department explainer and AI summaries first, then a runbook and a video per workflow.

Demo and test windows

The team tries each piece before it's real, and tells us what to fix.

Adoption sign-off gates each phase

Bryn signs off that the team is actually using it before we move on. No adoption, no advance.

The measure of success is the assistants and attorneys actually adopting it. That's the whole game, and it's what got under-served last time. Nick is on-site at every phase kickoff, not just this one.

9 • The roadmap

Seven phases under one agreement, appeals end to end first. Pricing is locked across all seven so you can plan the year, and you commit one phase at a time with Bryn's sign-off as the gate between them.



Wave 1A (now to mid-July): the foundation plus first-draft appeals and rebuttals.

Wave 1B (to mid-August): the rest of the appeals workflow automates around the draft. LLM usage is billed at cost, roughly \$200 to \$500 a month starting in Wave 1A.

10 · What happens next



The cadence

Week 1 is on-site shadowing on real appeals, so we capture the actual process, not the idealized one. Then the foundation and the summarizer, then the retrieval index and the appeal and rebuttal drafting.

Go-live demo to Bryn and Eddie around Week 6, then two weeks of tuning on the matters you select. You'll hear from us before any date slips, not after.