THE DODONAI GUIDE TO AI-NATIVE LAW

Going (AI) Native with Your Law Firm

How to Use AI to Run a Smarter, More Profitable Legal Practice

INSIDE: The 1% Problem · What AI-Native Looks Like · Architecture: Skills, Connections & Schedules · Document Processing Engine · From Documents to Strategy · and more

DODONAI

AI Services for Legal Teams

264 pages · 8 real agent outputs · prompts included

Dodonai, Inc. · dodon.ai · 2026

Contents

Contents 2
Chapter 1: The 1% Problem 9
7:32 PM on a Tuesday 9
The 1% Problem 10
Story 1: The Two-Person Firm With 200 Clients 11
Story 2: 32x Time Compression 11
Story 3: The Morning Briefing 12
The Pressure From Above 13
The Practice Without Systems 14
Your Effective Hourly Rate 15
From Chatbot to Operating System 16
What AI Won't Do for Your Practice 18
Try This Now: Your First Prompt With Real Data 19
The Time Audit 20
End-of-Chapter Checklist 21
Quick Reference Card: Chapter 1 22
Chapter 2: What AI-Native Looks Like 23
What Your AI Did While You Slept 23
A Morning at Chen Law Group 24
The Outputs: What You'd Actually See 25
Meeting Intel Briefing – Apr 2, 2026 25
Case Health Digest – Week of Mar 31, 2026 27
Inbox Digest – 2026-04-02 (Morning) 30
The Compound Effect 32
The Numbers Behind AI-Native 34
What This Isn't 36
The Shift in Your Day 37
Try This Now: Map Your Morning 37
What Comes Next 38

End-of-Chapter Checklist 38
Quick Reference Card: Chapter 2 39
Chapter 3: The Architecture – Skills, Connections, Schedules 40
What's Behind the Curtain 40
Layer 1: The Skill 41
Layer 2: The Connections 44
Layer 3: The Schedule 45
The Safety Layer 47
Putting the Layers Together 49
The Architecture is AI-Model Agnostic 50
What AI Won't Do in This Architecture 51
Try This Now: Write a Skill File 51
Quick Reference: The Architecture 52
Chapter 4: Dodonai – Your Document Processing Engine 53
Before Your Agents Can Think, Your Documents Need Structure 53
Medical Record Summaries & Chronologies 54
Deposition Summaries 55
Transcript Management 57
OCR – AI-Powered Text Extraction 58
eDiscovery – Semantic Document Search 59
Extract & Draft Agents (XD Agents) 61
Pricing 63
How Dodonai Feeds Your Agent Architecture 64
Security and Compliance 64
Try This Now: Process Your First Document 65
Quick Reference: Dodonai Products 66
Chapter 5: From Documents to Strategy 67
The Output That Changes the Conversation 67
Workflow 1: Medical Chronology to Demand Letter 68
Workflow 2: Deposition Summary to Cross-Examination Outline 70
Workflow 3: eDiscovery Results to Case Theory 73

Workflow 4: Intake Notes to Go/No-Go Case Evaluation 75
Workflow 5: Multiple Depositions to Trial Prep Matrix 78
The Manual vs. AI-Assisted Timeline..... 81
The Skills Approach: Encoding Your Methodology 81
What AI Won't Do Here..... 83
Try This Now: Your First Document-to-Strategy Workflow..... 83
Quick Reference: The Five Workflows 84
Chapter 6: Seven Practice Areas, One Architecture 86
The Same Architecture, Every Time 86
Personal Injury: From Records Chaos to Demand-Ready Cases..... 87
Real Estate: Closing Conditions That Track Themselves..... 89
Tax Disputes: Deadlines That Don't Forgive 91
Cybersecurity & Privacy: 50 States, 50 Deadlines, One Agent..... 94
Immigration: Where Every Milestone Is a Deadline..... 96
Estates & Probate: The 12-Month Lifecycle 98
Family & Collaborative Law: The Paper Trail After Every Session 100
The Unifying Pattern 103
What This Means for Multi-Practice Firms 104
What AI Won't Do Across These Practice Areas.....104
Try This Now.....105
Quick Reference: Seven Practice Areas, One Architecture 105
Chapter 7: The Agent Playbook.....107
Where Your 21.5 Hours Go.....107
The Seven Agents.....108
Build Your First Agent: The Deadline Monitor.....109
Agent 2: Inbox Triage.....116
Agent 3: Client Update Drafter.....117
Agent 4: Billing Reconciler.....119
Agent 5: Intake Pipeline.....120
Agent 6: Marketing Engine.....121
Agent 7: Case Health Dashboard.....121

The Compound Effect 122
Time Savings Estimates 123
Build Sequence: What to Do This Month 124
What AI Won't Do Here 124
Try This Now 125
Quick Reference: The Agent Playbook 125
Chapter 8: Marketing, Intake, and Growth 127
Sarah Chen's Marketing Department 127
The Content Engine 128
Review Management 130
Referral Nurturing 132
The Intake Pipeline 134
Building the System: Connections and Schedule 137
The Math 138
The Compound Effect 139
Getting Started This Week 140
Chapter 9: Operations – Billing, SOPs, Knowledge 142
The Revenue You're Leaving in the System 142
Time Entry Reconstruction 143
SOP Creation: From Tribal Knowledge to Written Procedures 145
Template Libraries: From Static Forms to Living Documents 148
Training and Onboarding: The New Hire's First Two Weeks 151
Connecting the Four Systems 153
The Connections Map 153
The Math 154
Getting Started This Week 155
What AI Won't Do Here 156
Quick Reference: Operations Systems 156
Chapter 10: The Safety Layer – Ethics, Risks, and Audit Trails 158
The Case That Changed the Conversation 158
The Hallucination Problem 159

Bar Ethics Rules: Where Things Stand 160
The Carlton Fields Position 163
Confidentiality: The Less Obvious Risk 163
The Embarrassment Prevention Layer 166
How the Safety Layer Maps to Ethics Rules 169
Building Your Firm's AI Policy 170
What AI Can't Protect You From 171
The Honest Assessment 172
The Verification Habit 173
Quick Reference: The Safety Layer 173
Chapter 11: Getting Started – Crawl, Walk, Run 175
7:14 AM on a Thursday 175
The Three Phases 176
Phase 1: Crawl 176
Phase 2: Walk 178
Phase 3: Run 180
Where Dodonai Fits 182
The Chapter Roadmap 184
Your First Week 185
The Firms That Wait 185
5:47 PM on a Tuesday 186
Start 187
Quick Reference Card: Chapter 11 188
Appendix A: Prompt Library 189
Master Index 189
General 191
Case Strategy 195
Medical Records 200
Depositions 203
Client Communication 206
Intake 206

Marketing 209
Billing/Operations 212
Operations (Practice-Area-Specific) 216
Prompts by Chapter 222
Appendix B: Dodonai Product Reference 223
Deposition Summaries 223
Medical Record Summaries & Chronologies 224
Transcript Management 225
OCR 226
eDiscovery 227
Extract & Draft Agents (XD Agents) 228
Pricing Tiers 229
Appendix C: Tools & Integrations 231
1. Clio 231
2. Gmail / Outlook (Email) 232
3. Google Calendar / Outlook Calendar 233
4. QuickBooks 235
5. Dodonai 236
6. Slack 237
7. Court Filing Systems 238
Getting Started: Your First API Connection 240
Appendix D: Sample Agent Outputs 242
1. Damages Narrative from Medical Chronology 242
2. Cross-Examination Outline from Deposition Summary 244
3. Case Evaluation Go/No-Go Memo 245
4. Client Status Update Letter 247
5. Billing Reconciliation Weekly Report 248
6. New PI File Opening SOP 249
7. Practice-Area Blog Post 251
What to Do in the First 72 Hours After a Car Accident in Texas 252
Get Medical Attention – Even If You Feel Fine 252

Document Everything at the Scene 252
Don't Talk to the Other Driver's Insurance Company 252
What Most People Don't Know 253
When to Talk to an Attorney 253
8. Breach Notification Compliance Matrix 253
How to Use These Outputs 255
Appendix E: AI Policy Template 256
[Firm Name] Artificial Intelligence Use Policy 256

Chapter 1: The 1% Problem

What You'll Walk Away With:

- Why single prompts are only 1% of what AI can do for your practice
- Three real stories of firms operating in the other 99%
- The math on what unstructured admin time actually costs a litigation practice
- A framework for seeing AI as an operating system, not a chatbot
- Your first prompt using real case data

7:32 PM on a Tuesday

You're at your desk at 7:32 PM. Not because something went wrong – because this is how Tuesdays work.

Your paralegal left at 5:00. Since then, you've reviewed a medical chronology for the Hernandez PI case, drafted responses to two discovery requests, returned four client calls you didn't get to during the day, and opened the demand package you've been meaning to finalize for the Kowalski file. There are 23 unread emails in your inbox. One is from opposing counsel on the Rivera matter – you can tell from the preview it's a motion to compel, which means you need to calendar a response deadline before you forget. Your Clio dashboard shows 14 matters with activity due this week. You billed 5.2 hours today but worked at least 9.

You are not disorganized. Your Google reviews are 4.9 stars. Your case results are strong. Your clients trust you. The problem is not competence. The problem is that there is one of you and there are roughly forty things that need to happen every day, and about half of them don't require your legal judgment at all – they require your time, your attention, and your memory.

This is the bottleneck that every solo and small-firm attorney hits. Not a skills problem. A bandwidth problem. You are simultaneously the trial lawyer, the office manager, the marketing department, the billing clerk, and the IT administrator. You're caught between practicing law and running a business, and most days, the urgent wins because it's louder than the important.

Here is what has changed: **AI can now handle the second category.** Not the legal judgment. Not the client relationships. Not the courtroom advocacy. The other half – the operational backbone of your practice that eats your evenings and weekends.

But not the way most attorneys are using it.

The 1% Problem

You've probably used AI by now. According to the ABA's 2026 TechReport, 79% of lawyers use AI daily. Wolters Kluwer's 2026 Future Ready Lawyer survey found that 69% use general-purpose AI like ChatGPT, Claude, or Gemini – doubled from 31% the year before.

So adoption isn't the issue. The issue is scope.

If you're like most attorneys, here's how you use AI: you open Claude or ChatGPT, type a question, get an answer, and close the window. Maybe you ask it to draft a demand letter section. Summarize a deposition transcript. Explain a statute. Research a narrow legal question. Each interaction is a single prompt – you ask, it answers, then it forgets you exist.

That's about 1% of what AI can do. It's the equivalent of hiring a brilliant associate and only asking them to fetch coffee.

The other 99% is what happens when AI remembers your instructions across sessions. When it connects to your calendar, your email, your case management system. When it runs on a schedule without being asked. When it checks on your cases at 7:00 AM, triages your inbox at 8:00 AM, flags approaching deadlines at 8:15, reconciles your time entries at 9:00, and sends you a briefing before you've finished your first cup of coffee.

That's not a chatbot. That's an operating system for your practice.

AI as Chatbot (the 1%)	AI as Operating System (the 99%)
One question, one answer	Remembers your instructions, case details, and preferences across sessions
Starts fresh every time	Connects to Clio, Gmail, your calendar, your billing system
You initiate every interaction	Runs on a schedule – alerts you proactively when something needs attention
A tool you pick up when you think of it	A team member who shows up every morning before you do

The shift from the left column to the right column is not theoretical. Firms are doing it right now. Three stories make it concrete.

Story 1: The Two-Person Firm With 200 Clients

In early 2026, Zack Shapiro published an article about how he runs Raines LLP – a two-person firm serving over 200 startup and investor clients. No enterprise legal AI tools. No Harvey subscription. No CoCounsel license. Just Claude with custom skill files that encode his analytical frameworks, preferred formats, and a decade of practice judgment.

The article hit 7.5 million views. Roughly 1.5 million of those readers were lawyers. Artificial Lawyer called it part of a "DIY Revolution in Law." Fortune covered it. agents.law interviewed him.

Why did it resonate? Because Shapiro wasn't describing a chatbot. He was describing something every litigator understands: a trained associate who knows how you think.

When Shapiro asks Claude to analyze a term sheet, it doesn't start from zero. His skill files tell Claude: here's how I evaluate liquidation preferences, here's the format my clients expect, here's what I flag as non-standard, here's the voice I use in client memos. The AI applies his analytical framework every time – consistently, at 2:00 AM if that's when the deal needs to close.

The secret, as Shapiro put it, isn't the model's training data. It's the quality and specificity of the input. And lawyers – trained to write with precision, detail every relevant fact, and anticipate counterarguments – are naturally equipped to give AI exceptional instructions.

Two people. Two hundred clients. No enterprise software. Just encoded expertise running on a capable AI model.

What Shapiro built isn't a technology product. It's a practice methodology – the same kind of thing a seasoned litigator does when they develop their own deposition outline template, their own demand letter structure, their own intake evaluation checklist. The difference is that those frameworks now execute autonomously instead of sitting in a Word document or in the attorney's head. Artificial Lawyer's coverage framed it as a "DIY Revolution" – lawyers building their own tools with general-purpose AI rather than waiting for legal tech vendors to build (and charge for) a purpose-built product. That revolution is already underway, and the firms participating in it are pulling ahead.

Story 2: 32x Time Compression

Legal Paradox, a firm based in Mexico, went further. They rebuilt every production process around AI agent orchestration – becoming what they describe as the world's first Claude Code-native law firm.

The results aren't incremental. They're structural.

Their benchmark case: a CNBV authorization (Mexico's financial regulatory approval, roughly equivalent to an SEC registration in complexity). The traditional approach required 8 lawyers working for 8 months. Legal Paradox completed it with 1 partner directing 500 AI agents in 1 week.

That's 32x time compression and 256x the human capacity. Not because the legal analysis got simpler – because every repeatable step in the process was encoded as an instruction set, connected to the firm's data, and executed autonomously. The partner's job shifted from doing the work to reviewing it.

You don't have to reorganize your entire firm around AI agents to benefit from this. But Legal Paradox proves the ceiling is far higher than most attorneys assume. The question is not whether AI can handle complex legal workflows. It can. The question is how much of your practice you're willing to systematize.

Story 3: The Morning Briefing

This one is closer to home.

At Dodonai, we built 12 automated agents that run our marketing, sales, and operations. Each one is a skill file written in plain English, connected to our systems, and triggered on a schedule.

Here is what a typical weekday morning looks like – before anyone on the team is awake:

- **7:00 AM** – An agent checks the calendar, identifies every external meeting for the day, and researches each attendee: their company, their role, their history with us, their last five emails, recent news. It compiles a briefing and sends it to Slack.
- **8:00 AM** – An inbox triage agent reads every new email, classifies it (client, vendor, newsletter, outreach reply), applies labels, archives the noise, and surfaces anything urgent in a prioritized digest.
- **8:15 AM** – A prospect pipeline agent cross-references active deals with recent emails, calendar activity, and engagement signals. It flags deals that have gone cold and recommends next steps.
- **9:00 AM** – A churn detection agent compares yesterday's subscriber data against today's, identifies cancellations, pulls each customer's full history, and drafts personalized save emails. All the operator does is open Gmail, scan the draft, and hit send.
- **9:30 AM** – A lifecycle agent checks every new user's onboarding progress, identifies who's stuck, and queues nudge emails for the ones who haven't hit key milestones.

None of this was asked for that morning. It was set up once and now it runs. The operator wakes up to a briefing, not a to-do list.

Now translate that to a litigation practice:

- **7:00 AM** – Agent checks your calendar, pulls the file history for every matter you have a meeting on today, reviews recent correspondence, surfaces approaching deadlines, and sends you a case prep briefing before you leave the house.
- **8:00 AM** – Inbox triage classifies overnight emails: client communications, opposing counsel filings, court notices, vendor invoices. Urgent items (a motion to compel, a new deposition notice) get flagged immediately. Newsletters get archived.
- **8:15 AM** – Deadline monitor scans all active matters against your jurisdiction's rules. Anything within 30 days gets an urgent alert. 60 days gets an action flag. 90 days gets a planning note.
- **9:00 AM** – Billing reconciler compares your calendar and email activity against yesterday's time entries. It finds the 2.3 hours you spent on the Patel file that never made it into Clio.
- **9:30 AM** – Client update drafter generates monthly status updates for your 12 most active matters, pulling case activity from the past 30 days.

Same architecture. Same three layers: a skill file (instructions), a connection to your systems (Clio, Gmail, calendar), and a schedule (cron job). Different inputs, identical pattern.

This is what the 99% looks like.

The Pressure From Above

These three stories would be interesting but not urgent if the market weren't shifting underneath you. It is.

According to the ACC/Everlaw 2026 GenAI Survey, 64% of in-house legal teams now expect to depend less on outside counsel because of AI. They're building internal capability – reviewing contracts, monitoring compliance, drafting standard agreements – with tools that didn't exist two years ago.

This doesn't mean corporate clients will stop hiring litigators. It means the bar for what outside counsel provides is rising. Clients who can draft their own NDAs and review their own leases will expect their litigators to be faster, leaner, and more efficient. A Harvard Law analysis put it plainly: "Clients expect the benefits passed directly to them."

The firms that build AI capability will deliver faster turnaround, more thorough analysis, and competitive billing. The firms that don't will lose work to the ones that do – not because the law is different, but because the economics are.

Here's the uncomfortable math. According to Wolters Kluwer's 2026 survey, 62% of firms using AI report 6-20% weekly time savings. At a $350/hour blended rate, 4 hours per week recovered is $72,800 per year per attorney (Rankings.io). That's not a rounding error. That's a junior associate's salary recovered in billable capacity – or, for a solo, the difference between working until 7:32 PM and leaving at 5:30.

And the trend is accelerating. Solo tech spending is growing 56% annually (Clio Legal Trends Report). Immigration practices lead AI adoption at 47% (Wolters Kluwer). Harvey, Thomson Reuters, LexisNexis, and a wave of startups are all shipping autonomous agent workflows. Artificial Lawyer declared a "DIY Revolution in Law" – attorneys building their own tools with general-purpose AI instead of waiting for vendors to build it for them.

The question is no longer whether to use AI. It's whether you'll use 1% of it or 99%.

The Practice Without Systems

To understand where the 99% lives, you need to see where your time actually goes.

Sarah Chen runs Chen Law Group, a solo PI practice in Austin. Three-person team: Sarah, one paralegal, and a part-time legal assistant. About $800K/year in revenue. Forty active cases at any given time – 70% personal injury, 15% medical malpractice, 15% workers' comp. She uses Clio for case management, QuickBooks for accounting, and Dodonai Pro for medical record summaries and deposition digests.

Sarah is good at her job. Her settlements are strong. Her clients recommend her. But like every solo, she carries the full operational load of the practice on top of the legal work. Here is where her week goes outside of active case work:

Task	Hours/Week	Requires Sarah's Legal Judgment?	Cost to Outsource
Reviewing and responding to emails (non-case-specific)	4 hrs	Partially	$35-50/hr VA = $160
Drafting client status updates	3 hrs	Partially	$45/hr paralegal = $135
Calendaring deadlines from court notices and filings	2 hrs	No	$45/hr paralegal = $90

Task	Hours/Week	Requires Sarah's Legal Judgment?	Cost to Outsource
Billing: reviewing entries, catching unbilled time	2.5 hrs	Partially	$75/hr billing specialist = $187
Marketing: Google review responses, referral follow-ups	2 hrs	No	$50/hr freelancer = $100
Intake screening: evaluating new leads, conflict checks	2 hrs	Partially	Can't easily outsource
SOP documentation and staff training	1.5 hrs	Partially	$100/hr consultant = $150
Vendor management, supply ordering, office admin	1.5 hrs	No	$25/hr VA = $37
Competitive research, CLE tracking, bar compliance	1 hr	No	$75/hr analyst = $75
Preparing for meetings (pulling files, reviewing history)	2 hrs	Partially	Can't easily outsource
Total	**21.5 hrs**		**$934+/week**

Twenty-one and a half hours per week. That's more than half a standard workweek on tasks that sit outside her core legal work – case strategy, depositions, motions, negotiation, trial prep.

Outsourcing all of it would cost roughly $4,000 per month. On $800K in revenue, that's 6% going to administrative overhead. Most solos absorb the hours instead.

Here's the alternative: a well-configured AI system handles 60-70% of that list – not perfectly, not without review, but well enough to cut 12-15 hours per week of administrative work down to 3-4 hours of review and approval.

At Sarah's effective hourly rate ($800K revenue / 2,200 billable hours = ~$364/hour), those 12 recovered hours per week represent $4,368 in capacity – every single week. That's time she can bill, or time she can spend on case development, or time she can spend not working at 7:32 PM on a Tuesday.

Your Effective Hourly Rate

This math matters, so take 30 seconds to do it for yourself.

Your annual take-home (or gross revenue if solo): $_____

Hours you actually work per week (be honest – not billable hours, total hours): _____

Weeks per year (subtract vacation): _____

Your effective hourly rate: $_____ / (____ hours x ____ weeks) = $____ /hr

Every hour you spend formatting a client update letter, calendaring a deadline from a court notice, or writing a Google review response is that number – your hourly rate – applied to a task that doesn't require a law license.

AI doesn't replace you. It replaces the version of you that's doing $40/hour administrative work at $364/hour.

From Chatbot to Operating System

If you've used AI at all, you already understand the basic interaction: you type a question, it generates a response. That model works. It saves time on individual tasks. But it has three structural limitations that prevent it from changing how your practice operates:

1. It forgets. Every conversation starts fresh. You explain your practice, your case load, your preferences – every single time. It's like briefing a temp paralegal who leaves at 5:00 and has amnesia the next morning.

2. It's reactive. You have to remember to ask. If you don't open the app and type a prompt, nothing happens. The motion deadline approaching in 12 days? The client who hasn't heard from you in 6 weeks? The 3 unbilled hours from last Thursday? AI won't flag any of it unless you think to ask.

3. It's isolated. A chatbot doesn't know what's in your Clio system, your Gmail inbox, your calendar, or your bank account. It works with whatever you paste into the conversation window. That means you're the integration layer – copying data between systems, providing context manually, stitching outputs together yourself.

The shift from chatbot to operating system solves all three:

Skills replace amnesia. A skill file is a set of plain-English instructions – written once – that tells AI how to handle a specific task every time. Your demand letter structure. Your client update format. Your intake evaluation criteria. Write it once, and it applies it consistently forever. If you can write a procedure manual for a paralegal, you can write a skill.

Schedules replace memory. Instead of remembering to check deadlines, you set a schedule: "Every weekday at 7:00 AM, scan all active matters for deadlines within 90 days." It runs whether you remember or not. Like a recurring calendar event – for your AI.

Connections replace isolation. AI can read your Clio data, scan your Gmail, check your calendar, and query your billing system. It doesn't need you to copy-paste. The information flows automatically, and the AI works with the full picture.

Skills, connections, schedules. Three layers. That's the entire architecture. Chapter 3 walks through each one in detail, but the concept is this simple: encode your instructions, connect your tools, set a trigger.

The Difference in Practice

Here's a concrete example. Sarah Chen needs to evaluate a new PI intake – a potential client who was injured in a rideshare accident.

The 1% approach (chatbot):

```
Can you help me evaluate a personal injury case involving
a rideshare accident?
```

She gets back a generic overview of rideshare liability law, some factors to consider, and a list of questions she could ask during the consult. Mildly useful. She could have pulled the same information from a practice guide in roughly the same amount of time.

The 99% approach (operating system):

Sarah has an intake evaluation skill that she wrote once – encoding her case acceptance criteria, her firm's minimum case value threshold, her checklist for rideshare-specific issues (insurance coverage tiers, platform TOS implications, driver classification), and her preferred format for go/no-go memos. When a new lead comes in through her website, the intake agent:

1. Pulls the lead's form submission
2. Applies her evaluation criteria against the facts provided
3. Runs a conflict check against her Clio contact database
4. Estimates the case value range based on the injury type and jurisdiction
5. Drafts a go/no-go memo in her format, with a recommended next step

The memo is waiting in her inbox when she sits down with her coffee. She reviews it in 3 minutes, decides to take the consult, and the agent queues a confirmation email to the lead with her available times pulled from her calendar.

Same AI. Same underlying model. The difference is that the second version knows Sarah's practice, connects to her systems, and runs without being asked. That's the shift from tool to team member.

What AI Won't Do for Your Practice

Honest expectations build trust. Hype destroys it. Here is where AI stops – and where you should be skeptical of anyone who says otherwise.

It won't exercise legal judgment. Whether to accept a case, how to frame a theory of liability, when to settle versus try a case, how to handle a difficult client – these require the pattern recognition, ethical reasoning, and professional judgment that come from years of practice. AI can analyze data to inform your judgment. It cannot replace it.

It won't file anything without your review. In a well-designed system, every external action – every email sent, every document filed, every client communication – goes through a human approval gate. AI drafts. You decide. This is non-negotiable, and any system that skips this step is a malpractice risk.

It won't verify its own citations. This is the hallucination problem, and it's real. As of 2026, over 600 AI hallucination cases are on record, with 128 lawyers implicated (National Law Review). Courts have disqualified firms and referred attorneys to state bars. The failure wasn't using AI – it was failing to verify the output. Every AI-generated citation must be checked against primary sources before filing. Every one.

It won't replace your best paralegal's judgment. Your senior paralegal who knows which judges have standing orders about formatting, who remembers that opposing counsel on the Davis file always files late, who catches the inconsistency in the medical records that the doctors missed – that institutional knowledge is irreplaceable. AI helps you capture and systematize it. It doesn't replicate it.

It won't handle your most sensitive client moments. The call where you tell a client their case isn't as strong as they hoped. The conversation with a grieving family about next steps. The negotiation where reading the room matters more than reading the brief. These are human interactions that require empathy, presence, and professional courage. AI can prepare you for them. It can't do them.

The right mental model: AI is like hiring a very fast, very knowledgeable junior associate who works 24 hours a day for the cost of a monthly bar association dues payment. They're not as good as you at anything that requires judgment. But they're better than you at the tasks you shouldn't be doing in the first place – and they free you to spend your time where it actually matters.

Try This Now: Your First Prompt With Real Data

Stop reading and do this. It takes 15 minutes and it will change your frame of reference for the rest of this book.

Step 1: Pick one task from this list – choose the one that's been sitting on your to-do list longest:

- A client status update you owe someone
- A Google review you've been meaning to respond to
- A demand letter section you've been drafting in your head
- An intake evaluation for a lead that came in this week
- A case summary for a referral source

Step 2: Open your AI assistant (Claude, ChatGPT, or any capable model) and use this prompt template. Fill in your details:

```
You are a litigation practice assistant for a [practice area] firm.
My firm is [name], located in [city]. I handle approximately [X]
active cases across [practice areas with rough percentages].

I use [case management system] for matter management, [email
provider] for correspondence, and [billing system] for time
tracking. I need help with the administrative and analytical
work that sits around my active caseload.

My team is [brief description]. My communication style with
clients is [professional/direct/warm  --  pick one].

Before responding, confirm your understanding and ask one
clarifying question if anything is ambiguous. Then provide
the output.

Today I need to: [specific task]
```

Here's what it looks like filled in for Sarah Chen:

```
You are a litigation practice assistant for a personal injury firm.
My firm is Chen Law Group, located in Austin, TX. I handle
approximately 40 active cases: 70% personal injury (auto accidents,
premises liability, slip and fall), 15% medical malpractice,
15% workers' compensation.

I use Clio for matter management, Gmail for correspondence, and
QuickBooks for accounting. I use Dodonai Pro for medical record
summaries and deposition digests. I need help with the
administrative and analytical work that sits around my active
caseload.

My team is myself (attorney), one full-time paralegal, and one
part-time legal assistant. My communication style with clients
is direct and empathetic  --  I explain things clearly without
legal jargon unless necessary.

Before responding, confirm your understanding and ask one
clarifying question if anything is ambiguous. Then provide
```

```
the output.

Today I need to: draft a 60-day status update letter to Maria
Hernandez regarding her PI case. She was rear-ended on I-35
on November 12, 2025. She's completed 4 months of chiropractic
treatment and 2 months of physical therapy. We're waiting on
the final treatment records from Austin Spine & Rehab and her
MRI results from St. David's imaging. The SOL is November 12,
2027. I want to reassure her that the case is progressing,
explain why we're waiting on records before making a demand,
and let her know the next milestone. Sign as "Sarah Chen."
```

Step 3: Time yourself. How long did it take to write the prompt, review the output, and edit it into something you'd actually send? Compare that to writing from scratch.

Step 4: Iterate once. Take the AI's output and give it one round of feedback:

```
Good start. A few adjustments:
- The opening paragraph is too formal  --  I want Maria to feel
  like I'm talking to her, not writing a brief
- Add a sentence about how the chiropractic records show
  consistent treatment compliance, which strengthens the case
- Cut the paragraph about the legal process  --  she doesn't
  need that level of detail
- End with a specific next step and timeline, not a vague
  "we'll be in touch"
```

Notice how the second version is closer to what you'd actually send. This draft-then-refine cycle is how you'll work with AI 90% of the time. First draft, feedback, refined draft. It's faster than writing from scratch and usually produces better output than either you or the AI would create alone.

Step 5: Assess the quality. Was the output 80% usable? 90%? That's the bar – not perfection, but a fast first draft you can refine in minutes instead of writing from a blank screen. If you rewrote more than half of it, your prompt needed more context. Go back and add detail. The difference between a vague prompt and a specific one is the difference between a generic form letter and a document that sounds like you wrote it.

The Time Audit

One more exercise. This one takes 10 minutes and produces a map of where AI fits into your practice.

Pull out a piece of paper or open a blank document. Write down every recurring task you handle each week that isn't active legal work – motions, depositions, court appearances, client counseling, negotiations. Focus on the operational layer: emails, billing, calendaring, admin, marketing, intake.

For each task, note three things:

1. **Hours per week** you spend on it
2. **Does it require your legal judgment?** (Honest answer – not "could a paralegal do it" but "does this actually require a JD?")
3. **What would you pay someone to do it?**

Total the hours. For most solos and small-firm attorneys, the number lands between 15 and 25 hours per week. That's the gap between what your case management system handles and what's still sitting on your shoulders. That gap is where the 99% lives.

Keep your list. You'll come back to it. In Chapter 7, we'll map your highest-impact tasks to specific agent designs. For now, just seeing the number is the point – most attorneys are surprised by how much of their week goes to work that doesn't require a bar card.

In Chapter 2, we'll walk through what your practice looks like when that gap is filled – not with a larger staff, but with AI agents that run on schedule, feed each other's outputs, and give you back your evenings. You'll see actual outputs: a meeting intel briefing, a case health digest, an inbox triage summary. The kind of documents you'd open at 8:00 AM and think, "I didn't ask for this, but I'm glad it's here."

The shift from 1% to 99% doesn't happen in a day. It happens one workflow at a time. The first one takes the longest. Each one after that is faster, because the pattern is always the same: encode your instructions, connect your tools, set a trigger.

Start with one. The one you keep putting off. Write the instructions for how a human would do it.

That becomes your first skill. Everything else builds from there.

End-of-Chapter Checklist

Before moving to Chapter 2, confirm you've:

- ☐ Calculated your effective hourly rate
- ☐ Completed the time audit – identified your weekly administrative hours
- ☐ Run at least one real practice task through an AI assistant using the prompt template
- ☐ Experienced the difference between a vague prompt and a context-rich prompt

- ☐ Identified 3-5 recurring tasks where AI could handle the first draft or analysis

---

Quick Reference Card: Chapter 1

Concept	Key Takeaway
The 1% Problem	Most lawyers use AI for single prompts. The other 99% is when it runs your practice on autopilot.
Three Stories	Shapiro (2-person firm, 200 clients), Legal Paradox (32x compression), Dodonai (12 agents before breakfast)
Client Pressure	64% of in-house teams building AI internally. Your clients expect the benefits passed to them.
The Admin Tax	15-25 hours/week of non-legal work at your hourly rate. That's where AI lives.
Chatbot vs. Operating System	Skills replace amnesia. Schedules replace memory. Connections replace isolation.
The Architecture	Three layers: skill files (instructions), connections (your tools), schedules (automation triggers).
Honest Limits	AI won't exercise legal judgment, verify citations, or replace your best paralegal's instincts. Always verify.
First Prompt	Context-rich prompts produce practice-ready output. Vague prompts produce generic filler.

Chapter 2: What AI-Native Looks Like

What You'll Walk Away With:

- A minute-by-minute walkthrough of an AI-native morning at a litigation practice
- Three real sample outputs – a meeting intel briefing, a case health digest, and an inbox triage summary – formatted exactly as you'd see them
- The compound effect: how three agents working together produce more value than three agents working alone
- The math on recovered revenue from automated time capture and proactive case management
- A clear picture of what "AI-native" means in practice, not in theory

---

What Your AI Did While You Slept

It's 6:47 AM on a Wednesday. Your alarm hasn't gone off yet. You're reaching for your phone the way you always do – half awake, one eye open, checking for fires.

But instead of 23 unread emails and a vague sense of dread, here's what's waiting on your phone:

A Slack message from your AI system, timestamped 7:02 AM:

> **Morning Briefing – Apr 2, 2026**
>
> 3 meetings today. Briefings ready for each.
>
> **Inbox**: 31 new messages triaged. 2 urgent (motion to compel on Rivera – response due Apr 16; court notice on Hernandez – hearing rescheduled to May 8). 6 action items. 23 archived.
>
> **Deadlines**: Kowalski SOL in 47 days – demand package still pending final medical records. Patel discovery responses due in 12 days – draft started, needs your review.
>
> **Billing**: 3.1 unbilled hours identified from yesterday – 1.4 hrs on Hernandez (phone call + record review), 0.9 hrs on Chen estate referral (emails), 0.8 hrs on Rivera (opposing counsel correspondence).
>
> **Client updates**: 4 matters past 30-day communication mark. Drafts queued.

You didn't ask for any of this. You set it up once – weeks ago – and now it runs every morning before you're awake. The meeting prep is already done. The deadlines are tracked. The billing gaps are caught. The client communications are drafted.

You get out of bed, pour your coffee, and start reviewing instead of researching.

This is what an AI-native practice looks like. Not a chatbot you open when you think of it. A system that works while you sleep and hands you a briefing when you wake up.

A Morning at Chen Law Group

Chapter 1 introduced Sarah Chen – solo PI practice in Austin, forty active cases, one paralegal, one part-time legal assistant. She uses Clio for case management, Gmail for correspondence, QuickBooks for accounting, and Dodonai Pro for medical record summaries and deposition digests.

Six weeks ago, Sarah configured seven AI agents. Each one is a skill file – plain-English instructions – connected to her existing tools and running on a schedule. Here's what her Wednesday morning looks like now.

Time	Agent	What It Produces
7:00 AM	**Meeting Intel**	Briefings for each meeting: case file summary, recent correspondence, opposing counsel history, approaching deadlines, client communication log
7:30 AM	**Deadline Monitor**	Scans all 40 active matters against Texas rules. Flags SOL dates, discovery cutoffs, motion deadlines, and hearing dates at 30/60/90-day intervals
8:00 AM	**Inbox Triage**	Classifies overnight emails: client, opposing counsel, court notice, vendor, newsletter. Urgent items flagged. Newsletters archived. Digest delivered
8:30 AM	**Billing Reconciler**	Cross-references yesterday's calendar entries, emails, and document activity against Clio time entries. Identifies unbilled hours
9:00 AM	**Case Health Dashboard**	Weekly rollup: case aging, settlement pipeline, matters with no activity in 14+ days, retainer balances running low
9:30 AM	**Client Update Drafter**	Generates status update drafts for any matter past its 30-day communication mark, pulling case activity from Clio
On demand	**Intake Evaluator**	When a new lead hits her website form, scores it against her case acceptance criteria and drafts a go/no-go memo

Sarah's paralegal, Maria, arrives at 8:30. By then, Sarah has already reviewed the meeting briefings, scanned the deadline alerts, and approved three time entries

the billing reconciler flagged. She didn't write a single email, pull a single file, or check a single deadline manually.

Her morning used to start with 45 minutes of triage – scanning emails, pulling up files for today's meetings, mentally running through her active cases to remember what's due. Now it starts with 10 minutes of review.

That difference – 35 minutes every morning, five days a week – adds up to nearly 150 hours per year. At her effective rate of $364/hour, that's $54,600 in recovered capacity. And that's just the morning.

The Outputs: What You'd Actually See

The agents are only as useful as what they produce. Generic summaries don't cut it. You need outputs specific enough to act on – with real case names, real deadlines, and real recommendations.

Here's what Sarah's agents actually produce. These are modeled on the real outputs that Dodonai's own 12-agent system generates daily for marketing, sales, and customer operations – translated into the legal equivalent.

Output 1: Meeting Intel Briefing

Sarah has a 10:00 AM meeting with Maria Hernandez – a client update on her PI case. At 7:02 AM, the meeting intel agent checked Sarah's calendar, pulled the case file from Clio, searched Gmail for recent correspondence, and compiled this briefing:

Meeting Intel Briefing – Apr 2, 2026

10:00 AM – Client Update: Hernandez v. Martinez

> **Duration**: 30 minutes | **Location**: Office (in-person)

Case Status

Field	Value
Matter	Hernandez v. Martinez (Auto PI)
Filed	Not yet – pre-litigation

Field	Value
SOL	**November 12, 2027** (589 days remaining)
Case Value Estimate	$175,000 - $250,000
Phase	Treatment complete. Waiting on final records.
Retainer Balance	$2,400 (of $5,000 original)
Last Client Contact	March 4, 2026 (29 days ago)

Treatment Summary (from Dodonai Chronology)

- **Chiropractic**: 4 months at Austin Spine & Rehab (48 visits, discharged Feb 28)
- **Physical therapy**: 8 weeks at Capital PT (16 sessions, discharged Mar 15)
- **Imaging**: MRI at St. David's (Jan 14) – L4-L5 disc herniation confirmed
- **Outstanding**: Final discharge summary from Austin Spine & Rehab (requested Mar 18, follow-up sent Mar 28). St. David's imaging report received Mar 30.

Recent Correspondence

- **Mar 28** – Maria sent follow-up to Austin Spine & Rehab requesting discharge summary (3rd request)
- **Mar 30** – St. David's imaging report received and uploaded to Clio
- **Mar 4** – Last call with Maria. She asked about timeline to settlement. Sarah explained we're waiting on final records before making demand.

Opposing Counsel Activity

- **Counsel**: Robert Fitzgerald, Fitzgerald & Associates (for State Farm)
- **Last contact**: Feb 12 – standard preservation letter exchange
- **Notes**: Fitzgerald typically responds within 10 business days. Known to lowball first offers on disc herniations. Usual pattern: first offer 40-50% of demand.

Talk Track

1. **Update on records**: St. David's imaging received. Still waiting on Austin Spine & Rehab discharge summary – follow up again if not received by Apr 7.
2. **Timeline**: Once final records arrive, demand package assembly takes ~2 weeks. Expect to send demand by late April if records arrive this week.
3. **Reassurance**: Treatment compliance is strong (48 chiropractic + 16 PT sessions, no gaps). MRI confirms disc herniation. This strengthens the case.
4. **Retainer**: Balance at $2,400. No additional retainer needed at this stage.
5. **Maria may ask**: "When will I get my money?" Be direct about the timeline: demand → 30 days for response → negotiation → settlement or litigation decision. Best case: 3-4 months. Be honest that it could be longer if State Farm lowballs.

That briefing took the agent about 90 seconds to compile. It would have taken Sarah 15-20 minutes to pull the same information manually – opening Clio, searching Gmail, reviewing the medical chronology, checking the calendar for last contact.

More importantly, the talk track at the bottom is specific to this meeting. It's not generic advice about client communication. It anticipates the question Maria is likely to ask, based on her last conversation. That's the difference between a tool that retrieves data and a system that prepares you for the meeting.

Output 2: Case Health Digest

Every Monday morning, Sarah's case health agent compiles a practice-wide overview. It connects to Clio for matter data, QuickBooks for financials, and cross-references the outputs from the deadline monitor and billing reconciler. Here's what it produces:

Case Health Digest – Week of Mar 31, 2026

> **Generated**: Monday, Mar 31, 2026 at 7:45 AM

Practice Dashboard

Metric	This Week	Last Week	Change
Active Matters	40	41	-1 (Orozco settled)
Pre-Litigation	24	25	-1
In Litigation	12	12	–
In Settlement Negotiation	4	4	–
New Intakes (7 days)	2	1	+1
Outstanding A/R	$47,200	$52,800	-$5,600 (Gutierrez paid)
Avg Case Age	8.2 months	8.0 months	+0.2

Deadline Alerts

Priority	Matter	Deadline	Days Out	Action
URGENT	Kowalski v. Apex Logistics	SOL expires Jun 18	**47 days**	Demand package pending final med records. File suit if no settlement by May 30.
URGENT	Patel v. St. David's Medical	Discovery responses due	**12 days**	Draft started Mar 25. Needs Sarah's review by Apr 7.
Action	Rivera v. Citywide Transport	Response to Motion to Compel due Apr 16	16 days	Motion received Apr 1. Calendar response deadline.
Action	Jackson v. Metro Rideshare	Deposition of plaintiff	22 days	Scheduled Apr 23. Prep outline needed by Apr 18.
Planning	Williams v. HEB	Mediation	58 days	Scheduled May 28. Mediation brief due May 21.
Planning	Torres v. Austin ISD	Expert designation deadline	71 days	Jun 10. Dr. Patel (biomechanics) retained. Report not yet received.

Stale Matters (no activity > 14 days)

Matter	Last Activity	Days Inactive	Recommended Action
Brooks v. Lonestar Insurance	Mar 12	19 days	Waiting on IME scheduling. Follow up with defense counsel.
Nguyen v. Amazon Logistics	Mar 14	17 days	Client treating. Check treatment compliance – next PT visit was scheduled Mar 21.
Foster v. Whataburger	Mar 10	21 days	Demand sent Feb 28. No response from adjuster. Follow up or escalate.

Billing Summary

Metric	Value
Hours billed last week	31.2
Unbilled hours caught by reconciler	4.7
Revenue from caught hours (at blended rate)	$1,711
YTD unbilled hours recovered	58.3
YTD revenue recovered	$21,222

Client Communication Status

Matter	Last Contact	Days Since	Draft Status
Hernandez v. Martinez	Mar 4	27 days	Update draft queued (meeting today)
Brooks v. Lonestar	Mar 8	23 days	Update draft queued
Nguyen v. Amazon	Mar 14	17 days	Update draft queued
Chen v. AT&T (workers' comp)	Mar 2	29 days	Update draft queued

Action Summary

Priority	Count	Details
Urgent deadlines	2	Kowalski SOL (47d), Patel discovery (12d)
Motion responses	1	Rivera motion to compel (16d)
Stale matters needing attention	3	Brooks, Nguyen, Foster
Client updates past 30 days	4	Drafts queued for review
Unbilled time to enter	4.7 hrs	$1,711 at blended rate

That digest would take 2-3 hours to compile manually – if Sarah did it at all. Most solos don't. They carry the case status in their heads, check deadlines reactively, and catch billing gaps only during month-end reconciliation (if then).

The billing reconciler line is worth pausing on. In just the first quarter, it caught 58.3 hours that Sarah worked but didn't bill – mostly phone calls, email correspondence, and record review sessions that didn't make it into Clio. At her blended rate, that's $21,222 in recovered revenue. Annually, that projects to over $80,000.

That number tracks with industry data. BigHand SmartTime reports that AI-assisted time capture recovers 8-10 extra billable hours per month per attorney. Billables AI found that firms using automated time reconstruction capture 10-30% more billable time than those relying on manual entry alone. One firm reported recouping their entire investment in two days.

Sarah isn't capturing more hours because she's working more. She's capturing hours she was already working but failing to record.

Output 3: Inbox Triage Summary

Twice daily – at 8:00 AM and 2:00 PM – Sarah's inbox agent reads every new email, classifies it, applies labels, archives the noise, and produces a prioritized digest. Here's what the morning triage looks like:

Inbox Digest – 2026-04-02 (Morning)

Summary

Metric	Count
New messages triaged	27
Urgent (respond today)	2
Action (respond this week)	4
FYI (archived)	18
Court notices (calendared)	1
Needs review	2

Urgent – Respond Today

1. Robert Fitzgerald rfitzgerald@fitzlaw.com

> **Subject**: Hernandez v. Martinez – Motion to Compel IME **Category**: Opposing Counsel | **Received**: Apr 1, 11:47 PM **Matter**: Hernandez v. Martinez **Action**: Motion to compel IME filed. Response deadline computed: **April 16, 2026** (15 days per TRCP Rule 196). Deadline added to Clio. **Snippet**: Please find attached Defendant's Motion to Compel Independent Medical Examination of Plaintiff Maria Hernandez pursuant to TRCP Rule 204.1…

2. Judge Patricia Morales clerk@traviscounty.courts

> **Subject**: Patel v. St. David's Medical – Hearing Rescheduled **Category**: Court Notice | **Received**: Apr 2, 6:12 AM **Matter**: Patel v. St. David's Medical **Action**: Status conference moved from Apr 18 to **May 8 at 2:00 PM**. Old date removed, new date calendared in Clio. **Snippet**: NOTICE: The status conference in Cause No. 2025-CI-48291 previously set for April 18, 2026, is hereby RESCHEDULED to May 8, 2026 at 2:00 PM…

Action – Respond This Week

1. Maria Hernandez maria.hernandez@gmail.com

> **Subject**: RE: Case Update **Category**: Client | **Received**: Apr 1, 4:22 PM **Matter**: Hernandez v. Martinez **Snippet**: Hi Sarah, thank you for the update. I wanted to ask – my employer is asking when I'll have information about the case outcome because they need it for their insurance. What should I tell them?

2. Dr. James Whitfield records@austinspinerehab.com

> **Subject**: RE: Records Request – Hernandez **Category**: Medical Provider | **Received**: Apr 1, 3:15 PM **Matter**: Hernandez v. Martinez **Snippet**: We received your request. The discharge summary is being finalized by Dr. Whitfield and should be available by April 7. We'll send it via the secure portal.

3. Maria Reyes maria@chenlawgroup.com

> **Subject**: Patel Discovery – Draft Ready for Review **Category**: Internal | **Received**: Apr 1, 5:30 PM **Matter**: Patel v. St. David's Medical **Snippet**: Sarah, I've completed the first draft of the discovery responses for Patel. They're in the Clio documents folder. Need your review by Thursday.

4. Tom Reeves (State Farm Adjuster) treeves@statefarm.com

> **Subject**: Foster v. Whataburger – Demand Response **Category**: Insurance/Adjuster | **Received**: Apr 2, 7:45 AM **Matter**: Foster v. Whataburger **Snippet**: Ms. Chen, we've reviewed your demand package dated February 28. We'd like to schedule a call to discuss our evaluation of the claim. Please advise on your availability next week.

FYI – Archived (18)

- Texas Bar Association – April CLE Calendar
- Clio – Product Update: New Billing Features
- LawPay – Monthly Statement
- Westlaw – Legal News Digest (x3)
- Austin Bar Association – Upcoming Events
- Google – New review on Google Business Profile (5 stars)
- *(+10 newsletters, vendor emails, and automated notifications)*

Needs Review (2)

1. **Unknown sender** legal@mccarthy-firm.com – "Referral from Judge Navarro" – possible new intake, needs manual classification
2. **USPS Informed Delivery** – certified mail notification – may be a court filing, needs manual check

Before this system, Sarah spent 30-40 minutes every morning scanning her inbox. She'd open each email, decide what it was, figure out which case it belonged to, determine whether it needed immediate action, and then either respond or add it to her mental to-do list. The court notice about Patel would sit unread until she got to it. The deadline from the motion to compel might not get calendared until later that day – or the next morning if she was in depositions.

Now the urgent items are surfaced immediately. The deadlines are computed and calendared before she reads the digest. The 18 newsletters and vendor emails never reach her attention. And the two ambiguous messages are flagged for manual review rather than silently buried in the middle of 27 emails.

The numbers add up. If inbox triage saves 30 minutes twice a day – which is conservative for a practice receiving 25-40 emails daily – that's 5 hours per week, or 250 hours per year. At $364/hour, that's $91,000 in capacity freed up for billable work or, more realistically, for going home on time.

The Compound Effect

Each of those agents is useful on its own. The meeting intel briefing saves 15 minutes of prep. The inbox triage saves 30 minutes of scanning. The deadline monitor prevents missed deadlines. The billing reconciler catches unbilled time.

But the real leverage shows up when agents feed each other. The total value exceeds the sum of the parts.

How Three Agents Become a Closed Loop

Watch what happens when the inbox triage agent processes that motion to compel from Robert Fitzgerald on the Hernandez case:

Step 1 – Inbox Triage Agent (8:00 AM): Reads the email. Classifies it as opposing counsel correspondence. Identifies the matter (Hernandez v. Martinez). Extracts the document type (Motion to Compel IME). Computes the response deadline (April 16, per TRCP Rule 196). Calendars the deadline in Clio. Flags it as urgent in Sarah's morning digest.

Step 2 – Deadline Monitor Agent (feeds from Clio, 7:30 AM next morning): Picks up the new April 16 deadline. Cross-references it against Sarah's existing calendar. Flags a conflict – she has a deposition on April 14 and 15. Recommends completing the response brief by April 11 to allow buffer. Adds a "Response draft due" task for April 11.

Step 3 – Client Update Agent (triggered by new motion filing): Detects that a substantive motion has been filed on Hernandez. Drafts a client notification email to Maria Hernandez explaining that opposing counsel has filed a motion requesting an independent medical examination, what that means, and what the next steps are. The draft is waiting in Sarah's Gmail drafts folder for review.

Step 4 – Billing Reconciler (next morning): Identifies that Sarah spent 45 minutes reviewing the motion to compel and researching IME case law that evening. She didn't enter it in Clio. The reconciler flags it: "0.8 hrs on Hernandez – motion review, Apr 2."

Step 5 – Case Health Dashboard (Monday rollup): The Hernandez matter now shows the new deadline, the pending response, the updated case status, and the time entry. The client communication log reflects the update email (once Sarah sends it). Everything is current.

Five steps. Four agents. Sarah's involvement: review the client email draft (2 minutes), approve the time entry (30 seconds), review the response brief deadline (already handled). Total active time: under 5 minutes for a sequence that would have taken 45 minutes to handle manually – if she caught every piece.

That's the compound effect. No single agent in that chain is particularly sophisticated. The inbox triage is classifying emails. The deadline monitor is checking dates. The client update drafter is writing a letter. The billing reconciler is comparing calendars to time sheets. But connected together, running on schedule, feeding each other's outputs, they create a closed loop that handles the entire lifecycle of an incoming filing – from receipt to calendaring to client communication to billing – without Sarah having to orchestrate any of it.

Three Agents vs. Three Times One Agent

The math is not linear. Here's why.

One agent (inbox triage alone) saves Sarah 30 minutes a day on email scanning. Useful, but limited – she still has to manually calendar deadlines, draft client updates, and enter time.

Add a second agent (deadline monitor). Now the deadlines extracted by the inbox triage agent are automatically calendared and tracked. Sarah no longer needs to read court notices for dates and manually enter them. The two agents together save more time than either one alone, because the second agent eliminates work that the first agent's output used to create.

Add a third agent (client update drafter). Now the case activity detected by the first two agents automatically triggers client communication. Sarah doesn't need to remember which clients are due for updates – the system tracks it and drafts the emails.

The third agent doesn't just save time on drafting. It eliminates the cognitive overhead of tracking which clients need to hear from you. That overhead is invisible – it's the background process running in your head while you're in a deposition, reminding you that Maria Hernandez hasn't heard from you in three weeks. When the system handles it, that mental bandwidth frees up for the work that actually requires your attention.

This is the same pattern that Dodonai uses internally. Twelve agents running on schedule, each one's output feeding the next. Churn detection flags a cancellation. The outreach composer drafts a save email. The reply monitor watches for a response. When the customer comes back, the weekly health digest reports the save and attributes it to the outreach. Closed loop. The operator spent 30 seconds reviewing a draft. The system handled everything else.

For a law firm, the pattern is identical: case event detected, deadline computed, client notified, time captured, dashboard updated. Same three layers – skills, connections, schedules. Different inputs. Identical architecture.

The Numbers Behind AI-Native

The sample outputs make the concept concrete. The numbers make the business case.

Time Recovery

The 2026 Wolters Kluwer Future Ready Lawyer survey found that 62% of firms using AI report 6-20% weekly time savings. For an attorney working 50 hours per week, 10% recovery is 5 hours. At a $350/hour blended rate, that's $9,100 per month – or $109,200 per year (Rankings.io puts the annual figure at $72,800 per attorney at slightly lower utilization assumptions).

Sarah's seven agents target the operational layer specifically:

Agent	Estimated Weekly Time Saved	Annual Value (at $364/hr)
Meeting Intel	1.5 hrs (15 min/meeting x 6 meetings)	$28,392
Deadline Monitor	1.0 hr (eliminated manual calendar checks)	$18,928
Inbox Triage	5.0 hrs (30 min x 2/day x 5 days)	$94,640
Billing Reconciler	1.5 hrs (catches ~5 hrs/week unbilled)	$28,392
Case Health Dashboard	2.0 hrs (replaced manual case reviews)	$37,856

Agent	Estimated Weekly Time Saved	Annual Value (at $364/hr)
Client Update Drafter	2.5 hrs (drafts for 8-10 matters/month)	$47,320
Intake Evaluator	1.0 hr (2-3 intakes/week)	$18,928
Total	**14.5 hrs/week**	**$274,456**

The $274,000 figure represents capacity, not cash. Sarah won't bill all of those hours. But even converting a third of that recovered time to billable work – 5 additional hours per week – adds $94,640 in annual revenue. For an $800K practice, that's an 11.8% revenue increase from operational automation alone, before any improvement in case outcomes or client service.

Revenue Capture

The billing reconciler deserves its own line item. Attorneys are notoriously bad at capturing all their time. The work happens – the phone call, the email exchange, the quick records review – but the time entry doesn't. It falls through the cracks because billing is a retrospective task competing with everything else.

BigHand SmartTime data shows that AI-assisted time capture recovers 8-10 extra billable hours per month per attorney. Billables AI and Smokeball report that firms using automated time reconstruction capture 10-30% more billable time than manual entry.

For Sarah, 8 hours per month at $364/hour is $34,944 per year in revenue that was previously earned but never billed. That's money she already worked for. The billing reconciler just makes sure it doesn't disappear.

The Quality Dividend

Not everything AI-native produces shows up in a spreadsheet. Some of it shows up in client satisfaction, case outcomes, and risk avoidance.

Deadlines don't get missed. The most common malpractice claim in litigation is a missed deadline – a blown SOL, a late discovery response, a forgotten motion deadline. When the deadline monitor scans every active matter every morning against jurisdiction-specific rules, and the inbox triage agent automatically calendars new deadlines from court notices, the risk drops from "human memory plus a tickler system" to "automated monitoring plus human verification."

Clients hear from you more often. The number one client complaint about attorneys (per Clio's Legal Trends data, year after year) is lack of communication. Not bad outcomes. Not high fees. Silence. When the client update drafter automatically generates status updates for any matter past its 30-day

communication mark, every client stays informed without Sarah having to remember who's due for an update.

Nothing falls through the cracks. The stale matter alerts on the case health digest caught the Foster v. Whataburger demand response that had gone unanswered for 33 days. Without the alert, Sarah might not have noticed until the adjuster called – or until the client called asking why nothing was happening. With the alert, she followed up within 24 hours, and the adjuster responded with a settlement offer two days later.

These are the outcomes that don't show up in a time-savings calculation but compound over years into a stronger practice, better client relationships, and fewer 3 AM anxiety spirals about what you might have forgotten.

What This Isn't

A necessary correction before we go further, because the sample outputs above might make it look too clean.

The outputs require review. Every meeting intel briefing, every client update draft, every deadline computation needs a human attorney to verify it before anything goes to a client, a court, or opposing counsel. The system produces a first draft and a best guess. You provide the judgment. Sarah spends 10 minutes reviewing the morning briefing, not 0 minutes. The motion to compel deadline was computed as April 16 – but Sarah confirms it against the local rules herself, because AI doesn't carry your bar license.

The setup takes work. Sarah's seven agents didn't appear overnight. Each one required writing a skill file, connecting to her systems, testing the output, refining the instructions, and running it for a week or two before trusting it. The first agent (deadline monitor) took a full afternoon. The seventh (intake evaluator) took two hours, because by then the pattern was familiar. Chapter 7 walks through building your first agent step by step.

It won't handle the hard parts. The meeting intel briefing gives Sarah everything she needs to walk into the Hernandez meeting prepared. It won't tell her how to handle Maria's anxiety about the timeline, whether to recommend accepting a low settlement offer, or how to frame the IME motion in a way that reassures rather than alarms. That's Sarah's job. The AI prepared the data. The attorney handles the client.

Some mornings will have errors. The inbox triage agent will occasionally miscategorize an email. The deadline monitor will flag a deadline that doesn't apply because the case settled last week. The billing reconciler will identify "unbilled

time" that was actually a personal appointment on Sarah's calendar. These are the rough edges of any automated system, and they're why human review is non-negotiable. The goal isn't perfection. The goal is that reviewing and correcting a mostly-right draft takes far less time than building it from scratch.

The Shift in Your Day

Here's the part that doesn't fit in a table. When you stop spending your mornings on triage and start spending them on review, the entire texture of your workday changes.

Sarah used to arrive at the office and immediately start reacting. Open email. Scan for fires. Pull files. Check deadlines. Return calls. Draft responses. By 10:00 AM, she'd handled a dozen urgent-but-not-important tasks and hadn't touched any of her actual legal work.

Now she arrives, opens the briefing, reviews the flagged items, approves or edits the drafts, and by 9:00 AM she's working on the Kowalski demand package – the $250,000 case that's been sitting on her desk for two weeks because there was never a clean block of time to focus on it.

That's the shift. Not "AI does everything." AI handles the operational backbone – the monitoring, the triaging, the drafting, the tracking – so you can spend your best hours on the work that requires your judgment, your expertise, and your license. The depositions. The negotiations. The case strategy. The client relationships.

The attorneys who've made this shift describe it the same way: "I feel like I have a chief of staff." Not a chatbot. Not a tool. A team member who shows up every morning before they do, handles the prep work, and hands them a briefing instead of a to-do list.

Try This Now: Map Your Morning

You don't need to build seven agents to start. You need to see your morning clearly.

Tomorrow morning, before you do anything else, write down every task you perform between arriving at the office and starting your first piece of substantive legal work. Be specific:

- Opened email, scanned for urgent items (12 minutes)
- Pulled the Johnson file to prep for 10:00 AM meeting (8 minutes)

- Checked Clio for upcoming deadlines this week (6 minutes)
- Reviewed paralegal's draft of discovery responses (15 minutes)
- Returned two client voicemails (14 minutes)
- Entered yesterday's time that I forgot to log (10 minutes)

Total those minutes. For most attorneys, the number falls between 45 minutes and 90 minutes – all before you touch any substantive work.

Now look at each item and ask: **Could an AI agent have done the first draft of this, or surfaced this information proactively, so I just had to review instead of build from scratch?**

Circle the ones where the answer is yes. Those are your agents. The biggest ones – the ones that eat the most time every single morning – are where you start.

Keep this list. In Chapter 7, you'll pick one and build it.

What Comes Next

You've seen what an AI-native morning produces. The meeting briefing that's waiting before your first cup of coffee. The case health digest that catches the deadline you would have missed. The inbox triage that turns 27 emails into 6 action items. The billing reconciler that finds $80,000 in unbilled time per year. The compound effect of agents feeding each other in closed loops.

This isn't theoretical. It's the same architecture – skills, connections, schedules – running in production right now for Dodonai's own operations, translated into the tools and workflows of a litigation practice.

But seeing the outputs is different from understanding the architecture. You know what these agents produce. You don't yet know how they're built – what a skill file actually looks like, how the connections to Clio and Gmail work, what it means to put an agent on a schedule.

Chapter 3 opens the hood. Three layers. No computer science required. If you can write instructions for a paralegal, you can write a skill.

End-of-Chapter Checklist

Before moving to Chapter 3, confirm you've:

- ☐ Reviewed the three sample outputs and identified which ones would be most valuable for your practice

- ☐ Mapped your morning – identified the pre-work triage tasks that eat 45-90 minutes daily
- ☐ Calculated the annual value of those triage hours at your effective hourly rate
- ☐ Identified 2-3 tasks where "review a draft" would be faster than "build from scratch"
- ☐ Understood the compound effect – how connected agents create more value than isolated ones

---

Quick Reference Card: Chapter 2

Concept	Key Takeaway
The Overnight Story	AI agents run before you wake up. You start the day with a briefing, not a to-do list.
The AI-Native Morning	7 agents, 7:00-9:30 AM: meeting intel, deadlines, inbox triage, billing reconciler, case health, client updates, intake evaluation.
Sample Outputs	Meeting intel briefings, case health digests, and inbox triage summaries – specific enough to act on, not generic summaries.
The Compound Effect	Connected agents create closed loops: filing received, deadline calendared, client notified, time captured, dashboard updated. Three agents > 3x one agent.
Time Recovery	62% of firms report 6-20% time savings. AI billing capture recovers 8-10 hours/month per attorney. $72,800/year in recovered capacity at $350/hr.
Revenue Capture	Automated time reconstruction catches 10-30% of previously unbilled work. Money you already earned but never billed.
The Quality Dividend	Fewer missed deadlines (malpractice prevention), more client communication (satisfaction), nothing falling through cracks (case outcomes).
What It Isn't	Outputs require review. Setup takes work. Won't handle judgment calls. Some mornings will have errors. Human approval is non-negotiable.

Chapter 3: The Architecture – Skills, Connections, Schedules

What You'll Walk Away With:

- How every AI agent works in three layers – and why none of them require coding
- A complete skill file you can adapt for your practice, annotated line by line
- The connections map: which tools plug in where, and how to swap in your own stack
- How schedules turn one-off prompts into autonomous systems
- The safety layer that prevents malpractice exposure: audit trails, human gates, and review queues
- Why this architecture is AI-model agnostic – and what that means for your firm

---

What's Behind the Curtain

Chapter 2 showed you the outputs: the meeting briefing waiting at 7:02 AM, the inbox already triaged, the billing gaps caught before you sat down. If you're like most attorneys, your first reaction was some version of "how does that actually work?"

The answer is simpler than you'd expect.

Every agent in Chapter 2 – the meeting intel, the inbox triage, the billing reconciler, the case health monitor – runs on the same three-layer architecture. Three layers, stacked the same way every time:

1. **The Skill** – a plain-English instruction file that tells the AI what to do
2. **The Connections** – APIs linking the AI to your existing tools (Clio, Gmail, Calendar, Dodonai, QuickBooks)
3. **The Schedule** – a trigger that makes it run automatically, like a recurring calendar event

That's it. Skill, connections, schedule. Every agent Sarah Chen runs is a variation on this pattern. Every agent we'll build in Chapters 7 through 9 follows it. Once you understand these three layers, you can design agents for any workflow in your practice.

There's a fourth layer underneath these three – the safety layer – which we'll cover at the end of this chapter and expand on in Chapter 10. For lawyers, it's the most important one.

Layer 1: The Skill

A skill file is a set of instructions written in plain English and saved as a markdown file. Not code. Not a programming language. English.

Here's the key insight: **if you can write a procedure manual for a paralegal, you can write a skill file for an AI agent.**

Think about how you'd onboard a new paralegal on their first week. You'd say something like: "Every morning, check the calendar for today's meetings. For each external meeting, look up the client in Clio, pull their case history, check if there's anything pending in their file, skim recent emails from them, and put a one-page briefing on my desk before 8:00 AM."

That's a skill file. You just wrote one.

The difference is that a paralegal would need three months to learn your systems, make mistakes, ask clarifying questions, and gradually internalize your preferences. A skill file encodes all of that on day one – and it executes the same way every time.

A Real Skill File: Morning Case Briefing

Here's a complete skill file adapted for a litigation practice. This is the kind of file that would produce the meeting briefing you saw in Chapter 2. Read it like you'd read a procedure manual – because that's exactly what it is.

```
# /case-briefing -- Morning Case Briefing

Generate intel briefings for today's external meetings.

## Context

Before calls and meetings with clients, opposing counsel, or experts,
the attorney needs context on each attendee: their case status,
upcoming
deadlines, recent correspondence, and any open items. This skill
automates that gathering from Calendar, Clio, Gmail, and Dodonai.

## Steps

### 1. Determine the date

Use today's date unless a different date is specified.
Format as YYYY-MM-DD for file naming.
```

```
### 2. Fetch calendar events

Pull today's events from the firm's Google Calendar.

### 3. Filter to external meetings

Skip the event if:
- It is an all-day event
- It was declined
- It has no external attendees
- The title matches internal patterns: "team meeting", "staff lunch",
  "internal review"

Identify external attendees:
- Remove anyone with a firm email domain
- Remove calendar room resources
- If no external attendees remain, skip the event

If no external meetings exist for the day:
- Save a note: "No external meetings on [date]."
- Stop here.

### 4. For each external attendee, gather intel

**From Clio:**
- Look up the contact by name or email
- Pull their associated matters (case name, status, practice area)
- Pull upcoming deadlines and tasks on those matters
- Pull recent time entries (last 14 days)
- Pull the current trust/retainer balance

**From Gmail:**
- Search for recent emails to/from this contact (last 30 days)
- Note the last email date, subject, and whether it needs a response
- Flag any emails from opposing counsel on their matters

**From Dodonai:**
- Check if any documents are being processed for their matters
- Pull recent job outputs (medical chronologies, deposition
summaries)

**From web search (if new contact):**
- Search for the attendee's name + firm name
- Note their practice area, bar status, and any recent news

### 5. Compile the briefing

For each meeting, create a section with:
- Meeting time and title
- Attendee name, role, and firm
- Case summary (one paragraph)
- Upcoming deadlines (bulleted list)
- Recent correspondence summary
- Open items requiring attorney attention
- Trust/retainer balance

### 6. Save the briefing

Save to: reports/case_briefings/briefing_[date].md
```

```
### 7. Send notification

Post a summary to Slack:
"Morning briefing ready -- [N] meetings today. [list meeting names]"
```

Read that again. Count the lines of code. There are zero. Every instruction is a sentence a human could follow. The AI reads these instructions, connects to the specified tools, and executes the steps in order.

A few things worth noting:

The structure is explicit. Step 1 through Step 7. The AI follows them sequentially. If Step 3 finds no external meetings, it stops. There's no ambiguity about when the process ends.

The filters are precise. "Skip the event if it was declined." "Remove anyone with a firm email domain." These are the same filters you'd teach a paralegal – the skill file just writes them down so they never get forgotten.

The data sources are named. Clio for case data, Gmail for correspondence, Dodonai for document outputs. The skill file tells the AI where to look. You don't need to explain how APIs work – you just name the source and the AI connects to it through the connections layer (next section).

The output format is defined. Each briefing gets a specific structure: meeting time, attendee, case summary, deadlines, correspondence, open items, trust balance. This means every briefing looks the same. Consistency is built in, not hoped for.

Writing Your First Skill

You don't need to start with something this comprehensive. Your first skill file could be four lines:

```
# /weekly-review -- Friday Case Review

## Steps

### 1. Pull all active matters from Clio with no activity in the last
14 days

### 2. For each matter, note the last activity date and next deadline

### 3. Save the list to reports/stale_matters_[date].md
```

That's a working skill. It does one thing – finds stale matters – and it does it consistently every Friday. You'll refine it over time, the same way you'd refine a paralegal's training after their first few weeks.

The point is that skill files aren't engineering. They're operational documentation. Every attorney already writes documents with this level of specificity –

engagement letters, discovery protocols, office procedure manuals. A skill file uses the same writing discipline for a different audience: your AI instead of your staff.

Layer 2: The Connections

A skill file tells the AI what to do. The connections layer tells it where to look.

Connections are API integrations – the bridges between your AI and the tools your practice already uses. When the Morning Case Briefing skill says "pull their associated matters from Clio," it's the connection to Clio's API that makes that possible.

You don't need to build these connections from scratch. AI model harnesses (we use Claude Code in our examples) come with pre-built connectors for hundreds of tools. You plug in your credentials, and the AI can read from and write to those systems.

Here's what a typical litigation firm's connection map looks like:

System	What the AI Accesses	Example Agent Use
Clio	Matters, contacts, tasks, deadlines, time entries, trust balances, billing	Case health monitor, deadline tracker, billing reconciler
Gmail / Outlook	Inbox messages, sent mail, drafts, labels	Inbox triage, client communication drafts, opposing counsel monitoring
Google Calendar / Outlook Calendar	Events, attendees, scheduling	Meeting intel briefings, deadline calendaring
Dodonai	Document processing outputs – medical chronologies, deposition summaries, eDiscovery results, OCR text	Case preparation, evidence analysis, document status checks
QuickBooks / Xero	Invoices, payments, accounts receivable, trust accounting	Revenue monitoring, collections follow-up, retainer alerts
Court Filing Services (varies)	E-filing status, case dockets, court notices	Deadline computation, filing confirmation, docket monitoring
Document Management (NetDocuments, etc.)	File storage, folder structure, document metadata	Auto-filing, document retrieval, naming conventions

System	What the AI Accesses	Example Agent Use
Slack / Teams	Notifications, channels	Agent alerts, morning briefings, urgent flags

You swap in your stack. The architecture doesn't care which specific tools you use. Clio or PracticePanther. Gmail or Outlook. QuickBooks or Xero. NetDocuments or a shared drive. If the tool has an API – and nearly every modern practice management tool does – your AI can connect to it.

Sarah Chen's setup uses six connections: Google Calendar, Gmail, Clio, Dodonai, QuickBooks, and Slack. That covers her entire practice stack. A larger firm might add NetDocuments, a court filing API, and Microsoft Teams. A simpler practice might start with just Calendar, Gmail, and Clio.

Read vs. Write Permissions

Connections have permission levels, and this matters for lawyers.

Most agent connections are **read-only**. The meeting briefing agent reads your calendar and reads Clio data. It doesn't modify anything. The inbox triage agent reads your inbox to classify messages. The billing reconciler reads your time entries.

Write permissions are more limited and more controlled. An agent might create a draft email (Gmail draft, not a sent email). It might create a calendar event for a computed deadline. It might add a time entry to Clio.

But every write action goes through the safety layer (more on that below). The agent drafts; you approve. The agent calendars; you confirm. No external action happens without human sign-off.

This is a deliberate architectural choice. AI agents should be aggressive about gathering information and conservative about taking action. Read everything, write nothing without approval.

Layer 3: The Schedule

Skills tell the AI what to do. Connections tell it where to look. Schedules tell it when to run.

A schedule is a trigger – typically a cron job – that fires an agent at a set time. If you've ever set a recurring event in Google Calendar, you understand the concept. The only difference is that instead of reminding you to do something, it tells your AI to do something.

Here are the schedules from Sarah Chen's practice (adapted from the setup in Chapter 2):

Schedule	Agent	What It Does
Weekday 7:00 AM	Morning Case Briefing	Pulls calendar, researches attendees, compiles briefing
Weekday 7:30 AM	Inbox Triage	Classifies messages, labels, archives noise, flags urgent
Weekday 8:00 AM	Deadline Monitor	Checks Clio for deadlines within 14/30/60-day windows
Monday 8:30 AM	Case Health Digest	Scans all 40 active matters for stale cases, low trust balances, overdue tasks
Daily 9:00 AM	Billing Reconciler	Cross-references calendar/email activity with Clio time entries
Daily 9:30 AM	Client Communication Check	Flags matters past 30-day client contact mark, drafts follow-ups
Friday 4:00 PM	Weekly Summary	Rolls up the week's agent outputs into one digest

Notice the pattern: every agent runs before Sarah starts her day or at a natural checkpoint. By the time she opens her laptop at 8:15, five agents have already completed their work. The morning briefing is done. The inbox is sorted. The deadlines are checked. The billing gaps are identified.

She didn't ask for any of it. She set the schedules once, and they run until she changes them.

The "Recurring Calendar Event" Mental Model

If schedules feel abstract, think of it this way. You already have a recurring calendar event that says "Review open cases" every Monday. You probably ignore it half the time because you're busy. Now imagine that calendar event didn't remind you to review cases – it reviewed them for you, compiled the results, and sent you the summary.

That's a schedule. Same recurring trigger, different executor. Instead of you doing the work, your AI does it and hands you the result.

Schedule Granularity

Not every agent needs to run daily. Some run weekly (case health digest), some run multiple times a day (inbox triage – morning and afternoon), some run only on specific events (a new filing arrives, a payment posts). You design the cadence around the workflow.

The general principle: **run the agent as often as the underlying data changes in ways that matter.** Deadlines don't change hourly, so a daily check is enough. Your

inbox gets new messages constantly, so twice-daily triage makes sense. Your case portfolio shifts weekly, so a Monday health digest captures the meaningful movement.

The Safety Layer

This is where most legal technology conversations start. For good reason.

AI agents that connect to your case management system, read client communications, and generate documents create obvious risk. Malpractice carriers, bar ethics committees, and your clients all have legitimate concerns about AI operating inside a law practice.

The safety layer addresses those concerns with four principles. These aren't optional features. They're architectural requirements – baked into every agent, every workflow, every schedule.

Principle 1: AI Drafts, Humans Decide

No agent sends an email. No agent files a document. No agent communicates with a client, opposing counsel, a court, or any third party.

Agents draft. Humans approve.

When the client communication agent identifies a matter that hasn't had contact in 30 days, it drafts a follow-up email and saves it to your Gmail drafts folder. You review the draft, edit it, and hit send – or delete it. When the deadline monitor computes a response deadline from a court order, it creates a proposed calendar event. You confirm it or adjust it.

Every external action passes through a human approval gate. No exceptions.

This isn't a limitation. It's the design. The agent's job is to do the 90% of work that doesn't require legal judgment – gathering data, identifying patterns, computing dates, drafting routine language. Your job is the 10% that does: deciding whether the draft is right, whether the deadline computation accounts for local rules, whether this particular client needs a phone call instead of an email.

Principle 2: Full Audit Trail

Every agent run is logged with timestamps, inputs, outputs, and decisions. When the deadline monitor computes a response date, the log shows which document triggered the computation, which jurisdiction rules were applied, and what date was produced. When the inbox triage agent classifies a message as "FYI" and archives it, the log records the classification and the reasoning.

This matters for three reasons:

Defensibility. If a malpractice claim alleges a missed deadline, you can show exactly when the agent flagged the deadline, what rule it applied, when you were notified, and what action you took. The audit trail is your evidence that the system worked and that you exercised supervision.

Debugging. If an agent produces an unexpected result, the log shows you why. You can trace the issue to a specific step, a specific data source, or a specific instruction in the skill file – and fix it.

Ethics compliance. ABA Formal Opinion 512 requires competence, communication, and supervision when using AI tools. An audit trail demonstrates all three. You understand how the tool works (competence). You can explain its role to clients (communication). You review every output before it reaches the outside world (supervision).

Principle 3: Uncertain Items Route to a Review Queue

Not every piece of information fits neatly into a category. A court filing might be ambiguous about whether a deadline applies. An email might be unclear about whether it requires a response or is purely informational. A medical record might contain a note that could be relevant to causation – or might not.

When an agent encounters ambiguity, it doesn't guess. It routes the item to a review queue with a flag: "Low confidence – needs attorney review."

This is the malpractice prevention layer. A paralegal who's unsure about a deadline should escalate it to the attorney. Your AI does the same thing. The difference is that the AI escalates consistently every time, at the same confidence threshold, without the social pressure to "just figure it out" that sometimes keeps a paralegal from asking.

Principle 4: Access Controls and Scope Limits

Each agent has defined permissions. The inbox triage agent can read emails and apply labels, but it can't delete messages or send replies. The billing reconciler can read time entries and suggest additions, but it can't modify existing entries or generate invoices. The deadline monitor can read court filings and create proposed calendar events, but it can't accept or decline hearing dates.

Scope limits prevent cascading errors. If an agent malfunctions, the worst it can do is generate an incorrect draft or a wrong calendar suggestion – both of which go through human review before any action is taken. It can't send an email to opposing counsel, file a document with the court, or transfer funds from a trust account.

How This Maps to Ethics Rules

Bar ethics rules for AI are still evolving, but the framework is converging. ABA Formal Opinion 512 (2024) established the baseline: competence, communication, and supervision. California's proposed amendments to Rules 1.1, 1.4, 1.6, 3.3, 5.1, and 5.3 extend these obligations specifically to AI use. As of early 2026, every state bar that has issued AI guidance follows some version of this framework.

The safety layer maps directly:

Ethics Obligation	How the Architecture Addresses It
Competence (Rule 1.1)	You understand the three layers – skill, connections, schedule. You wrote or approved the skill file. You know what data the agent accesses.
Communication (Rule 1.4)	You can explain to clients how AI is used in their matter. "Our system checks for upcoming deadlines daily and flags them for my review."
Supervision (Rule 5.1 / 5.3)	Every agent output passes through human review. No external action without attorney approval. Full audit trail of every run.
Confidentiality (Rule 1.6)	Data stays within your tool stack. AI providers under zero-retention policies. No client data used for model training.
Candor (Rule 3.3)	AI-generated content is reviewed before submission. Citation verification catches hallucination risk before anything reaches a court.

We'll go deeper on ethics in Chapter 10, including hallucination risk (600+ documented cases and 128 lawyers implicated, per National Law Review), disclosure obligations, and the 43% of firms that still have no AI policy (Wolters Kluwer 2026). For now, the key takeaway: the architecture has the safety layer built in, not bolted on.

Putting the Layers Together

Here's how the three layers plus the safety layer combine for a single agent – the Morning Case Briefing from earlier in this chapter:

The Skill: A markdown file with seven steps. Check the calendar, filter to external meetings, look up each attendee in Clio, pull recent emails, check Dodonai for document status, compile the briefing, send a notification.

The Connections: Google Calendar (read events), Clio (read contacts, matters, deadlines, time entries), Gmail (read messages), Dodonai (read job outputs), Slack (write notification).

The Schedule: Weekday 7:00 AM.

The Safety Layer: The agent reads data but never modifies it. The briefing is saved as a file and a Slack notification. No emails sent, no calendar events created, no client contact. The attorney reviews the briefing and decides what action to take.

When Sarah Chen wakes up and reads her morning briefing, she's looking at the product of all four layers working together. The skill defined what to do. The connections retrieved the data. The schedule triggered it at 7:00 AM. The safety layer ensured nothing happened that required her judgment until she was ready to exercise it.

The same four layers drive every agent in the rest of this book: the document workflow agents in Chapter 5, the practice area agents in Chapter 6, the operational agents in Chapter 7 through 9, and the advanced configurations in Chapter 11.

The Architecture is AI-Model Agnostic

One more thing worth stating explicitly.

We use Claude Code in these examples because it's the model harness we currently work with. Claude Code can read files, run scripts, connect to APIs through what Anthropic calls MCP (Model Context Protocol) servers, and execute skill files as structured workflows.

But the architecture – skills, connections, schedules, plus the safety layer – doesn't depend on Claude. It works with any AI model that can follow multi-step instructions and connect to external tools.

The skill file you saw earlier in this chapter is written in English. It's not Claude-specific syntax. Any capable AI system that can process natural language instructions, call APIs, and generate structured output can execute it. If you're using GPT, Gemini, or another model, you'd adapt the connection method, but the skill file itself wouldn't change.

This matters because AI models improve rapidly. The model you use today won't be the model you use in 18 months. The architecture survives model changes because the intelligence lives in the skill files and connections – not in the model. When a better model arrives, you swap it in and your agents keep running.

Think of it like a law practice's operations manual. If you hire a new associate, you don't rewrite your firm procedures. You hand them the manual and they follow it. The skill file is the manual. The AI model is the associate. The manual stays the same even when the associate changes.

What AI Won't Do in This Architecture

Honesty about limitations builds more trust than overselling capabilities.

It won't exercise legal judgment. The deadline monitor computes dates from court rules, but it doesn't decide whether to request an extension. The case health agent flags a stale matter, but it doesn't determine the right strategy for re-engaging the client. Judgment calls remain yours.

It won't handle novel situations gracefully. Skill files encode known workflows. When something genuinely unprecedented happens – a novel procedural question, a first-impression issue, a client situation that doesn't fit any pattern – the agent will either route it to the review queue or produce a generic response. You need to recognize those moments and take over.

It will sometimes get things wrong. AI hallucinates. It misreads documents. It misclassifies emails. The safety layer catches most errors before they reach the outside world, but you need to review outputs with the same critical eye you'd apply to a first-year associate's work product. Trust but verify – emphasis on verify.

It won't replace your staff. Maria the paralegal still reviews every medical chronology. David still handles client phone calls. The agents handle the gathering, sorting, computing, and drafting. Your team handles the judgment, relationship, and quality control. The agents free your team from production work so they can do the work that actually requires their expertise.

Try This Now: Write a Skill File

Pick one task you do every week that follows a repeatable pattern. The Monday case review. The Friday billing check. The daily email triage. Whatever makes you groan when it shows up on your calendar.

Now write the procedure for how you'd teach someone to do it. Not the legal judgment parts – the gathering, checking, compiling parts. Write it as numbered steps. Be specific about where the information comes from and what the output should look like.

Here's a template:

```
# /[name] -- [What it does]

## Context
```

```
[One sentence: why this task matters and when it runs]

## Steps

### 1. [First action]
[Where to look, what to pull]

### 2. [Second action]
[What to check, what to filter]

### 3. [Third action]
[How to compile the output]

### 4. Save and notify
[Where to save the result, who to alert]
```

Fill that in for your chosen task. You just wrote your first skill file.

If you want to test it, paste it into an AI assistant (any one – Claude, GPT, Gemini) with the underlying data and see what it produces. The skill file works as a prompt, even before you've set up connections or schedules. That's the crawl stage. Walk and run come in Chapter 11.

Quick Reference: The Architecture

Layer	What It Is	Law Firm Analogy
Skill	Plain-English instruction file	Procedure manual for a paralegal
Connections	API links to your tools (Clio, Gmail, Calendar, etc.)	The paralegal's access to your systems
Schedule	Automated trigger (daily, weekly, on-event)	The recurring calendar reminder – except it does the work
Safety Layer	Human approval gates, audit trail, review queue	Supervising attorney reviews before anything goes out

Connections Map	What the AI Accesses
Clio / PracticePanther	Cases, contacts, deadlines, billing, trust
Gmail / Outlook	Inbox, sent mail, drafts
Calendar	Events, attendees
Dodonai	Medical chronologies, deposition summaries, eDiscovery, OCR
QuickBooks / Xero	Invoices, payments, AR, trust accounting
Slack / Teams	Notifications, briefing delivery

Next: Chapter 4 walks through Dodonai – the document processing engine that turns raw medical records, deposition transcripts, and evidence into the structured data your agents work with. Before your AI can analyze a case, it needs the documents in usable form. That's what Dodonai does.

Chapter 4: Dodonai – Your Document Processing Engine

What You'll Walk Away With:

- What each of Dodonai's six products does, the specific problem it solves, and what the output looks like
- Side-by-side cost and time comparisons: manual, outsourced, and Dodonai for medical records and depositions
- How document processing fits into the agent architecture from Chapter 3 – agents need structured data; Dodonai produces it
- The pricing math: what it actually costs per case, per page, per month
- Which product to start with based on your practice area

Before Your Agents Can Think, Your Documents Need Structure

In Chapter 3, the Morning Case Briefing agent checked Dodonai for "recent job outputs – medical chronologies, deposition summaries." That step takes two seconds in the agent's workflow, but it depends on something important: the documents were already processed.

Raw medical records are PDFs. Scanned faxes. Handwritten notes. Photocopies of photocopies. A 400-page stack from four different providers with no consistent formatting, no index, and no table of contents. Your AI can't pull a treatment timeline from that. Neither can your paralegal, at least not without spending half a day reading through it.

Deposition transcripts are slightly better – they're text, at least – but a 200-page transcript is still a wall of Q&A that takes two hours to distill into something usable for motion practice.

This is the document processing layer. Before AI agents can analyze case data, flag deadlines, cross-reference testimony, or draft demand sections, someone – or something – needs to turn raw documents into structured output. Date-ordered chronologies. Searchable transcript libraries. Extracted data fields. Text pulled from scanned pages.

Dodonai is the engine that handles this layer. Six products, each solving a specific document processing problem that every litigation practice faces. This chapter walks through each one.

Medical Record Summaries & Chronologies

The Problem

A personal injury case with four treating providers generates 200 to 500 pages of medical records. A med mal case can hit thousands. The records arrive as PDFs – sometimes scanned, sometimes faxed, often with pages out of order, duplicate records from different requests, and handwritten notes that are barely legible.

Someone has to read through those records, extract every relevant date, diagnosis, treatment, procedure, and provider visit, and organize them into a chronology the attorney can use for case evaluation, demand preparation, and trial.

The manual approach: a paralegal spends 4 to 8 hours per case reading records, creating a chronology in Word, cross-referencing dates, and flagging gaps in treatment. At $50/hour, that's $200 to $400 in paralegal time per case.

The outsourced approach: send the records to a medical record summary service. Cost: $500 to $800 per case. Turnaround: 3 to 7 business days. Quality varies. Revisions add time and cost.

The Dodonai Approach

Upload the medical records PDF to Dodonai. Select the Medical Record summary type. Enter any key topics, diagnoses, or providers you want highlighted. Hit Summarize.

Dodonai processes the records and produces two outputs:

1. **A general summary** of the patient's healthcare history – the narrative overview
2. **A detailed medical chronology** – every provider visit, diagnosis, treatment, procedure, and medication organized in date order with source citations

A 500-page file processes in under 15 minutes. Output includes an interactive viewer with side-by-side source linking – click any chronology entry and see the exact page in the source records it came from.

You can sort and filter entries by provider or date, customize the executive summary, and download in PDF, Word, or TXT.

The Math

Method	Cost per Case (300 pages)	Turnaround	Attorney/Paralegal Time
Manual (paralegal)	$200-400 (4-8 hrs at $50/hr)	Same day (but consumes the day)	4-8 hours
Outsourced (summary service)	$500-800	3-7 business days	30 min review
Dodonai (Pro plan)	$6 (300 pages x $0.02/page)	Under 15 minutes	30 min review

Sarah Chen processes roughly 8 new cases per month. At an average of 300 pages of medical records per case, she was spending $4,000 to $6,400 per month on outsourced summaries – or losing 32 to 64 paralegal hours. With Dodonai Pro ($100/month for 5,000 credits), the same volume costs $48 in credits. Maria the paralegal now spends 30 minutes reviewing each chronology instead of 6 hours creating it. That's 44 hours per month of paralegal capacity returned to the practice.

One PI attorney told us his Dodonai medical record summaries are "the same quality as what I used to pay $800 for elsewhere." That's not a marginal improvement. It's a structural change in case economics.

What It Won't Do

Dodonai extracts and organizes what's in the records. It doesn't interpret causation, assess damages, or render opinions about standard of care. Those are attorney and expert judgment calls. The chronology gives you the foundation – every date, every provider, every treatment – and you build the legal analysis from there.

Always verify AI-generated medical chronologies against the source records before relying on them in filings or proceedings.

Deposition Summaries

The Problem

A 100-page deposition transcript takes a paralegal roughly 2 hours to summarize manually. A senior paralegal billing at $50-75/hour produces a summary for $100 to $150 in labor cost. Outsourcing to a deposition summary service runs $3 to $5 per page – $300 to $500 for that same 100-page transcript, with a multi-day turnaround.

Most litigation practices process dozens of depositions per case. A complex commercial dispute might have 30 to 50 transcripts. A PI case with multiple defendants could have 8 to 12. Multiply the per-transcript cost by the case volume and deposition summaries become one of the largest production expenses in litigation.

Three Formats

Dodonai produces deposition summaries in three formats, each suited to a different use:

Page-Line Summary – The most common format for litigation. A table with two columns: a summary of each testimony section and the corresponding page-line citation from the transcript. Organized chronologically. This is what you'd use for motion practice, trial prep, and witness cross-examination binders.

Narrative Summary – A flowing overview of the testimony. Useful for quick case assessment, team briefings, and client updates where the reader needs the gist without citation-level detail.

Issue-Based Summary – Testimony organized by topic rather than chronological order. Useful for complex commercial disputes where you need all testimony on a specific issue (damages, notice, breach) collected in one section.

The Numbers

Upload a PDF, TXT, or Word transcript. Select the format. Enter any key topics for emphasis. Dodonai produces the summary.

Processing time on real transcripts, benchmarked against manual summaries by a reviewing attorney:

Transcript	Pages	Manual Time	Dodonai Time
Corporate Representative	64	1 hr 48 min	84 seconds
Insurance Investigator	44	1 hr 50 min	79 seconds
Truck Driver (PI)	37	1 hr 26 min	61 seconds
Petitioner (Family)	52	1 hr 52 min	90 seconds
Damages Expert	40	1 hr 45 min	74 seconds

The reviewing attorney's assessment: "I'm a little embarrassed to say that the AI deposition summaries created by Dodonai were even better summaries than the ones I wrote."

Cost Comparison

Method	Cost per Transcript (100 pages)	Turnaround
Manual (paralegal)	$100-150 (2 hrs at $50-75/hr)	Same day

Method	Cost per Transcript (100 pages)	Turnaround
Outsourced (summary service)	$300-500 ($3-5/page)	2-5 business days
Dodonai (Pro plan)	$2 (100 pages x $0.02/page)	~3 minutes

For Sarah Chen's PI practice, the impact is straightforward. She takes roughly 8 new cases per month with an average of 4 depositions per case. At an outsourced rate of $400 per summary, that's $12,800/month in deposition summary costs. With Dodonai Pro, the same volume runs through 3,200 credits at $0.02/page – roughly $64/month in credit usage. The $100/month Pro subscription pays for itself on the first transcript.

What It Won't Do

Deposition summaries capture what the deponent said. They don't assess credibility, flag inconsistencies with other testimony, or identify impeachment opportunities. That analysis is the attorney's work – and it gets done faster when you're reading a structured summary instead of a 200-page transcript.

Transcript Management

The Problem

A litigation practice accumulates transcripts. Multiple depositions per case, multiple cases per year. Without a dedicated system, transcripts live as individual PDFs scattered across folders. Finding specific testimony means opening files one at a time and using Ctrl-F for keyword searches – which misses relevant passages that use different language.

When you're preparing for trial and need every instance where a witness discussed treatment dates, you're searching "treatment," "appointment," "visit," "saw the doctor," "went to the hospital" – and you're still not sure you found everything.

The Solution

Dodonai's Transcript Management is a library for your deposition transcripts with three capabilities that individual PDFs don't offer:

Semantic Search. Search by concept, not keywords. Ask for testimony about "causation of injuries" and the system finds relevant passages even when the word "causation" never appears in the transcript. This is the same semantic search technology behind the eDiscovery product, applied to your transcript library.

AI Chat / Copilot. Ask natural language questions across your transcript library. “What did Dr. Martinez say about the plaintiff’s pre-existing condition?” The AI pulls the relevant testimony, cites the source, and highlights the passages.

Blue Book Citations. Highlight any passage in the transcript viewer and Dodonai formats the citation in Blue Book format, ready to copy and paste into a brief or motion. One click from highlighted testimony to properly formatted citation.

Transcript Management is included with all Dodonai plans. No additional credit cost for uploading, searching, or citing.

Why This Matters for Trial Prep

Most trial lawyers have a deposition management problem they’ve accepted as normal. Transcripts live in case folders – individual PDFs named something like “Smith_Depo_03-14-26.pdf.” Finding testimony means remembering which witness said what, opening the right file, and running keyword searches that may or may not find the passage you need.

Transcript Management changes the unit of work. Instead of searching within a single transcript, you search across all transcripts on a matter simultaneously. “What did any witness testify about the plaintiff’s activities after the accident?” One query, every relevant passage from every deposition, with citations formatted for your brief.

For Sarah Chen’s PI practice with 40 active matters and an average of 4 depositions per case, that’s 160 transcripts. Finding every mention of a specific treating provider across all transcripts in a case – work that used to take an afternoon of Ctrl-F searches – takes seconds.

OCR – AI-Powered Text Extraction

The Problem

Not every document arrives as clean, searchable text. Medical records come as scanned PDFs. Older court filings are photocopied images. Handwritten notes from physicians are barely legible to humans, let alone traditional OCR software.

Legacy OCR tools – the kind built into Adobe Acrobat or standalone scanning software – struggle with complex layouts, mixed formatting, handwritten content, and poor scan quality. The result: garbled text, missed pages, manual cleanup that takes longer than the original scan.

The Solution

Dodonai's OCR uses AI to extract text from PDFs, including:

- **Handwritten documents** – physician notes, handwritten correspondence, field notes
- **Tables** – medical records with tabular data, billing statements, lab results
- **Images within documents** – embedded charts, diagrams, photographs with captions
- **Poor-quality scans** – faded faxes, low-resolution copies, multi-generation photocopies

The AI doesn't just extract characters. It understands context – recognizing that a series of numbers in a medical record is a lab result, not a phone number. Tables are preserved as structured data, not flattened into a text stream.

OCR is free with all Dodonai plans. No credit cost, no per-page charge. Process your scanned documents first, then run them through eDiscovery or summarization for complete coverage.

Why This Matters for Litigation

Every document you can't search is a document you might miss in discovery review. OCR is the prerequisite for eDiscovery, summarization, and transcript management. Run OCR on your scanned records first, and every downstream tool works with clean, searchable text instead of image files.

---

eDiscovery – Semantic Document Search

The Problem

Traditional eDiscovery platforms charge per-gigabyte hosting fees plus per-user licensing. A mid-size firm on Relativity or Everlaw can spend $50,000 or more per year before processing a single document. For solo practitioners and small firms, that pricing makes eDiscovery functionally inaccessible. The result: keyword searches across individual PDFs, hoping you found everything.

Even with a traditional eDiscovery platform, keyword search has a fundamental limitation. Searching for "negligence" misses documents that discuss "failure to exercise reasonable care." Searching for "breach of contract" misses communications about "not meeting the terms of the agreement." You run dozens of keyword variations and still can't be certain you've captured every relevant document.

The Solution

Dodonai's eDiscovery searches by meaning, not keywords.

Upload your document set – PDFs, Word files, emails, text files. Describe what you're looking for in plain language. "Communications about the defendant's knowledge of the product defect before the recall." The AI finds relevant documents even when the exact terms don't appear.

Ranked Results. The most relevant documents surface first, with highlighted passages showing why each document matched.

Audit Trail. Every search is logged – the query, the results, the relevance rankings. This gives you a defensible review methodology. If opposing counsel challenges your discovery process, the audit trail shows exactly what you searched and how results were ranked.

Source Verification. Click any result and see the exact passage in the original document. No black-box relevance scores. You can verify every match.

Dodonai eDiscovery is included with all plans. Combined with OCR (also free), even scanned document productions become fully searchable by concept.

The Cost Comparison

Platform	Annual Cost (mid-size firm)	Search Type
Relativity	$50,000+	Keyword + TAR
Everlaw	$50,000+	Keyword + TAR
Dodonai	$360-1,200/year	Semantic (concept-based)

For Sarah Chen's practice, eDiscovery was never economically viable before Dodonai. At $50,000/year, the platform would have cost more than 6% of her annual revenue. At $1,200/year on the Pro plan, she has access to the same core capability – semantic search across document sets – as firms spending 40 times more.

Practical Example: PI Document Review

Sarah receives a production of 2,000 pages of the defendant's internal emails in a premises liability case. The core question: did the property owner know about the hazardous condition before the plaintiff's fall?

With keyword search, she'd run queries for "slip," "fall," "hazard," "repair," "maintenance," "complaint," "warning," "danger," "wet floor" – and still worry she missed the email where someone wrote "we need to fix that walkway before somebody gets hurt."

With Dodonai eDiscovery, she searches: “communications about knowledge of unsafe conditions on the property.” The AI finds that email – and the three others that discuss the condition using language she never would have included in a keyword list.

The audit trail logs the search, the query, and every result returned. If opposing counsel challenges her review methodology, the log demonstrates a thorough, systematic search.

What It Won’t Do

eDiscovery surfaces relevant documents. It doesn’t make privilege determinations, assess responsiveness under specific discovery requests, or produce documents to opposing counsel. The attorney reviews every flagged document and makes those judgment calls. The AI’s role is to ensure no relevant document gets missed in the review – not to replace the reviewer.

Extract & Draft Agents (XD Agents)

The Problem

Every litigation practice has documents it processes repeatedly. Intake forms where the same 15 fields need to be extracted and entered into Clio. Insurance policies where you need to find the coverage limits, exclusions, and reservation of rights language. Employment contracts where you need the termination provisions, non-compete terms, and compensation structure.

These aren’t novel analyses. They’re data extraction tasks that follow the same pattern every time. But they still require a human to read the document, find the relevant sections, and manually transfer the data to where it needs to go.

The Solution

Dodonai XD Agents let you create custom extraction templates. Define the fields you need – once – and the AI extracts them from any source document.

How it works:

1. **Create a template.** Define the structure, sections, data types, and formats you need. Example: “For an intake form, extract: client name, date of incident, type of injury, treating providers, insurance carrier, policy number, prior attorney.”

2. **Select source documents.** Upload the intake form, insurance policy, medical record, or whatever document contains the data.
3. **Run the extraction.** The AI extracts the specified fields from the source documents and populates your template.
4. **Review and adjust.** Edit any extracted data directly in the platform before exporting.

Templates are reusable. Build an intake extraction template once, and it runs against every new intake form with zero reconfiguration. Build a policy review template, and it pulls the same fields from every insurance policy.

Use Cases for Litigation

Template	Fields Extracted	Source Documents
PI Intake	Client name, DOI, injury type, providers, insurance, policy number	Intake forms
Insurance Policy Review	Coverage limits, exclusions, reservations, deductible, named insureds	Policy PDFs
Medical Bill Audit	Provider, service date, CPT codes, billed amount, paid amount, balance	Medical billing records
Employment Review	Termination provisions, non-compete scope, compensation, change of control	Employment agreements
Discovery Response Organizer	Request number, objection, response, documents referenced	Discovery responses

XD Agents use 2 credits per page – double the standard rate – reflecting the custom extraction logic. For a 10-page insurance policy, that's 20 credits or $0.40 on the Pro plan.

The Compound Value

The real power of XD Agents emerges when you combine them with other Dodonai products. Process medical records with OCR first, then run them through an XD Agent template that extracts treating providers, diagnoses, and procedure dates. Feed that extracted data into your medical chronology. Use the chronology to identify gaps in treatment. Cross-reference those gaps against deposition testimony using Transcript Management.

Each product handles one step. Chained together, they produce a case file that would have taken a paralegal two days to compile manually.

For Sarah Chen's workers' comp cases, she built an XD Agent template that extracts the employer, date of injury, job title, body parts injured, treating providers, and work restrictions from the initial claim form. That template runs on

every new intake. The extracted fields go directly into Clio – no manual data entry, no transcription errors.

Pricing

Dodonai's pricing is credit-based. One credit equals approximately one page (roughly 400 words). Credits are used for document processing – summaries, chronologies, and XD Agent extractions. Transcript Management, eDiscovery, and OCR are included with all plans at no additional credit cost.

Plan	Monthly	Annual	Monthly Credits	Annual Credits	Effective Cost/Page
Lite	$30	$300	200	4,800	$0.15/page (monthly), $0.06/page (annual)
Basic	$60	$600	1,000	24,000	$0.06/page (monthly), $0.025/page (annual)
Pro	$100	$1,000	5,000	120,000	$0.02/page (monthly and annual)
Enterprise	Custom	Custom	Unlimited	Unlimited	Contact for pricing

Key details:

- Annual credits roll over – unused credits carry to the next month
- No per-seat fees – everyone at the firm uses one account
- Free trial available – process your first documents before committing
- SOC 2 certified, HIPAA compliant, zero data retention on AI processing
- BAA (Business Associate Agreement) available for HIPAA-covered entities

Which Plan Fits?

Lite ($30/month) – Solo practitioners handling a few cases per month. 200 credits covers roughly two medical record summaries or four short deposition summaries.

Basic ($60/month) – Small firms with steady caseloads. 1,000 credits covers roughly 10 cases per month at 100 pages each.

Pro ($100/month) – Active litigation practices. Sarah Chen's 40-case portfolio runs comfortably on Pro's 5,000 monthly credits. At $0.02/page, this is where the economics get compelling.

Enterprise (custom) – Multi-attorney firms or high-volume practices that need unlimited processing, dedicated support, and custom integrations.

How Dodonai Feeds Your Agent Architecture

This is where Chapter 3 and Chapter 4 connect.

In the agent architecture, Dodonai sits in the connections layer. Your AI agents connect to Dodonai the same way they connect to Clio or Gmail – as a data source. When the Morning Case Briefing agent checks for "recent document outputs," it's reading Dodonai's processed results: the medical chronology that finished last night, the deposition summary you ran yesterday, the eDiscovery search results from this morning.

Without Dodonai, your agents would be reading raw PDFs – unstructured, unindexed, unsearchable. With Dodonai, they're reading structured chronologies with source citations, searchable transcript libraries, and extracted data fields.

The document processing layer is the foundation. Everything else – the case analysis, the deadline tracking, the strategic planning that your agents will handle in Chapters 5 through 9 – builds on top of structured, processed documents.

The Document Processing Workflow

Here's the typical flow for a new PI case at Chen Law Group:

1. **Intake form arrives** – XD Agent extracts client details and populates Clio
2. **Medical records received** – OCR processes scanned pages, then Medical Record Summary generates the chronology
3. **Depositions taken** – Transcripts uploaded to Transcript Management library, summaries generated in page-line format
4. **Discovery production received** – OCR processes scanned documents, eDiscovery indexes everything for semantic search
5. **Case prep begins** – Attorney searches transcripts for specific testimony, cross-references medical chronology, runs eDiscovery queries

At every step, the processed output is available to the AI agents from Chapter 3. The Morning Case Briefing agent can report: "Medical chronology completed for Hernandez – 14 months of treatment documented across 4 providers. 2 depositions summarized. eDiscovery index ready with 1,847 documents." That status update took zero attorney time to compile.

Security and Compliance

A brief note on the question every attorney asks: where does my client's data go?

Dodonai is SOC 2 certified and HIPAA compliant. All documents are encrypted in transit and at rest using AES-256 encryption. BAAs are available for firms handling protected health information.

The critical detail for client confidentiality: Dodonai's AI providers operate under zero data retention policies. Documents are processed and the results are returned. No client data is stored by the AI model after processing completes. No client data is used to train AI models. You control when your data is deleted.

This matters because AI-processed legal documents are potentially discoverable. Having a clear data retention policy – and the ability to document exactly how client data was processed and where it resides – is the baseline for responsible AI use in a litigation practice.

Try This Now: Process Your First Document

Pick a document from a current case – a medical record set, a deposition transcript, or a stack of scanned PDFs that need OCR. Sign up for a Dodonai free trial and process it.

Here's what to try based on your practice area:

Personal injury: Upload a set of medical records (200+ pages). Generate a medical chronology. Compare it to a chronology you or your paralegal created manually. Note what the AI captured that you missed – and what you'd add that it didn't.

Any litigation practice: Upload a recent deposition transcript. Generate a page-line summary. Time how long it takes. Compare the output to what a paralegal would have produced. Ask yourself: if this takes 90 seconds instead of 2 hours, how does that change your case prep workflow?

Document-heavy cases: Upload a set of scanned documents and run OCR. Then upload the OCR output to eDiscovery and try a semantic search. Search for a concept you know is in the documents, using language that doesn't match the exact words on the page. See what comes back.

The goal isn't to evaluate whether AI document processing is "good enough." It's to experience the time compression firsthand. A 300-page medical record chronology in under 15 minutes changes how you think about case prep economics.

Quick Reference: Dodonai Products

Product	What It Does	Cost	Speed
Medical Record Summaries & Chronologies	Date-ordered chronology + narrative summary from medical records	1 credit/page	500 pages in <15 min
Deposition Summaries	Page-line, narrative, or issue-based digest of testimony	1 credit/page	100 pages in ~3 min
Transcript Management	Searchable library with AI Chat and Blue Book citations	Included	Instant search
OCR	Text extraction from scans, handwriting, tables, poor PDFs	Free	Minutes
eDiscovery	Semantic document search across large document sets	Included	Seconds per query
XD Agents	Custom extraction templates for repeatable document processing	2 credits/page	Minutes

Plan	Monthly Cost	Credits	Best For
Lite	$30	200	Solo, light caseload
Basic	$60	1,000	Small firm, steady volume
Pro	$100	5,000	Active litigation practice
Enterprise	Custom	Unlimited	Multi-attorney, high volume

Next: Chapter 5 shows how to move from processed documents to strategic outputs – demand letter sections built from medical chronologies, deposition designations compiled from transcript searches, and case evaluations assembled from structured data.

Chapter 5: From Documents to Strategy

What You'll Walk Away With:

- The three-step workflow that turns processed documents into strategic outputs: Process (Dodonai) -> Analyze (AI) -> Act (You)
- Five complete worked examples, each with real prompts you can copy, paste, and adapt for your cases
- How to build a demand letter from a medical chronology in 45 minutes instead of a full day
- How to turn 10 deposition summaries into an impeachment-ready cross-exam outline
- The "skills" approach: encoding your analytical framework so AI applies YOUR methodology every time
- Where AI analysis ends and attorney judgment begins – for every workflow

---

The Output That Changes the Conversation

Sarah Chen finishes reviewing a medical chronology from Dodonai on the Hernandez PI case. Fourteen months of treatment across four providers, 412 pages of records, organized into a date-ordered chronology with source citations in under 15 minutes.

That chronology is useful on its own. But it's not a demand letter. It's not a damages narrative. It's not a case evaluation. It's structured data – the foundation for strategic work that still needs to happen.

What used to happen next: Sarah or Maria would spend the afternoon extracting treatment milestones, organizing them into a damages narrative, cross-referencing the treatment timeline against pre-existing conditions, and drafting the medical specials section. Four to six hours of work that required attention but not much legal judgment.

What happens now: Sarah pastes the chronology into her AI assistant, runs a prompt she's used dozens of times, and gets a structured damages narrative in eight minutes. She adjusts two paragraphs where the AI understated the severity of a lumbar fusion and drops the narrative into her demand template. Total time from chronology to demand section: 45 minutes.

That's the workflow this chapter teaches. Chapter 4 covered document processing. Chapter 3 covered the agent architecture. This chapter is about what you DO with Dodonai's structured outputs once they exist.

The pattern:

1. **Process** (Dodonai) – Raw documents become structured data
2. **Analyze** (AI) – Structured data becomes strategic output
3. **Act** (You) – Strategic output becomes attorney work product after your review

Five workflows. Each one takes a Dodonai output you already know how to produce, feeds it through an AI prompt, and delivers something you'd normally spend hours creating manually.

Workflow 1: Medical Chronology to Demand Letter

The Input

A Dodonai Medical Record Summaries & Chronologies output for a PI auto accident case. The plaintiff is Maria Gonzalez, rear-ended at a stoplight on IH-35 in Austin. Four treating providers over 14 months. The chronology is 18 pages of date-ordered entries with source citations.

The Prompt Chain

This workflow uses two prompts in sequence. The first extracts and organizes the damages narrative. The second builds the medical specials section for the demand package.

Prompt 1: Damages Narrative

```
I'm preparing a demand letter for a personal injury case. Below is
the
medical chronology generated from the plaintiff's medical records.

CASE FACTS:
- Plaintiff: Maria Gonzalez, age 34, administrative assistant
- Incident: Rear-end collision on IH-35, Austin, TX on 03/15/2025
- Liability: Clear -- defendant rear-ended plaintiff at stoplight,
  police report confirms
- Pre-existing conditions: None documented in records

TASK:
Analyze the medical chronology and produce a damages narrative
suitable
for inclusion in a demand letter. Structure it as follows:
```

```
1. INJURY SUMMARY: One paragraph describing all diagnosed injuries,
   severity, and prognosis
2. TREATMENT TIMELINE: Narrative of treatment progression from
   initial ER visit through most recent appointment, emphasizing:
   - Escalation of treatment (conservative to interventional)
   - Duration and frequency of care
   - Any surgical interventions or recommendations
   - Ongoing symptoms and functional limitations
3. IMPACT ON DAILY LIFE: How injuries affected plaintiff's ability
   to work, perform daily activities, and maintain quality of life
   (draw from any functional assessments or physician notes in the
   chronology)
4. PROGNOSIS: Future treatment recommendations, permanent
   impairment ratings if documented, and ongoing care needs
5. GAPS AND FLAGS: Note any treatment gaps longer than 30 days,
   any pre-existing conditions referenced, or any entries that
   the attorney should verify before sending the demand

Use specific dates, provider names, and treatment details from the
chronology. Cite the chronology entry dates so I can verify each
statement against the source records.

Do not editorialize or overstate. Let the medical evidence speak. The
tone should be factual and precise -- this will be reviewed by an
insurance adjuster and possibly defense counsel.

[PASTE DODONAI MEDICAL CHRONOLOGY HERE]
```

Prompt 2: Medical Specials Calculation

```
Using the damages narrative above and the medical chronology, prepare
the medical specials section of the demand package.

Structure as a table with these columns:
- Provider name
- Date range of treatment
- Description of services
- Total billed amount (if documented in records)
- Status (paid, outstanding, liened)

After the table, include:
1. TOTAL MEDICAL SPECIALS (sum of all billed amounts)
2. FUTURE MEDICAL COSTS (estimated based on prognosis section --
   note these are projections and flag if an economist or life care
   planner should be retained)
3. LOST WAGES CALCULATION (if employment records were in the
   chronology -- if not, flag that these need to be obtained
   separately)

Flag any billing entries where the amount seems inconsistent with the
treatment described, or where records from a provider appear
incomplete.
```

What the AI Produces

A 2-3 page damages narrative organized in your demand letter structure. The AI pulls specific dates ("On March 15, 2025, Ms. Gonzalez presented to St. David's

Medical Center ER with complaints of cervical and lumbar pain, headache, and left shoulder radiculopathy"), references providers by name, and tracks treatment escalation from conservative care through injections to a recommended surgical consultation.

The "Gaps and Flags" section is where this gets strategically valuable. The AI identifies a 45-day gap between the plaintiff's last physical therapy session and her follow-up with the orthopedist – the gap defense counsel will cite to argue the injuries resolved. Sarah now knows to address it proactively in the demand before the adjuster asks.

The second prompt produces a specials table that Maria the paralegal would normally spend 90 minutes compiling from billing records.

What the Attorney Does

Sarah reviews the narrative against the chronology. She strengthens the language around the lumbar fusion recommendation – the AI was conservative, and the operative consultation note documented a clear surgical indication. She adds a sentence about the plaintiff's inability to lift her two-year-old daughter, noted in a functional capacity evaluation the AI flagged but didn't emphasize. She adjusts two entries in the specials table where PIP payments need to be reflected and adds lost wages manually.

Total time: 45 minutes from chronology to completed demand section. Previous workflow: 4-6 hours.

Verification note: Always verify AI-generated medical summaries and damages calculations against the source records before including them in any demand letter or filing. The AI extracts and organizes; the attorney confirms accuracy and completeness.

Workflow 2: Deposition Summary to Cross-Examination Outline

The Input

A Dodonai Deposition Summaries output in page-line format for the defendant's corporate representative in a premises liability case. The deponent testified about the property's inspection schedule, maintenance records, and prior incident reports. The summary is 12 pages with page-line citations.

The Prompt

```
I'm preparing cross-examination questions for the defendant's
corporate
representative in a premises liability case. Below is the page-line
deposition summary generated from the transcript.

CASE CONTEXT:
- Plaintiff slipped on a wet floor in the defendant's grocery store
- Defendant claims regular inspection and maintenance schedule
- Key issue: whether defendant had actual or constructive notice of
  the hazardous condition
- Deposition taken: 02/10/2026
- Deponent: James Whitfield, Store Manager and Corporate Rep

TASK:
Analyze this deposition summary and produce:

1. KEY ADMISSIONS: Testimony that supports the plaintiff's case.
   For each admission, quote the relevant summary entry with its
   page-line citation and explain why it matters.

2. INTERNAL CONTRADICTIONS: Places where the deponent contradicts
   himself within this testimony. For each contradiction, cite both
   statements with page-line references and frame the inconsistency.

3. CONTRADICTIONS WITH KNOWN FACTS: Based on the testimony, flag
   statements that likely conflict with:
   - Store surveillance footage (if it exists)
   - Maintenance logs or inspection records
   - Prior incident reports
   - Employee witness statements
   Note: I'll need to verify these against actual documents. Flag
   them as "areas to investigate."

4. CROSS-EXAMINATION QUESTIONS: For each admission and
   contradiction, draft 3-5 leading questions suitable for cross.
   Follow standard cross-examination form:
   - One fact per question
   - Leading (answerable with "yes" or "no")
   - Build from undeniable facts toward the key admission
   - Pin the witness to the prior testimony before revealing the
     contradiction
   Include the page-line citation after each question so I can
   impeach directly from the transcript.

5. AREAS FOR FOLLOW-UP DISCOVERY: Identify any documents,
   witnesses, or records referenced in the testimony that haven't
   been produced and should be requested.

[PASTE DODONAI PAGE-LINE DEPOSITION SUMMARY HERE]
```

What the AI Produces

The output is a structured cross-examination preparation package. A real example of what this produces:

Key Admission (example from output): > Whitfield testified that he personally conducts floor inspections “every hour on the hour” (p. 47, ln. 8-12). However, he also testified that on the day of the incident, he was in a management meeting from 2:00 PM to 3:30 PM (p. 83, ln. 4-9). The plaintiff fell at 2:45 PM. This means either the hourly inspection didn’t happen during the window when the hazard existed, or someone else was responsible for inspections – and Whitfield didn’t identify who.

Cross-Examination Sequence (example): > Q: Mr. Whitfield, you testified that floor inspections happen every hour. (p. 47, ln. 8-12) > Q: You are the person responsible for conducting those inspections. (p. 47, ln. 14-16) > Q: On January 12, 2026, you were in a management meeting. (p. 83, ln. 4-9) > Q: That meeting started at 2:00 PM. (p. 83, ln. 6) > Q: The meeting lasted until approximately 3:30 PM. (p. 83, ln. 9) > Q: My client fell at 2:45 PM. > Q: You were not on the floor at 2:45 PM. > Q: No one conducted a floor inspection between 2:00 PM and 2:45 PM.

That question sequence would take an experienced litigator 30-45 minutes to build manually after carefully reading the full summary and cross-referencing the timeline. The AI built it in two minutes.

What the Attorney Does

Sarah makes three changes. She reorders the contradiction sequence to establish the inspection policy firmly before revealing the meeting conflict. She adds a question about the maintenance log – which she knows from the production shows no entry between 1:00 PM and 4:00 PM. And she cuts two questions that are argumentative rather than leading.

She cross-references the “Contradictions with Known Facts” section against the actual document production. Two of the four flags turn out to be legitimate inconsistencies she hadn’t caught. One becomes the centerpiece of her cross.

Verification note: AI-generated cross-examination outlines are analytical aids, not finished work product. The attorney must verify every page-line citation against the original transcript and confirm that the logical inferences drawn from the testimony are legally sound.

Workflow 3: eDiscovery Results to Case Theory

The Input

Sarah is handling a premises liability case against a commercial property management company. Opposing counsel produced 3,200 pages of internal emails, maintenance records, and incident reports. Sarah ran them through Dodonai OCR (the scanned maintenance logs) and then indexed everything in Dodonai eDiscovery.

She's already run several semantic searches: - "Knowledge of unsafe conditions on the property" - "Prior complaints about the stairwell" - "Maintenance requests that were deferred or delayed" - "Budget decisions affecting property maintenance"

Each search returned ranked results with highlighted passages and source citations. She exported the top 20 results from each query – 62 unique documents after deduplication.

The Prompt

```
I'm developing a case theory for a premises liability case. Below are
the relevant documents identified through eDiscovery review. Each
document includes the source file name, the highlighted passage, and
the relevance ranking.

CASE CONTEXT:
- Plaintiff fell on a deteriorated exterior stairwell at a commercial
  office building
- Defendant: Meridian Property Management (property manager)
- Building owner: separate entity, not yet a party
- Key legal issue: Did Meridian have actual or constructive knowledge
  of the stairwell condition? Did they have a duty to repair?
- Jurisdiction: Texas -- premises liability standard requires proof
  that defendant knew or should have known of the condition

TASK:
Analyze these 62 documents and produce:

1. TIMELINE OF KNOWLEDGE: Build a chronological timeline showing
   when Meridian first learned of (or should have learned of)
   the stairwell condition. Cite each document by filename and
   relevant passage.

2. PATTERN ANALYSIS: Identify any pattern of deferred maintenance,
   budget cuts to property upkeep, ignored complaints, or similar
   conduct across the document set. Look for:
   - Repeated complaints about the same issue
   - Maintenance requests that were acknowledged but not completed
   - Budget communications that prioritized cost savings over safety
```

```
   - Communications showing awareness of risk

3. HOT DOCUMENTS: Identify the 5-10 most damaging documents for
   the defendant. For each, explain:
   - What it says
   - Why it's damaging
   - How it fits the case theory
   - Which witness should be questioned about it

4. MISSING PIECES: What evidence would strengthen this case theory
   but isn't in the current document set? What additional discovery
   should I pursue?

5. CASE THEORY MEMO: In 2-3 paragraphs, articulate the strongest
   case theory supported by this evidence. Write it as if you're
   explaining the theory to co-counsel joining the case.

[PASTE EDISCOVERY RESULTS HERE]
```

What the AI Produces

The timeline is the most immediately useful output. It shows that five tenants emailed Meridian about the stairwell over a seven-month period before the plaintiff's fall. The first complaint was eight months before the incident. Three of the five received a response acknowledging the issue and promising repairs. No repair was completed.

The pattern analysis reveals something Sarah hadn't connected across individual documents: Meridian's regional director sent a company-wide email four months before the incident directing property managers to "defer all non-emergency exterior maintenance through Q4 to meet budget targets." The stairwell repair request was classified as "non-emergency exterior" one week later.

The hot documents list prioritizes the regional director's budget email, the three acknowledgment responses to tenant complaints, and an internal inspection report from two months before the fall that rated the stairwell as "requires immediate attention."

What the Attorney Does

Sarah reads the case theory memo and recognizes the shape of her case: this isn't just "they knew and didn't fix it." It's "corporate policy caused them to defer a known safety hazard to meet quarterly budget targets." That reframing changes the trial narrative and potentially supports a gross negligence argument.

She uses the hot documents list to prepare deposition outlines for the property manager and the regional director. She drafts a supplemental discovery request targeting Q4 budget directives for other Meridian properties. The "Missing Pieces" section flags that no inspection records from the building owner were produced – Sarah files a motion to compel.

Step	Manual Approach	AI-Assisted Approach
Document review (3,200 pages)	2-3 days (paralegal)	2 hours (OCR + eDiscovery indexing via Dodonai)
Pattern identification	Half-day (attorney reads flagged docs)	8 minutes (AI analysis of exported results)
Hot document identification	Ongoing through review	Produced with pattern analysis
Case theory development	Attorney synthesizes after full review	Draft theory in 8 minutes, attorney refines in 30
Total	**3-4 days**	**Half a day**

Verification note: AI-identified "hot documents" and pattern analyses must be verified against the source documents. The AI may misinterpret passage context, miss privileged material, or draw inferences that don't hold up under closer reading. Every document flagged as strategically significant should be read in full by the attorney.

Workflow 4: Intake Notes to Go/No-Go Case Evaluation

The Input

Sarah gets roughly 15 new intake leads per month and signs about 8. The intake evaluation is one of the highest-leverage decisions in a PI practice – taking a weak case costs months of unrecoverable time, and declining a strong case leaves money on the table.

David, Sarah's part-time legal assistant, conducts the initial intake call and enters notes into Clio. The intake form captures: caller name, incident date, incident type, injuries described, treatment status, insurance information, defendant information, and David's notes on the call.

Sarah also has a Dodonai XD Agents template that extracts structured fields from any attached documents – police reports, initial medical records, insurance correspondence – so by the time she reviews an intake, the key facts are organized.

The Prompt

This prompt encodes Sarah's actual case acceptance criteria – the factors she weighs on every intake. This is a prime candidate for the "skills" approach from

Chapter 3: save it as a skill file and it becomes her firm's case evaluation methodology, applied consistently to every intake.

```
You are a case evaluation assistant for a PI practice in Austin, TX.
Analyze the intake information below and produce a go/no-go
recommendation memo.

FIRM'S CASE ACCEPTANCE CRITERIA:
- Minimum viable case value: $25,000 (we don't take cases below this)
- Clear or probable liability preferred (we evaluate contested
  liability cases individually)
- Treatment must be ongoing or recently completed (gaps >90 days
  from incident with no treatment are a red flag)
- Statute of limitations: 2 years in Texas for PI (flag if <6
  months remaining)
- Red flags: prior claims for same body part, attorney shopping
  (multiple prior attorneys), pre-existing conditions not
  disclosed on intake, incident on plaintiff's own property,
  government entity defendant (sovereign immunity / notice
  requirements)

INTAKE INFORMATION:
[PASTE INTAKE NOTES FROM CLIO AND XD AGENT EXTRACTION HERE]

PRODUCE:
1. CASE SUMMARY: 3-4 sentences summarizing the facts
2. LIABILITY ASSESSMENT: Clear / Probable / Contested / Unclear
   -- with reasoning
3. DAMAGES ESTIMATE: Conservative range based on injury type,
   treatment level, and comparable cases in Texas. Note: this is
   a rough screening estimate, not a formal valuation.
   - Medical specials (current and estimated future)
   - General damages multiplier range (1.5x-5x depending on
     severity, surgery, permanence)
   - Total estimated case value range
4. RED FLAGS: Any of the criteria above that are triggered
5. STATUTE CHECK: SOL deadline and time remaining
6. RECOMMENDATION: ACCEPT / DECLINE / INVESTIGATE FURTHER
   -- with a one-paragraph explanation
7. IF ACCEPT: Suggested next steps (records to request, witnesses
   to identify, experts to retain)
8. IF INVESTIGATE: What specific information would change the
   recommendation

Be conservative in your estimates. I'd rather decline a marginal
case and be wrong than accept one and lose six months. A "no" that
should have been a "yes" costs me one case fee. A "yes" that should
have been a "no" costs me 200 hours of unrecoverable time.
```

What the AI Produces

A structured memo that takes about 90 seconds to generate. For a typical intake – say, a rear-end collision with documented cervical and lumbar injuries, three months of chiropractic care, and a pending orthopedic referral – the output looks like this:

Case Summary: Caller reports rear-end collision on 01/22/2026 at the intersection of Lamar and 5th, Austin, TX. Police report confirms defendant at fault (citation issued). Caller describes neck and lower back pain with radiating symptoms into left arm. Currently treating with Dr. Reyes (chiropractic, 3x/week since 02/01/2026) with referral to Austin Spine Associates for orthopedic evaluation scheduled 04/15/2026.

Liability Assessment: Clear. Police report confirms defendant at fault. Rear-end collision with citation issued.

Damages Estimate: - Current medical specials: ~$4,800 (12 weeks chiropractic at ~$400/week) - Estimated future medical: $8,000-$25,000 (depending on orthopedic findings – if injections or surgical recommendation, specials increase significantly) - General damages multiplier: 2x-3.5x (soft tissue with radiating symptoms, no surgery yet) - **Estimated case value range: $25,000-$105,000**

Red Flags: None identified.

Statute Check: SOL expires 01/22/2028. 21 months remaining. No urgency.

Recommendation: ACCEPT. Clear liability, active treatment, orthopedic referral suggests potential escalation beyond conservative care. Low end of value estimate meets minimum threshold; if orthopedic evaluation leads to interventional treatment, value increases substantially.

What the Attorney Does

Sarah reads the memo in two minutes. Clean case. She accepts it, and the suggested next steps become Maria's task list for the new file.

On a different intake – one where the caller had a prior workers' comp claim for the same body part and hasn't treated in 60 days – the AI flags both red flags and recommends INVESTIGATE FURTHER: "Obtain prior workers' comp medical records to assess whether current complaints are distinguishable from the prior injury." That gives David a clear list of follow-up questions for the potential client.

Decision	Manual Process	AI-Assisted Process
Initial screening	15-20 min per intake (Sarah reviews notes, estimates value, checks SOL)	2 min (AI memo + Sarah's review)
Monthly intake volume	15 leads x 20 min = 5 hours	15 leads x 2 min = 30 minutes
Time saved per month	–	4.5 hours
Consistency	Varies by day (tired Fridays vs. fresh Mondays)	Same criteria applied every time

Verification note: AI-generated case valuations are screening estimates based on the information provided. They are not formal case evaluations and should not be

relied upon for client counseling, settlement authority, or fee agreements. The attorney's judgment, informed by jurisdiction-specific experience and the complete case file, determines case acceptance.

Workflow 5: Multiple Depositions to Trial Prep Matrix

The Input

Sarah is preparing for trial on a PI case involving a multi-vehicle accident. The case has 11 deposition transcripts: the plaintiff, three defendants, two eyewitnesses, the responding police officer, two treating physicians, an accident reconstructionist, and the defendant's biomechanical expert. Each transcript has been processed through Dodonai Deposition Summaries in page-line format and uploaded to Transcript Management.

The Prompt

This is the most complex workflow because it cross-references multiple documents. The AI needs all 11 summaries. For larger cases, you may need to feed them in batches or use a model with a longer context window.

```
I'm preparing for trial in a multi-vehicle PI case. Below are page-
line
deposition summaries for all 11 witnesses. I need a comprehensive
trial preparation matrix.

CASE SUMMARY:
- Three-vehicle chain collision on IH-35, Austin, TX
- Plaintiff (our client): rear passenger in Vehicle 2
- Defendant A: driver of Vehicle 1 (stopped suddenly)
- Defendant B: driver of Vehicle 3 (rear-ended Vehicle 2, pushing
  it into Vehicle 1)
- Key dispute: which impact caused plaintiff's injuries -- the
  initial rear-end from Vehicle 3 or the secondary front-end
  collision with Vehicle 1
- Each defendant blames the other for the mechanism of injury

WITNESSES:
1. Plaintiff - Angela Torres
2. Defendant A - Michael Park (Vehicle 1 driver)
3. Defendant B - Lisa Drummond (Vehicle 3 driver)
4. Defendant B's passenger - Ryan Drummond
5. Eyewitness 1 - Sandra Okafor (adjacent lane)
6. Eyewitness 2 - Thomas Briggs (behind Vehicle 3)
7. Officer Patricia Vega (responding officer)
8. Dr. James Okonkwo (treating orthopedist)
9. Dr. Helen Strauss (treating neurologist)
```

```
10. Dr. Andrew Kim (defense biomechanical expert)
11. Dr. Rebecca Liang (accident reconstruction, plaintiff's expert)

PRODUCE:

1. WITNESS-BY-WITNESS SUMMARY GRID:
   Create a table with one row per witness and columns for:
   - Witness name and role
   - Position on mechanism of injury (which impact caused injuries)
   - Key admissions favorable to plaintiff
   - Key admissions favorable to defense
   - Credibility issues identified
   - Page-line citations for critical testimony

2. CONTRADICTION MATRIX:
   For each material fact in dispute, create a grid showing what
   each witness testified. Focus on:
   - Sequence of impacts
   - Speed of vehicles
   - Time between impacts
   - Whether plaintiff's vehicle moved after initial impact
   - Plaintiff's complaints immediately after the collision
   - Whether plaintiff was bracing or unaware before impact

3. IMPEACHMENT OPPORTUNITIES:
   For each witness, identify:
   - Internal contradictions (within their own testimony)
   - Contradictions with other witnesses
   - Contradictions with documentary evidence (police report,
     medical records, accident reconstruction)
   - Prior inconsistent statements (if referenced in any deposition)
   Cite page-line for every identified inconsistency.

4. CROSS-EXAMINATION PRIORITIES:
   Rank the defense witnesses by importance and recommend the
   top 3-5 cross-examination objectives for each, with the
   supporting testimony citations.

5. TRIAL THEMES:
   Based on the testimony as a whole, identify the 2-3 strongest
   trial themes for the plaintiff. For each theme, list the
   witnesses and testimony that support it.

[PASTE ALL 11 PAGE-LINE DEPOSITION SUMMARIES HERE]
```

What the AI Produces

The witness summary grid alone would take a paralegal a full day to compile manually. The AI produces it in five to eight minutes.

The contradiction matrix is where this workflow delivers the most trial value. In a multi-vehicle case, the sequence of events is everything. The matrix might reveal:

Disputed Fact	Torres (Plaintiff)	Park (Def A)	Drummond (Def B)	Okafor (Eyewitness)	Briggs (Eyewitness)
Time between impacts	"Almost instant – one right after the other" (p.34, ln.8)	"Several seconds apart" (p.22, ln.14)	"I don't know, it all happened fast" (p.41, ln.3)	"Two separate hits, maybe 2-3 seconds apart" (p.18, ln.7)	"Simultaneous" (p.15, ln.22)
Vehicle 2 movement	"We got pushed forward into the car ahead" (p.35, ln.2)	"The car behind me rolled into me slowly" (p.23, ln.9)	"I barely tapped them" (p.42, ln.18)	"Vehicle 2 lurched forward significantly" (p.19, ln.1)	"Hard impact from behind, then they hit the car in front" (p.16, ln.5)

When you lay out six witnesses' accounts side by side, inconsistencies become visible in a way they never are reading summaries one at a time. Defendant B says she "barely tapped" Vehicle 2, but two independent eyewitnesses describe a significant impact. That discrepancy becomes the foundation for cross-examination.

The trial themes section synthesizes across the entire testimony record: (1) Defendant B was following too closely, supported by eyewitness testimony and the accident reconstructionist; (2) the rear impact was the primary mechanism of injury, supported by the treating orthopedist's testimony about posterior disc herniation consistent with rear-end forces; (3) Defendant B's minimization is contradicted by the physical evidence and every other witness.

What the Attorney Does

Sarah uses the witness summary grid as her trial binder index – one page that tells her where every key piece of testimony lives across 11 depositions. She reorganizes the cross-examination priorities: the AI recommended leading with Defendant B's speed, but Sarah knows the jury will respond more to the "barely tapped" minimization. She'll establish impact severity through the eyewitnesses first, then confront Defendant B with the contradiction.

She uses Transcript Management in Dodonai to verify every page-line citation. For a trial prep matrix, citation accuracy is non-negotiable – misquoting deposition testimony at trial is the fastest way to lose credibility with the judge and jury.

Verification note: Every page-line citation in a trial prep matrix must be verified against the original transcript before use in court. AI-generated contradiction analyses may identify apparent inconsistencies that resolve when the full

testimony context is read. Always read the cited passages in the original transcript, not just the summary.

The Manual vs. AI-Assisted Timeline

Across all five workflows, the pattern is consistent: the strategic analysis that used to consume most of the time becomes a review-and-refine task instead of a build-from-scratch task.

Workflow	Manual Time	AI-Assisted Time	Attorney Time Saved
Medical chronology to demand letter section	4-6 hours	45 minutes	3-5 hours
Deposition summary to cross-exam outline	2-3 hours	30 minutes	1.5-2.5 hours
eDiscovery results to case theory	3-4 days	Half a day	2-3 days
Intake evaluation	5 hours/month (15 intakes)	30 minutes/month	4.5 hours/month
11 depositions to trial prep matrix	2-3 days	3-4 hours	1.5-2.5 days

For Sarah Chen's practice at an effective rate of $364/hour, the demand letter workflow alone saves $1,100-$1,800 per case. Across 8 new cases per month, that's $8,800-$14,400 in recovered capacity – time Sarah can spend on case strategy, client communication, depositions, and trial prep. Or, if she prefers, it's 8-15 hours per month she's no longer working evenings and weekends.

The point is not that AI replaces attorney analysis. In every workflow, Sarah's review improved the output – catching nuances the AI missed, reordering arguments for strategic effect, adding context from experience. The point is that she's starting from a structured draft instead of a blank page. Review is faster than creation. And the strategic quality of the final work product is at least as good, because Sarah is spending her time on judgment calls instead of data organization.

The Skills Approach: Encoding Your Methodology

You've probably noticed something about the prompts in this chapter. They're long. They're specific. They include case acceptance criteria, demand letter

structure, cross-examination form rules, and jurisdiction-specific standards. That level of detail is why they work.

As Zack Shapiro put it when describing how he runs Raines LLP with two people and 200 clients: "The secret isn't the model's training data – it's the quality and specificity of the input. And lawyers' training in precise, detailed communication is a natural superpower for prompting."

But copying and pasting a 400-word prompt every time you evaluate an intake is not a sustainable workflow. Here's where the skills approach from Chapter 3 changes the equation.

Take the intake evaluation prompt from Workflow 4. It contains Sarah's case acceptance criteria, her damages estimation methodology, her red flag checklist, and her preferred output format. That's not a one-off prompt. That's her firm's case evaluation methodology, encoded in text.

In Chapter 3, we covered how skill files work: a plain-English instruction file your AI assistant references every time you run that task. The intake evaluation prompt becomes a skill called /case-eval. Sarah types /case-eval, pastes the intake notes, and the AI applies her criteria, her methodology, her preferred output format – automatically.

The same principle applies to every workflow in this chapter:

Workflow	Skill File	What It Encodes
Demand letter	/demand-narrative	Damages narrative structure, tone, specials table format, gap-flagging criteria
Cross-exam	/cross-prep	Impeachment analysis framework, leading question format, contradiction identification methodology
Case theory	/case-theory	Pattern analysis categories, hot document criteria, case theory memo structure
Intake eval	/case-eval	Case acceptance criteria, damages multiplier ranges, red flag checklist, SOL tracking
Trial prep	/trial-matrix	Witness grid format, contradiction matrix structure, cross-exam priority ranking

Each skill file is a 1-2 page markdown document. You write it once. It captures your analytical framework – not generic legal analysis, but the specific way your firm approaches these tasks. Your damages multiplier ranges. Your red flag criteria. Your preferred demand letter structure.

Then the skill file improves over time. After 30 intakes, Sarah notices the AI consistently underestimates future medical costs when surgery is recommended. She adds one line: "For intakes where surgery is recommended or likely, include surgical cost estimates of $15,000-$80,000 depending on procedure type." Every subsequent evaluation applies the updated methodology automatically.

This is what Shapiro means by "encoded expertise." It's not that the AI knows how to evaluate a PI case. It's that YOUR methodology is written down, formalized, and applied consistently – something most solo practitioners never have time to do, even though it's exactly what separates a well-run practice from one that operates on instinct and memory.

Dodonai Consulting helps firms build these skill files. If you know what your methodology should be but don't want to spend the time encoding it, that's the consulting engagement – translating your practice expertise into skill files, connecting them to your tools, scheduling them to run autonomously. That's the Dodonai Approach: encode expertise as skills, connect to existing tools, schedule for autonomy, layer in human safety checks.

What AI Won't Do Here

Every workflow in this chapter involves AI generating analysis from structured documents. Here's what requires your judgment, not the AI's:

Legal strategy decisions. The AI can identify contradictions between witnesses. It cannot decide which contradiction to lead with at trial. It can draft a damages narrative. It cannot decide whether to demand $150,000 or $250,000. These are attorney judgment calls informed by experience, jurisdiction-specific knowledge, and the specific judge, jury pool, and opposing counsel you're facing.

Credibility assessments. The AI can note that a witness changed their account between page 34 and page 78. It cannot assess whether that change reflects deception, confusion, or a clarification prompted by counsel.

Ethical compliance. The AI doesn't know your bar's rules on AI-generated work product or privilege implications. Chapter 10 covers the ethics framework in detail.

Client communication. The AI can draft a case evaluation memo. It cannot deliver the news that a case doesn't meet your criteria or counsel a client on whether to accept a settlement offer. The human relationship is yours.

Try This Now: Your First Document-to-Strategy Workflow

Pick one of the five workflows and try it on a real case.

Easiest starting point: Workflow 4 (intake evaluation). Take your most recent intake notes, paste them with the case evaluation prompt, and see what the AI

produces. You don't need Dodonai for this one – just your intake notes and an AI assistant.

If you're already using Dodonai: Try Workflow 1 or 2. Process a medical record set or deposition transcript, then paste the output into your AI assistant with the corresponding prompt. Note where the AI gets it right, where it misses, and what you'd change in the prompt.

The real exercise: Turn the prompt into a skill file. Save it as a markdown file with your firm's specific criteria, your preferred format, your analytical framework. That's the step that turns a one-off experiment into a practice methodology.

Here's a template to get started:

```
# /[your-skill-name] -- [One-line description]

## Context
[Why this workflow exists. What problem it solves. When to run it.]

## Inputs Required
[What documents or data the AI needs to do this analysis.]

## Instructions
[Your full prompt -- the analysis methodology, output format,
specific criteria, and verification requirements.]

## Output Format
[Exactly what the AI should produce -- sections, tables, structure.]

## Verification Checklist
[What the attorney must check before using the output.]
```

Once you've saved your first skill file, you've formalized your analytical methodology – something most solo and small-firm attorneys never have time to do. Unlike a process that lives in your head, a skill file applies consistently, improves incrementally, and runs whether you're in the office or not.

Quick Reference: The Five Workflows

#	Workflow	Input (Dodonai)	AI Output	Attorney Action
1	Medical chronology to demand letter	Medical Record Summaries & Chronologies output	Damages narrative + medical specials table	Review, adjust severity language, verify specials, add lost wages
2	Deposition to cross-exam	Deposition Summaries (page-line format)	Admissions, contradictions, cross-exam questions with citations	Reorder sequence, add documentary evidence, cut argumentative questions

#	Workflow	Input (Dodonai)	AI Output	Attorney Action
3	eDiscovery to case theory	eDiscovery search results (exported)	Knowledge timeline, pattern analysis, hot documents, case theory memo	Verify hot docs, develop deposition outlines, draft supplemental discovery
4	Intake to go/no-go	XD Agents extraction + intake notes	Case summary, liability/damages assessment, red flags, recommendation	Review valuation, accept/decline, assign follow-up tasks
5	Multiple depositions to trial prep	Multiple Deposition Summaries + Transcript Management library	Witness grid, contradiction matrix, impeachment map, trial themes	Verify citations, reorder cross-exam strategy, annotate for trial binder

The pattern: Process (Dodonai) -> Analyze (AI) -> Act (You). Every time.

Chapter 6: Seven Practice Areas, One Architecture

What You'll Walk Away With:

- How the three-layer architecture from Chapter 3 applies to seven distinct practice areas – with no structural changes
- Practice-specific agent workflows for personal injury, real estate, tax disputes, cybersecurity, immigration, estates, and family law
- Prompts and automation patterns using each area's native terminology
- A summary table proving the underlying pattern is always the same: documents in, deadlines computed, things filed, people notified
- Competing tools in each area – what the market is building, and why you can build it yourself

---

The Same Architecture, Every Time

Chapters 3 through 5 built an architecture and demonstrated it through personal injury workflows. If you're a PI attorney, you might be thinking this book was written for you. If you practice tax law or immigration or estates, you might be wondering whether any of this applies.

It all applies. Every word.

The architecture from Chapter 3 – skill files, connections, schedules, plus the safety layer – doesn't care about practice area. It cares about a pattern. And the pattern is the same in every area of law:

Documents come in. Deadlines get computed. Things get filed. People get notified.

A PI firm receives medical records from eight providers and computes a statute of limitations. An immigration firm receives an IRCC status update and computes a biometrics window. An estates firm receives an executor's document package and computes beneficiary notification deadlines. A cybersecurity firm receives breach details and computes notification deadlines across 50 states.

Different documents. Different deadlines. Different filing requirements. Different people to notify. Same architecture.

This chapter walks through seven practice areas – each one built from the same three layers you already understand. By the end, you'll see that the question isn't "does this work for my practice area?" It's "which workflow do I automate first?"

Personal Injury: From Records Chaos to Demand-Ready Cases

The Pain

Every PI attorney knows this math. A single auto accident case involves medical records from 4 to 8 providers. Each provider sends 50 to 150 pages. A moderately complex case lands you 200 to 500 pages of records, bills, imaging reports, and correspondence – arriving at different times, in different formats, from different facilities.

If you outsource the medical chronology, you're paying $500 to $800 per case. If your paralegal builds it in-house, it takes 6 to 10 hours of manual review. Multiply that across 30 to 40 active cases and you're either bleeding cash to vendors or burying your best paralegal in records organization instead of case strategy.

Meanwhile, the statute of limitations is ticking. Treatment gaps need to be identified before they become defense ammunition. And when the records are finally organized, someone still has to draft the demand.

The Agent Workflow

The PI records agent follows the same three-layer pattern from Chapter 3:

Skill: When new medical records arrive (via email attachment, client portal upload, or provider fax-to-email), classify each document by type – treatment records, billing statements, imaging reports, pharmacy records, correspondence. Upload to Dodonai for chronology processing. Once the chronology is complete, scan for treatment gaps longer than 30 days and flag them. Track the statute of limitations for each case. When all expected providers have submitted records, draft a demand package outline.

Connections: Gmail or Outlook (incoming records), Dodonai (Medical Record Summaries & Chronologies), Clio or your case management system (case details, provider list, SOL dates), Slack or Teams (alerts).

Schedule: Check for new records daily. SOL monitoring runs weekly with escalating alerts at 180, 90, 60, and 30 days.

What This Looks Like in Practice

Gordon Herrin runs a PI firm on the Northshore of Louisiana – St. Tammany Parish, where hurricane season doesn't just bring storms, it brings a wave of insurance claims and personal injury cases. When a major storm hits, intake can double overnight. Every claimant needs an insurance demand, and if the carrier stonewalls, a FEMA appeal or a lawsuit. Same fact pattern repeated dozens of times, but every detail has to be right.

Here's the agent at work during storm season:

```
You are a medical records intake agent for a personal injury
practice.

WHEN: New email attachments arrive from medical providers
(match sender domains against the provider list in
[CASE_MANAGEMENT_SYSTEM])

DO:
1. Classify each attachment:
   - Treatment records (progress notes, operative reports, discharge
summaries)
   - Billing statements (itemized charges, EOBs, lien letters)
   - Imaging (radiology reports, MRI/CT scan reports)
   - Pharmacy records
   - Correspondence (IME notices, peer review letters)

2. Match to the correct case by patient name and DOI

3. Upload treatment records and imaging to Dodonai for chronology
processing

4. After chronology is complete, scan for:
   - Treatment gaps > 30 days (flag for attorney review -- defense
will exploit these)
   - Missing providers (compare against the known treater list)
   - Inconsistencies between billing and treatment dates

5. Update the case status: "[X] of [Y] providers received.
   Missing: [provider names]. SOL: [date] ([days] remaining)."

6. When all providers are received, alert:
   "Records complete for [case name]. Chronology ready.
   Ready for demand preparation."
```

During hurricane season, this agent processes 50 claimants as easily as 5. The skill file doesn't change. The volume just scales.

```
You are a demand preparation agent for a personal injury practice.

CONTEXT: All medical records have been received and processed through
Dodonai Medical Record Summaries & Chronologies for the following
case:

[Paste Dodonai chronology output here]

CASE DETAILS:
```

```
- Plaintiff: [name]
- DOI: [date]
- Liability: [brief liability summary]
- Policy limits: [amount]
- Pre-existing conditions: [if any]

TASK:
1. Extract total medical specials from the chronology
   (separate by provider, by date, by type)
2. Identify the treatment arc: initial ER/urgent care → specialist
referrals
   → ongoing treatment → MMI or current status
3. Flag any gaps, inconsistencies, or missing records that need
resolution
   before sending the demand
4. Draft a damages narrative (3-5 paragraphs) that walks the adjuster
   through the injury, treatment, and impact on daily life

OUTPUT: A demand preparation memo with specials table, treatment
timeline summary, flagged issues, and draft damages narrative.

NOTE: Attorney must verify all medical specials against source
records,
review the damages narrative for accuracy, and add lost wages
and non-economic damages analysis before finalizing.
```

The Competition

This workflow isn't theoretical. Supio – acquired by Thomson Reuters for $91 million in 2025 – built its entire business on AI-powered medical record review for PI firms. Tavrn offers five modules covering document processing through case analysis. Anytime AI focuses on PI document automation. These companies are building purpose-built products for the exact workflow described above.

The difference: those products lock you into their platform and their pricing. The agent approach gives you the same workflow using your existing tools, at a fraction of the cost, with skill files you control and can modify as your practice evolves.

Real Estate: Closing Conditions That Track Themselves

The Pain

A commercial real estate closing involves dozens of conditions scattered across 20 or more emails from title companies, lenders, surveyors, environmental consultants, and opposing counsel. Title docs trickle in over weeks. Someone has to manually match each incoming document against the closing checklist, update

the timeline, and flag what's still missing – usually the day before closing, when the missing estoppel certificate becomes a crisis.

For cross-border transactions, add compliance requirements that vary by jurisdiction. A deal spanning Hawaii and Asia requires different disclosure obligations, tax withholding rules, and regulatory filings on each side. Miss a compliance requirement and the deal doesn't close – or worse, it closes and the compliance problem surfaces later.

The Agent Workflow

Skill: Monitor the attorney's inbox for emails related to active closings (match by deal name, property address, or counterparty). When a document arrives, classify it against the closing checklist. Update the checklist status. If a condition is now satisfied, mark it complete. If a deadline is approaching with an outstanding condition, escalate. For cross-border deals, cross-reference incoming documents against the jurisdiction-specific compliance checklist. Generate a daily closing status report for each active deal.

Connections: Gmail or Outlook (incoming documents and correspondence), Clio or deal management system (closing checklists, counterparty contacts), document management system (filed documents), Slack or Teams (alerts).

Schedule: Check for incoming closing documents twice daily (morning and afternoon). Generate closing status reports daily at 7:00 AM.

What This Looks Like in Practice

David Austin practices commercial real estate out of Honolulu, handling acquisitions, sales, leasing, and financing across Hawaii, the mainland US, and Asia. He already thinks in workflows – he built Nova Lease, a technology-driven commercial leasing platform with an online portal for real-time lease progress tracking, version control, and document management.

For someone who's already built a tech-enabled leasing portal, the closing checklist agent is a natural extension. The pattern David built for Nova Lease – tracking conditions, flagging what's outstanding, keeping all parties informed – is exactly what an AI agent does for every closing:

```
You are a closing conditions tracker for a commercial real estate
practice.

ACTIVE DEAL: [Deal name / property address]
CLOSING CHECKLIST:
[Paste or reference the closing conditions list from your deal
management system]

WHEN: New emails arrive matching this deal
(match by property address, deal name, or counterparty email domains)
```

```
DO:
1. Classify the attachment or email content against the closing
checklist:
   - Title commitment / title insurance
   - Survey / ALTA survey
   - Environmental (Phase I, Phase II)
   - Estoppel certificates
   - Tenant subordination agreements
   - Lender approval / commitment letter
   - Government approvals / permits
   - Entity formation documents

2. If the document satisfies a checklist condition, mark it RECEIVED
   with the date and source

3. If a condition deadline is within 5 business days and status is
   still OUTSTANDING, escalate: "[Deal name]: [condition] due in
   [N] days -- not yet received from [responsible party]."

4. For cross-border deals, flag any jurisdiction-specific
requirements:
   - FIRPTA withholding obligations (foreign seller)
   - State/country-specific disclosure requirements
   - Tax treaty considerations
   - Regulatory approvals required

5. Generate daily status: "[Deal name]  --  [X] of [Y] conditions
met.
   Outstanding: [list]. Next deadline: [condition] due [date]."
```

The Competition

Gavel offers template-based document automation that covers parts of this workflow. Clio Manage tracks deal timelines and deadlines. But neither provides the inbox monitoring, automatic document-to-checklist matching, or cross-border compliance flagging that the agent approach delivers. The closing checklist agent sits on top of whatever deal management tools you already use – it doesn't replace them, it watches the inbox and keeps them current.

Tax Disputes: Deadlines That Don't Forgive

The Pain

Tax dispute work runs on deadlines that are both inflexible and consequential. In Canada, a taxpayer has 90 days from the date of a Notice of Assessment to file a Notice of Objection with CRA. Miss it, and your client's right to object is gone – they're looking at a Tax Court application for extension, which is discretionary and far from guaranteed. In the United States, the IRS imposes similar windows: 90

days to petition the Tax Court after a Notice of Deficiency, 30 days to respond to proposed adjustments.

These deadlines arrive buried in dense correspondence. A single practitioner managing 100 or more active CRA files is tracking dozens of overlapping deadline windows, each computed from a different Notice of Assessment date. The computation itself is straightforward – 90 days from the NOA date. The risk is that a notice arrives on a Thursday, sits in a pile through a busy trial week, and by the time someone calendars it, three weeks of the window have already elapsed.

The Agent Workflow

Skill: Monitor incoming correspondence (email, scanned mail) for CRA Notices of Assessment, IRS Notices of Deficiency, or other regulatory correspondence with embedded deadlines. Extract the assessment details: taxpayer name, assessment date, assessment type, amount assessed, and the deadline trigger date. Compute the objection or response deadline. Draft the response document from the firm's precedent template. Set a reminder cascade at 60, 30, 14, and 7 days.

Connections: Gmail or Outlook (incoming CRA/IRS correspondence), Clio or case management (client file, deadline tracking), document management (precedent templates, filed objections), Slack or Teams (deadline alerts).

Schedule: Check for new correspondence daily. Deadline monitoring runs daily with escalating alerts.

What This Looks Like in Practice

Brandon Siegal runs a boutique tax dispute firm in Toronto. Before founding his practice, he spent five years as Department of Justice counsel representing CRA on over 600 litigation files. He has seen both sides of a missed deadline – as the government lawyer who benefits when a taxpayer's objection is late, and now as the private practitioner whose clients depend on timely filings.

For a practice managing hundreds of active CRA files, the correspondence monitoring agent is the difference between systematic deadline tracking and a spreadsheet that someone forgot to update:

```
You are a tax correspondence monitoring agent for a tax dispute
practice.

WHEN: New email or scanned document arrives from:
- Canada Revenue Agency (CRA) domains or letterhead
- Internal Revenue Service (IRS) domains or letterhead
- Tax Court of Canada / US Tax Court

DO:
1. Classify the document:
   - Notice of Assessment (NOA) / Notice of Reassessment
```

```
   - Notice of Deficiency (IRS)
   - Notice of Confirmation
   - Proposed adjustment letter
   - Requirement to Pay / Collection notice
   - Tax Court hearing notice / scheduling order

2. Extract:
   - Taxpayer name and file number
   - Assessment date and tax year(s)
   - Amount assessed or in dispute
   - Deadline trigger: the date from which response periods are
computed

3. Compute deadlines:
   - NOA → 90-day objection window (CRA) or 90-day Tax Court petition
(IRS)
   - Proposed adjustment → 30-day response window
   - Collection notice → flag for immediate review
   - Hearing notice → appearance date + pretrial deadlines

4. Match to client file in [CASE_MANAGEMENT_SYSTEM]

5. Set reminder cascade: 60 days, 30 days, 14 days, 7 days, 3 days

6. Alert: "[Client name]  --  [Document type] received [date].
   Objection deadline: [date] ([days] remaining).
   Tax year: [year]. Amount: [amount]."

IMPORTANT: All computed deadlines must be verified by the attorney
against the source document. Holiday and weekend adjustments
vary by jurisdiction and must be confirmed manually.
```

The OECD has recognized AI in tax administration as a major use case – governments are already using it to flag audit targets and compute assessments. This is the taxpayer side of the same trend: using AI to monitor, extract, compute, and respond to the correspondence that tax authorities are increasingly generating through their own automated systems.

The Competition

Thomson Reuters ONESOURCE covers tax compliance and research but isn't built for the dispute-side workflow of monitoring correspondence, computing response deadlines, and drafting objections. The agent approach fills a gap that enterprise tax software doesn't address because it's designed for compliance, not litigation.

Cybersecurity & Privacy: 50 States, 50 Deadlines, One Agent

The Pain

When a client reports a data breach, the clock starts running in every affected jurisdiction simultaneously. And the clocks aren't synchronized.

As of 2026, all 50 US states, the District of Columbia, Guam, Puerto Rico, and the US Virgin Islands have their own breach notification statutes. The deadlines range from 24 hours (certain incidents under some state laws) to 90 days. Content requirements vary – some states require specific language about the type of data exposed, others mandate offering credit monitoring, others require notifying the state Attorney General before notifying affected individuals. Some states have a separate notification obligation to the AG's office. Some require notification to consumer reporting agencies if the breach exceeds a threshold number of affected residents.

A single breach affecting customers in 15 states means 15 different deadlines, 15 different content requirements, 15 different filing procedures. Managing this manually is error-prone, time-intensive, and the cost of getting it wrong is regulatory enforcement and civil liability.

This is, as one practitioner put it, the most jurisdiction-complex problem in legal practice.

The Agent Workflow

Skill: When a breach incident report is received, extract the breach details: type of data exposed, number of affected individuals, states where affected individuals reside, date of breach discovery. Generate a compliance matrix showing, for each affected state: notification deadline (computed from discovery date), content requirements, AG notification requirements, threshold triggers (number of affected residents), and filing procedures. Draft notification letters for each jurisdiction. Track filing status. Set deadline cascades.

Connections: Email or intake form (breach incident details), compliance database or jurisdiction reference (state-by-state notification requirements), document management (draft notifications, filed notices), Clio or case management (client file, deadline tracking), Slack or Teams (deadline alerts).

Schedule: Once triggered by a breach incident, the agent runs daily to track filing progress and escalate approaching deadlines.

What This Looks Like in Practice

Kellam Parks practices cybersecurity and data privacy out of Virginia Beach. He chaired the Virginia State Bar's Special Committee on Technology and the Future Practice of Law. He built a fully paperless, cloud-based firm from day one. He teaches CLE courses on technology in law practice nationally. If anyone in a room of lawyers understands the architecture behind AI agents, it's the person who's been building technology-driven legal workflows for over a decade.

For Parks, the breach notification matrix is the killer use case – not because the legal analysis is complex (each state's statute is clear on its face), but because the operational challenge of tracking 15 or 30 jurisdictions simultaneously is where mistakes happen:

```
You are a breach notification compliance agent for a cybersecurity
and data privacy practice.

BREACH INCIDENT REPORT:
- Client: [company name]
- Date of breach discovery: [date]
- Type of data exposed: [PII categories  --  SSN, financial, health,
etc.]
- Number of affected individuals by state:
  [State]: [count]
  [State]: [count]
  ...

FOR EACH AFFECTED STATE, GENERATE:
1. NOTIFICATION DEADLINE
   - Statute citation
   - Deadline computation from discovery date
   - Any "without unreasonable delay" language
     (note: flag these  --  they require attorney judgment on
reasonableness)

2. CONTENT REQUIREMENTS
   - Required elements in notification letter
   - Credit monitoring obligation (yes/no, duration if required)
   - Specific language mandated by statute

3. AG / REGULATORY NOTIFICATION
   - Is separate AG notification required?
   - AG notification deadline (before or concurrent with individual
notice?)
   - Filing method (online portal, email, mail)
   - Threshold trigger (e.g., notify AG if > 500 residents affected)

4. CONSUMER REPORTING AGENCY NOTIFICATION
   - Required if breach exceeds [threshold] residents in the state?

OUTPUT: Compliance matrix (table format), draft notification letters
per jurisdiction, filing tracker with status column.

IMPORTANT: State breach notification laws are amended frequently.
All deadlines and content requirements must be verified against
```

```
current statutes before any notification is sent or filed.
Attorney review is mandatory before any external communication.
```

The Competition

BigID, OneTrust, and Exterro all offer breach response automation as part of larger privacy management platforms. These are enterprise-grade tools with enterprise-grade pricing. For a mid-size or boutique cybersecurity practice, the agent approach delivers the same compliance matrix and tracking workflow at a fraction of the cost, using the firm's existing tools and the attorney's own jurisdiction-specific knowledge encoded in the skill file.

Immigration: Where Every Milestone Is a Deadline

The Pain

Immigration law is deadline-driven in a way that few other practice areas match. A single client file can have a dozen concurrent deadlines: biometrics appointment windows that expire, medical certificates with 12-month validity that must still be valid at the time of landing, hearing dates with strict filing deadlines for supporting materials, work permit renewals that must be filed before the current permit lapses, and status changes that trigger new deadline cascades.

In Canada, IRCC portal updates arrive without notification – the practitioner has to log in and check. In the United States, USCIS case status updates are slightly more automated but still require monitoring. A solo practitioner managing 50 or more active files is logging into portals daily, manually tracking which biometrics windows are closing, which medical certs are expiring, and which hearing dates are approaching.

It's not surprising that immigration leads AI adoption at 47% – the highest of any practice area, according to Wolters Kluwer's 2026 survey. The practice is perfectly structured for automation: repetitive monitoring tasks, date-driven workflows, standardized forms, and milestone-based client communication.

The Agent Workflow

Skill: Check immigration portals daily for status changes. For each active client file, track: current application stage, biometrics status (scheduled, completed, expired), medical certificate validity, hearing dates, work permit expiry, and any pending requests for additional documentation. When a status change occurs,

compute downstream deadlines. Draft a client update when a milestone is reached. Flag urgently when a deadline window is closing.

Connections: IRCC/USCIS portals (status monitoring – via browser automation or API where available), Gmail or Outlook (client communication, portal notification emails), Clio or immigration case management (client files, deadlines), document management (forms, supporting documents), Slack or Teams (alerts).

Schedule: Portal monitoring daily at 7:00 AM. Deadline cascade check daily. Client update drafts generated on milestone triggers.

What This Looks Like in Practice

Debbie Rachlis practices immigration and refugee law as a solo in Toronto. She teaches immigration law at Queen's University's Graduate Diploma program and organizes the Junior Refugee & Immigration Lawyers Network. As a solo managing active IRCC files, every milestone is tracked personally – biometrics windows, medical cert expiry, hearing dates, BOC submissions.

For a practice built on milestone tracking, the immigration monitoring agent replaces the manual portal checks and spreadsheet deadline tracking that consume the first hour of every morning:

```
You are an immigration case monitoring agent for an immigration
and refugee law practice.

DAILY CHECK (run at 7:00 AM):

FOR EACH ACTIVE CLIENT FILE:
1. Check IRCC portal / USCIS case status for updates
   - Flag any status change (new stage, request for docs, decision)

2. Compute deadline windows:
   - Biometrics: appointment scheduled? Completed? Window expiring?
   - Medical certificate: exam date + 12-month validity
     → flag at 60 days before expiry if application still pending
   - Hearing date: days remaining, filing deadlines for supporting
materials
   - Work permit: current expiry date, renewal filing deadline
   - BOC (Basis of Claim): 15-day deadline from referral

3. If status changed:
   - Draft client update: "[Client name]  --  your application status
     has changed to [new status]. Here's what this means and
     what happens next: [plain language explanation]."
   - Compute any new deadlines triggered by the status change
   - Update case management system

4. Daily summary:
   - Applications with status changes: [list]
   - Deadlines within 30 days: [list with days remaining]
   - Biometrics/medical expiring: [list]
   - No changes (but monitored): [count] files checked
```

```
IMPORTANT: Portal status information should be verified by the
practitioner. Deadline computations for hearing materials and
BOC submissions must be confirmed against the specific
tribunal's scheduling order or referral notice.
```

The Competition

DocketWise and VisaLaw are building exactly this – purpose-built immigration case management with automated status tracking and deadline computation. They're proving the market demand for this workflow. The agent approach offers the same functionality but integrates with whatever tools the practitioner already uses, without migrating to a new platform. For a solo practitioner like Rachlis who already has a system that works, the agent adds the monitoring and deadline layer without replacing everything underneath.

Estates & Probate: The 12-Month Lifecycle

The Pain

A probate file is a 12-month lifecycle with sequential deadlines. Unlike litigation – where deadlines are triggered by opposing counsel's filings and can arrive unpredictably – estate administration follows a largely predictable sequence: executor appointment, beneficiary notification, asset inventory, creditor notification, court filings at specific intervals, interim accountings, and final distribution.

The predictability is a double-edged sword. Because the deadlines are known in advance, there's an assumption that they'll be tracked. When a firm is managing 20 or 30 active estate files, each at a different stage of the lifecycle, the "known" deadlines start overlapping. The beneficiary notification deadline for the Smith estate falls in the same week as the court accounting for the Jones estate and the distribution date for the Patel estate. None of these are surprises. All of them require action. And if anyone drops the ball, the executor – your client – is personally liable.

The Agent Workflow

Skill: When a new estate file opens, generate the complete lifecycle timeline based on jurisdiction-specific rules: notification deadlines, filing dates, accounting periods, distribution windows. Create notification letter drafts for beneficiaries. Generate filing checklists for each stage. Track executor document intake (will, death certificate, asset statements, insurance policies). Monitor the lifecycle and alert as each stage approaches.

Connections: Clio or estate-specific case management (estate files, beneficiary lists, asset inventories), Gmail or Outlook (executor correspondence, beneficiary notifications), document management (will, death certificate, filed documents), accounting system (estate accounting, distribution tracking), Slack or Teams (deadline alerts).

Schedule: New estate setup runs on file opening. Lifecycle monitoring runs weekly. Escalating alerts at 30, 14, and 7 days before each filing deadline.

What This Looks Like in Practice

Alex Bainov practices estates and business law from Winnipeg. He's a member of STEP – the Society of Trust and Estates Practitioners – an organization whose members think in exactly the kind of lifecycle terms that make AI agent automation a natural fit. Estate administration isn't about reacting to unexpected filings. It's about executing a known sequence on time, every time, across every active file.

For a STEP practitioner, the estate lifecycle agent is less of a conceptual leap and more of an obvious next step – encoding the lifecycle they already manage mentally into a system that tracks it automatically:

```
You are an estate administration lifecycle agent.

NEW ESTATE FILE OPENED:
- Deceased: [name]
- Date of death: [date]
- Jurisdiction: [province/state]
- Executor: [name, contact]
- Estimated estate value: [amount]
- Will located: [yes/no]

GENERATE LIFECYCLE TIMELINE:

1. IMMEDIATE (Days 1-14):
   - [ ] Obtain death certificate (original + [N] certified copies)
   - [ ] Locate and secure the will
   - [ ] Notify beneficiaries (deadline: [jurisdiction-specific,
typically 10-30 days])
   - [ ] Draft beneficiary notification letters
   - [ ] Secure estate assets (change locks, redirect mail, freeze
accounts)

2. PROBATE APPLICATION (Days 14-60):
   - [ ] Prepare estate inventory (assets + liabilities)
   - [ ] File probate application / Letters of Administration
   - [ ] Publish creditor notification (deadline: [jurisdiction-
specific])
   - [ ] Notify known creditors directly

3. ADMINISTRATION (Days 60-270):
   - [ ] Collect estate assets
   - [ ] Pay debts and administration expenses
   - [ ] File tax returns (terminal return deadline: [date])
   - [ ] Obtain tax clearance certificate
```

```
   - [ ] Prepare interim accounting if required

4. DISTRIBUTION (Days 270-365):
   - [ ] Prepare final estate accounting
   - [ ] Obtain beneficiary approval of accounting
   - [ ] File court accounting if required
   - [ ] Distribute estate assets per will / intestacy rules
   - [ ] Close estate accounts

FOR EACH ITEM: Track status (pending/in-progress/complete),
responsible party, deadline, and any dependencies.

IMPORTANT: Estate administration timelines and requirements vary
significantly by jurisdiction. All deadlines must be verified
against the applicable provincial/state rules. Tax return
deadlines should be confirmed with the estate's accountant.
```

The Competition

Gavel offers document automation templates for estate planning documents. EstateMin provides estate administration tracking with task management and beneficiary portals. These tools validate the demand for lifecycle-based estate automation. The agent approach provides the same lifecycle tracking while connecting to the firm's existing case management and document systems – no migration, no new platform to learn.

Family & Collaborative Law: The Paper Trail After Every Session

The Pain

Collaborative family law is one of the most process-heavy areas of legal practice. Every four-way meeting – both spouses, both attorneys, sometimes a coach, sometimes a financial neutral – generates a cascade of follow-up work. Session summaries need to be drafted and distributed to all parties. Financial disclosure trackers need to be updated. Revised agreements need to incorporate what was discussed. Disclosure reminders need to go to both sides. And the next session needs to be scheduled with an agenda that reflects where things stand.

The collaborative process is designed to keep people out of court. But the administrative burden of maintaining the process – the paper trail that makes it work – falls on the lawyers. After a two-hour four-way, the attorney might spend another two hours on session notes, tracker updates, draft revisions, and distribution emails. That's not legal work. That's production work.

And it's not just volume. It's coordination. The session summary goes to four or more people. Each person needs different pieces – the coach needs the emotional dynamics section, the financial neutral needs the updated asset table, each attorney needs the revised terms and their client's action items. Getting the right information to the right people after every session is a distribution problem that compounds with every meeting.

The Agent Workflow

Skill: After each collaborative session, generate: (1) a session summary organized by topic discussed, agreements reached, and open items; (2) updated financial disclosure tracker with any new information exchanged; (3) revised agreement sections incorporating new terms discussed; (4) action items for each party; (5) disclosure reminders for outstanding items. Draft distribution emails to each participant with their relevant sections. Prepare the next session agenda from open items.

Connections: Meeting notes or transcription (session input), document management (agreements, trackers, prior session summaries), Gmail or Outlook (distribution drafts), Clio or case management (matter details, participant contacts), Slack or Teams (internal alerts).

Schedule: Triggered after each collaborative session (manual trigger or transcription completion trigger). Disclosure reminder checks run weekly.

What This Looks Like in Practice

In collaborative practice, the attorney's value is in the room – reading the dynamics, managing the negotiation, keeping the process constructive when emotions escalate. The paper trail afterward is essential but mechanical. AI handles the paper trail so the attorney can focus on the people.

```
You are a collaborative law session follow-up agent.

SESSION COMPLETED:
- Case: [case name]
- Date: [session date]
- Participants:
  - [Spouse 1 name] + [Attorney 1 name]
  - [Spouse 2 name] + [Attorney 2 name]
  - [Coach name, if applicable]
  - [Financial neutral name, if applicable]
- Session notes / transcript: [attached or referenced]

GENERATE:

1. SESSION SUMMARY (distribute to all participants):
   - Topics discussed (bulleted, factual, neutral tone)
   - Agreements reached (specific terms, not generalities)
   - Open items requiring further discussion
```

```
   - Next steps for each party

2. UPDATED DISCLOSURE TRACKER:
   - New financial information exchanged this session
   - Outstanding disclosure items by party
   - Deadlines for outstanding items

3. REVISED AGREEMENT SECTIONS:
   - Draft updated language for any terms agreed upon
   - Track changes from prior version
   - Flag any terms that were discussed but not finalized
     (mark as "UNDER DISCUSSION  --  NOT AGREED")

4. ACTION ITEMS BY PARTY:
   - [Spouse 1]: [list]
   - [Spouse 2]: [list]
   - [Attorney 1]: [list]
   - [Attorney 2]: [list]
   - [Coach]: [list]
   - [Financial neutral]: [list]

5. DISTRIBUTION DRAFTS:
   - Draft email to each participant with their relevant sections
   - Each email includes: session summary, their action items,
     disclosure reminders specific to them

6. NEXT SESSION AGENDA (draft):
   - Open items carried forward
   - New items flagged during this session
   - Disclosure items due before next session

IMPORTANT: Session summaries in collaborative practice are shared
documents. Tone must be neutral and factual. Do not characterize
either party's position or emotional state. Agreements must be
clearly distinguished from items still under discussion.
Attorney must review all outputs before distribution.
```

The Competition

Deliberately.ai is building AI-powered tools specifically for collaborative and family law workflows. Gavel offers document automation templates for family law agreements. These purpose-built tools reflect the reality that collaborative practice has unique process requirements that general-purpose case management software doesn’t address well. The agent approach delivers similar automation while keeping the attorney in control of the workflow design – particularly important in collaborative practice, where every practitioner has a slightly different process for managing sessions and distribution.

The Unifying Pattern

Seven practice areas. Seven different sets of documents, deadlines, filing requirements, and people to notify. Seven different vocabularies – treaters and SOL in PI, Notices of Assessment and objection windows in tax, biometrics and BOC in immigration, estoppel certificates and FIRPTA in real estate, breach notification matrices in cybersecurity, probate lifecycles in estates, four-way session summaries in family law.

But strip away the vocabulary and the pattern is identical:

Architecture Layer	PI	Real Estate	Tax	Cybersecurity	Immigration	Estates	Family Law
Documents In	Medical records from providers	Title docs, surveys, lender letters	CRA/IRS notices, assessments	Breach incident reports	Portal updates, IRCC correspondence	Executor documents, asset statements	Session notes, financial disclosures
Classify & Extract	Records vs. bills vs. imaging	Match to closing checklist	Assessment type, amount, trigger date	Data type, affected states, affected count	Status change type, new deadlines	Will, death cert, asset inventory	Topics, agreements, open items
Compute Deadlines	SOL, treatment gap flags	Closing condition deadlines	90-day objection, 30-day response	State-by-state notification deadlines	Biometrics window, medical cert expiry, hearing dates	Beneficiary notification, creditor notice, court filing	Disclosure deadlines, next session prep
Draft Outputs	Demand package outline	Closing status report	Notice of Objection	Notification letters per jurisdiction	Client milestone update	Beneficiary letters, filing checklists	Session summary, revised agreements
Notify People	Attorney: records complete	Attorney: condition missing	Attorney: deadline cascade	Attorney: filing tracker	Client: status change; Attorney: deadline alert	Executor: next steps; Attorney: filing due	All participants: summary + action items
Safety Gate	Attorney reviews demand	Attorney review	Attorney verifies deadline	Attorney verifies statutes current	Practitioner confirms portal data	Attorney verifies jurisdiction rules	Attorney reviews before

Architecture Layer	PI	Real Estate	Tax	Cybersecurity	Immigration	Estates	Family Law
		s status	computation				distribution

Six rows. Seven columns. Same architecture in every cell.

The skill file changes. The connections swap in a different case management system or portal. The schedule adjusts to the workflow's cadence. But the three layers plus the safety layer – the architecture from Chapter 3 – remain structurally identical.

This is the point. You don't need seven different AI products for seven practice areas. You need one architecture and seven skill files.

What This Means for Multi-Practice Firms

If your firm handles more than one practice area – and most small firms do – the architecture advantage compounds. A firm doing PI and family law uses the same connections (Clio, Gmail, Slack), the same scheduling system, and the same safety layer. The only difference is the skill files.

Your paralegal doesn't need to learn a new software platform when the firm takes on an estates matter. The case management system is the same. The AI architecture is the same. Only the workflow instructions change.

This is also why the skills approach from Chapter 3 matters so much. When your firm's expertise is encoded in skill files – not locked inside a vendor's platform – you can add new practice areas by writing new skills, not by buying new software.

What AI Won't Do Across These Practice Areas

The limitations from Chapter 3 apply everywhere, and a few are worth restating in context:

Jurisdiction-specific rules change. The breach notification deadlines are current as of this writing but will change. The CRA objection window is statutory but holiday adjustments vary. The probate lifecycle differs by province and state. Every computed deadline must be verified against current rules. AI computes dates. The attorney verifies them.

Novel fact patterns break the pattern. The agent workflows above handle the repeatable 80% of each practice area. When a case presents a genuinely novel issue – an ambiguous assessment notice, a cross-border breach involving GDPR

and state law simultaneously, a contested will with unusual provisions – the agent routes it to the review queue. The attorney handles it. That's the safety layer doing its job.

AI-generated legal documents must be reviewed. Every demand letter, Notice of Objection, notification letter, and session summary produced by these agents is a draft. The attorney reviews, edits, and approves before anything reaches a client, opposing counsel, a court, or a regulatory body. No exceptions. No shortcuts.

Try This Now

Pick your practice area from the seven above (or adapt the pattern to your own). Take the prompt example from that section and modify it for one active case in your practice right now:

1. Replace the placeholders with real case details
2. Paste it into your AI assistant
3. Review the output

You'll see two things: first, the output is genuinely useful – a starting point that saves real time. Second, it needs your judgment – adjustments, corrections, additions that only you can make. That's the pattern working exactly as designed: AI does the production work, you do the legal work.

For the next step – turning these one-off prompts into automated agents that run on a schedule – Chapter 7 walks through the agent playbook: how to design, build, test, and deploy your first agent.

Quick Reference: Seven Practice Areas, One Architecture

Practice Area	Documents In	Key Deadline	Agent Output	Competing Tools
Personal Injury	Medical records (200-500 pp/case)	SOL + treatment gap flags	Chronology, demand prep memo	Supio (Thomson Reuters), Tavrn, Anytime AI
Real Estate	Title docs, surveys, lender letters	Closing condition deadlines	Closing status report, compliance flags	Gavel, Clio Manage

Practice Area	Documents In	Key Deadline	Agent Output	Competing Tools
Tax Disputes	CRA/IRS assessments, correspondence	90-day objection / 30-day response	Notice of Objection draft, deadline cascade	Thomson Reuters ONESOURCE
Cybersecurity	Breach incident reports	50-state notification deadlines	Compliance matrix, draft letters, filing tracker	BigID, OneTrust, Exterro
Immigration	Portal updates, IRCC/USCIS correspondence	Biometrics, medical cert, hearing dates	Client updates, deadline dashboard	DocketWise, VisaLaw
Estates & Probate	Executor documents, asset statements	Beneficiary notice, creditor notice, court filings	Lifecycle timeline, notification letters, checklists	Gavel, EstateMin
Family & Collaborative	Session notes, financial disclosures	Disclosure deadlines, next session prep	Session summary, revised agreements, action items	Deliberately.ai, Gavel

The pattern: Documents in –> deadlines computed –> things filed –> people notified. The specifics change. The architecture doesn't.

Next: Chapter 7 builds on this foundation with the agent playbook – how to design your first agent, test it, refine the skill file, and deploy it on a schedule. Whether you picked PI, estates, immigration, or any other area from this chapter, the build process is the same.

Chapter 7: The Agent Playbook

What You'll Walk Away With:

- Seven agents that work for any practice area – PI, real estate, tax, immigration, estates, or anything else
- A complete walkthrough for building Agent #1 (Deadline Monitor) in an afternoon, including the skill file, the cron schedule, and the sample output
- Brief but concrete descriptions of Agents #2 through #7, with connection maps, sample outputs, and effort estimates
- The compound effect: how three or more agents feeding each other produce results none of them could alone
- A mapping of the 21.5-hour admin tax from Chapter 1 to specific agents that reclaim that time

---

Where Your 21.5 Hours Go

Remember Sarah Chen's time audit from Chapter 1? She mapped 21.5 hours per week of administrative work – tasks that consume attorney-rate time without requiring a law license.

Here's that audit again. But this time, every line maps to a specific agent:

Task	Hours/Week	Agent
Reviewing and responding to emails (non-case-specific)	4 hrs	**Agent 2: Inbox Triage**
Drafting client status updates	3 hrs	**Agent 3: Client Update Drafter**
Calendaring deadlines from court notices and filings	2 hrs	**Agent 1: Deadline Monitor**
Billing: reviewing entries, catching unbilled time	2.5 hrs	**Agent 4: Billing Reconciler**
Marketing: Google review responses, referral follow-ups	2 hrs	**Agent 6: Marketing Engine**
Intake screening: evaluating new leads, conflict checks	2 hrs	**Agent 5: Intake Pipeline**
SOP documentation and staff training	1.5 hrs	Partially: Agent 7 + encoded skill files
Vendor management, supply ordering, office admin	1.5 hrs	Not automated (human judgment)

Task	Hours/Week	Agent
Competitive research, CLE tracking, bar compliance	1 hr	Partially: Agent 7 reporting
Preparing for meetings (pulling files, reviewing history)	2 hrs	**Agent 1: Deadline Monitor** + Morning Briefing (Ch 3)
Total	**21.5 hrs**	**~17.5 hrs addressable by agents**

Seventeen and a half hours per week. At Sarah's effective rate of $364/hour, that's $6,370 in weekly capacity trapped in tasks that AI handles well. Not all of it disappears – you still review every output, approve every draft, verify every deadline. But the 17.5 hours of production work compresses into roughly 4 hours of review work. The net recovery is 13+ hours per week.

Seven agents. One architecture. Here's the playbook.

The Seven Agents

#	Agent	What It Does	Connections	Runs	Effort
1	**Deadline Monitor**	Scans all active matters for SOL dates, discovery deadlines, motion deadlines, hearing dates. Alerts at 90 days (planning), 60 days (action), 30 days (urgent).	Clio, Calendar, Slack	Daily 7:00 AM	Afternoon build
2	**Inbox Triage**	Classifies every email (client, court, opposing counsel, vendor, newsletter). Labels, archives noise, surfaces urgent items in a prioritized digest.	Gmail/Outlook, Clio, Slack	2x daily (8 AM, 2 PM)	Afternoon build
3	**Client Update Drafter**	Generates monthly status update drafts for every active matter by pulling 30 days of case activity, correspondence, and upcoming milestones.	Clio, Gmail, Dodonai	Monthly (1st of month)	Weekend project
4	**Billing Reconciler**	Cross-references calendar events and email activity against Clio time entries. Flags unbilled work, suggests missing	Clio, Calendar, Gmail, QuickBooks	Weekly (Friday 5 PM)	Weekend project

#	Agent	What It Does	Connections	Runs	Effort
		entries, catches underbilled time.			
5	**Intake Pipeline**	Scores web leads against your case acceptance criteria, runs conflict checks, drafts follow-up emails, prepares consult memos with case value estimates.	Web form/email, Clio, Calendar, Gmail	On new lead arrival	Weekend project
6	**Marketing Engine**	Drafts blog posts from practice expertise, generates Google review responses, queues referral nurture emails, tracks content calendar.	Gmail, Google My Business, Website CMS, Slack	Weekly (Monday 7 AM)	Weekend project
7	**Case Health Dashboard**	Rolls up all active matters into a single digest: case aging, pipeline value, profitability by practice area, stale matters, trust balance alerts.	Clio, QuickBooks, Slack	Weekly (Monday 8 AM)	Multi-day build

Every agent follows the three-layer architecture from Chapter 3: skill file, connections, schedule. The safety layer applies to all of them – AI drafts, you decide. The first agent takes the longest because you're also setting up the infrastructure. Each subsequent one reuses the same connections and scheduling system.

Build Your First Agent: The Deadline Monitor

This is the walkthrough. By the end of it, you'll have an agent that scans every active matter in your case management system, computes deadline windows, and sends you a prioritized alert every morning at 7:00 AM – before you've left the house.

Total build time: 2-3 hours. One afternoon.

Why Start Here

Missed deadlines are the number one source of legal malpractice claims. The ABA's profile of malpractice claims consistently shows calendar-related errors as the leading cause. An agent that checks every deadline every single day – without relying on your memory or Maria's – is the most defensible thing you can build.

It's also the lowest risk. This agent only reads data. It doesn't modify anything in Clio, doesn't send emails, doesn't create calendar events. If the report is wrong, you catch it during your morning review. The worst-case failure mode is a false alarm – infinitely better than a missed deadline.

From day one, you'll open your laptop to a prioritized list of every deadline within 90 days. No more mental load. No more "did I calendar the response deadline from that order last Thursday?" anxiety at 11 PM.

Step 1: Set Up the Connection to Clio

Your AI needs read-only access to your case management system. In Clio, this means creating an API key with read permissions for matters, tasks, calendar entries, and contacts.

Here's what you're connecting:

Data	Where It Lives in Clio	What the Agent Reads
Active matters	Matters list (status: open)	Matter name, practice area, client name, responsible attorney
Deadlines	Tasks and calendar entries	Due dates, task types (SOL, discovery, motion, hearing), associated matter
SOL dates	Custom fields or matter notes	Statute of limitations date per matter
Court dates	Calendar entries	Hearing dates, trial dates, mediation dates

If you use PracticePanther, MyCase, or another case management system, the data is the same – only the API endpoint changes. The architecture from Chapter 3 doesn't care which specific system you use.

Read-only access. This is critical. The Deadline Monitor never writes to Clio. It reads matters, reads deadlines, reads calendar entries, and produces a report. Set the API permissions to read-only and leave them there.

Step 2: Write the Skill File

This is the complete skill file for the Deadline Monitor. It's a markdown file – plain English, no code. You can adapt it for your practice by changing the deadline windows, the matter filters, or the output format.

```
# /deadline-monitor -- Daily Deadline & SOL Monitor

Check all active matters for approaching deadlines and statute of
limitations dates. Generate a prioritized morning alert.

## Context

This agent runs every weekday at 7:00 AM. It scans all active matters
in the firm's case management system, identifies every deadline
```

```
within
90 days, and produces a tiered alert. The attorney reviews the alert
before starting the workday.

## Steps

### 1. Pull all active matters

Fetch every matter with status "Open" or "Active" from Clio.
For each matter, retrieve:
- Matter name and case number
- Client name
- Practice area (PI, med mal, workers' comp, etc.)
- Responsible attorney
- Date opened
- Statute of limitations date (from custom field or matter notes)

### 2. Pull all deadlines and calendar entries

For each active matter, fetch:
- All tasks with due dates in the next 90 days
- All calendar entries (hearings, mediations, depositions, trial
dates)
  in the next 90 days
- Any court-imposed deadlines (discovery cutoffs, motion deadlines,
  expert disclosure dates)

### 3. Classify by urgency

Sort every deadline into three tiers:

**URGENT (within 30 days):**
- Statute of limitations expiring
- Discovery responses due
- Motion filing deadlines
- Expert disclosure deadlines
- Hearing or trial dates

**ACTION REQUIRED (31-60 days):**
- Upcoming depositions
- Mediation dates
- Demand deadlines (policy limits, pre-suit requirements)
- Client document requests due

**PLANNING (61-90 days):**
- Trial dates approaching
- Expert retention deadlines
- Settlement conference dates
- SOL dates entering the action window next month

### 4. Flag anomalies

Check for:
- Any active matter with NO deadline or SOL date on file
  (flag: "No SOL recorded -- verify immediately")
- Any SOL date within 30 days where no demand or suit has been filed
  (flag: "CRITICAL -- SOL imminent, no filing on record")
- Any matter with no activity in the last 60 days and an SOL
  within 180 days (flag: "Stale matter approaching SOL")
- Duplicate or conflicting deadlines on the same matter
```

```
### 5. Generate the morning alert

Format the output as a daily digest:

**Header:** "Deadline Monitor -- [date] -- [total deadlines flagged]"

**Section 1: URGENT (30 days)**
For each item:
- Matter name | Client | Deadline type | Due date | Days remaining
- One-line note on what action is needed

**Section 2: ACTION REQUIRED (60 days)**
Same format as above.

**Section 3: PLANNING (90 days)**
Same format as above.

**Section 4: ANOMALIES**
- Matters with no SOL recorded
- Stale matters approaching SOL
- Any other flags from Step 4

**Footer:** "Total active matters checked: [N].
Next run: tomorrow at 7:00 AM."

### 6. Save and notify

Save the digest to: reports/deadlines/deadline_[date].md
Send a summary to Slack:
"Deadline alert ready -- [N] urgent, [N] action, [N] planning."

If any URGENT items exist, also send a separate Slack message
for each one: "URGENT: [matter name] -- [deadline type]
due [date] ([N] days)."

## Safety Notes

- All computed deadlines must be verified against source documents.
  Holiday and weekend adjustments vary by jurisdiction.
- SOL computation is based on the date recorded in the case
  management system. If the recorded date is wrong, the alert
  will be wrong. Garbage in, garbage out.
- This agent reads data. It does not modify matters, create
  calendar events, or send external communications. Read-only.
```

It's a procedure manual. If you wanted Maria to do this check every morning, you'd give her roughly these same instructions.

Three things to note. **The anomaly check is the safety net.** Step 4 catches the cases that fall through cracks: the matter where nobody recorded the SOL, the case that went quiet while a deadline crept closer. These are the malpractice claim scenarios. The agent catches them every single day.

The tiers match how you actually prioritize. Thirty days is "drop everything." Sixty days is "put it on the calendar." Ninety days is "start thinking about it." Adjust the windows for your practice.

The output is designed for a 2-minute scan. The format – tier, matter name, deadline type, days remaining – lets you triage 40 matters in the time it takes to drink half a cup of coffee.

Step 3: Set Up the Schedule

A schedule is a trigger that runs the agent at a set time. On macOS, this is a launchd plist file. On Linux or Windows, it's a cron job. The concept is identical to a recurring calendar event – except instead of reminding you to do something, it tells your AI to do something.

Here's the actual pattern, adapted from production systems running 12+ agents on automated schedules:

The shell wrapper – a simple script that activates your environment and runs the skill:

```
#!/bin/bash
# Daily Deadline Monitor -- runs via launchd every weekday at 7am
source "$(dirname "$0")/common.sh"
wait_for_network || exit 1

log_start "Deadline Monitor"
cd "$SCRIPT_DIR" && claude -p "/deadline-monitor" \
    --permission-mode auto >> "$LOG" 2>&1
EXIT_CODE=$?
echo "=== Exit code: $EXIT_CODE ===" >> "$LOG"
alert_on_failure "Deadline Monitor" $EXIT_CODE
```

The schedule file (macOS launchd plist):

```
<?xml version="1.0" encoding="UTF-8"?>
<!DOCTYPE plist PUBLIC "-//Apple//DTD PLIST 1.0//EN"
    "http://www.apple.com/DTDs/PropertyList-1.0.dtd">
<plist version="1.0">
<dict>
    <key>Label</key>
    <string>com.chenlawgroup.deadline-monitor</string>
    <key>ProgramArguments</key>
    <array>
        <string>/bin/bash</string>
        <string>/path/to/your/cron/run-deadline-monitor.sh</string>
    </array>
    <key>StartCalendarInterval</key>
    <dict>
        <key>Hour</key>
        <integer>7</integer>
        <key>Minute</key>
        <integer>0</integer>
    </dict>
    <key>StandardOutPath</key>
    <string>/path/to/your/reports/cron.log</string>
    <key>StandardErrorPath</key>
    <string>/path/to/your/reports/cron.log</string>
</dict>
</plist>
```

To install it:

```
cp com.chenlawgroup.deadline-monitor.plist ~/Library/LaunchAgents/
launchctl load ~/Library/LaunchAgents/com.chenlawgroup.deadline-
monitor.plist
```

That's it. The agent runs every day at 7:00 AM. If the machine is asleep, it runs when the machine wakes up. If the network is unavailable (the wait_for_network function), it retries for two minutes before logging a failure.

On Windows or Linux, the schedule is a standard cron entry instead of a launchd plist. The pattern is the same: trigger a shell script at a set time, log the output, alert on failure.

Step 4: Test It

Before you set the schedule, run the agent manually:

```
claude -p "/deadline-monitor"
```

Check the output against your actual Clio data. Verify:

- Does it pull all active matters? (Compare the count against your Clio dashboard.)
- Are the SOL dates correct? (Spot-check three or four.)
- Are the urgency tiers right? (Is the discovery response due in 12 days showing up under URGENT?)
- Do the anomaly flags make sense? (Does it catch the matter where you forgot to enter the SOL?)

If something is wrong, edit the skill file. Add more specificity. If the agent is pulling closed matters, add a filter instruction. If it's missing a deadline type, add it to Step 2. Skill files are living documents – you refine them the same way you'd refine a paralegal's training after their first week.

What the Morning Alert Looks Like

Here's a sample output from the Deadline Monitor running against Sarah Chen's 40-matter caseload:

```
================================================================
DEADLINE MONITOR -- April 4, 2026 -- 11 deadlines flagged
================================================================

----- URGENT (within 30 days) -----

Hernandez v. Lyft | PI-Auto | Discovery responses due
  Due: April 14, 2026 (10 days)
  Action: Responses drafted, pending attorney review and signature

Rivera v. St. David's | Med Mal | Expert disclosure deadline
  Due: April 22, 2026 (18 days)
```

```
  Action: Dr. Kaplan retained, need to serve report

Thompson v. Meridian Apts | PI-Premises | SOL expires
  Due: April 28, 2026 (24 days)
  Action: CRITICAL -- suit not yet filed. SOL imminent.

----- ACTION REQUIRED (31-60 days) -----

Kowalski v. Target Corp | PI-Premises | Mediation scheduled
  Due: May 15, 2026 (41 days)
  Action: Mediation brief due 7 days prior (May 8)

Patel v. Meridian Trucking | PI-Auto | Deposition of plaintiff
  Due: May 22, 2026 (48 days)
  Action: Prep session needed, coordinate with client

Nguyen v. ABC Roofing | Workers' Comp | Hearing date
  Due: June 2, 2026 (59 days)
  Action: Compile exhibits, prepare witness list

----- PLANNING (61-90 days) -----

Davis v. HCA Healthcare | Med Mal | Trial date
  Due: June 20, 2026 (77 days)
  Action: Trial prep timeline begins -- see case health dashboard

[3 more planning items]

Chang v. Unitrin Insurance | PI-Auto | SOL enters action window
  Due: July 1, 2026 (88 days)
  Note: SOL is September 1, 2026. Entering 60-day action window
  next month. Demand package should be finalized before then.

----- ANOMALIES -----

WARNING: Martinez v. City of Austin | PI-Premises
  No SOL date recorded. Matter opened January 12, 2026.
  Verify SOL and enter in Clio immediately.

WARNING: Wilson v. Allstate | PI-Auto
  No case activity in 74 days. SOL is October 3, 2026 (182 days).
  Stale matter approaching SOL window.

================================================================
Total active matters checked: 40
Next run: tomorrow at 7:00 AM
================================================================
```

That took 90 seconds. Compiling it manually – pulling up each matter in Clio, checking the calendar, cross-referencing deadlines, sorting by urgency – takes 45-60 minutes.

The Thompson SOL alert prevents a malpractice claim. The Martinez anomaly catches a data entry gap before it becomes a problem. The Wilson stale matter flag surfaces the case you forgot about because nothing was happening on it – which is exactly when things fall through cracks.

Sarah scans it in two minutes over coffee. She texts Maria about the Thompson filing priority, makes a mental note to check the Martinez SOL, and opens the Wilson file. The agent didn't make any of those decisions. It gave her the information to make them – reliably, every single morning.

What You Just Built

Skill: The markdown file with seven steps. **Connection:** Read-only Clio API + Slack notifications. **Schedule:** A launchd plist firing every weekday at 7:00 AM. **Safety layer:** The agent reads and reports. Every action – filing the Thompson suit, checking the Martinez SOL, re-engaging the Wilson file – is decided by Sarah.

Total build time: 2-3 hours. The agent runs indefinitely from that point forward.

Agent 2: Inbox Triage

The Problem

Sarah's inbox gets 60-80 emails per day. About 15 need her attention. The court notice buried between a Lexis digest and a Clio activity alert is the one that matters. Four hours per week on email is conservative.

What It Does

The Inbox Triage agent classifies every email into six categories (Court/Filing, Opposing Counsel, Client, Vendor/Admin, Newsletter/CLE, Outreach Reply), applies labels, archives the noise, and surfaces urgent items in a prioritized digest.

Schedule: Twice daily – 8:00 AM and 2:00 PM. **Connections:** Gmail or Outlook, Clio (sender-to-matter matching), Slack.

The Skill File (Key Section)

```
### Classify each message

For each unread message in the inbox, classify by examining:
- Sender domain (court domains, known opposing counsel, client
emails)
- Subject line patterns ("ORDER", "NOTICE", "RE: [matter name]")
- Message content (filing references, case numbers, deadline
language)

Priority tiers:
- **Urgent:** Court orders, motions requiring response, client
  emergencies, messages from judges
- **Action:** Client questions, opposing counsel correspondence,
  scheduling requests, items requiring a reply within 48 hours
```

```
- **FYI:** Newsletters, vendor updates, automated notifications,
  marketing emails, CLE announcements

Apply labels: "Triage/Urgent", "Triage/Action", "Triage/FYI"
Archive all FYI items. Leave Urgent and Action in inbox.
```

Sample Output

```
INBOX TRIAGE -- April 4, 2026 (8:00 AM run)
47 messages processed | 4 Urgent | 8 Action | 35 FYI (archived)

URGENT:
  [Court] Rivera v. St. David's -- Order Granting Motion to Compel
    Response deadline: April 18 (14 days). Calendar immediately.
  [Client] Maria Hernandez -- voicemail follow-up, requesting
    case status update. Last contact: March 2.
  [Opposing] Davis v. HCA -- Settlement demand response from
    defense counsel. Time-sensitive.
  [Court] Thompson v. Meridian -- Scheduling order issued.
    Trial date set June 20. Pre-trial deadlines embedded.

ACTION:
  [Client] Patel -- question about deposition prep
  [Scheduling] Kowalski mediation -- mediator proposing new dates
  [Vendor] Clio -- annual renewal notice (due April 30)
  [5 more Action items]
```

Effort: Afternoon build. You'll refine the classification rules over the first week as you see which emails get miscategorized. **Handoff:** When Inbox Triage detects a court order with a deadline, it feeds that to the Deadline Monitor. When it identifies a client communication requiring a response, it flags it for the Client Update Drafter. Agents feeding agents.

Agent 3: Client Update Drafter

The Problem

Forty active matters. Bar ethics rules – and common sense – require regular client communication. But drafting 40 status updates per month takes 3+ hours that always get deprioritized.

What It Does

On the 1st of each month, the agent pulls 30 days of case activity for every active matter and generates a draft status email for each client. **Connections:** Clio, Gmail, Dodonai. **Schedule:** Monthly, 1st of month at 7:00 AM.

The Skill File (Key Section)

```
### For each active matter, generate a client update draft

Pull from Clio:
- Activity in the last 30 days (filings, correspondence, depositions)
- Next upcoming deadline or milestone
- Trust/retainer balance

Pull from Gmail:
- Date of last client communication
- Any unanswered client emails

Draft the update email:
- Opening: acknowledge the client by name, reference the matter
- Body: summarize recent activity in plain language (no legal jargon
  unless the client is sophisticated). What happened, what it means,
  what happens next.
- Close: next milestone with estimated timeline, invitation to call
  with questions
- Tone: direct, warm, reassuring. The client should feel informed,
  not alarmed.

Save as Gmail draft, addressed to the client. Do NOT send.
Flag any matter where the last client contact was over 45 days ago:
"OVERDUE -- last contact [date]."
```

Sample Output (Single Draft)

```
To: maria.hernandez@email.com
Subject: Case Update -- Hernandez v. Lyft (March 2026)

Maria,

Here's where things stand on your case as of the end of March.

We received the final treatment records from Austin Spine & Rehab
last week, and your MRI results from St. David's arrived on March
22. Your medical chronology is now complete -- 14 months of
consistent treatment following the November 2025 accident. That
consistency is a strength for your case.

Next steps: we're finalizing the demand package this month. That
means compiling your total medical expenses, documenting your
treatment timeline, and drafting the damages narrative. I expect
to have the demand ready to send to Lyft's insurance carrier by
mid-April.

After we send the demand, the insurance company typically has
30-45 days to respond. I'll keep you posted as soon as we hear
back.

Your statute of limitations is November 12, 2027 -- we have
plenty of time, and we're well ahead of schedule.

If you have any questions or if anything has changed with your
treatment, give me a call.
```

```
Sarah Chen
Chen Law Group
```

Effort: Weekend project. You'll tune the tone after the first batch – some updates will be too formal, some too detailed. That's the refinement cycle from Chapter 3.

Agent 4: Billing Reconciler

The Problem

Sarah billed 5.2 hours on a day she worked at least 9. That 3.8-hour gap is the norm, not the exception. BigHand SmartTime data shows attorneys recover 8-10 extra billable hours per month with AI-assisted time tracking. Billables AI reports 10-30% more billable time captured. At Sarah's rate, 10 extra hours per month is $43,680 per year in recovered revenue.

What It Does

Every Friday at 5:00 PM, the Billing Reconciler cross-references three data sources to find unbilled work:

1. **Calendar events** – meetings, calls, depositions, court appearances that should have corresponding time entries
2. **Email activity** – substantive emails sent or received per matter that suggest research, drafting, or review time
3. **Clio time entries** – what was actually billed

It produces a reconciliation report: here's what the calendar and email show you did, here's what Clio shows you billed, and here's the gap.

Connections: Clio (time entries, matters), Google Calendar (events with matter names or client names), Gmail (emails associated with matters), Slack (weekly summary).

Schedule: Weekly, Friday 5:00 PM.

Sample Output (Excerpt)

```
BILLING RECONCILER -- Week of March 30 - April 3, 2026

UNBILLED GAPS DETECTED:

  Hernandez v. Lyft | PI-Auto
    Calendar: 45-min call with client (March 31, 2:15 PM)
    Clio entry: None found
    Suggested entry: 0.8 hrs -- Client conference call,
    case status discussion
```

```
  Davis v. HCA | Med Mal
    Email: 6 substantive emails exchanged with expert
    witness Dr. Kaplan (April 1-2)
    Clio entries: 0.3 hrs total
    Suggested additional: 1.2 hrs -- Expert correspondence,
    report review and feedback

  Patel v. Meridian Trucking | PI-Auto
    Calendar: 2-hr deposition prep session (April 3)
    Clio entry: 1.0 hr billed
    Gap: 1.0 hr underbilled
    Suggested correction: 2.0 hrs -- Deposition preparation

WEEKLY SUMMARY:
  Total calendar/email activity identified: 38.2 hours
  Total Clio time entries: 31.6 hours
  Identified gap: 6.6 hours ($2,402 at $364/hr)
  Suggested new entries: 14 (review and approve in Clio)
```

Effort: Weekend project. The trickiest part is teaching the agent which emails represent billable work versus administrative correspondence. You'll tune this over the first month.

Agent 5: Intake Pipeline

The Problem

Sarah gets 15 leads per month, signs 8. The other 7 still consume time. Speed matters – Rankings.io found that firms responding within 5 minutes dramatically outperform on conversion. Sarah's leads sit in her inbox until she has a gap, which is sometimes end of day.

What It Does

On new lead arrival, the agent extracts case facts, scores the lead against Sarah's acceptance criteria, runs a Clio conflict check, drafts an acknowledgment email, and prepares a go/no-go consult memo. **Connections:** Web form/email, Clio, Gmail, Calendar. **Schedule:** Event-based or every 2 hours during business hours.

The Skill File (Key Section)

```
### Score the lead

Evaluate against these criteria:

| Factor | Weight | Scoring |
|--------|--------|---------|
| Injury severity | 30% | Catastrophic (10), Serious (7), Moderate
(5), Minor (2) |
```

```
| Liability clarity | 25% | Clear (10), Probable (7), Disputed (4),
Weak (1) |
| Insurance coverage | 20% | Commercial policy (10), Standard auto
(7), Minimum/uninsured (3) |
| Case value estimate | 15% | >$100K (10), $50-100K (7), $25-50K (5),
<$25K (2) |
| Practice fit | 10% | Core PI (10), Adjacent (6), Outside practice
areas (1) |

Score > 70: Strong intake -- draft consult confirmation immediately
Score 50-70: Review recommended -- flag for attorney evaluation
Score < 50: Likely decline -- draft polite referral response
```

Effort: Weekend project. The scoring weights are where you encode your case acceptance judgment. You'll adjust them after the first 10 leads.

Agent 6: Marketing Engine

Sarah knows marketing matters – her 4.9-star Google rating didn't happen by accident. But reviews, blog posts, and referral follow-ups always lose the priority battle to active case work.

The Marketing Engine runs a weekly cycle: drafts one blog post from your topic queue, monitors Google My Business for new reviews and drafts responses, and queues quarterly referral nurture emails. **Connections:** Google My Business, Gmail, website CMS, Clio, Slack. **Schedule:** Weekly, Monday 7:00 AM. **Effort:** Weekend project.

Chapter 8 covers this agent in full detail – the content engine, review management, referral nurturing, and intake pipeline working together as a growth system.

Agent 7: Case Health Dashboard

Sarah knows the status of the 10 matters she's actively working on. The other 30? General sense, no systematic view. Which cases are aging past profitability? Which trust balances need replenishment? Which practice areas are generating the best revenue per hour?

The Case Health Dashboard treats your caseload like a portfolio, answering three questions weekly: **What needs attention?** (stale matters, low trust balances, deadline-activity mismatches). **What's the pipeline worth?** (estimated values by practice area and stage). **What's working?** (revenue per hour, time-to-resolution trends, settlement rates).

Connections: Clio, QuickBooks, Slack. **Schedule:** Weekly, Monday 8:00 AM (runs after the Deadline Monitor).

Sample Output (Excerpt)

```
CASE HEALTH DASHBOARD -- Week of March 30, 2026

PORTFOLIO SUMMARY:
  Active matters: 40
  Total estimated pipeline value: $1.8M
  Matters in demand stage: 8 ($640K)
  Matters in active litigation: 6 ($720K)

ATTENTION REQUIRED:
  3 matters with no activity > 30 days:
    Wilson v. Allstate (74 days -- also flagged by Deadline Monitor)
    Barnes v. Dominos (45 days)
    Lee v. Metro Transit (32 days)

  2 matters with trust balance < $500:
    Kowalski v. Target -- $220 remaining (mediation upcoming)
    Nguyen v. ABC Roofing -- $0 balance (hearing in 59 days)

PRACTICE AREA PERFORMANCE (last 90 days):
  PI-Auto: 24 matters, $14.2K avg value, $412/hr effective rate
  Med Mal: 6 matters, $118K avg value, $380/hr effective rate
  Workers' Comp: 6 matters, $22K avg value, $285/hr effective rate

RESOLVED THIS MONTH:
  Ortiz v. Amazon Logistics -- settled $67,500 (PI-Auto)
    Time invested: 142 hrs. Effective rate: $475/hr.
    Resolution time: 11 months (below 14-month avg).
```

Effort: Multi-day build. This is the most complex agent because it pulls from the most data sources and requires the most nuanced output formatting. Build it after you've had Agents 1-4 running for at least a month.

The Compound Effect

Individual agents save time. Multiple agents feeding each other change how your practice operates. Here's a Thursday in April – nothing has gone wrong, but without agents, several things would have gone unnoticed:

7:00 AM – Deadline Monitor flags that the Davis v. HCA trial date is 77 days away and the expert disclosure deadline is 18 days out.

8:00 AM – Inbox Triage classifies a new email from defense counsel on the Davis matter as Urgent. It's a settlement offer – $180K on a case Sarah has valued at $250K+.

8:05 AM – Case Health Dashboard (from last Monday's run) shows Davis has $720 in trust remaining with a trial date approaching. If the case doesn't settle, she'll need a trust replenishment before trial prep begins.

8:10 AM – Billing Reconciler (from last Friday) shows 4.2 unbilled hours on Davis from the prior week – expert correspondence and records review that hadn't been entered.

Now Sarah opens her laptop. In two minutes of scanning, she knows:

- Expert disclosure is due in 18 days (Deadline Monitor)
- Defense has made a settlement offer (Inbox Triage)
- Trust needs replenishment if this goes to trial (Case Health Dashboard)
- There's $1,528 in unbilled time to enter (Billing Reconciler)

No agent made a decision. Each one surfaced a piece of the picture. Sarah sees the full picture – not because she remembered to check four systems, but because four agents checked them for her. One agent is a convenience. Three agents feeding each other create situational awareness that would take a full-time office manager to replicate.

The system doesn't replace Sarah's judgment on whether to accept the $180K offer. It ensures she has every relevant fact in front of her when she makes the call.

Time Savings Estimates

Conservative estimates. Actual savings depend on caseload and current systems.

Agent	Weekly Time Without	Weekly Time With (Review Only)	Net Savings	Annual Value ($364/hr)
Deadline Monitor	2.0 hrs	0.3 hrs (scan digest)	1.7 hrs	$32,177
Inbox Triage	4.0 hrs	1.0 hrs (review flagged items)	3.0 hrs	$56,784
Client Update Drafter	3.0 hrs (monthly, ~0.75/wk)	0.5 hrs (review drafts)	0.25 hrs	$4,732
Billing Reconciler	2.5 hrs	0.5 hrs (approve entries)	2.0 hrs	$37,856
Intake Pipeline	2.0 hrs	0.5 hrs (review memos)	1.5 hrs	$28,392
Marketing Engine	2.0 hrs	0.5 hrs (review drafts)	1.5 hrs	$28,392
Case Health Dashboard	2.0 hrs	0.3 hrs (scan digest)	1.7 hrs	$32,177

Agent	Weekly Time Without	Weekly Time With (Review Only)	Net Savings	Annual Value ($364/hr)
Total	**17.5 hrs**	**3.6 hrs**	**~13.9 hrs**	**$220,511**

That $220K is capacity, not guaranteed revenue. How much converts depends on you. But the ceiling is real, and it's supported from another angle: BigHand SmartTime shows 8-10 extra billable hours per month per attorney from AI time-tracking alone – and the Billing Reconciler is one of seven agents.

Build Sequence: What to Do This Month

Don't build all seven at once. Each agent works better when the ones before it are running.

Week	Build	Why This Order
Week 1	Agent 1: Deadline Monitor	Highest impact, lowest risk. Sets up your connections to Clio and Slack. Proves the architecture works.
Week 2	Agent 2: Inbox Triage	Second-highest daily time savings. Reuses your Gmail connection. Starts feeding deadlines to Agent 1.
Week 3	Agent 4: Billing Reconciler	Immediate financial impact. Uses connections already established in Weeks 1-2.
Week 4	Agent 3: Client Update Drafter	Monthly cadence means you set it up once and it runs on the 1st. Uses Clio + Gmail connections from prior agents.
Month 2	Agent 5: Intake Pipeline	Builds on the lead scoring judgment you've been refining informally.
Month 2	Agent 6: Marketing Engine	See Chapter 8 for the full build guide.
Month 3	Agent 7: Case Health Dashboard	This agent is most valuable once the others have been producing data for a few weeks.

Each agent after the first takes less time because the infrastructure is already in place. The Deadline Monitor takes an afternoon because you're building the foundation. The Client Update Drafter takes an hour because it's just a new skill file on the same foundation.

What AI Won't Do Here

It won't catch every deadline. If a deadline isn't in Clio, the agent can't find it. Entering deadlines correctly is still your responsibility – the agent checks what's there, not what should be there.

It won't draft a perfect client update every time. Some matters have nuances the agent won't capture – a difficult client relationship, a strategic tone adjustment. You'll edit roughly 30% of drafts significantly. Still faster than writing from scratch.

It won't replace your billing judgment. The Reconciler flags gaps, but that 45-minute call might have been a courtesy call you'd never bill for. The agent suggests; you decide.

It won't eliminate your marketing effort. Blog posts need your review for accuracy and voice. Referral emails need your personal touch. The agent handles the 80% production work; you supply the 20% judgment.

It will sometimes get things wrong. An email miscategorized. A deadline computed without accounting for a local holiday rule. The safety layer catches these at the review stage – which is why every agent produces drafts, not final outputs.

Try This Now

Three options:

Write the skill file. Adapt the Deadline Monitor for your practice – change the windows, add your deadline types, modify the format. Save it as a markdown file. Time: 30 minutes.

Test it manually. Paste the skill file into your AI assistant with a list of your active matters and deadlines exported from Clio. Compare the output against your calendar. Time: 45 minutes.

Build it. Follow the walkthrough above. Clio API, skill file, schedule, test. Working Deadline Monitor by end of afternoon. Time: 2-3 hours.

If you want help with the build – setting up connections, refining skill files, designing agents for your workflows – that's what Dodonai Consulting does. Agent design, workflow automation, tool integration, ongoing maintenance.

Quick Reference: The Agent Playbook

Agent	What It Does	Connections	Schedule	Build Time
1. Deadline Monitor	SOL, discovery, motion, hearing alerts at 30/60/90 days	Clio, Calendar, Slack	Daily 7 AM	Afternoon

Agent	What It Does	Connections	Schedule	Build Time
2. Inbox Triage	Classify, label, archive, surface urgent	Gmail, Clio, Slack	2x daily	Afternoon
3. Client Update Drafter	Monthly status emails per active matter	Clio, Gmail, Dodonai	Monthly	Weekend
4. Billing Reconciler	Calendar + email vs time entries, flag gaps	Clio, Calendar, Gmail, QuickBooks	Weekly Fri	Weekend
5. Intake Pipeline	Score leads, conflict check, follow-up, consult memo	Web form, Clio, Gmail, Calendar	On intake	Weekend
6. Marketing Engine	Blog posts, review responses, referral nurture	Gmail, GMB, CMS, Clio	Weekly Mon	Weekend
7. Case Health Dashboard	Case aging, pipeline, profitability, stale matters	Clio, QuickBooks, Slack	Weekly Mon	Multi-day

Compound Effect	Agent A Flags	Agent B Acts
Deadline approaching trial	Deadline Monitor	Case Health Dashboard prioritizes, Client Update Drafter sends prep communication
Court order arrives with deadline	Inbox Triage extracts deadline	Deadline Monitor adds to tracking, Billing Reconciler accounts for response time
Stale matter detected	Case Health Dashboard flags	Client Update Drafter queues outreach, Billing Reconciler checks for unbilled work on file
New lead arrives	Intake Pipeline scores and preps	Marketing Engine tracks conversion for content strategy refinement

The pattern: Skill (instructions) + Connections (your tools) + Schedule (trigger) + Safety layer (human review). Same three layers from Chapter 3. Seven different skill files. One architecture.

Next: Chapter 8 goes deep on Agent 6 – the Marketing Engine. Blog content, review management, referral nurturing, and the intake pipeline working together to turn your practice expertise into a growth system that runs on autopilot.

Chapter 8: Marketing, Intake, and Growth

What You'll Walk Away With:

- A content engine that drafts weekly blog posts from your practice expertise – no writing habit required
- A review management system that monitors your online reputation and drafts responses while you sleep
- A referral nurturing pipeline that maintains relationships with your top referral sources on autopilot
- A scored intake pipeline that evaluates, follows up, and schedules consultations before you touch the file
- The math: how a $20/month AI assistant replaces a $55K/year marketing coordinator
- Four prompts you can deploy this week – blog post generation, review response, intake evaluation, and referral follow-up

---

Sarah Chen's Marketing Department

Sarah Chen has a 4.9-star Google rating, strong case results, and a referral network that keeps her at 40 active cases. She has no blog. Her website hasn't been updated in two years. Her "intake process" is Maria answering the phone and writing the caller's name on a sticky note. Her "marketing strategy" is hoping that the three attorneys who send her PI referrals keep sending them.

This works until it doesn't. One of those referring attorneys retires. Another takes on a partner who went to law school with a different PI lawyer. The phone gets quieter by two calls a week, and suddenly Sarah's pipeline is thinner than it's been in three years. She has no content ranking in search, no review response history, no systematic follow-up with the 40% of leads who call but don't immediately sign.

She doesn't need a marketing degree. She needs a system.

The firms that dominate local search in PI, med mal, and workers' comp aren't necessarily better lawyers. They're firms that publish content consistently, respond to every review within hours, nurture referral relationships quarterly, and follow up with every lead who expressed interest. They do this because they have

a marketing coordinator – someone making $50-60K a year whose entire job is keeping these systems running.

Sarah can build the same system with the agent architecture from Chapter 3. One skill file for each function. Connected to the tools she already has. Running on a schedule. Total cost: her AI subscription (Claude Pro, in her case) she's already paying for, plus the time to set it up.

We know because we built it. Dodonai's own marketing runs on this exact pattern – 12 agents handling content, outreach, lead scoring, inbox triage, reputation monitoring, and referral nurturing. The same skills-connections-schedules architecture from Chapter 3, applied to business development instead of case management. The workflows in this chapter are adapted from that system.

The Content Engine

The Problem

Rankings.io's data shows that law firms publishing weekly blog content generate 3-4x more organic leads than firms with static websites. But "write a weekly blog post" has been on Sarah's to-do list for two years. It never happens because a blog post takes a litigator 3-4 hours to write – hours that could be spent billing at $364/hour. That's over $1,000 in opportunity cost per post.

The result: her website ranks for "Chen Law Group" (her own name) and nothing else. Potential clients searching "Austin car accident lawyer" or "Texas slip and fall attorney" never find her.

The System

An AI content agent works like having a marketing writer who knows your practice area inside out. It drafts blog posts from topics you'd naturally discuss with clients – common questions, case outcome patterns, legal process explanations – and formats them for SEO without you having to think about keyword research.

The workflow:

1. **Topic generation (monthly):** The agent reviews your practice areas, recent case types, and seasonal trends to suggest 4-5 topics for the month.
2. **Draft generation (weekly):** Each week, the agent takes a topic and produces a complete blog post – 800-1,200 words, naturally incorporating search terms, written in your voice.
3. **Your review (15 minutes):** You read it, adjust anything that doesn't sound like you, and approve it for posting.

4. **Publication:** The post goes live on your website. Over time, these posts compound – 52 posts per year covering the questions your potential clients are typing into Google.

The Prompt: Blog Post Generation

```
You are a content writer for a personal injury law firm in Austin,
TX.
The firm handles auto accidents, slip and fall, premises liability,
medical malpractice, and workers' compensation cases.

TASK:
Write a blog post on the following topic:

TOPIC: [TOPIC -- e.g., "What to do in the first 72 hours after a
car accident in Texas"]

REQUIREMENTS:
- Length: 800-1,200 words
- Audience: Potential clients who have been injured and are
  researching their options. They are not lawyers. Write in plain
  English, but don't condescend.
- Tone: Informative, confident, empathetic. Not salesy. Not
  "call us now for a free consultation" energy. Think: helpful
  attorney explaining the process over coffee.
- Structure:
  1. Opening that addresses the reader's situation directly
     (2-3 sentences)
  2. Clear section headers (H2s) breaking the topic into
     actionable steps or key points
  3. Specific Texas law references where relevant (cite the
     statute or rule by name -- e.g., Texas Civil Practice &
     Remedies Code Section 33.001 for comparative fault)
  4. A "what most people don't know" section with one non-obvious
     insight that demonstrates expertise
  5. Brief closing that encourages the reader to consult an
     attorney without being pushy
- SEO: Naturally incorporate these search terms where they fit
  (don't force them): [PRIMARY KEYWORD], [2-3 SECONDARY KEYWORDS]
- Do NOT include a disclaimer or "this is not legal advice"
  boilerplate -- the firm adds that sitewide
- Do NOT use phrases like "navigating the legal landscape" or
  "in today's world" or "it's important to note"

VOICE SAMPLE (match this tone):
[PASTE 2-3 PARAGRAPHS FROM AN EXISTING PAGE ON YOUR WEBSITE
OR A CLIENT EMAIL YOU'VE WRITTEN -- THIS TEACHES THE AI YOUR
NATURAL WRITING STYLE]
```

Making It a Skill

To run this weekly on a schedule (as described in Chapter 3), encode the prompt as a skill file and connect it to a simple topic list. The agent pulls the next topic each week, generates the draft, and saves it for your review. You spend 15 minutes approving instead of 4 hours writing.

The compound effect matters here. After six months, you have 26 posts covering common PI questions in your jurisdiction. After a year, 52. Each one is a permanent page on your website that can rank in search and bring in leads for years. Firms that started publishing AI-assisted content in early 2025 are already seeing organic traffic increases of 30-50% within 12 months – because they're covering topics they never had time to write about before.

What AI Won't Do Here

AI doesn't know your local opposing counsel's tendencies, your jurisdiction's unwritten courtroom norms, or the specific fact patterns that make your practice unique. The posts it drafts will be competent and accurate on general Texas personal injury law. The 15 minutes you spend reviewing is where you add the specifics that make the content genuinely authoritative – the detail about how Travis County judges handle comparative fault arguments differently than Williamson County, or why the first 72 hours matter more in Texas than in states without a modified comparative fault rule.

Review Management

The Problem

92% of consumers read online reviews before choosing a service provider (BrightLocal 2024). For local service businesses like law firms, Google reviews are the single strongest trust signal after referrals. Sarah has 23 Google reviews – all five stars except one four-star. But she hasn't responded to any of them. The last review was posted three months ago and sits unanswered.

Her competitor down the street has 87 reviews, each one with a personalized response posted within 24 hours. Google's algorithm rewards engagement, and potential clients reading reviews notice when the attorney takes the time to respond.

The System

A review monitoring agent checks your Google Business Profile daily (and Avvo, if you have a profile there). When a new review appears:

- **Positive reviews:** The agent drafts a personalized thank-you response that references specific details from the review without revealing case information. You approve and post it.

- **Negative or mixed reviews:** The agent flags it for your personal attention, drafts a professional response template, and surfaces it as an urgent item in your morning briefing.
- **Review requests:** After a case resolves favorably, the agent drafts a personalized review request email to the client. Timing matters – the request goes out 2-3 days after settlement or verdict, when satisfaction is highest.

The Prompt: Review Response

```
You are responding to an online review for a personal injury law
firm in Austin, TX. The attorney is Sarah Chen.

REVIEW PLATFORM: [Google / Avvo]
REVIEWER NAME: [NAME]
STAR RATING: [1-5]
REVIEW TEXT:
[PASTE THE REVIEW TEXT]

TASK:
Draft a response following these rules:

FOR POSITIVE REVIEWS (4-5 stars):
- Thank the reviewer by first name
- Reference something specific they mentioned (e.g., "I'm glad
  the team kept you informed throughout the process") WITHOUT
  revealing any case details, injury types, or legal strategy
- Keep it to 2-3 sentences. Warm but professional.
- Do NOT say "we appreciate your business" or anything that
  sounds like a retail store
- Do NOT reference specific case outcomes, settlement amounts,
  or legal proceedings -- this protects client confidentiality
  and complies with bar advertising rules

FOR NEGATIVE REVIEWS (1-3 stars):
- Acknowledge the reviewer's frustration without admitting fault
- Do NOT reference any case specifics
- Invite them to contact the office directly to discuss their
  concerns: "I'd welcome the chance to discuss this with you
  directly -- please call our office at [PHONE]"
- Keep it to 2-3 sentences. Professional, empathetic, not
  defensive
- Flag this for the attorney's personal review before posting

GENERAL RULES:
- Texas State Bar advertising rules apply -- no guarantees of
  results, no specific case details, no client testimonials
  that imply a guaranteed outcome
- Sign as "Sarah Chen" (not "The Team" or "Chen Law Group")
- Match the tone of the review -- if the reviewer is emotional,
  be warm. If they're matter-of-fact, be direct.
```

The Operational Reality

This isn't glamorous work. Responding to reviews is a task that's easy to do, easy to skip, and compounds over time. The firms with 100+ well-managed reviews didn't get there through a single effort – they got there by responding to every review within 24 hours for three years straight. An AI agent makes that consistency automatic.

Set the review check to run daily. Every morning, your briefing tells you whether there's a new review waiting for your approval. Most days there isn't. On the days there is, you spend 30 seconds reading the draft, maybe adjust a word, and approve it. Over 12 months, you've responded to every review within hours and built the kind of online reputation that ranks in local search and wins trust before a prospect ever picks up the phone.

Referral Nurturing

The Problem

Referrals are the lifeblood of most PI practices. Sarah gets roughly 60% of her cases from three sources: two plaintiff's attorneys who don't handle PI, and a chiropractor who treats accident victims. Those relationships produce $480K of her $800K annual revenue.

But she hasn't sent any of them a case update in four months. She doesn't track which referral source sent which client. She has no systematic touchpoint – no quarterly thank-you, no case result notification, no "here's something interesting happening in PI law" email. The relationships work because they're personal, but personal relationships atrophy without maintenance.

The System

A referral nurturing agent tracks your referral sources and maintains regular touchpoints:

1. **Referral tracking:** Log every referral source in a simple spreadsheet or Clio contact field. The agent maintains a running record: who referred whom, case outcome, total value of referred cases per source.
2. **Quarterly touchpoints:** Every 90 days, the agent drafts a personalized email to each referral source – not a form letter, but a note referencing your relationship, recent case activity (anonymized), and something genuinely useful.

3. **Case result notifications:** When a referred case resolves, the agent drafts a thank-you email to the referring source with the outcome (with client permission). This closes the loop and reinforces the relationship.
4. **Educational content sharing:** When your content engine produces a blog post relevant to a referral source's practice, the agent includes it in the next touchpoint.

The Prompt: Referral Follow-Up

```
You are drafting a quarterly referral nurture email for Sarah Chen,
a personal injury attorney in Austin, TX.

REFERRAL SOURCE:
- Name: [NAME]
- Type: [Attorney / Chiropractor / Other Professional]
- Relationship duration: [e.g., "5 years"]
- Cases referred in last 12 months: [NUMBER]
- Last communication: [DATE and brief description]
- Notable context: [e.g., "Dr. Martinez's patients are often
  auto accident victims; he refers clients who need legal
  representation for their injury claims"]

TASK:
Draft a short email (150-250 words) from Sarah to this referral
source. The email should:

1. Open with a genuine, specific reference to the relationship --
   NOT "I hope this email finds you well" or "Just checking in."
   Reference something real: a recent referral, a shared
   professional connection, a development in their field.

2. Share ONE piece of useful information. Options:
   - A recent legal development relevant to their practice
     (e.g., a new Texas statute affecting auto accident claims
     that a chiropractor's patients should know about)
   - An anonymized case result that demonstrates the value of
     early referral
   - A blog post or resource from the firm's website

3. Close with a specific, low-pressure next step -- coffee,
   a brief call, an invitation to a CLE event. NOT "let me
   know if you have any referrals."

TONE: This is a genuine professional relationship, not a sales
email. Sarah values this person. The email should read like it
was written by Sarah personally, not generated by a marketing
system.

DO NOT:
- Use phrases like "valued referral partner" or "mutually
  beneficial relationship"
- Reference specific client names or case details without
  [PLACEHOLDER] tags for Sarah to fill in
- Include any marketing language, CTAs, or promotional content
- Use em-dashes
```

The Math on Referral Nurturing

Sarah's three primary referral sources send her roughly 20 cases per year at an average value of $45K. That's $900K in case value from relationships she maintains informally. If one source stops referring – retirement, a new relationship, relocation – that's $300K in annual case value gone with no pipeline to replace it.

Systematic nurturing does two things: it protects existing relationships by keeping them warm, and it creates data that helps you build new ones. When you know that Dr. Martinez's referrals convert at 80% and average $52K in case value, you can identify three more chiropractors in your area and apply the same relationship-building approach. The agent drafts the initial outreach. You provide the personal connection.

The Intake Pipeline

The Problem

Sarah gets roughly 15 new leads per month. Some call the office. Some fill out the contact form on her website (which she checks twice a day, sometimes). Some email directly. Of those 15, she signs about 8. The other 7 don't all disappear because the case was bad. Some disappear because nobody followed up within 24 hours, or because they called on a Friday afternoon when Maria was at lunch, or because they filled out the website form and got a generic auto-reply that said "we'll be in touch" and then nobody was in touch for three days.

According to Clio's 2024 Legal Trends Report, 79% of clients expect to hear from an attorney within 24 hours of reaching out, and 10% expect a response within an hour. The firms that respond fastest win the client. In PI, where the same injured person is often contacting three firms simultaneously, speed to first contact is the single biggest determinant of who signs the case.

The System

An AI-powered intake pipeline turns every incoming lead into a scored, evaluated, and followed-up prospect – automatically.

Step 1: Capture Every inquiry – phone call (logged by Maria or voicemail transcription), web form submission, email – enters the pipeline. The agent consolidates them into a single intake queue regardless of source.

Step 2: Score The agent evaluates each lead against your practice criteria:

Scoring Factor	What the Agent Checks
Practice area fit	Does this match PI, med mal, or workers' comp? Or is it a family law question you don't handle?
Jurisdiction	Is the incident in Texas? Which county?
Estimated case value	Based on injury type, treatment described, and liability indicators
Statute of limitations	How much time remains? (2 years for PI in Texas, with exceptions)
Liability indicators	Clear liability (rear-end collision) vs. disputed (comparative fault scenario)
Urgency signals	Medical treatment ongoing? Evidence preservation needed? SOL approaching?

Step 3: Follow Up Based on the score, the agent triggers an immediate response:

- **High-value, high-fit leads** (score 80+): Draft a personalized email within the hour acknowledging their inquiry, confirming you handle their case type, and offering two specific consultation time slots from your calendar.
- **Medium-fit leads** (score 50-79): Draft an email within 4 hours with a brief response and an invitation to schedule a call.
- **Low-fit leads** (score below 50): Draft a polite decline with a referral to an appropriate attorney if the case type doesn't match your practice.

Step 4: Conflict Check Before any consultation is scheduled, the agent runs the prospective client's name and opposing party (if known) against your existing client list in Clio. Conflicts get flagged immediately – no consultation is booked until you clear it.

Step 5: Consultation Prep For leads that schedule a consultation, the agent prepares a brief: the inquiry details, the scoring analysis, relevant case law for their fact pattern, and a draft fee agreement template pre-filled with the client's information.

The Prompt: Intake Evaluation

```
You are evaluating a new intake inquiry for a personal injury
law firm in Austin, TX. The firm handles auto accidents, slip
and fall, premises liability, medical malpractice, and workers'
compensation.

INQUIRY DETAILS:
- Source: [Phone / Web form / Email / Referral from ___]
- Date received: [DATE]
- Name: [NAME]
- Contact: [PHONE / EMAIL]
- Description of incident: [PASTE WHATEVER THE PROSPECT
  DESCRIBED -- this may be a voicemail transcript, a web form
  submission, or an email. It will often be informal and
  incomplete.]
```

```
TASK:
Evaluate this inquiry and produce a structured intake assessment:

1. PRACTICE AREA FIT (Yes / Partial / No)
   - Does this fall within the firm's practice areas?
   - If partial or no, what practice area does it actually
     involve?

2. CASE VALUE ESTIMATE (High / Medium / Low / Cannot Assess)
   - Based on the injury described, treatment mentioned, and
     liability indicators, estimate the likely case value range
   - Note: This is a preliminary estimate for triage purposes
     only. It is not a case evaluation.

3. LIABILITY ASSESSMENT (Strong / Moderate / Weak / Insufficient
Info)
   - What liability indicators are present?
   - What additional facts would strengthen or weaken the claim?

4. STATUTE OF LIMITATIONS
   - Based on the incident date and case type, calculate the
     SOL deadline under Texas law
   - Flag if less than 6 months remain

5. URGENCY FLAGS
   - Evidence preservation needed? (surveillance footage,
     vehicle damage, scene conditions)
   - Medical treatment ongoing or recommended?
   - SOL pressure?
   - Other time-sensitive factors?

6. CONFLICTS CHECK NEEDED
   - List the names that need to be run through the conflicts
     database: prospective client, any opposing parties
     mentioned, any businesses or entities involved

7. RECOMMENDED ACTION
   - Priority: High / Medium / Low / Decline
   - Suggested next step (schedule consultation, request more
     information, refer out, decline with explanation)
   - Draft response email (2-3 sentences appropriate for the
     priority level)

8. MISSING INFORMATION
   - What critical facts are missing from this inquiry?
   - Suggest 3-5 specific questions for the initial callback

IMPORTANT:
- This is a screening tool, not a legal opinion. Flag
  uncertainty rather than guessing.
- If the injury description is too vague to assess, say so
  and recommend a phone screen to gather more detail.
- Do not assess damages. That requires medical records and
  is beyond the scope of intake triage.
```

What the Pipeline Actually Looks Like

Here’s a Tuesday at Chen Law Group with the intake pipeline running:

9:15 AM – A web form submission comes in. Auto accident on MoPac, rear-end collision three weeks ago, the prospect describes neck pain and mentions they've been to the ER and a follow-up with an orthopedist. The agent scores this: PI auto accident (practice area fit: yes), clear liability (rear-end), active medical treatment (case development positive), 23 months remaining on SOL (no urgency), estimated medium-high value. It drafts a response within the hour offering a consultation slot at 2:00 PM Thursday.

11:30 AM – A voicemail from someone asking about a slip and fall at an apartment complex. The description is vague – they fell, they're sore, they don't mention medical treatment. The agent scores this: premises liability (fit: yes), liability unclear (need to determine property owner negligence), no documented treatment (case development concern), estimated low-medium value. It drafts a response asking three specific follow-up questions: Have you sought medical treatment? Do you know who owns the property? Were there any witnesses or was an incident report filed?

2:45 PM – An email from a prospect asking about a contract dispute with a home contractor. The agent scores this: practice area fit: no. It drafts a polite response explaining that Chen Law Group focuses on personal injury cases and recommending the Travis County Bar Association's lawyer referral service for contract disputes.

Sarah reviews the three intake assessments at the end of the day. The first one is a strong lead – she confirms the Thursday consultation and asks Maria to run the conflict check. The second needs more information before she decides. The third is already handled. Total time Sarah spent on intake that day: 5 minutes reviewing three assessments instead of 45 minutes fielding calls, researching case types, and drafting individual responses.

Building the System: Connections and Schedule

All four agents in this chapter follow the same architecture from Chapter 3. Here's how they connect:

Agent	Connections (Read)	Connections (Write)	Schedule
Content Engine	Website analytics, topic list, voice sample file	Blog draft folder (for your review)	Weekly (Monday morning)
Review Monitor	Google Business Profile, Avvo profile	Draft response queue, morning briefing	Daily (7:00 AM)

Agent	Connections (Read)	Connections (Write)	Schedule
Referral Nurture	Clio contacts, referral tracking spreadsheet, case results log	Email draft folder	Quarterly (1st of month) + on case resolution
Intake Pipeline	Web form submissions, Gmail, voicemail transcriptions, Clio (conflict check)	Intake assessment queue, draft response emails, calendar (consultation slots)	Continuous (checks every 30 minutes during business hours)

The safety rails from Chapter 3 apply here with particular force:

- **AI drafts, you send.** Every email, review response, and blog post sits in a draft queue until you approve it. The agent never sends anything on your behalf. This isn't just a best practice – Texas State Bar advertising rules require attorney review of client-facing communications.
- **No case details in public responses.** The review response agent is specifically instructed to never reference case types, outcomes, or legal strategy. Bar advertising rules and client confidentiality require it.
- **Conflict check before consultation.** The intake pipeline flags every prospective client and adverse party for conflict screening before any consultation is booked. This is a hard gate – no override, no bypass.

The Math

Here's what the full marketing system costs versus the alternatives:

	AI Marketing System	Part-Time Marketing Coordinator	Doing Nothing
Monthly cost	$20 (AI subscription, e.g. Claude Pro)	$2,500-3,000 (20 hrs/week)	$0
Annual cost	$240	$30,000-36,000	$0
Blog posts/year	52 (weekly)	24-36 (depends on skill)	0
Review responses	Same day, every review	Next business day, most reviews	None
Referral touchpoints	Every 90 days, every source	Quarterly if remembered	Sporadic
Intake response time	Under 1 hour	Same business day	When someone checks
Setup time	8-12 hours (one-time)	2-4 weeks onboarding	None
Ongoing attorney time	2-3 hours/week (review and approve)	1-2 hours/week (manage coordinator)	0 (but $0 pipeline growth)

	AI Marketing System	Part-Time Marketing Coordinator	Doing Nothing
Scales with practice?	Yes – same system handles 5 or 50 leads/month	Hits capacity at ~30 leads/month	N/A

The real comparison isn't cost – it's what happens when you don't have the system.

Sarah's practice generates $800K on referrals alone. If her content engine brings in even 2 additional cases per month from organic search – conservative for a firm publishing weekly in a metro market – that's 24 new cases per year. At her average PI case value of $45K, that's $1,080,000 in additional case value from a system that costs $240/year to run.

Even if only a quarter of those leads convert and the average value is lower than her referral cases, the math is overwhelming. The cost of the AI system is a rounding error compared to the cost of not having it.

The Compound Effect

These four agents don't just run in parallel – they feed each other.

The **content engine** produces blog posts that rank in search. Those posts bring in **intake leads** that the pipeline scores and follows up. Clients who have a good experience leave **reviews** that the monitor responds to, which improves local search rankings, which brings in more leads. Referral sources see your **online presence** growing and feel more confident sending clients your way, which the **referral nurture agent** tracks and reinforces.

This is the same compound loop described in Chapter 2, applied to business development. Each agent makes the others more effective. After six months, the system is generating leads from three channels – referrals, organic search, and reputation – instead of one.

At Dodonai, we watched this happen with our own marketing automation. Our content engine feeds our SEO rankings. Our inbox triage agent detects replies to outreach emails and updates our tracking. Our lead scoring system prioritizes prospects based on engagement signals from multiple agents. No single agent is responsible for revenue growth – the system is.

Getting Started This Week

You don't need all four agents running by Friday. Here's the sequence:

This week: Deploy the intake evaluation prompt. You don't need automation for this – just paste every new inquiry into your AI assistant with the intake evaluation prompt from this chapter. Score it, draft a response, send it. You'll immediately notice two things: faster response times and more consistent evaluation. This is the highest-impact starting point because it directly affects case signups.

Week 2-3: Set up the content engine. Pick four blog topics from questions clients ask you most frequently. Generate one post per week. Review it, publish it. After a month, you'll have four pages of content ranking for terms you weren't visible for before.

Week 4: Deploy review monitoring. This one is the simplest to set up and maintain. One daily check, occasional response drafts. Low effort, high compound value over time.

Month 2: Build the referral nurture system. This requires the most personalization but the least frequency. Set up your referral tracking, draft your first round of quarterly emails, and let the agent maintain the cadence going forward.

By month three, you have a marketing department. Not a person in a seat – a system that runs on the architecture from Chapter 3, reviewed by you, producing the same outputs a marketing coordinator would, at a fraction of the cost and with perfect consistency.

What AI Won't Do Here

AI won't build your reputation. It won't try a case, win a verdict, or earn a client's trust. The content engine drafts posts, but the expertise behind them is yours. The intake pipeline scores leads, but the judgment call on whether to take a case is yours. The referral nurture agent drafts emails, but the relationships are yours.

What AI does is eliminate the gap between having a marketing strategy and executing it. Every solo practitioner knows they should publish content, respond to reviews, nurture referral sources, and follow up with leads quickly. Almost none of them do it consistently because the day fills up with billable work and the marketing falls off the list.

The system closes that gap. You provide the expertise, the judgment, and the relationships. The agents provide the consistency, the speed, and the follow-through. That's the division of labor that lets a solo practitioner compete with firms that have dedicated marketing staff – not by working more hours, but by building the system once and letting it run.

Quick Reference: Chapter 8

Agent	Function	Frequency	Attorney Time
Content Engine	Weekly SEO blog posts from practice expertise	Weekly	15 min review per post
Review Monitor	Google/Avvo monitoring + response drafts	Daily	30 sec per review (approve draft)
Referral Nurture	Quarterly touchpoints to referral sources	Quarterly + on case resolution	5 min per email (review + personalize)
Intake Pipeline	Lead scoring, follow-up, conflict check, consult prep	Continuous during business hours	5 min per lead (review assessment)

Prompt	Use Case	Page
Blog Post Generation	Weekly content drafts from practice topics	This chapter
Review Response	Google/Avvo review reply drafts	This chapter
Intake Evaluation	New lead scoring and triage	This chapter
Referral Follow-Up	Quarterly nurture emails to referral sources	This chapter

Chapter 9: Operations – Billing, SOPs, Knowledge

What You'll Walk Away With:

- How to recover 10-30% more billable time by reconstructing unbilled work from calendar, email, and document activity
- A system for turning tribal knowledge ("ask Maria") into written SOPs in under an hour
- AI-ready template libraries that auto-populate from case data – not static Word forms
- An onboarding workflow that turns a new paralegal's first two weeks from hand-holding into a structured plan
- The math: $72,800/year in recovered capacity per attorney, and what that number actually means for your practice

---

The Revenue You're Leaving in the System

Sarah Chen billed 5.2 hours on a Tuesday she worked at least 9.

That gap – 3.8 hours of substantive legal work that never became a time entry – is the norm for most litigators. Not because Sarah is careless about billing. Because by 5:30 PM, when she finally opens Clio to enter time, the 45-minute call with Dr. Kaplan about the Hernandez medical records has blurred into the hour she spent reviewing the expert's revised report. The three emails to opposing counsel on Rivera somehow aren't in her mental ledger. The 20 minutes she spent preparing for the deposition on Patel feels too small to log, so she rounds down or skips it entirely.

Multiply that across a week, and you get the billing leakage that every practice management consultant warns about. BigHand SmartTime data puts the recovery opportunity at 8-10 extra billable hours per month per attorney when AI-assisted time tracking is in place. Billables AI reports 10-30% more billable time captured. At a $350/hour blended rate, those recovered hours add up to $72,800 per year (Rankings.io).

That's not theoretical revenue. It's work you already did – depositions you prepared for, expert correspondence you managed, client calls you took – that never made it from your calendar into your billing system.

Billing leakage is the most visible example, but it's part of a broader pattern. Maria knows how to open a new PI file because she's done it 300 times – but that process lives in her head, not in a written SOP. Your engagement letter template was last updated in 2019. Your new paralegal's onboarding plan is "shadow Maria for two weeks."

Four systems fix it: time entry reconstruction, SOP creation, template libraries, and onboarding automation.

Time Entry Reconstruction

The Problem

In Chapter 1, Sarah Chen's time audit showed 2.5 hours per week on billing administration. But the real cost isn't the reconciliation time – it's the hours that never get reconciled at all. The 45-minute call between two meetings? Gone. The email chain that required 30 minutes of research? Billed as 0.1 hours or skipped entirely.

Chapter 7 introduced the Billing Reconciler as Agent #4. Here, we go deeper: the complete workflow for time entry reconstruction, the prompt that drives it, and the operational changes that make it stick.

How It Works

Time entry reconstruction is a three-source comparison:

Source	What It Shows	Example
Calendar	Meetings, calls, depositions, court appearances, deadlines with time blocks	"45-min call with Dr. Kaplan – Hernandez" on March 31 at 2:15 PM
Email	Substantive correspondence per matter – replies, forwards, attachments sent/received	6 emails exchanged with expert witness Dr. Kaplan on April 1-2, including attachment reviews
Document activity	Files opened, edited, or created in your DMS or case management system	Medical chronology for Hernandez opened 3 times on April 1, edited twice
Clio time entries	What was actually billed	0.3 hrs total for Hernandez on April 1-2

The agent cross-references these four sources, identifies gaps, and generates suggested time entries with matter names, descriptions, and recommended durations. It doesn't enter them for you – it drafts them for review.

The Prompt

```
You are a billing reconciler for a litigation practice. Your job is
to
identify unbilled work by comparing three data sources and flagging
gaps.

DATA SOURCES:
1. Calendar events for the past [WEEK/DAY] (provided below)
2. Email activity by matter for the past [WEEK/DAY] (provided below)
3. Current time entries in Clio for the same period (provided below)

RECONCILIATION RULES:
- A calendar event with a client or matter name that has no
corresponding
  Clio time entry is an UNBILLED GAP
- 3+ substantive emails on a single matter in one day with less than
  0.5 hrs billed suggests missing correspondence/review time
- A calendar block of 30+ minutes with a partial Clio entry (billed
  duration < 75% of calendar duration) is UNDERBILLED
- Exclude: internal team meetings with no client/matter reference,
  personal calendar events, automated notifications
- Exclude: emails that are clearly administrative (scheduling,
  forwarding without comment, out-of-office replies)

FOR EACH GAP, PROVIDE:
- Matter name and case type
- Source of the evidence (calendar event, email chain, or both)
- Date and approximate time
- Suggested time entry description (in the style of a Clio narrative
--
  professional, specific, using standard billing language)
- Suggested duration (in tenths of an hour)
- Confidence level: HIGH (calendar event with no entry at all),
  MEDIUM (email activity suggests unbilled time), LOW (indirect
  evidence only)

WEEKLY SUMMARY:
- Total activity identified across all sources (hours)
- Total billed in Clio (hours)
- Total gap (hours and dollar value at $[RATE]/hr)
- Number of suggested entries by confidence level

Sort gaps by dollar value, highest first. Group by matter.
```

This prompt runs against real data from your systems – the calendar events, email logs, and Clio exports the agent pulls through its connections (Chapter 3). The output is what the Chapter 7 Billing Reconciler produces every Friday at 5:00 PM: a reconciliation report with specific, approvable time entries.

What the Output Looks Like

Chapter 7 showed the full sample output from the Billing Reconciler. Here's the summary line that matters:

```
WEEKLY SUMMARY:
  Total activity identified: 38.2 hours
```

```
Total Clio time entries: 31.6 hours
Identified gap: 6.6 hours ($2,402.40 at $364/hr)
Suggested entries: 14 (8 high confidence, 4 medium, 2 low)
```

Each suggested entry includes the matter name, the evidence source (calendar event, email chain, or both), a narrative description in standard billing language, and a recommended duration in tenths of an hour. Sarah's review takes 15-20 minutes. She approves 11 of the 14 suggestions, adjusts two descriptions, and deletes one that flagged an internal call incorrectly. Those 15 minutes recovered $2,100 in billable time that would have otherwise evaporated.

The Compounding Effect

Time entry reconstruction changes billing behavior. After a month of seeing the same patterns – expert correspondence consistently underbilled, client calls routinely missed – you start catching them in real time. The agent trains you to bill more accurately by showing you what you keep missing. Most firms see a spike in the first month (the backlog effect) and then a sustained improvement as billing habits shift.

What AI Won't Do Here

The agent flags gaps. It doesn't decide what's billable. Whether a particular email exchange constitutes substantive legal work, whether a research detour was necessary for the matter, whether the engagement letter caps certain work categories – that judgment stays with you. Every suggested entry needs human review before it enters Clio.

The agent also can't reconstruct work that left no digital trace. Hallway conversations with co-counsel, legal theories worked through on a yellow pad, Westlaw research without a matter tag – none of that shows up in calendar or email data. Time entry reconstruction captures digitally-evidenced work. The rest still depends on contemporaneous time tracking.

SOP Creation: From Tribal Knowledge to Written Procedures

The Problem

"Ask Maria, she knows."

Every firm has a Maria. The paralegal who has been there long enough to answer every procedural question. How do we open a new PI file? Ask Maria. What's the

process for ordering records from Memorial Hermann? Maria handles that. How do we prepare the demand package? Maria has a system.

Maria's knowledge is an asset – until she takes a two-week vacation, calls in sick on a filing deadline, or gives notice. At that point, the firm discovers that its operational procedures live in one person's memory. Nothing is written down. Nothing is transferable.

This is a documentation problem that technology now solves in under an hour.

The Workflow

SOP creation with AI follows a simple loop: **describe, generate, refine.** You talk through the process – or have Maria talk through it – and the AI turns it into a structured, step-by-step procedure with checklists.

Here's the prompt:

```
I'm going to describe a process we follow at our law firm. Turn it
into a formal Standard Operating Procedure (SOP) with the following
structure:

SOP FORMAT:
- Title and purpose (one sentence)
- Roles involved (who does what)
- Trigger (what initiates this process)
- Step-by-step procedure with numbered steps
- For each step: the action, who performs it, expected timeframe,
  and any system where the action is recorded
- Checklist version (checkbox format for daily use)
- Common exceptions and how to handle them
- Quality check: what does "done correctly" look like?

RULES:
- Use specific system names (Clio, Gmail, QuickBooks) where relevant
- Include the actual form names, template names, or document titles
  we reference
- Flag any step where I'm vague and ask me to clarify
- Write for someone who has never done this before -- assume no
  prior knowledge of our firm's systems

Here is the process:

[DESCRIBE THE PROCESS -- speak naturally, as if you're training
a new hire. Include the exceptions and edge cases you've learned
from experience. The messier and more detailed, the better.]
```

What This Looks Like in Practice

Sarah sits down with Maria for 10 minutes and records Maria describing how she opens a new PI file. Maria talks through it naturally:

"OK, so when we get a new PI case, first I create the matter in Clio. I use the client's last name and a short description – like 'Martinez - Auto Accident I-35.'

Then I set up the custom fields: date of loss, SOL date, insurance carrier, policy limits if we have them. I create the standard folder structure in the matter – Correspondence, Medical Records, Billing Records, Pleadings, Discovery, Expert, and Settlement. Then I send the engagement letter from the template – but I always check whether it's a contingency case or hybrid, because the language is different. After the engagement letter goes out, I calendar the SOL date and set a 90-day reminder. Then I order the crash report from DPS and request the first set of medical records – I usually start with the ER and the primary care doc. Oh, and I add the insurance adjuster's contact info if we have it, and I send the LOR to the carrier within 48 hours."

That 10-minute description becomes a two-page SOP with 14 numbered steps, a checklist version, exception handling for contingency vs. hybrid fee structures, and a quality check section that specifies what a properly opened file looks like in Clio. The AI flags two places where Maria was vague ("I usually start with the ER" – which ER? The one from the crash report or the client's reported ER?) and asks for clarification.

Total time: 10 minutes of Maria talking, 5 minutes of review and clarification, 15 minutes for the AI to generate. A written SOP that would have taken 2-3 hours to draft from scratch.

Building the SOP Library

One SOP isn't a system. A library of SOPs is. The high-value procedures to document first:

Process	Why It's High Priority	Typical Complexity
New PI file opening	Done 8x/month at Sarah's volume; every step affects downstream case management	Medium – 12-15 steps
Medical records request and tracking	Errors here delay the entire case timeline	High – multiple providers, follow-up cycles
Demand package preparation	High-dollar output, quality-sensitive, multi-step	High – 20+ steps with conditional branches
Discovery response workflow	Deadline-driven, malpractice risk if missed	High – varies by request type
Client communication cadence	Affects retention and satisfaction; often inconsistent	Low – 5-8 steps
Engagement letter processing	Compliance-critical, fee structure variations	Medium – 8-10 steps with conditionals
Settlement disbursement	Trust account rules, bar compliance requirements	High – compliance-critical

Process	Why It's High Priority	Typical Complexity
Expert retention process	Timing matters for depositions and trial prep	Medium – 10-12 steps

Start with the process your firm runs most frequently. For Sarah, that's new PI file opening (8 per month) and medical records tracking (ongoing across all 40 matters). Document those two first. Add one new SOP per week. Within two months, you have the core operational manual that most firms never build.

The Compound Value of Written SOPs

Written SOPs do more than protect against the "Maria leaves" scenario. They become the foundation for three other systems in this chapter:

1. **Template libraries** – SOPs define when and how templates are used, turning isolated documents into workflow components
2. **Onboarding plans** – new hires learn from written procedures instead of oral tradition
3. **Agent instructions** – an AI agent that processes intake or prepares demand packages needs a procedure to follow. Your SOPs become skill files (Chapter 3)

That last point deserves emphasis. In Chapter 3, we described a skill file as "a set of plain-English instructions your AI follows every time." An SOP is a skill file for humans. Convert your SOPs into skill files and you've automated the procedure entirely – with a human review step at the end.

Template Libraries: From Static Forms to Living Documents

The Problem

Your engagement letter template was created four years ago. It's a Word document with yellow-highlighted fields: [CLIENT NAME], [DATE OF LOSS], [FEE PERCENTAGE]. Your paralegal opens it, manually replaces each field, saves it with the client's name, and files it. Sometimes a field gets missed. Sometimes the fee language reflects your old fee structure instead of your current one. Sometimes the demand letter template references a section of the insurance code that was renumbered last legislative session.

Static templates are better than no templates. But they degrade over time, they require manual population, and they don't connect to your case data. The result:

your paralegal spends 20-30 minutes per template customizing a document that should take 2 minutes.

AI-Ready Templates

An AI-ready template is a document structure paired with a data mapping and a generation prompt. Instead of a Word document with highlighted fields, it's an instruction set that tells your AI: "Here's the template structure. Here's where the case data comes from. Generate the document."

The key difference: AI-ready templates pull data directly from your case management system. The client name, date of loss, insurance carrier, medical providers, treatment dates, and special damages figures come from Clio. The AI populates the template, formats it to your specifications, and produces a draft for review.

The Prompt: Engagement Letter Generator

```
Generate an engagement letter for a new personal injury case using
the
following template structure and case data.

TEMPLATE STRUCTURE:
1. Opening paragraph: Confirm retention, identify the specific
incident
   (date, location, type of accident), name the adverse party or
   parties
2. Scope of representation: PI claim arising from [INCIDENT TYPE],
   including negotiation, litigation if necessary, and settlement
3. Fee arrangement: [CONTINGENCY/HYBRID/HOURLY] -- use the
   appropriate language block below
   - Contingency: [FEE_PERCENTAGE]% of gross recovery, advancing
     costs, client responsible for costs if no recovery
   - Hybrid: $[RETAINER] retainer + [FEE_PERCENTAGE]% of recovery
     above $[THRESHOLD]
   - Hourly: $[RATE]/hour, $[RETAINER] retainer, monthly billing
4. Client obligations: cooperation, truthfulness, medical treatment
   compliance, no direct contact with adverse parties or insurers
5. Termination provisions: either party, written notice, client
   responsible for costs incurred to date
6. Closing: signature blocks for attorney and client, date line

CASE DATA (from Clio):
- Client: [CLIENT_NAME]
- Matter: [MATTER_NAME]
- Date of loss: [DATE_OF_LOSS]
- Incident type: [INCIDENT_TYPE]
- Adverse party: [ADVERSE_PARTY]
- Fee structure: [FEE_TYPE]
- Fee percentage: [FEE_PERCENTAGE]
- Retainer amount: [RETAINER] (if applicable)

STYLE:
- Professional but accessible -- the client is not a lawyer
```

```
- Short paragraphs, no legalese where plain English works
- Texas-compliant engagement letter format
- Match the tone and structure of [ATTACH: your most recent
  engagement letter as a style reference]

Generate the complete letter. Flag any missing data fields that
prevent completion.
```

Scaling to a Template Library

The engagement letter is one template. A PI practice needs a library:

Template	Frequency	Data Sources	Complexity
Engagement letter	Every new case (~8/month)	Clio intake fields	Medium
Letter of representation (LOR)	Every new case	Clio, insurance carrier info	Low
Demand letter	Every case pre-suit (~3-4/month)	Clio, medical records, billing records, Dodonai medical chronology	High
Discovery responses (interrogatories)	Per case in litigation	Clio, case notes, client interview	High
Discovery responses (RFPs)	Per case in litigation	Clio, document index	Medium
Client status update	Monthly per active case (~40/month)	Clio activity log, calendar, upcoming deadlines	Medium
Expert retention letter	As needed (~2-3/month)	Clio, expert CV, case theory	Medium
Settlement statement	At resolution (~4-5/month)	Clio, medical billing, lien amounts, fee calculation	High
Referral thank-you letter	Per referral (~3-4/month)	Clio intake source, referrer contact	Low

Each template follows the same pattern: a structure definition, a data mapping to your case management system, and a generation prompt. Build them one at a time, starting with the highest-frequency templates. The engagement letter and LOR handle every new intake. The client status update covers your largest recurring output (40 per month for Sarah). The demand letter is the highest-value document and deserves the most detailed template.

Living Documents, Not Frozen Forms

When Texas updates its insurance code provisions, you update the reference in one template prompt – not in 15 saved Word documents. When you refine your demand letter structure after a strong settlement, the improvement propagates to every future demand. This is the skill file concept from Chapter 3 applied to

documents. Your templates are instructions your AI follows. When you improve the instructions, every output improves.

Training and Onboarding: The New Hire's First Two Weeks

The Problem

A new paralegal starts Monday. In most small firms: she shadows Maria for two weeks. She watches, asks questions, and picks up some habits correctly and some incorrectly. By month three, she's developed her own version of the firm's procedures – slightly different from Maria's, which was slightly different from the written version that doesn't exist.

If you've built the SOP library described in the previous section, onboarding changes from a knowledge transfer problem to a scheduling problem. The procedures exist. The question is: in what order does the new hire learn them, and how do you verify she's learned them correctly?

The Prompt: Onboarding Plan Generator

```
Generate a 2-week onboarding plan for a new paralegal at a personal
injury law firm (solo practitioner, 40 active cases, using Clio,
Gmail/Google Workspace, QuickBooks, and Dodonai for document
processing).

AVAILABLE SOPs (the new hire should learn these during onboarding):
[LIST YOUR SOPs -- paste the titles and one-sentence descriptions]

ONBOARDING PLAN REQUIREMENTS:
- Day-by-day schedule for 10 business days
- Each day: 2-3 learning objectives, specific SOPs to review,
  hands-on tasks to practice, and a verification checkpoint
- Week 1 focus: systems access, core procedures (file opening,
  calendaring, basic correspondence), observation of active matters
- Week 2 focus: independent execution under review (open a real
  file, draft a real LOR, process real medical records), exception
  handling, quality standards
- Include "shadow" sessions where the new hire observes the
  paralegal or attorney handling specific tasks
- Include a Day 10 assessment: checklist of competencies the new
  hire should demonstrate independently

PRACTICAL CONSTRAINTS:
- The existing paralegal (Maria) can dedicate 1 hour/day to
  training. Remaining learning is self-directed from SOPs.
- The new hire has prior paralegal experience but is new to PI
  and new to Clio
- Systems access (Clio login, email, Dodonai account) should be
```

```
  set up on Day 1

Format as a table: Day | Morning | Afternoon | SOPs Referenced |
Checkpoint
```

What the Output Looks Like

The AI generates a structured plan that maps your SOPs to a learning sequence:

Week 1: Foundation

Day	Morning	Afternoon	SOPs Referenced	Checkpoint
1	Systems setup (Clio, Gmail, Dodonai). Firm overview with Sarah.	Shadow Maria: observe 2 client calls, 1 file review. Read SOPs 1-3.	Firm Overview, Clio Navigation, File Structure	Can log into all systems, navigate Clio matter list
2	Review PI case lifecycle SOP. Walk through 3 active matters in Clio with Maria.	Practice: open a test matter in Clio, create folder structure.	New PI File Opening, Clio Matter Setup	Test matter created correctly per SOP checklist
3	Medical records request SOP. Review 5 completed requests in active files.	Shadow Maria: observe medical records follow-up calls. Draft 2 practice requests.	Medical Records Request, Provider Contact Protocol	2 practice requests match SOP format
4	Calendaring SOP. Review deadline tracking across 5 active matters.	Practice: calendar deadlines from 3 sample court notices.	Deadline Calendaring, SOL Tracking	All 3 deadlines calendared correctly with reminders
5	Client communication SOP. Review recent client update letters.	Draft 2 practice client status updates from Clio data. Maria reviews.	Client Communication Cadence, Status Update Template	Updates match firm style, all required fields populated

Week 2: Supervised Execution

The second week shifts from practice exercises to real matters under review. The new paralegal opens an actual new file (Maria checks every step against the SOP), drafts a real LOR, processes incoming medical records for an active case, and by Day 10, demonstrates independent competency on the core procedures.

Why This Matters Beyond Onboarding

The onboarding plan isn't a one-time artifact. It becomes a performance reference (check the SOP, not Maria's memory), a hiring tool (the plan exists before the next hire starts), a cross-training framework (David can learn demand package preparation from the same SOP library), and – critically – an agent training resource. The SOPs that train your paralegal are the same instructions that train your AI agents. This is the bridge between Chapter 3's skill files and your firm's operational knowledge.

Connecting the Four Systems

These four workflows reinforce each other. SOPs define when templates are used. Templates generate trackable document activity that the billing reconciler captures. The billing reconciler reveals undocumented procedures (expert correspondence that's consistently underbilled because no one documented it as a billable workflow). New hire questions during onboarding expose gaps in existing SOPs. And SOPs written for humans convert directly into AI skill files (Chapter 3).

The virtuous cycle: you write an SOP for demand package preparation. That SOP defines the template structure for the demand letter. The template generates the demand letter from case data. The billing reconciler catches the 3 hours of preparation work that went unbilled. The next hire learns the entire process from the SOP and template, not from watching over someone's shoulder.

The Connections Map

Here's what these systems need access to:

System	Clio	Calendar	Gmail	QuickBooks	Dodonai	File System
Billing Reconciler	Time entries, matters	Events, time blocks	Sent/received per matter	Invoice history	–	–
SOP Library	Matter types, workflow data	–	–	–	–	SOP storage (drive/wiki)
Template Library	All case data fields	Upcoming dates	Prior correspondence (style ref)	Fee/cost data	Medical chronologies, deposition summaries	Template storage

System	Clio	Calendar	Gmail	QuickBooks	Dodonai	File System
Onboarding	User permissions	Training schedule	New hire communications	–	Account setup	SOP access

If you built the Deadline Monitor and Inbox Triage agents from Chapter 7, the Clio, Calendar, and Gmail connections already exist. The Billing Reconciler adds QuickBooks. The Template Library adds Dodonai as a data source for medical chronologies and deposition summaries that feed into demand letters and case evaluations.

The Math

Here's what these four systems recover for a practice like Sarah Chen's:

System	Weekly Time Saved	How	Annual Value (at $364/hr)
Billing Reconciler	2.0 hrs admin time recovered + unbilled hours captured	Friday reconciliation replaces daily billing anxiety; 6+ hrs/week in unbilled work captured	$37,856 admin + $72,800 recovered billing
SOP Creation	1.0 hr/week (ongoing maintenance) vs. 1.5 hrs/week (current ad hoc training)	Written procedures replace repeated verbal explanations	$26,208 (Maria's time freed)
Template Library	1.5 hrs/week	Auto-populated templates replace manual document assembly	$28,392
Onboarding	Variable (2-3 weeks saved per hire)	Structured plan vs. unstructured shadowing	~$5,000 per new hire in reduced ramp time

The billing reconciler line deserves explanation. The $37,856 represents Sarah's own time saved on billing administration (2 hours/week she no longer spends manually reconciling). The $72,800 represents the additional revenue from unbilled work the agent catches – work she already performed but wasn't capturing. These are different pools: one is efficiency, the other is revenue recovery.

Combined, the four systems recover roughly $90,000-$110,000 in annual value for a solo practice. Scale that to a 5-attorney firm and you're looking at $400,000+ in recovered capacity and revenue.

What “Recovered Capacity” Actually Means

A note on the $72,800 figure, because it appears throughout this book and throughout legal AI marketing. Recovered capacity is not the same as recovered revenue. Those 10 additional billable hours per month are only worth $72,800/year if you can bill them. For a solo at capacity, that’s one more case per quarter. For a solo who’s burning out, it’s leaving the office an hour earlier. For a firm with excess demand, it converts almost dollar-for-dollar.

The honest framing: time entry reconstruction ensures you get paid for work you already do. Whether the recovered capacity becomes revenue, work-life balance, or case quality improvement depends on how you deploy it.

Getting Started This Week

You don’t build all four systems at once. Here’s the sequence:

Week 1: Billing Reconciler

This has the most immediate financial impact. Export your Clio time entries, calendar events, and email activity for the past week. Run the time entry reconstruction prompt against the data. See what comes back. If it identifies even 3-4 hours of unbilled work, you’ve justified the setup time in the first run.

Week 2: First Three SOPs

Pick the three procedures your team runs most frequently. Sit with Maria (or whoever holds the knowledge) for 10 minutes each. Record the description, run it through the SOP prompt, and review the output. You’ll have three written procedures by Friday.

Week 3: Template Library (Start with Two)

Convert your engagement letter and client status update into AI-ready templates. These are your highest-frequency documents. Test them against 2-3 recent cases to verify the output matches your standards.

Month 2: Expand and Connect

Add SOPs weekly. Convert each into both a template (where applicable) and a potential skill file. Set up the Billing Reconciler to run automatically every Friday (Chapter 7’s Agent #4 schedule). When the next hire comes, the onboarding plan generates itself from the SOP library you’ve built.

What AI Won't Do Here

It won't decide what's billable. The billing reconciler flags gaps. You decide whether a particular activity constitutes billable work under your engagement letter and fee structure. Ethical billing is a judgment call, not an algorithm.

It won't replace Maria's judgment. SOPs capture procedures, not the experience that informs when to deviate from them. Maria knows that Memorial Hermann takes 6 weeks to produce records but Methodist responds in 10 days. That knowledge belongs in the SOP's notes section, but the judgment of how to sequence record requests across a complex case still requires a human who understands the matter.

It won't maintain itself. SOPs go stale. Templates drift from current practice. Plan for a monthly 30-minute review: audit the SOP library, update changed procedures, refine template prompts based on the outputs you've been editing.

It won't handle trust account compliance. Settlement disbursements, IOLTA requirements, and client fund management involve bar compliance rules that require attorney oversight at every step. AI can generate the calculations. A human must verify them.

Quick Reference: Operations Systems

System	What It Does	Connections	Schedule	Effort to Build
Billing Reconciler	Cross-references calendar + email + docs against time entries; flags unbilled work	Clio, Calendar, Gmail, QuickBooks	Weekly (Friday PM)	Weekend project
SOP Library	Converts verbal process descriptions into structured, written procedures	Firm knowledge + AI	On-demand (10 min per SOP)	30 min per SOP
Template Library	AI-ready document templates that auto-populate from case data	Clio, Dodonai, file system	On-demand (per document)	1-2 hrs per template
Onboarding Plan	Generates structured training schedule from SOP library	SOP Library, Clio, Calendar	Per new hire	1 hr to generate

Prompt	Section	Purpose
Billing reconciliation	Time Entry Reconstruction	Weekly unbilled work detection
SOP generator	SOP Creation	Convert verbal procedures to written SOPs
Engagement letter generator	Template Libraries	Auto-populated document from case data
Onboarding plan generator	Training and Onboarding	2-week structured training schedule

Chapter 10: The Safety Layer – Ethics, Risks, and Audit Trails

---

What You'll Walk Away With:

- The real hallucination numbers: 600+ cases, 128 lawyers, courts issuing sanctions – and what actually went wrong in each pattern
- A clear map of your ethical obligations under ABA Formal Opinion 512 and the emerging state-by-state rules
- Why 43% of firms with no AI policy are carrying uninsured risk
- The confidentiality framework: zero-retention policies, privilege landmines, and the discoverability of your chat logs
- The four-principle safety architecture that prevents AI errors from reaching clients, courts, or opposing counsel
- A ready-to-adopt AI policy template (Appendix E) you can customize for your firm this week

---

The Case That Changed the Conversation

In 2023, a New York attorney submitted a brief containing six case citations generated by ChatGPT. None of the cases existed. The court discovered the fabrications, sanctioned the attorney, sanctioned his firm, and referred the matter to the state bar. The attorney told the court he didn't know ChatGPT could generate fictitious citations. The court was unpersuaded.

That case – *Mata v. Avianca* – became the opening anecdote in every legal AI presentation for a year. But it wasn't the end of the story. It was the beginning.

As of early 2026, the National Law Review has documented 600+ cases involving AI hallucinations, implicating 128 attorneys. Courts have issued monetary sanctions, disqualified firms from ongoing representations, referred lawyers to state bar disciplinary authorities, and in at least one case struck an entire filing from the record.

The pattern in these cases is remarkably consistent. It's not that the attorneys used AI. It's that they treated AI output as verified work product. They skipped the step that every lawyer applies to every other source of information in their practice: checking it.

A first-year associate who hands you a brief with unchecked citations gets a lecture. An AI tool that generates unchecked citations isn't the one who faces the bar complaint. You are.

This chapter is the honest assessment of what can go wrong when AI operates inside a law practice – and the specific safeguards that prevent it.

The Hallucination Problem

AI hallucination is a technical term for a simple failure: the model generates text that looks authoritative but is factually wrong. In legal work, this most often manifests as fabricated case citations – complete with realistic-sounding case names, accurate-looking reporter citations, and plausible holdings that don't correspond to any real case.

Why does this happen? Large language models predict the next word in a sequence based on patterns in their training data. When asked for a case supporting a legal proposition, the model doesn't search Westlaw. It generates text that looks like a citation because it has seen millions of citations during training. Sometimes the generated citation corresponds to a real case. Sometimes it doesn't. The model has no internal mechanism to distinguish between the two.

The Three Patterns

The 600+ documented hallucination cases fall into three categories:

Pattern 1: Fabricated citations. This is the *Mata v. Avianca* pattern. The attorney asks AI to find cases supporting a legal argument. The AI generates citations that don't exist. The attorney includes them in a filing without checking. The court or opposing counsel discovers the fabrication. This pattern accounts for the majority of documented cases.

Pattern 2: Real case, wrong holding. The AI cites a real case, but mischaracterizes what the court held. The case name and reporter citation are accurate; the description of the holding is not. This is harder to catch because the citation passes a superficial check – the case exists – but fails a substantive one. An attorney who Shepardizes the citation but doesn't re-read the case may miss the error.

Pattern 3: Outdated or overruled authority. The AI cites a case that was good law when the model's training data was collected but has since been overruled, distinguished, or superseded by statute. The model's training data has a cutoff; the law doesn't.

Each pattern has a different detection method:

Pattern	What Goes Wrong	How to Catch It
Fabricated citation	Case doesn't exist	Verify every citation in Westlaw, Lexis, or the primary source
Real case, wrong holding	Case exists but holding is mischaracterized	Read the actual case, not just the AI's summary of it
Outdated authority	Case was good law at training cutoff but has since been overruled	Shepardize or KeyCite every citation; check for subsequent history

The common thread across all three: **verification catches the error before it reaches a court.** The attorneys who got sanctioned didn't fail because they used AI. They failed because they skipped verification. The distinction matters, because it means AI-assisted legal work is not inherently risky – unverified AI-assisted legal work is.

The Rule

Every attorney in your firm, every paralegal, every contract attorney needs to internalize one rule:

Never file, send, or rely on AI-generated work without verifying every citation against the primary source and every factual claim against the case file.

This is not a new professional obligation. It's the same duty of competence you already apply to work produced by associates, contract attorneys, and legal research services. AI doesn't get a lower standard. It gets the same one.

Bar Ethics Rules: Where Things Stand

The legal profession's regulatory response to AI has been faster and more substantive than most attorneys realize. The framework is converging around five obligations, and your practice needs to address all of them.

ABA Formal Opinion 512 (2024)

ABA Formal Opinion 512 is the baseline. Issued in 2024, it establishes that existing Model Rules of Professional Conduct apply to attorney use of generative AI tools. The opinion doesn't create new obligations – it maps existing ones onto AI use:

Competence (Rule 1.1). You must understand the AI tools you use well enough to recognize their limitations. This doesn't mean understanding the technical architecture of neural networks. It means knowing that AI hallucinates, understanding what zero-retention means, and being able to evaluate whether an

AI output is reliable for the task at hand. If you don't understand the tool well enough to supervise its output, you don't understand it well enough to use it on client matters.

Communication (Rule 1.4). You must keep clients informed about how AI is used in their representation. The degree of communication depends on the context. Using AI to draft an internal case analysis memo is different from using AI to generate a section of a brief filed with the court. The opinion emphasizes that clients have the right to know about material aspects of their representation, and AI use on their matter is material.

Supervision (Rules 5.1 and 5.3). The duty to supervise extends to AI-generated work product the same way it extends to work product from associates and non-lawyer assistants. A supervising partner who signs off on an AI-generated brief without review bears the same responsibility as one who signs off on an associate's brief without review. The AI doesn't reduce the supervision obligation. As the opinion notes, it may increase it, because AI errors can be subtler than human errors.

Confidentiality (Rule 1.6). Inputting client information into an AI tool constitutes a disclosure. If the tool stores, retains, or uses that data for training, you have a confidentiality problem. The opinion requires attorneys to understand their tools' data handling practices before using them with client information.

Candor (Rule 3.3). Submitting AI-generated content to a tribunal without verification risks violating the duty of candor. If the AI fabricates a citation or mischaracterizes a holding, and you present it to the court as accurate, you've made a misrepresentation – regardless of whether you knew it was false.

California's Proposed Amendments

California has gone further than the ABA by proposing specific amendments to six rules of professional conduct. Public comment closed in May 2026, and adoption is expected. The proposed amendments touch Rules 1.1, 1.4, 1.6, 3.3, 5.1, and 5.3 – the same rules mapped in ABA Formal Opinion 512, but with AI-specific language baked into the rule text rather than addressed through an interpretive opinion.

The California approach signals a direction other states will likely follow: explicit AI provisions in the rules themselves, not just guidance opinions.

State-by-State Variation

This is where it gets complicated. As of early 2026:

- Multiple states have issued formal ethics opinions on AI use, most tracking ABA Formal Opinion 512
- Some state bars require affirmative disclosure of AI use in court filings

- Several federal judges have entered standing orders requiring AI disclosure
- Some jurisdictions have no AI-specific guidance at all

The variation creates a practical problem for any firm practicing across state lines – and in federal courts, where individual judges may impose their own requirements. You cannot assume that the rules in your home jurisdiction apply everywhere you practice.

The practical response: Check the AI disclosure requirements in every jurisdiction where you file. Maintain a current list. When in doubt, disclose. Overcommunication builds trust with courts; concealment creates risk.

The 43% Problem

Here's a number that should concern every managing partner: 43% of law firms have no AI policy and no plans to create one (Wolters Kluwer 2026 Future Ready Lawyer).

That's not a technology gap. It's a liability gap. Consider what that means in practice:

- No approved tool list, so personnel use whatever AI they find convenient – including free-tier tools that train on user input
- No verification requirements, so AI-generated citations may be filed without checks
- No data classification, so client data may be entered into tools with unknown data handling practices
- No supervision standards, so AI outputs may bypass attorney review
- No disclosure policy, so the firm has no consistent approach when courts or clients ask about AI use
- No incident response plan, so when something goes wrong (and at scale, it will), there's no playbook

Every one of these gaps is a potential disciplinary complaint, malpractice claim, or disqualification motion waiting to happen.

The fix is not complicated. It's a written policy. Appendix E provides a complete, ready-to-customize template that covers approved tools, data handling, verification, disclosure, supervision, training, and incident response. A motivated partner can adapt it for their firm in an afternoon. The cost of not having one is measured in bar complaints and malpractice premiums.

The Carlton Fields Position

Not every firm has landed where we have on AI adoption. Carlton Fields – an Am Law 200 firm with roughly 350 attorneys – has prohibited the use of generative AI for legal research and written advocacy.

Their reasoning deserves a fair hearing, because it reflects genuine concerns:

Liability concentration. When an AI hallucination makes it into a filed brief, the liability falls entirely on the firm and the signing attorney. For a firm with deep pockets and a strong reputation, the reputational damage of a *Mata v. Avianca*-type incident may outweigh the efficiency gains. Their calculus is that the risk-reward ratio doesn't favor generative AI for court-facing work – yet.

Quality control at scale. A 350-attorney firm can't monitor every attorney's AI usage the way a solo practitioner can. The verification step that prevents hallucinations is a human behavior, not a technical control. At scale, ensuring that every attorney applies that step every time is a management challenge that grows with headcount.

Client expectations. Some clients – particularly in regulated industries – may have their own restrictions on AI use by outside counsel. A blanket prohibition simplifies compliance with those client-specific requirements.

You can disagree with their conclusion and still learn from their analysis. Every firm that adopts AI needs to answer the questions Carlton Fields is raising:

- What is your verification protocol, and how do you enforce it?
- What happens when verification fails and an error reaches a court?
- Can you demonstrate to your malpractice carrier that your AI safeguards are adequate?
- Do your clients know you're using AI, and are they comfortable with it?

If you can answer those questions with specifics – not generalities, but specific protocols, specific verification workflows, specific audit trails – then you've built the safety layer that makes AI adoption defensible. If you can't, Carlton Fields might have a point.

Confidentiality: The Less Obvious Risk

Hallucinations get the headlines. Confidentiality failures are the quieter, potentially more damaging risk.

Every time you type a prompt into an AI tool, you're transmitting data to a server. If that prompt contains client information – case facts, strategy discussions, witness

names, settlement figures, privileged communications – you've made a disclosure. The question is whether that disclosure is protected.

Zero-Retention Policies

The first line of defense is your AI provider's data handling policy. The two critical questions:

1. **Does the provider retain your inputs after processing?** A zero data retention (ZDR) policy means your prompts and the AI's responses are deleted after the session ends. They are not stored in a database, not used to improve the model, not accessible to the provider's employees.
2. **Does the provider use your inputs to train or fine-tune models?** Even if the provider deletes the raw data, some may use it in aggregate to improve their models. For client data, this is a confidentiality problem.

Dodonai operates under a zero data retention policy – documents uploaded for processing are not stored after the job completes, and no client data is used for model training. Claude's commercial plans (Pro, Team, Enterprise) also operate under zero data retention: conversations are not used for model training, and data is deleted according to the stated retention policy.

Free-tier AI tools often do not provide these protections. Most consumer-grade AI tools explicitly state in their terms of service that user inputs may be used for model improvement. If a paralegal enters case facts into a free ChatGPT account, those facts may become part of the model's training data – a confidentiality breach that may be impossible to remediate.

The rule: Before any client information enters any AI tool, confirm that the tool has a zero data retention policy or an enterprise agreement with appropriate data handling terms. Appendix E's data classification tiers (Tier 1, Tier 2, Tier 3) provide a framework for this.

Privilege Considerations

AI-generated work product raises privilege questions that don't yet have settled answers in most jurisdictions.

The *Heppner* case raised the issue directly: when an attorney uses AI to analyze case strategy, are the prompts and responses protected as attorney work product? The argument for protection is that the prompts reflect the attorney's mental impressions, legal theories, and analytical framework – the classic work product elements. The argument against is that the attorney shared those impressions with a third-party technology provider, potentially waiving the protection.

The safest position: treat AI prompts containing privileged information the same way you'd treat a phone call with a consulting expert. Use a tool with zero data retention. Don't include more privileged information than the task requires. And be aware that opposing counsel may argue for disclosure.

Chat Logs Are Discoverable

This is the risk that catches attorneys off guard. Your AI conversation history – every prompt, every response, every iteration – is a document. If your AI tool stores that history (and many do, for the user's convenience), those logs are potentially discoverable in litigation.

Think about what's in a typical AI conversation: case strategy discussions, privileged analysis of your client's weaknesses, draft work product, settlement figures, witness preparation notes. In a malpractice action, those logs could be produced in discovery. In a fee dispute, they could reveal how much of the work was AI-assisted. In a disciplinary proceeding, they could show whether you verified AI output before filing.

Practical safeguards:

- **Use tools with auto-deletion or zero-retention.** If the logs don't exist, they can't be produced.
- **Don't use AI tools as a case strategy journal.** Keep your analytical thinking in your case file, where attorney-client privilege and work product protections are well-established.
- **Frame prompts carefully.** "Analyze the following deposition excerpt for impeachment opportunities" is safer than "Here's our entire case strategy – tell me what's wrong with it."
- **Know your tool's retention settings.** Some tools let you disable conversation history. If yours does, consider whether to enable it.

Know Your Tool's Data Policy

Before you upload a single document or type a single case fact, you should be able to answer these questions about every AI tool your firm uses:

Question	Why It Matters
Does the tool have a zero data retention policy?	Determines whether client data is stored after processing
Does the tool use inputs for model training?	Determines whether client data could appear in model outputs for other users
Where are the servers located?	May affect compliance with data residency requirements

Question	Why It Matters
Does the tool have SOC 2 or similar certification?	Independent verification of security practices
Does the tool offer a BAA for HIPAA compliance?	Required if processing medical records or health information
What happens to data if the vendor is acquired or goes bankrupt?	Your data handling policy is only as durable as the vendor

If a vendor can't answer these questions clearly, that tool doesn't belong in your practice.

The Embarrassment Prevention Layer

Chapter 3 introduced the safety layer as the fourth element of the agent architecture: AI drafts, humans decide. Full audit trail. Uncertain items route to a review queue. Access controls limit scope.

This chapter goes deeper on why those principles aren't optional and how they work in practice.

The term comes from a pattern used in high-stakes automated systems – nuclear power plants, air traffic control, financial trading. The pattern is the same everywhere: automated systems handle the volume, computation, and monitoring. Humans authorize every action that reaches the outside world.

A nuclear reactor's automated systems monitor thousands of variables every second. They adjust coolant flow, control rod positions, and turbine speeds in real time. But they don't vent radiation to the atmosphere. That requires a human authorization. Not because the human is faster at monitoring – the automated system is better at that. But because the consequences of an error are severe enough that a human checkpoint is worth the latency.

Your law practice operates on the same principle. Filing a fabricated citation, sending a privileged document to opposing counsel, or miscalendaring a statute of limitations – these consequences are severe enough that every external action deserves a human checkpoint.

Principle 1: AI Drafts, Humans Decide

This is the foundational principle. It's worth restating with precision.

No AI agent in the architecture described in this book sends an email, files a document, communicates with a client, contacts opposing counsel, modifies a court filing, transfers funds, or takes any action that reaches a person outside your firm.

Agents draft. You decide.

When the client communication agent identifies a matter that hasn't had contact in 30 days, it drafts a follow-up email and saves it to your Gmail drafts folder. You open the draft, read it, edit it if needed, and hit send – or delete it. The agent never sends the email. The architecture doesn't allow it.

When the deadline monitor computes a response deadline from a court order, it creates a proposed calendar event with the computed date, the underlying rule, and the source document reference. You confirm the date against local rules, any stipulated extensions, and your own professional judgment – then accept, adjust, or delete it.

This isn't a transitional measure until AI gets "good enough." It's a permanent architectural decision. AI accelerates the preparation; the judgment remains yours.

Principle 2: Full Audit Trail

Every agent run produces a log. The log records:

- **When** the agent ran (timestamp)
- **What data** it accessed (which calendar events, which emails, which Clio matters)
- **What it produced** (the briefing, the draft email, the deadline computation)
- **What decisions it made** (why it classified an email as "urgent" rather than "FYI," why it flagged a matter as stale, why it computed a specific deadline)
- **What it routed to review** (ambiguous items it couldn't classify with high confidence)

Why does this matter? Three reasons.

Malpractice defense. If a client alleges a missed deadline, the audit trail shows when the deadline was flagged, what rule was applied, when you were notified, and what action you took. Either way, you have a record – not a reconstruction from memory.

Ethics compliance. ABA Formal Opinion 512 requires competence, communication, and supervision. The audit trail demonstrates all three: you understand the tool (you wrote the skill file), you can show clients how AI is used (the trail shows exactly what the agent did), and you reviewed every output before it reached anyone outside the firm (the trail logs your review).

Continuous improvement. When an agent produces an unexpected result, the log shows exactly where the process diverged. You trace it to a specific skill file instruction, a specific data source, or a specific edge case – and you fix it.

Principle 3: Uncertain Items Route to a Review Queue

Not every input fits a clean category. An email from opposing counsel might be routine correspondence or it might contain a settlement offer buried in the third paragraph. A court notice might trigger a deadline or might be informational. A medical record note might be relevant to causation or might be incidental.

When an agent's confidence in a classification falls below a threshold, it doesn't guess. It routes the item to a review queue with a flag: "Low confidence – needs attorney review."

This is the malpractice prevention mechanism. Consider the alternative: an inbox triage agent that classifies a motion to compel as "FYI" and archives it. You don't see it for three days. The response deadline passes. That's a malpractice claim.

The review queue prevents that failure mode. If the agent isn't sure whether an email is routine or urgent, it doesn't archive it. It flags it and leaves it for your review. The cost is a few extra items in your review queue each day. The benefit is that ambiguous items never fall through the cracks.

A paralegal who's unsure about a deadline should escalate it to the attorney. Your AI does the same thing. The difference is that the AI escalates consistently, every time, at the same confidence threshold, without the social pressure to "just figure it out" that sometimes keeps a junior team member from asking.

Principle 4: Access Controls and Scope Limits

Each agent operates within defined permissions. These permissions are set when the agent is built, not when it runs.

Agent	Can Read	Can Write	Cannot Do
Meeting Briefing	Calendar, Clio, Gmail, Dodonai	Save briefing file, send Slack notification	Send emails, modify calendar, edit Clio data
Inbox Triage	Gmail inbox	Apply labels, archive FYI messages	Delete emails, send replies, forward messages
Deadline Monitor	Clio matters, court filings	Create proposed calendar events	Accept hearing dates, modify deadlines, file documents
Billing Reconciler	Calendar, email, Clio time entries	Suggest additions to time entries	Create invoices, modify existing entries, process payments
Client Update Drafter	Clio matters, Gmail correspondence	Save draft emails	Send emails, modify case records

The principle: **read aggressively, write conservatively.** Agents should gather every piece of information that's relevant to their task. But their write permissions should be limited to drafts, proposals, and notifications – never to final actions.

Scope limits prevent cascading failures. If an agent malfunctions, the worst it can do is generate an incorrect draft or a wrong calendar suggestion. Both go through human review. It can't send an email to opposing counsel. It can't file a document with the court. It can't disburse trust funds. The blast radius of any agent error is contained to a draft that sits in a queue until you review it.

How the Safety Layer Maps to Ethics Rules

Chapter 3 introduced a mapping between the safety layer and bar ethics obligations. Here's the expanded version, with specific implementation details:

Ethics Obligation	Rule	Safety Layer Implementation	Evidence You Can Point To
Competence	Rule 1.1 / ABA Op. 512	You wrote or approved every skill file. You understand what data each agent accesses and what it produces.	The skill files themselves; your review history; training records
Communication	Rule 1.4 / ABA Op. 512	Your engagement letter discloses AI use. You can explain each agent's role in plain English: "Our system checks for upcoming deadlines daily and flags them for my review."	Engagement letter AI disclosure; client communication records
Supervision	Rules 5.1, 5.3 / ABA Op. 512	Every agent output passes through human review. No external action without attorney approval. Full audit trail of every run.	Audit logs showing agent output, attorney review, and final action
Confidentiality	Rule 1.6 / ABA Op. 512	Data stays within your approved tool stack. AI providers operate under zero-retention policies. No client data used for model training. Approved tool list with data classification tiers.	Vendor ZDR confirmations; approved tool list; data handling matrix
Candor	Rule 3.3 / ABA Op. 512	AI-generated content is reviewed before submission. Citation verification is mandatory. No AI output is filed without attorney confirmation.	Verification checklist in case file; audit trail showing review step

Ethics Obligation	Rule	Safety Layer Implementation	Evidence You Can Point To
Disclosure	Varies by jurisdiction	AI use disclosed in engagement letter. Court-specific disclosure requirements tracked and complied with.	Engagement letter language; jurisdiction disclosure checklist

The point isn't that the safety layer was designed to satisfy ethics rules. It's that the safety layer was designed to prevent mistakes – and preventing mistakes is what ethics rules require.

Building Your Firm's AI Policy

If your firm is in the 43% without an AI policy, this is the section that fixes it.

An AI policy doesn't need to be a 50-page document drafted by a committee over six months. It needs to answer seven questions:

1. **Which AI tools are approved for use on client matters?** Name them. Classify them by data handling tier. If it's not on the approved list, it's not approved.
2. **What data can go into which tools?** Define what types of information can be entered into each tier of tool. Client names and case facts are different from privileged strategy discussions, which are different from sealed court records.
3. **What verification is required before AI output is used?** Every citation checked against primary sources. Every factual claim checked against the case file. No exceptions.
4. **What disclosure do you provide to clients?** Add an AI disclosure to your engagement letter. Decide whether you disclose proactively or only when asked. (Proactive disclosure is the safer position.)
5. **What disclosure do you provide to courts?** Track the disclosure requirements in every jurisdiction where you file. When in doubt, disclose.
6. **Who supervises AI-assisted work product?** Define the review chain. AI outputs get the same review as a junior associate's first draft.
7. **What happens when something goes wrong?** An incident response plan: stop, assess, notify, preserve, correct.

Appendix E provides a complete, ready-to-customize template that covers all seven questions. It includes sections on approved tools, data classification tiers, verification workflows, disclosure obligations, supervision standards, training

requirements, and incident response. A managing partner or office administrator can adapt it for their firm in a single sitting.

What the Policy Looks Like in Practice

A policy that lives in a binder doesn't protect anyone. Here's how the policy translates into daily operations:

Morning triage. Your inbox triage agent runs at 7:30 AM. It classifies messages, applies labels, and archives newsletters. Before you act on any flagged message, you glance at the audit log to confirm the classifications make sense. This takes two minutes. It's your supervision step.

Draft review. Your client update drafter generates status updates for active matters. You open each draft, confirm the case facts match the file, check that the tone is appropriate for the client, and hit send or edit. This is the same review you'd do for a paralegal's draft. It's your verification step.

Deadline check. Your deadline monitor flags three deadlines within 30 days. You confirm each date against the court order, check local rules for any modifications, and accept the calendar events. If one looks off, you compute it yourself. This is your competence step – you understand the tool well enough to catch its errors.

End of week. You spend ten minutes reviewing the week's audit logs. Any anomalies? Any classifications that seemed off? Any agents that didn't run when expected? This is your quality control step – the same weekly review you'd do for any system your practice depends on.

The total time investment for maintaining the safety layer is roughly 20-30 minutes per day. That's a fraction of the time the agents save – and it's the time that makes the entire system defensible.

What AI Can't Protect You From

The safety layer catches AI errors. It doesn't catch human errors.

Overreliance. If you stop reading your agents' outputs critically – if you rubber-stamp every draft, accept every deadline computation without checking, file every AI-assisted brief without reading it through – the safety layer fails. It's a human gate, and gates only work when the human actually gates.

Scope creep. An agent that starts as an inbox triage tool and gradually gets expanded to draft substantive legal analysis without a corresponding increase in review rigor is a liability. Every time you expand an agent's scope, revisit the review protocol.

Tool drift. AI vendors update their models, their terms of service, and their data handling practices. A tool that had zero data retention six months ago may have changed its policy. Review your approved tool list quarterly, and re-confirm data handling terms annually.

The competence floor. ABA Formal Opinion 512 says you must understand the tools you use. If you can't explain to a client – in plain, non-technical language – what your AI agents do and how you supervise them, you haven't met the competence standard. Understanding doesn't mean you need to explain transformer architectures. It means you can say: "I use an automated system that checks my calendar every morning, pulls the case file for each meeting, and sends me a briefing. I review the briefing before every meeting. The system doesn't contact anyone on my behalf."

Staff training. Your policy is only as strong as the weakest user. If a paralegal enters client data into a free-tier AI tool because nobody told them not to, your written policy doesn't help. Training is not optional. Every person with access to client data needs to understand the approved tool list, the data classification tiers, and the verification requirements.

The Honest Assessment

This book advocates for AI adoption in law practice. This chapter balances that advocacy with candor about the risks.

AI will hallucinate. Not sometimes. Regularly. The rate varies by model, by task, and by how the prompt is structured – but no current AI model can guarantee factual accuracy. Every AI-generated output is a draft that requires verification. This is not a flaw that will be fixed in the next model version. It's an inherent characteristic of how large language models work.

Bar rules will tighten. The trend is clear: more states will adopt AI-specific ethics provisions. Disclosure requirements will expand. Competence standards will become more specific. Firms that build their AI practices around strong safety protocols now will be ahead of the compliance curve. Firms that don't will be retrofitting under pressure.

Malpractice carriers are watching. Insurers are beginning to ask about AI use in applications and renewals. Firms with documented AI policies, verification workflows, and audit trails will be in a stronger position than firms that can't describe their safeguards. An AI policy isn't just an ethics document. It's a malpractice insurance document.

The strongest argument for AI is also an argument for caution. AI catches things humans miss: unbilled time, approaching deadlines, stale matters, client contacts that have gone too long without a follow-up. That same capability means AI is operating on your client data, inside your case management system, with access to privileged information. The power that makes it useful is the same power that makes it risky. The safety layer exists because both things are true simultaneously.

The Verification Habit

The single highest-value behavior change this chapter recommends is not a policy, a technology, or an architecture. It's a habit.

Verify everything.

Treat every AI output the way you'd treat a brief from a brilliant but unreliable first-year associate. Assume the research is thorough and the reasoning is sound – but check every citation, confirm every date, validate every case fact against the file. The associate is fast, hardworking, and available at 2:00 AM. But they make mistakes. So does AI.

The attorneys in the 600+ hallucination cases didn't get sanctioned because they used AI. They got sanctioned because they broke a habit that predates AI by centuries: verifying your sources before you stake your reputation on them.

Build the habit. Apply it to every AI output. Apply it to your agents' work product every morning. Apply it to every draft, every deadline computation, every case analysis.

The safety layer – the architecture, the audit trail, the review queue, the access controls – exists to make verification efficient. It puts the right outputs in front of you at the right time, with the right context, so you can apply your judgment quickly. But it doesn't substitute for your judgment. Nothing does.

Quick Reference: The Safety Layer

Principle	What It Means	What It Prevents
AI Drafts, Humans Decide	No agent takes external action. All outputs go to human review.	Fabricated citations filed with courts, privileged data sent to opposing counsel, incorrect communications to clients

Principle	What It Means	What It Prevents
Full Audit Trail	Every agent run is logged with timestamp, inputs, outputs, and decisions	Unverifiable claims about system behavior; inability to reconstruct what happened when something goes wrong
Review Queue	Ambiguous items are flagged, not guessed at	Misclassified urgent items getting archived; edge cases falling through the cracks
Access Controls	Each agent has defined read/write permissions; write is limited to drafts and proposals	Cascading errors; agents taking actions beyond their intended scope

Ethics Obligation	How You Meet It
Competence (Rule 1.1)	Understand your tools; write or approve skill files; catch agent errors
Communication (Rule 1.4)	Disclose AI use to clients in engagement letter; explain agent roles
Supervision (Rules 5.1, 5.3)	Review every output; maintain audit trail; human gate on all external actions
Confidentiality (Rule 1.6)	Approved tool list; data classification tiers; zero-retention confirmation
Candor (Rule 3.3)	Verify every citation; review every filing; never submit unverified AI output

Your immediate next step: If your firm doesn't have an AI policy, turn to Appendix E. Customize the template. Adopt it. The 43% of firms without one are carrying risk that a written policy eliminates.

Next: Chapter 11 lays out the implementation path – crawl, walk, run. You've seen the architecture, the workflows, and the safety layer. Now we'll build the plan that takes your firm from "I use AI sometimes" to "my practice runs on AI," starting with week one.

Chapter 11: Getting Started – Crawl, Walk, Run

What You'll Walk Away With:

- A three-phase roadmap that meets you where you are – whether you've never typed a prompt or you're ready to wire up a full agent system
- Specific chapter references for every phase, so you know exactly where to go next
- A "first skill" template you can copy, fill in, and have running this week
- The Dodonai engagement model – self-serve, done-with-you, and managed paths
- A vision of what your practice looks like six months from now if you start today

---

7:14 AM on a Thursday

You're at your desk, but something is different.

Your phone buzzed once at 7:00 AM – a Slack notification with your meeting intel briefing. Today's 10:30 mediation with opposing counsel on the Hernandez case: their firm's recent settlements in similar PI matters, the adjuster's authority range based on comparable claims, and three points from your last email exchange you'll want to reference. You didn't ask for it. It ran while you were pouring coffee.

At 7:15, your inbox triage digest arrives. Fourteen new emails overnight, already classified. Two flagged urgent: opposing counsel on the Rivera matter filed a motion to compel (deadline calendared automatically), and a client on the Davis file replied with medical records you've been waiting for. Four action items queued for your review this morning. Eight newsletters and vendor emails archived. Zero unread threads you need to manually sort.

By 7:30, your case health dashboard is open. Forty active matters. No deadlines within 30 days that don't have a next action assigned. The Kowalski demand package you've been meaning to finalize? Your AI flagged it three days ago. Maria already has the final treatment records uploaded. The draft demand narrative has been waiting since Tuesday.

You billed 6.8 hours yesterday and worked 7.5. Not because you're working harder. Because the other 3-4 hours that used to go to administrative overhead –

the email triage, the deadline checking, the client update drafting, the billing reconciliation – happen before you sit down.

It's 7:32 AM. You're at your desk. But this time, you're not drowning. You're deciding what matters most.

This is what the book has been building toward. Not a theoretical future. A Thursday morning that's different from last Thursday because you built one system, then another, then a third, and now they compound.

The question is how to get here. The answer is smaller than you think.

The Three Phases

Every firm that goes AI-native follows the same progression. Not because someone designed it that way, but because the learning curve naturally breaks into three stages. Each one builds on the last, and each one delivers value before you move to the next.

Phase	What It Looks Like	Time Investment	Weekly Time Saved	Chapter Reference
Crawl	AI Desktop for one-off tasks. Single prompts, copy-paste workflows.	15 minutes to start	4-6 hours/week	Ch 1, Ch 5
Walk	First skill file, first scheduled agent. One system running without you.	2-3 hours to build	8-12 hours/week	Ch 3, Ch 7
Run	Connected systems, compound workflows, agents feeding agents.	Ongoing (1 hr/week maintenance)	15-21 hours/week	Ch 2, Ch 6, Ch 8, Ch 9

These aren't arbitrary labels. They map to the three structural limitations we identified in Chapter 1: the chatbot forgets (Crawl fixes that with better prompts), it's reactive (Walk fixes that with schedules), and it's isolated (Run fixes that with connections). Each phase removes one limitation.

Phase 1: Crawl

Where most firms are right now. You open Claude or ChatGPT, type a question, get a response. Maybe you use it a few times a week, maybe a few times a day. This is the 1% from Chapter 1 – but it's a real 1%, and it's worth getting right before moving on.

The Crawl phase has one objective: **prove to yourself that AI produces practice-ready output when given sufficient context.** That's it. No installation, no configuration, no new software. Just better prompts.

What to Do This Week

Pick one task from this list – the one that's been sitting on your desk longest:

- A client status update letter (Ch 5 has the demand letter workflow that extends to any client communication)
- A deposition summary analysis – upload the transcript and ask for a cross-examination outline (Ch 5, Workflow 2)
- An intake evaluation for a lead that came in this week (Ch 5, Workflow 4, has the go/no-go evaluation prompt)
- A medical chronology review – have AI flag gaps in treatment, missing providers, or inconsistencies (Ch 4 covers how Dodonai structures the records first)
- A case theory memo synthesizing your eDiscovery results (Ch 5, Workflow 3)

Time yourself. Write down how long the task takes without AI. Then do it with AI, using the context-rich prompt approach from Chapter 1's "Try This Now" exercise. Compare.

Most attorneys report 40-60% time savings on their first structured attempt. The Wolters Kluwer 2026 Future Ready Lawyer survey found 62% of respondents reporting 6-20% weekly time savings. Those numbers come from firms still in Crawl phase – single prompts, no automation. At 4-6 hours saved per week and a blended rate of $350/hour, that's $72,800-$109,200 in recovered capacity per year (Rankings.io) before you've installed anything.

The Crawl Mindset

Three rules for this phase:

1. Use real case data, not test scenarios. The quality gap between a generic prompt and a context-rich prompt is the difference between a form letter and a document that sounds like you wrote it. Chapter 1 demonstrated this with the Sarah Chen status update. Chapter 5 extended it across five complete workflows. Use your actual cases.

2. Iterate once. The first output is a rough draft. Give one round of feedback: "Too formal in the opening," "Add the treatment compliance angle," "Cut the legal process explanation." The second output is usually 85-90% ready to send. This draft-then-refine cycle is how you'll work with AI permanently.

3. Notice where you're repeating yourself. Every time you give AI the same context – your practice areas, your communication style, your firm details, your case evaluation criteria – you're doing work that a skill file would eliminate. That repetition is the signal that you're ready for Walk.

What Crawl Won't Do

Crawl is limited by the three constraints from Chapter 1. The AI still forgets between sessions. It still waits for you to open it. It still can't see your Clio data or your inbox. You're the integration layer, manually providing context every time.

That's fine. Crawl's purpose is proof of concept – showing you what AI produces when you give it the right instructions. Once you see that consistently, the motivation to systematize it is natural.

Phase 2: Walk

The shift from tool to system. Walk is where most of this book lives, and it's where the largest jump in value happens. You go from using AI when you remember to having AI work while you sleep.

Walk requires one thing Crawl didn't: a model harness. In our examples, we use Claude Code, but any environment that supports skill files (persistent instructions), file access, and scheduled execution works. The concept is universal; the implementation details differ by platform.

Your First Agent: The Deadline Monitor

Chapter 7 walked through building this step by step – the skill file, the Clio connection, the shell wrapper, the launchd schedule. Here's why we recommend starting here:

- **It prevents malpractice.** Calendar errors are the leading cause of legal malpractice claims (ABA malpractice profiles). An agent that scans every active matter for approaching deadlines every morning is the highest-ROI system you can build.
- **It's simple.** One connection (Clio), one schedule (daily at 7:00 AM), one output (a sorted list of upcoming deadlines with next actions).
- **You'll trust it fast.** Within a week, you'll check the output against what you already know and confirm it's catching everything. That trust is what makes you willing to build the next agent.

If deadlines aren't your pain point, Chapter 7 offers six other starting agents: inbox triage, client update drafter, billing reconciler, intake pipeline, marketing engine,

and case health dashboard. Pick the one that makes you groan every week. That's your first skill.

Writing Your First Skill File

Chapter 3 introduced the three-layer architecture: skills, connections, schedules. The skill file is where it starts. Here's a template you can adapt for any recurring task:

```
# Agent: [NAME]
# Schedule: [WHEN IT RUNS]
# Purpose: [ONE SENTENCE  --  what this agent does and why]

## Context
You are an AI assistant for [FIRM NAME], a [PRACTICE AREA] firm
in [CITY]. The attorney is [NAME]. The team includes [STAFF].

The firm uses [CASE MANAGEMENT SYSTEM] for matter management,
[EMAIL] for correspondence, and [BILLING SYSTEM] for accounting.

## Task
Every time this agent runs, perform these steps in order:

1. [FIRST STEP  --  e.g., "Query Clio for all active matters"]
2. [SECOND STEP  --  e.g., "For each matter, check the next
   deadline within 90 days"]
3. [THIRD STEP  --  e.g., "Sort by urgency: within 30 days =
   URGENT, 31-60 = ACTION, 61-90 = MONITOR"]
4. [FOURTH STEP  --  e.g., "For any matter with no next action
   assigned, flag it as STALE"]

## Output Format
[DESCRIBE EXACTLY WHAT THE OUTPUT SHOULD LOOK LIKE]

- Group by urgency tier
- Include: matter name, client name, deadline date, days
  remaining, next action, assigned to
- If no deadlines within 90 days, state "All clear" with
  the total active matter count

## Rules
- Never send anything externally. Draft only.
- If uncertain about a deadline date, flag it for human
  review rather than guessing.
- [ANY FIRM-SPECIFIC RULES]
```

This is the same structure demonstrated in Chapter 3's Morning Case Briefing skill file – context, task, output format, rules. The analogy holds: if you can write a procedure manual for a paralegal, you can write a skill file.

Putting It on a Schedule

A skill file sitting in a folder is just a document. A skill file on a schedule is an agent.

Chapter 7 walked through the mechanics: a shell wrapper script that loads your credentials, a launchd plist (on Mac) or cron job (on Linux) that triggers it, and a log file for troubleshooting. The technical steps are straightforward. The conceptual shift is the important part.

Once you schedule your first agent, you have crossed the line from using AI to running AI. It works while you sleep. It checks things you'd forget. It's consistent in ways you can't be at 9:00 PM on a Tuesday after a full day of depositions.

You now have one agent. It runs without you.

The Walk Checklist

Before moving to Run, confirm you've done these:

- ☐ Installed a model harness (Claude Code or equivalent)
- ☐ Written one skill file for your most painful recurring task
- ☐ Connected it to at least one data source (Clio, Gmail, or your calendar)
- ☐ Put it on a daily or weekly schedule
- ☐ Verified its output against your own knowledge for at least one full week
- ☐ Identified the second agent you want to build

That last item matters. Once the first agent is running reliably, the second one takes half the time to build. The pattern is always the same: encode the instructions, connect the tools, set the trigger.

Phase 3: Run

Where the compound effect takes over. Run is the vision from Chapter 2 – the practice where AI agents feed each other's outputs, where the morning briefing is waiting before you arrive, where your effective hourly rate climbs because the administrative tax drops to near zero.

Run doesn't happen in a weekend. It's a progression: 1 agent to 3 to 7. Each new agent makes the existing ones more valuable because they share data.

The Compound Architecture

Chapter 2 demonstrated this with Sarah Chen's morning: the meeting intel agent uses data from the case health agent, which uses data from the inbox triage agent, which flags items that feed back into the case health agent. Chapter 7 called this the "compound effect" and walked through the Thursday scenario where four agents surface interconnected information about the Davis v. HCA case.

Here's how the compound effect builds in practice:

Agents Running	What Compounds	Example
1 (Deadline Monitor)	Nothing yet – standalone value	You catch a SOL approaching in 45 days that wasn't on your calendar
3 (+ Inbox Triage, Client Updates)	Inbox triage flags a client email about a case the deadline monitor is tracking. Client update drafter uses the deadline data to include timeline info.	Monday morning: "Mrs. Hernandez replied about her treatment records. Her demand deadline is in 60 days. Draft update includes the timeline."
5 (+ Billing Reconciler, Intake Pipeline)	Billing catches 3 unbilled hours on a matter the case health agent flagged as stale. Intake scoring uses outcome data from closed cases to calibrate case value estimates.	"The Kowalski file has 4.2 unbilled hours from last Thursday. Also: new lead scored 84/100 – similar profile to your last three successful med mal cases."
7 (+ Marketing Engine, Case Health)	Content engine writes blog posts from practice expertise (Ch 8). Reviews and referral agents drive intake. The full loop closes.	Your blog post on rideshare accident claims (written by AI, edited by you) generates 3 leads this month. The intake agent scored all three. One became a client.

Chapter 6 showed this architecture working across seven practice areas – PI, real estate, tax, cybersecurity, immigration, estates, family law. The pattern is identical every time: documents in, deadlines computed, things filed, people notified. Only the skill files change.

The Connections Map

Run requires connecting AI to your firm's systems. Chapter 3 covered this as "Layer 2: The Connections" and established the critical distinction between read and write permissions. Chapter 10 expanded this into the safety layer with specific access controls.

The essential connections for a litigation practice:

System	What AI Reads	What AI Writes	Chapter Reference
Clio	Matters, contacts, deadlines, activity history	Draft tasks, draft notes (never filed directly)	Ch 3, Ch 7
Gmail	Inbox, sent mail, threads	Draft emails only (never auto-sends)	Ch 3, Ch 8
Calendar	Events, attendees, scheduling conflicts	Draft events only	Ch 2, Ch 7
QuickBooks	Invoices, payments, time entries	Nothing (read-only)	Ch 9

System	What AI Reads	What AI Writes	Chapter Reference
Dodonai	Processed documents: chronologies, deposition summaries, eDiscovery results	Upload new documents for processing	Ch 4, Ch 5

The safety principle from Chapter 10 applies everywhere: **AI drafts, humans decide.** No agent sends an email, files a document, or creates a calendar event without your explicit approval. This isn't a limitation – it's the design. Chapter 10's embarrassment prevention layer (the nuclear power plant analogy) ensures that every external action routes through your review.

What Run Produces

Chapter 9 laid out the operational math. With seven agents running:

Category	Weekly Hours Recovered	Annual Value (at $350/hr)
Deadline monitoring + case health	3-4 hours	$54,600-$72,800
Inbox triage + client updates	4-5 hours	$72,800-$91,000
Billing reconciliation	2-3 hours	$36,400-$54,600
Intake pipeline	2-3 hours	$36,400-$54,600
Marketing + content	2-3 hours	$36,400-$54,600
SOP maintenance + training	1-2 hours	$18,200-$36,400
Total	**14-20 hours/week**	**$254,800-$364,000**

These aren't projections. They're the same time savings estimates from Chapter 7's analysis, derived from the time audit in Chapter 1 that showed most solo and small-firm attorneys spending 15-25 hours per week on non-legal administrative work.

The honest caveat from Chapter 9 applies: recovered capacity is not recovered revenue unless you fill that time with billable work or business development. But the capacity is real. What you do with it – take more cases, invest in trial preparation, leave the office at 5:30, build the practice you actually wanted when you went to law school – is your decision.

Where Dodonai Fits

This book has presented three ways to work with Dodonai, reflecting three levels of engagement. Each maps to a phase.

Engagement	What You Get	Best For	Starting Point
Dodonai App	Self-serve document processing: Medical Record Summaries & Chronologies, Deposition Summaries, Transcript Management, OCR, eDiscovery, XD Agents. Upload documents, get structured output.	Firms in Crawl phase who want immediate ROI from document processing. Also the foundation layer for Walk and Run – your agents need structured data to work with.	Sign up at app.dodon.ai. Lite plan ($30/month) for small caseloads, Pro ($100/month) for 40+ active matters. Ch 4 has the full product walkthrough.
Dodonai Consulting	Done-with-you agent design. We help you map your workflows, write your skill files, build your connections, deploy your schedules. You own the system. We help you build it.	Firms ready for Walk-to-Run who want to move faster than DIY allows. Attorneys who know what they want automated but don't want to spend weekends configuring it.	Start a conversation about your practice. We scope the engagement based on your case volume, your systems, and your pain points. Ch 7's agent playbook is the menu we work from.
Managed Automation	Ongoing support. We maintain your agents, update skill files as your practice evolves, add new integrations, troubleshoot issues. Your AI system, professionally maintained.	Firms in Run phase who want to focus on practicing law, not managing infrastructure. Multi-attorney firms where the agent system serves several practitioners.	Follows a Consulting engagement. Monthly retainer based on system complexity.

The Dodonai Approach – the methodology this entire book teaches – works regardless of whether you engage Dodonai or build everything yourself. Chapter 3 laid out the architecture. Chapters 5-9 provided the workflows and prompts. Chapter 10 covered the safety layer. Everything you need is in these pages.

Dodonai Consulting accelerates the timeline. Instead of spending 20 hours over two months figuring out Clio API connections and launchd scheduling, you spend that time practicing law while someone who's built these systems before handles the configuration. That's the value proposition – not a proprietary product you're locked into, but implementation expertise applied to your practice.

The Chapter Roadmap

Every chapter in this book was designed for a specific moment in your journey. Here's when to read (or re-read) each one:

When You're Ready To...	Go To	What You'll Find
Understand why single prompts aren't enough	**Ch 1: The 1% Problem**	The 1%/99% framing, three real firm stories, the time audit that maps your administrative overhead
See what an AI-native practice actually produces	**Ch 2: What AI-Native Looks Like**	A full morning walkthrough with sample outputs – meeting briefing, case health digest, inbox triage summary
Understand the three-layer architecture	**Ch 3: The Architecture**	Skills, connections, schedules explained with annotated skill files and the safety layer
Process your documents into structured data	**Ch 4: Dodonai**	All six products: Medical Records, Depositions, Transcripts, OCR, eDiscovery, XD Agents with pricing and benchmarks
Turn documents into case strategy	**Ch 5: From Documents to Strategy**	Five complete workflows: medical chronology to demand letter, deposition to cross-exam outline, eDiscovery to case theory, intake to go/no-go, multi-deposition to trial prep matrix
See the architecture in your practice area	**Ch 6: Seven Practice Areas**	PI, real estate, tax, cybersecurity, immigration, estates, family law – each with pain points, workflows, and practice-area-specific prompts
Build your first agent	**Ch 7: The Agent Playbook**	Seven agents, step-by-step. Start with the Deadline Monitor. Complete skill files, schedules, and sample outputs
Automate marketing and intake	**Ch 8: Marketing, Intake, and Growth**	Content engine, review management, referral nurturing, scored intake pipeline
Systematize billing and operations	**Ch 9: Operations**	Time entry reconstruction, SOP creation, template libraries, onboarding automation
Ensure your system is ethical and defensible	**Ch 10: The Safety Layer**	Hallucination patterns, bar ethics rules (ABA Op. 512), AI policy framework, the verification habit

This isn't a book you read once and shelve. It's a reference you return to as your system grows. Week one, you're in Chapters 1 and 5. Month two, you're deep in Chapters 3 and 7. Six months from now, you're pulling prompts from Chapters 6 and 8. The structure was built for that.

Your First Week

If you've read this far but haven't started yet, here's the plan. Five days, one task per day, 30-60 minutes each.

```
FIRST WEEK PLAN  --  FROM READER TO PRACTITIONER

Day 1 (Monday): The Proof
- Pick one real task from your current caseload
- Use the prompt template from Ch 1 ("Try This Now" section) with
your actual
  case data
- Time yourself: how long with AI vs. your estimate without
- Goal: see that AI produces 80-90% usable output with the
  right context

Day 2 (Tuesday): The Audit
- Complete the time audit from Ch 1
- List every recurring non-legal task you do each week
- Mark the top 3 by hours spent
- Goal: know where your 15-25 administrative hours go

Day 3 (Wednesday): The Document Test
- Upload one document to Dodonai (medical records, a deposition
  transcript, or a scanned file for OCR)
- Review the structured output
- Goal: see the difference between raw documents and structured
  data your AI can actually work with (Ch 4)

Day 4 (Thursday): The Workflow
- Take the document output from Day 3 and run it through one
  of Ch 5's workflows
- Medical chronology → demand narrative, or deposition summary
  → cross-exam outline
- Goal: experience the Process → Analyze → Act pipeline with
  your own case data

Day 5 (Friday): The Decision
- Review your time audit (Day 2) and your AI results (Days 1,
  3, 4)
- Decide: which task becomes your first skill file?
- Write the skill file using the template from this chapter
- Goal: have a written skill file ready to schedule
```

That's the Crawl-to-Walk transition in five days. Not five months. Not after a conference or a CLE or a committee report. This week.

The Firms That Wait

There's a version of this chapter that ends with gentle encouragement and a reassurance that there's no rush. This isn't that chapter.

The numbers tell a specific story. 79% of lawyers use AI daily (ABA 2026 TechReport). Solo tech spending is growing 56% annually – twice the industry average (Clio Legal Trends). 64% of in-house legal teams are building AI capabilities internally and expect the efficiency gains passed through to outside counsel (ACC/Everlaw GenAI Survey). The competitive landscape is shifting, and the shift rewards early movers.

The firms that wait aren't standing still. They're falling behind relative to firms that have moved from Crawl to Walk. Zack Shapiro serves 200 clients with two people. Legal Paradox achieved 32x time compression. These aren't edge cases anymore – they're the new pace of the profession. The ABA and state bars are building ethics frameworks (Chapter 10) not to slow adoption, but because adoption is already widespread and accelerating.

You don't need to be Legal Paradox. You don't need 500 AI agents or a complete practice rebuild. You need one skill file, running on a schedule, producing output you trust. That's the minimum viable version of AI-native, and it's enough to change how your Tuesdays end.

5:47 PM on a Tuesday

Let's go back to where this book started.

It's Tuesday. Not 7:32 PM – 5:47 PM. You're at your desk, but you're not staying late. You're wrapping up.

The Hernandez demand package went out at 2:00 PM. Your AI drafted the narrative from Dodonai's medical chronology on Monday (Chapter 5's demand letter workflow). You spent 40 minutes editing it – refining the causation language, adjusting the settlement range, adding the paragraph about treatment compliance that you know this particular adjuster responds to. That's judgment. That's what you went to law school for.

The two discovery responses that used to eat your post-5:00 PM hours? Your client update agent drafted preliminary responses Tuesday morning, cross-referencing each request against your Clio matter notes and document index. You reviewed and refined them between meetings. Done by 3:00 PM.

The 23 unread emails? Triaged at 8:00 AM. The motion to compel on the Rivera matter was flagged urgent, the response deadline was calendared, and a draft motion in opposition outline was waiting in your review queue by the time you'd finished your first meeting. The client calls? Your AI drafted callback prep notes for each one – who they are, their case status, the last three communications, the

likely reason for the call. You returned all four calls between 1:00 and 2:00 PM, each one prepared and productive.

The 14 matters with activity due this week? Your case health dashboard showed them sorted by urgency at 7:30 AM. Eight had next actions already assigned. Four needed your attention – you handled them by noon. Two were status quo, no action needed.

You billed 6.4 hours today and worked 7.5. The gap between worked and billed dropped from 42% to 15%, not because you started tracking harder, but because the billing reconciler (Chapter 9) catches the time your manual entries miss.

Maria is still your paralegal. David still works part-time. You didn't hire anyone. You didn't buy an enterprise software suite. You wrote skill files in plain English, connected them to the tools you already use, and set them on schedules.

Your practice is the same practice. Your clients are the same clients. Your judgment is the same judgment.

The difference is where your time goes.

Start

This book has covered a lot of ground. Ten chapters. Five document-to-strategy workflows. Seven agent designs across seven practice areas. Four operational systems. An ethics framework. Dozens of prompts.

But every firm that has gone AI-native – the ones in Chapter 1, the Inner Circle attorneys in Chapter 6, the practices that inspired every workflow in this book – started the same way.

One task. The one they kept putting off.

They wrote the instructions for how a human would do it. They were specific – the way lawyers are trained to be specific, with the facts that matter and the format they expect. They gave those instructions to an AI and saw that the output was 80% there on the first try and 90% there after one round of feedback.

Then they did it again the next week. And the next. And eventually they thought: why am I giving these instructions every time? Why doesn't this just run?

That's the moment the skill file gets written. That's the moment the schedule gets set. That's the moment the practice starts to change.

You don't need to believe that AI will reshape the legal profession. You don't need a technology strategy or a committee or a pilot program. You need 15 minutes, one real task from your desk, and the prompt template from Chapter 1.

Start there. Time yourself. See what happens.

Everything else builds from there.

Quick Reference Card: Chapter 11

Phase	What You Do	What Changes	Time to Value
Crawl	AI Desktop for one-off tasks. Context-rich prompts with real case data.	4-6 hrs/week recovered. Proof that AI produces practice-ready output.	15 minutes
Walk	First skill file + first scheduled agent. One system running autonomously.	8-12 hrs/week recovered. You stop being the reminder system.	2-3 hours
Run	Connected systems, compound workflows, agents feeding agents.	15-21 hrs/week recovered. AI handles the operational layer of your practice.	Ongoing build

Dodonai Engagement	What It Is	When to Start
App	Self-serve document processing (6 products, from $30/month)	Crawl phase – get structured data from your documents
Consulting	Done-with-you agent design and deployment	Walk-to-Run transition – accelerate the build
Managed Automation	Ongoing system maintenance and evolution	Run phase – focus on law, not infrastructure

Chapter	Best For	Key Deliverable
Ch 1	Starting the journey	Time audit + first prompt
Ch 2	Seeing the destination	Sample outputs from an AI-native morning
Ch 3	Understanding the architecture	Skills + connections + schedules
Ch 4	Processing documents	Dodonai product guide + pricing
Ch 5	Building case strategy	5 document-to-strategy workflows
Ch 6	Your practice area	7 practice areas, same architecture
Ch 7	Building agents	7 agents, step-by-step
Ch 8	Growing the practice	Marketing, intake, reviews, referrals
Ch 9	Running operations	Billing, SOPs, templates, onboarding
Ch 10	Staying safe	Ethics, hallucinations, AI policy
Ch 11	Getting started	You are here

Appendix A: Prompt Library

This appendix collects every prompt from *Going (AI) Native with Your Law Firm* into a single, scannable reference. Each prompt is copy-paste ready – fill in the bracketed placeholders with your case data and paste into any capable AI assistant (Claude, GPT, Gemini, or equivalent). Prompts are grouped by category and listed in chapter order within each category. Use the master index below to find what you need quickly.

Master Index

#	Prompt Title	Category	Chapter	Section
1	Practice Context Template	General	Ch 1	Try This Now: Your First Prompt
2	Practice Context Template (Filled Example)	General	Ch 1	Try This Now: Your First Prompt
3	Iterative Feedback Example	General	Ch 1	Try This Now: Your First Prompt
4	Morning Case Briefing Skill	Case Strategy	Ch 3	Layer 1: The Skill
5	Friday Case Review Skill (Minimal)	Case Strategy	Ch 3	Writing Your First Skill
6	Skill File Template	General	Ch 3	Try This Now: Write a Skill File
7	Damages Narrative from Medical Chronology	Medical Records	Ch 5	Workflow 1
8	Medical Specials Calculation	Medical Records	Ch 5	Workflow 1
9	Cross-Examination Outline from Deposition	Depositions	Ch 5	Workflow 2
10	Case Theory from eDiscovery Results	Case Strategy	Ch 5	Workflow 3
11	Intake Go/No-Go Evaluation	Intake	Ch 5	Workflow 4
12	Trial Prep Matrix from Multiple Depositions	Depositions	Ch 5	Workflow 5
13	Document-to-Strategy Skill File Template	General	Ch 5	The Skills Approach
14	PI Medical Records Intake Agent	Medical Records	Ch 6	Personal Injury

#	Prompt Title	Category	Chapter	Section
15	PI Demand Preparation Agent	Medical Records	Ch 6	Personal Injury
16	Real Estate Closing Conditions Tracker	Operations	Ch 6	Real Estate
17	Tax Correspondence Monitoring Agent	Operations	Ch 6	Tax Disputes
18	Breach Notification Compliance Agent	Operations	Ch 6	Cybersecurity & Privacy
19	Immigration Case Monitoring Agent	Operations	Ch 6	Immigration
20	Estate Administration Lifecycle Agent	Operations	Ch 6	Estates & Probate
21	Collaborative Law Session Follow-Up Agent	Operations	Ch 6	Family & Collaborative Law
22	Deadline Monitor Skill File	Case Strategy	Ch 7	Agent 1: Deadline Monitor
23	Inbox Triage Classification Rules	Operations	Ch 7	Agent 2: Inbox Triage
24	Client Update Drafter Skill	Client Communication	Ch 7	Agent 3: Client Update Drafter
25	Intake Lead Scoring Rules	Intake	Ch 7	Agent 5: Intake Pipeline
26	Blog Post Generation	Marketing	Ch 8	The Content Engine
27	Review Response Drafter	Marketing	Ch 8	Review Management
28	Referral Follow-Up Email	Marketing	Ch 8	Referral Nurturing
29	Intake Evaluation and Triage	Intake	Ch 8	The Intake Pipeline
30	Billing Reconciler	Billing/Operations	Ch 9	Time Entry Reconstruction
31	SOP Generator	Billing/Operations	Ch 9	SOP Creation
32	Engagement Letter Generator	Billing/Operations	Ch 9	Template Libraries
33	Onboarding Plan Generator	Billing/Operations	Ch 9	Training and Onboarding
34	First Agent Skill File Template	General	Ch 11	Phase 2: Walk
35	First Week Implementation Plan	General	Ch 11	Your First Week

General

1. Practice Context Template

Source: Chapter 1, “Try This Now: Your First Prompt” **Category:** General

Use this as the foundation for any one-off AI interaction. It sets your practice context so the AI produces firm-specific output instead of generic filler.

```
You are a litigation practice assistant for a [practice area] firm.
My firm is [name], located in [city]. I handle approximately [X]
active cases across [practice areas with rough percentages].

I use [case management system] for matter management, [email
provider] for correspondence, and [billing system] for time
tracking. I need help with the administrative and analytical
work that sits around my active caseload.

My team is [brief description]. My communication style with
clients is [professional/direct/warm -- pick one].

Before responding, confirm your understanding and ask one
clarifying question if anything is ambiguous. Then provide
the output.

Today I need to: [specific task]
```

---

2. Practice Context Template (Filled Example)

Source: Chapter 1, “Try This Now: Your First Prompt” **Category:** General

A filled-in version of the practice context template showing the level of detail that produces practice-ready output.

```
You are a litigation practice assistant for a personal injury firm.
My firm is Chen Law Group, located in Austin, TX. I handle
approximately 40 active cases: 70% personal injury (auto accidents,
premises liability, slip and fall), 15% medical malpractice,
15% workers' compensation.

I use Clio for matter management, Gmail for correspondence, and
QuickBooks for accounting. I use Dodonai Pro for medical record
summaries and deposition digests. I need help with the
administrative and analytical work that sits around my active
caseload.

My team is myself (attorney), one full-time paralegal, and one
part-time legal assistant. My communication style with clients
is direct and empathetic -- I explain things clearly without
legal jargon unless necessary.

Before responding, confirm your understanding and ask one
clarifying question if anything is ambiguous. Then provide
the output.
```

```
Today I need to: draft a 60-day status update letter to Maria
Hernandez regarding her PI case. She was rear-ended on I-35
on November 12, 2025. She's completed 4 months of chiropractic
treatment and 2 months of physical therapy. We're waiting on
the final treatment records from Austin Spine & Rehab and her
MRI results from St. David's imaging. The SOL is November 12,
2027. I want to reassure her that the case is progressing,
explain why we're waiting on records before making a demand,
and let her know the next milestone. Sign as "Sarah Chen."
```

---

3. Iterative Feedback Example

Source: Chapter 1, “Try This Now: Your First Prompt” **Category:** General

Demonstrates the draft-then-refine cycle. After receiving the AI’s first output, use targeted feedback like this to get the output 90% ready to send.

```
Good start. A few adjustments:
- The opening paragraph is too formal -- I want Maria to feel
  like I'm talking to her, not writing a brief
- Add a sentence about how the chiropractic records show
  consistent treatment compliance, which strengthens the case
- Cut the paragraph about the legal process -- she doesn't
  need that level of detail
- End with a specific next step and timeline, not a vague
  "we'll be in touch"
```

---

6. Skill File Template (Blank)

Source: Chapter 3, “Try This Now: Write a Skill File” **Category:** General

A minimal template for writing your first skill file. Fill in the brackets and you have a working agent instruction set.

```
# /[name] -- [What it does]

## Context
[One sentence: why this task matters and when it runs]

## Steps

### 1. [First action]
[Where to look, what to pull]

### 2. [Second action]
[What to check, what to filter]

### 3. [Third action]
[How to compile the output]

### 4. Save and notify
[Where to save the result, who to alert]
```

---

13. Document-to-Strategy Skill File Template

Source: Chapter 5, “The Skills Approach” **Category:** General

A template for encoding any document-to-strategy workflow as a reusable skill file. Save your analytical methodology once and apply it consistently.

```
# /[your-skill-name] -- [One-line description]

## Context
[Why this workflow exists. What problem it solves. When to run it.]

## Inputs Required
[What documents or data the AI needs to do this analysis.]

## Instructions
[Your full prompt -- the analysis methodology, output format,
specific criteria, and verification requirements.]

## Output Format
[Exactly what the AI should produce -- sections, tables, structure.]

## Verification Checklist
[What the attorney must check before using the output.]
```

---

34. First Agent Skill File Template

Source: Chapter 11, “Phase 2: Walk” **Category:** General

A complete template for your first scheduled agent. Includes context, task steps, output format, and safety rules.

```
# Agent: [NAME]
# Schedule: [WHEN IT RUNS]
# Purpose: [ONE SENTENCE -- what this agent does and why]

## Context
You are an AI assistant for [FIRM NAME], a [PRACTICE AREA] firm
in [CITY]. The attorney is [NAME]. The team includes [STAFF].

The firm uses [CASE MANAGEMENT SYSTEM] for matter management,
[EMAIL] for correspondence, and [BILLING SYSTEM] for accounting.

## Task
Every time this agent runs, perform these steps in order:

1. [FIRST STEP -- e.g., "Query Clio for all active matters"]
2. [SECOND STEP -- e.g., "For each matter, check the next
   deadline within 90 days"]
3. [THIRD STEP -- e.g., "Sort by urgency: within 30 days =
   URGENT, 31-60 = ACTION, 61-90 = MONITOR"]
4. [FOURTH STEP -- e.g., "For any matter with no next action
   assigned, flag it as STALE"]
```

```
## Output Format
[DESCRIBE EXACTLY WHAT THE OUTPUT SHOULD LOOK LIKE]

- Group by urgency tier
- Include: matter name, client name, deadline date, days
  remaining, next action, assigned to
- If no deadlines within 90 days, state "All clear" with
  the total active matter count

## Rules
- Never send anything externally. Draft only.
- If uncertain about a deadline date, flag it for human
  review rather than guessing.
- [ANY FIRM-SPECIFIC RULES]
```

---

35. First Week Implementation Plan

Source: Chapter 11, “Your First Week” **Category:** General

A structured five-day plan for transitioning from reader to practitioner. Follow this sequence to go from your first prompt to a written skill file in one week.

```
FIRST WEEK PLAN -- FROM READER TO PRACTITIONER

Day 1 (Monday): The Proof
- Pick one real task from your current caseload
- Use the prompt template from Ch 1 ("Try This Now" section) with
your actual
  case data
- Time yourself: how long with AI vs. your estimate without
- Goal: see that AI produces 80-90% usable output with the
  right context

Day 2 (Tuesday): The Audit
- Complete the time audit from Ch 1
- List every recurring non-legal task you do each week
- Mark the top 3 by hours spent
- Goal: know where your 15-25 administrative hours go

Day 3 (Wednesday): The Document Test
- Upload one document to Dodonai (medical records, a deposition
  transcript, or a scanned file for OCR)
- Review the structured output
- Goal: see the difference between raw documents and structured
  data your AI can actually work with (Ch 4)

Day 4 (Thursday): The Workflow
- Take the document output from Day 3 and run it through one
  of Ch 5's workflows
- Medical chronology -> demand narrative, or deposition summary
  -> cross-exam outline
- Goal: experience the Process -> Analyze -> Act pipeline with
  your own case data

Day 5 (Friday): The Decision
```

```
- Review your time audit (Day 2) and your AI results (Days 1,
  3, 4)
- Decide: which task becomes your first skill file?
- Write the skill file using the template from this chapter
- Goal: have a written skill file ready to schedule
```

Case Strategy

4. Morning Case Briefing Skill

Source: Chapter 3, “Layer 1: The Skill” **Category:** Case Strategy

A complete, annotated skill file that produces meeting prep briefings by pulling data from Calendar, Clio, Gmail, and Dodonai. This is the reference example for how to write any skill file.

```
# /case-briefing -- Morning Case Briefing

Generate intel briefings for today's external meetings.

## Context

Before calls and meetings with clients, opposing counsel, or experts,
the attorney needs context on each attendee: their case status,
upcoming
deadlines, recent correspondence, and any open items. This skill
automates that gathering from Calendar, Clio, Gmail, and Dodonai.

## Steps

### 1. Determine the date

Use today's date unless a different date is specified.
Format as YYYY-MM-DD for file naming.

### 2. Fetch calendar events

Pull today's events from the firm's Google Calendar.

### 3. Filter to external meetings

Skip the event if:
- It is an all-day event
- It was declined
- It has no external attendees
- The title matches internal patterns: "team meeting", "staff lunch",
  "internal review"

Identify external attendees:
- Remove anyone with a firm email domain
- Remove calendar room resources
- If no external attendees remain, skip the event

If no external meetings exist for the day:
```

```
- Save a note: "No external meetings on [date]."
- Stop here.

### 4. For each external attendee, gather intel

**From Clio:**
- Look up the contact by name or email
- Pull their associated matters (case name, status, practice area)
- Pull upcoming deadlines and tasks on those matters
- Pull recent time entries (last 14 days)
- Pull the current trust/retainer balance

**From Gmail:**
- Search for recent emails to/from this contact (last 30 days)
- Note the last email date, subject, and whether it needs a response
- Flag any emails from opposing counsel on their matters

**From Dodonai:**
- Check if any documents are being processed for their matters
- Pull recent job outputs (medical chronologies, deposition
summaries)

**From web search (if new contact):**
- Search for the attendee's name + firm name
- Note their practice area, bar status, and any recent news

### 5. Compile the briefing

For each meeting, create a section with:
- Meeting time and title
- Attendee name, role, and firm
- Case summary (one paragraph)
- Upcoming deadlines (bulleted list)
- Recent correspondence summary
- Open items requiring attorney attention
- Trust/retainer balance

### 6. Save the briefing

Save to: reports/case_briefings/briefing_[date].md

### 7. Send notification

Post a summary to Slack:
"Morning briefing ready -- [N] meetings today. [list meeting names]"
```

---

5. Friday Case Review Skill (Minimal)

Source: Chapter 3, “Writing Your First Skill” **Category:** Case Strategy

A minimal four-line skill file showing you can start simple. Finds stale matters with no activity in 14 days.

```
# /weekly-review -- Friday Case Review

## Steps
```

```
### 1. Pull all active matters from Clio with no activity in the last
14 days

### 2. For each matter, note the last activity date and next deadline

### 3. Save the list to reports/stale_matters_[date].md
```

---

10. Case Theory from eDiscovery Results

Source: Chapter 5, “Workflow 3: eDiscovery Results to Case Theory” **Category:** Case Strategy

Analyzes exported eDiscovery search results to build a timeline of knowledge, identify patterns, surface hot documents, and articulate a case theory. Use after running semantic searches in Dodonai eDiscovery.

```
I'm developing a case theory for a premises liability case. Below are
the relevant documents identified through eDiscovery review. Each
document includes the source file name, the highlighted passage, and
the relevance ranking.

CASE CONTEXT:
- Plaintiff fell on a deteriorated exterior stairwell at a commercial
  office building
- Defendant: Meridian Property Management (property manager)
- Building owner: separate entity, not yet a party
- Key legal issue: Did Meridian have actual or constructive knowledge
  of the stairwell condition? Did they have a duty to repair?
- Jurisdiction: Texas -- premises liability standard requires proof
  that defendant knew or should have known of the condition

TASK:
Analyze these 62 documents and produce:

1. TIMELINE OF KNOWLEDGE: Build a chronological timeline showing
   when Meridian first learned of (or should have learned of)
   the stairwell condition. Cite each document by filename and
   relevant passage.

2. PATTERN ANALYSIS: Identify any pattern of deferred maintenance,
   budget cuts to property upkeep, ignored complaints, or similar
   conduct across the document set. Look for:
   - Repeated complaints about the same issue
   - Maintenance requests that were acknowledged but not completed
   - Budget communications that prioritized cost savings over safety
   - Communications showing awareness of risk

3. HOT DOCUMENTS: Identify the 5-10 most damaging documents for
   the defendant. For each, explain:
   - What it says
   - Why it's damaging
   - How it fits the case theory
   - Which witness should be questioned about it

4. MISSING PIECES: What evidence would strengthen this case theory
```

```
   but isn't in the current document set? What additional discovery
   should I pursue?

5. CASE THEORY MEMO: In 2-3 paragraphs, articulate the strongest
   case theory supported by this evidence. Write it as if you're
   explaining the theory to co-counsel joining the case.

[PASTE EDISCOVERY RESULTS HERE]
```

---

22. Deadline Monitor Skill File

Source: Chapter 7, “Build Your First Agent: The Deadline Monitor” **Category:** Case Strategy

A complete, production-ready skill file for daily deadline monitoring across all active matters. The most important agent to build first – prevents malpractice from missed deadlines.

```
# /deadline-monitor -- Daily Deadline & SOL Monitor

Check all active matters for approaching deadlines and statute of
limitations dates. Generate a prioritized morning alert.

## Context

This agent runs every weekday at 7:00 AM. It scans all active matters
in the firm's case management system, identifies every deadline
within
90 days, and produces a tiered alert. The attorney reviews the alert
before starting the workday.

## Steps

### 1. Pull all active matters

Fetch every matter with status "Open" or "Active" from Clio.
For each matter, retrieve:
- Matter name and case number
- Client name
- Practice area (PI, med mal, workers' comp, etc.)
- Responsible attorney
- Date opened
- Statute of limitations date (from custom field or matter notes)

### 2. Pull all deadlines and calendar entries

For each active matter, fetch:
- All tasks with due dates in the next 90 days
- All calendar entries (hearings, mediations, depositions, trial
dates)
  in the next 90 days
- Any court-imposed deadlines (discovery cutoffs, motion deadlines,
  expert disclosure dates)

### 3. Classify by urgency
```

```
Sort every deadline into three tiers:

**URGENT (within 30 days):**
- Statute of limitations expiring
- Discovery responses due
- Motion filing deadlines
- Expert disclosure deadlines
- Hearing or trial dates

**ACTION REQUIRED (31-60 days):**
- Upcoming depositions
- Mediation dates
- Demand deadlines (policy limits, pre-suit requirements)
- Client document requests due

**PLANNING (61-90 days):**
- Trial dates approaching
- Expert retention deadlines
- Settlement conference dates
- SOL dates entering the action window next month

### 4. Flag anomalies

Check for:
- Any active matter with NO deadline or SOL date on file
  (flag: "No SOL recorded -- verify immediately")
- Any SOL date within 30 days where no demand or suit has been filed
  (flag: "CRITICAL -- SOL imminent, no filing on record")
- Any matter with no activity in the last 60 days and an SOL
  within 180 days (flag: "Stale matter approaching SOL")
- Duplicate or conflicting deadlines on the same matter

### 5. Generate the morning alert

Format the output as a daily digest:

**Header:** "Deadline Monitor -- [date] -- [total deadlines flagged]"

**Section 1: URGENT (30 days)**
For each item:
- Matter name | Client | Deadline type | Due date | Days remaining
- One-line note on what action is needed

**Section 2: ACTION REQUIRED (60 days)**
Same format as above.

**Section 3: PLANNING (90 days)**
Same format as above.

**Section 4: ANOMALIES**
- Matters with no SOL recorded
- Stale matters approaching SOL
- Any other flags from Step 4

**Footer:** "Total active matters checked: [N].
Next run: tomorrow at 7:00 AM."

### 6. Save and notify

Save the digest to: reports/deadlines/deadline_[date].md
```

```
Send a summary to Slack:
"Deadline alert ready -- [N] urgent, [N] action, [N] planning."

If any URGENT items exist, also send a separate Slack message
for each one: "URGENT: [matter name] -- [deadline type]
due [date] ([N] days)."

## Safety Notes

- All computed deadlines must be verified against source documents.
  Holiday and weekend adjustments vary by jurisdiction.
- SOL computation is based on the date recorded in the case
  management system. If the recorded date is wrong, the alert
  will be wrong. Garbage in, garbage out.
- This agent reads data. It does not modify matters, create
  calendar events, or send external communications. Read-only.
```

Medical Records

7. Damages Narrative from Medical Chronology

Source: Chapter 5, “Workflow 1: Medical Chronology to Demand Letter” **Category:** Medical Records

The first prompt in a two-prompt chain. Takes a Dodonai medical chronology output and produces a structured damages narrative for a demand letter. Includes a “Gaps and Flags” section that identifies issues defense counsel will exploit.

```
I'm preparing a demand letter for a personal injury case. Below is
the
medical chronology generated from the plaintiff's medical records.

CASE FACTS:
- Plaintiff: Maria Gonzalez, age 34, administrative assistant
- Incident: Rear-end collision on IH-35, Austin, TX on 03/15/2025
- Liability: Clear -- defendant rear-ended plaintiff at stoplight,
  police report confirms
- Pre-existing conditions: None documented in records

TASK:
Analyze the medical chronology and produce a damages narrative
suitable
for inclusion in a demand letter. Structure it as follows:

1. INJURY SUMMARY: One paragraph describing all diagnosed injuries,
   severity, and prognosis
2. TREATMENT TIMELINE: Narrative of treatment progression from
   initial ER visit through most recent appointment, emphasizing:
   - Escalation of treatment (conservative to interventional)
   - Duration and frequency of care
   - Any surgical interventions or recommendations
   - Ongoing symptoms and functional limitations
3. IMPACT ON DAILY LIFE: How injuries affected plaintiff's ability
```

```
   to work, perform daily activities, and maintain quality of life
   (draw from any functional assessments or physician notes in the
   chronology)
4. PROGNOSIS: Future treatment recommendations, permanent
   impairment ratings if documented, and ongoing care needs
5. GAPS AND FLAGS: Note any treatment gaps longer than 30 days,
   any pre-existing conditions referenced, or any entries that
   the attorney should verify before sending the demand

Use specific dates, provider names, and treatment details from the
chronology. Cite the chronology entry dates so I can verify each
statement against the source records.

Do not editorialize or overstate. Let the medical evidence speak. The
tone should be factual and precise -- this will be reviewed by an
insurance adjuster and possibly defense counsel.

[PASTE DODONAI MEDICAL CHRONOLOGY HERE]
```

---

8. Medical Specials Calculation

Source: Chapter 5, “Workflow 1: Medical Chronology to Demand Letter” **Category:** Medical Records

The second prompt in the demand letter chain. Run after the Damages Narrative prompt. Produces the medical specials table and flags incomplete billing data.

```
Using the damages narrative above and the medical chronology, prepare
the medical specials section of the demand package.

Structure as a table with these columns:
- Provider name
- Date range of treatment
- Description of services
- Total billed amount (if documented in records)
- Status (paid, outstanding, liened)

After the table, include:
1. TOTAL MEDICAL SPECIALS (sum of all billed amounts)
2. FUTURE MEDICAL COSTS (estimated based on prognosis section --
   note these are projections and flag if an economist or life care
   planner should be retained)
3. LOST WAGES CALCULATION (if employment records were in the
   chronology -- if not, flag that these need to be obtained
   separately)

Flag any billing entries where the amount seems inconsistent with the
treatment described, or where records from a provider appear
incomplete.
```

---

14. PI Medical Records Intake Agent

Source: Chapter 6, “Personal Injury” **Category:** Medical Records

An agent skill for classifying incoming medical records, uploading them for chronology processing, scanning for treatment gaps, and tracking provider completeness. Scales from 5 cases to 50 with no changes.

```
You are a medical records intake agent for a personal injury
practice.

WHEN: New email attachments arrive from medical providers
(match sender domains against the provider list in
[CASE_MANAGEMENT_SYSTEM])

DO:
1. Classify each attachment:
   - Treatment records (progress notes, operative reports, discharge
summaries)
   - Billing statements (itemized charges, EOBs, lien letters)
   - Imaging (radiology reports, MRI/CT scan reports)
   - Pharmacy records
   - Correspondence (IME notices, peer review letters)

2. Match to the correct case by patient name and DOI

3. Upload treatment records and imaging to Dodonai for chronology
processing

4. After chronology is complete, scan for:
   - Treatment gaps > 30 days (flag for attorney review -- defense
will exploit these)
   - Missing providers (compare against the known treater list)
   - Inconsistencies between billing and treatment dates

5. Update the case status: "[X] of [Y] providers received.
   Missing: [provider names]. SOL: [date] ([days] remaining)."

6. When all providers are received, alert:
   "Records complete for [case name]. Chronology ready.
   Ready for demand preparation."
```

---

15. PI Demand Preparation Agent

Source: Chapter 6, “Personal Injury” **Category:** Medical Records

Produces a demand preparation memo from a completed Dodonai medical chronology – extracting specials, mapping the treatment arc, flagging issues, and drafting the damages narrative.

```
You are a demand preparation agent for a personal injury practice.

CONTEXT: All medical records have been received and processed through
Dodonai Medical Record Summaries & Chronologies for the following
case:

[Paste Dodonai chronology output here]

CASE DETAILS:
```

```
- Plaintiff: [name]
- DOI: [date]
- Liability: [brief liability summary]
- Policy limits: [amount]
- Pre-existing conditions: [if any]

TASK:
1. Extract total medical specials from the chronology
   (separate by provider, by date, by type)
2. Identify the treatment arc: initial ER/urgent care -> specialist
referrals
   -> ongoing treatment -> MMI or current status
3. Flag any gaps, inconsistencies, or missing records that need
resolution
   before sending the demand
4. Draft a damages narrative (3-5 paragraphs) that walks the adjuster
   through the injury, treatment, and impact on daily life

OUTPUT: A demand preparation memo with specials table, treatment
timeline summary, flagged issues, and draft damages narrative.

NOTE: Attorney must verify all medical specials against source
records,
review the damages narrative for accuracy, and add lost wages
and non-economic damages analysis before finalizing.
```

Depositions

9. Cross-Examination Outline from Deposition

Source: Chapter 5, “Workflow 2: Deposition Summary to Cross-Examination Outline” **Category:** Depositions

Analyzes a page-line deposition summary to extract key admissions, internal contradictions, and draft leading cross-examination questions with transcript citations.

```
I'm preparing cross-examination questions for the defendant's
corporate
representative in a premises liability case. Below is the page-line
deposition summary generated from the transcript.

CASE CONTEXT:
- Plaintiff slipped on a wet floor in the defendant's grocery store
- Defendant claims regular inspection and maintenance schedule
- Key issue: whether defendant had actual or constructive notice of
  the hazardous condition
- Deposition taken: 02/10/2026
- Deponent: James Whitfield, Store Manager and Corporate Rep

TASK:
Analyze this deposition summary and produce:
```

```
1. KEY ADMISSIONS: Testimony that supports the plaintiff's case.
   For each admission, quote the relevant summary entry with its
   page-line citation and explain why it matters.

2. INTERNAL CONTRADICTIONS: Places where the deponent contradicts
   himself within this testimony. For each contradiction, cite both
   statements with page-line references and frame the inconsistency.

3. CONTRADICTIONS WITH KNOWN FACTS: Based on the testimony, flag
   statements that likely conflict with:
   - Store surveillance footage (if it exists)
   - Maintenance logs or inspection records
   - Prior incident reports
   - Employee witness statements
   Note: I'll need to verify these against actual documents. Flag
   them as "areas to investigate."

4. CROSS-EXAMINATION QUESTIONS: For each admission and
   contradiction, draft 3-5 leading questions suitable for cross.
   Follow standard cross-examination form:
   - One fact per question
   - Leading (answerable with "yes" or "no")
   - Build from undeniable facts toward the key admission
   - Pin the witness to the prior testimony before revealing the
     contradiction
   Include the page-line citation after each question so I can
   impeach directly from the transcript.

5. AREAS FOR FOLLOW-UP DISCOVERY: Identify any documents,
   witnesses, or records referenced in the testimony that haven't
   been produced and should be requested.

[PASTE DODONAI PAGE-LINE DEPOSITION SUMMARY HERE]
```

---

12. Trial Prep Matrix from Multiple Depositions

Source: Chapter 5, “Workflow 5: Multiple Depositions to Trial Prep Matrix”
Category: Depositions

The most complex workflow in the book. Cross-references multiple deposition summaries to produce a witness grid, contradiction matrix, impeachment map, and trial themes. Essential for multi-party litigation.

```
I'm preparing for trial in a multi-vehicle PI case. Below are page-
line
deposition summaries for all 11 witnesses. I need a comprehensive
trial preparation matrix.

CASE SUMMARY:
- Three-vehicle chain collision on IH-35, Austin, TX
- Plaintiff (our client): rear passenger in Vehicle 2
- Defendant A: driver of Vehicle 1 (stopped suddenly)
- Defendant B: driver of Vehicle 3 (rear-ended Vehicle 2, pushing
  it into Vehicle 1)
- Key dispute: which impact caused plaintiff's injuries -- the
  initial rear-end from Vehicle 3 or the secondary front-end
```

```
  collision with Vehicle 1
- Each defendant blames the other for the mechanism of injury

WITNESSES:
1. Plaintiff - Angela Torres
2. Defendant A - Michael Park (Vehicle 1 driver)
3. Defendant B - Lisa Drummond (Vehicle 3 driver)
4. Defendant B's passenger - Ryan Drummond
5. Eyewitness 1 - Sandra Okafor (adjacent lane)
6. Eyewitness 2 - Thomas Briggs (behind Vehicle 3)
7. Officer Patricia Vega (responding officer)
8. Dr. James Okonkwo (treating orthopedist)
9. Dr. Helen Strauss (treating neurologist)
10. Dr. Andrew Kim (defense biomechanical expert)
11. Dr. Rebecca Liang (accident reconstruction, plaintiff's expert)

PRODUCE:

1. WITNESS-BY-WITNESS SUMMARY GRID:
   Create a table with one row per witness and columns for:
   - Witness name and role
   - Position on mechanism of injury (which impact caused injuries)
   - Key admissions favorable to plaintiff
   - Key admissions favorable to defense
   - Credibility issues identified
   - Page-line citations for critical testimony

2. CONTRADICTION MATRIX:
   For each material fact in dispute, create a grid showing what
   each witness testified. Focus on:
   - Sequence of impacts
   - Speed of vehicles
   - Time between impacts
   - Whether plaintiff's vehicle moved after initial impact
   - Plaintiff's complaints immediately after the collision
   - Whether plaintiff was bracing or unaware before impact

3. IMPEACHMENT OPPORTUNITIES:
   For each witness, identify:
   - Internal contradictions (within their own testimony)
   - Contradictions with other witnesses
   - Contradictions with documentary evidence (police report,
     medical records, accident reconstruction)
   - Prior inconsistent statements (if referenced in any deposition)
   Cite page-line for every identified inconsistency.

4. CROSS-EXAMINATION PRIORITIES:
   Rank the defense witnesses by importance and recommend the
   top 3-5 cross-examination objectives for each, with the
   supporting testimony citations.

5. TRIAL THEMES:
   Based on the testimony as a whole, identify the 2-3 strongest
   trial themes for the plaintiff. For each theme, list the
   witnesses and testimony that support it.

[PASTE ALL 11 PAGE-LINE DEPOSITION SUMMARIES HERE]
```

Client Communication

24. Client Update Drafter Skill

Source: Chapter 7, “Agent 3: Client Update Drafter” **Category:** Client Communication

A skill file section for generating monthly status update emails for every active matter. Pulls case activity from Clio and correspondence history from Gmail, drafts in plain language, and flags overdue contacts.

```
### For each active matter, generate a client update draft

Pull from Clio:
- Activity in the last 30 days (filings, correspondence, depositions)
- Next upcoming deadline or milestone
- Trust/retainer balance

Pull from Gmail:
- Date of last client communication
- Any unanswered client emails

Draft the update email:
- Opening: acknowledge the client by name, reference the matter
- Body: summarize recent activity in plain language (no legal jargon
  unless the client is sophisticated). What happened, what it means,
  what happens next.
- Close: next milestone with estimated timeline, invitation to call
  with questions
- Tone: direct, warm, reassuring. The client should feel informed,
  not alarmed.

Save as Gmail draft, addressed to the client. Do NOT send.
Flag any matter where the last client contact was over 45 days ago:
"OVERDUE -- last contact [date]."
```

Intake

11. Intake Go/No-Go Evaluation

Source: Chapter 5, “Workflow 4: Intake Notes to Go/No-Go Case Evaluation” **Category:** Intake

Encodes a firm’s case acceptance criteria into a structured evaluation framework. Produces a go/no-go recommendation with case value estimate, red flag identification, and suggested next steps. A prime candidate for conversion to a permanent skill file.

```
You are a case evaluation assistant for a PI practice in Austin, TX.
Analyze the intake information below and produce a go/no-go
```

```
recommendation memo.

FIRM'S CASE ACCEPTANCE CRITERIA:
- Minimum viable case value: $25,000 (we don't take cases below this)
- Clear or probable liability preferred (we evaluate contested
  liability cases individually)
- Treatment must be ongoing or recently completed (gaps >90 days
  from incident with no treatment are a red flag)
- Statute of limitations: 2 years in Texas for PI (flag if <6
  months remaining)
- Red flags: prior claims for same body part, attorney shopping
  (multiple prior attorneys), pre-existing conditions not
  disclosed on intake, incident on plaintiff's own property,
  government entity defendant (sovereign immunity / notice
  requirements)

INTAKE INFORMATION:
[PASTE INTAKE NOTES FROM CLIO AND XD AGENT EXTRACTION HERE]

PRODUCE:
1. CASE SUMMARY: 3-4 sentences summarizing the facts
2. LIABILITY ASSESSMENT: Clear / Probable / Contested / Unclear
   -- with reasoning
3. DAMAGES ESTIMATE: Conservative range based on injury type,
   treatment level, and comparable cases in Texas. Note: this is
   a rough screening estimate, not a formal valuation.
   - Medical specials (current and estimated future)
   - General damages multiplier range (1.5x-5x depending on
     severity, surgery, permanence)
   - Total estimated case value range
4. RED FLAGS: Any of the criteria above that are triggered
5. STATUTE CHECK: SOL deadline and time remaining
6. RECOMMENDATION: ACCEPT / DECLINE / INVESTIGATE FURTHER
   -- with a one-paragraph explanation
7. IF ACCEPT: Suggested next steps (records to request, witnesses
   to identify, experts to retain)
8. IF INVESTIGATE: What specific information would change the
   recommendation

Be conservative in your estimates. I'd rather decline a marginal
case and be wrong than accept one and lose six months. A "no" that
should have been a "yes" costs me one case fee. A "yes" that should
have been a "no" costs me 200 hours of unrecoverable time.
```

---

25. Intake Lead Scoring Rules

Source: Chapter 7, “Agent 5: Intake Pipeline” **Category:** Intake

A weighted scoring framework for triaging new leads. Encode your case acceptance judgment as weights and thresholds, then apply consistently to every intake.

```
### Score the lead

Evaluate against these criteria:
```

```
| Factor | Weight | Scoring |
|--------|--------|---------|
| Injury severity | 30% | Catastrophic (10), Serious (7), Moderate
(5), Minor (2) |
| Liability clarity | 25% | Clear (10), Probable (7), Disputed (4),
Weak (1) |
| Insurance coverage | 20% | Commercial policy (10), Standard auto
(7), Minimum/uninsured (3) |
| Case value estimate | 15% | >$100K (10), $50-100K (7), $25-50K (5),
<$25K (2) |
| Practice fit | 10% | Core PI (10), Adjacent (6), Outside practice
areas (1) |

Score > 70: Strong intake -- draft consult confirmation immediately
Score 50-70: Review recommended -- flag for attorney evaluation
Score < 50: Likely decline -- draft polite referral response
```

---

29. Intake Evaluation and Triage

Source: Chapter 8, “The Intake Pipeline” **Category:** Intake

A comprehensive intake evaluation prompt that scores new leads, assesses liability, checks statute of limitations, flags urgency signals, identifies conflicts to check, and recommends next steps. More detailed than the Ch 5 version, covering the full intake pipeline workflow.

```
You are evaluating a new intake inquiry for a personal injury
law firm in Austin, TX. The firm handles auto accidents, slip
and fall, premises liability, medical malpractice, and workers'
compensation.

INQUIRY DETAILS:
- Source: [Phone / Web form / Email / Referral from ___]
- Date received: [DATE]
- Name: [NAME]
- Contact: [PHONE / EMAIL]
- Description of incident: [PASTE WHATEVER THE PROSPECT
  DESCRIBED -- this may be a voicemail transcript, a web form
  submission, or an email. It will often be informal and
  incomplete.]

TASK:
Evaluate this inquiry and produce a structured intake assessment:

1. PRACTICE AREA FIT (Yes / Partial / No)
   - Does this fall within the firm's practice areas?
   - If partial or no, what practice area does it actually
     involve?

2. CASE VALUE ESTIMATE (High / Medium / Low / Cannot Assess)
   - Based on the injury described, treatment mentioned, and
     liability indicators, estimate the likely case value range
   - Note: This is a preliminary estimate for triage purposes
     only. It is not a case evaluation.

3. LIABILITY ASSESSMENT (Strong / Moderate / Weak / Insufficient
```

```
Info)
   - What liability indicators are present?
   - What additional facts would strengthen or weaken the claim?

4. STATUTE OF LIMITATIONS
   - Based on the incident date and case type, calculate the
     SOL deadline under Texas law
   - Flag if less than 6 months remain

5. URGENCY FLAGS
   - Evidence preservation needed? (surveillance footage,
     vehicle damage, scene conditions)
   - Medical treatment ongoing or recommended?
   - SOL pressure?
   - Other time-sensitive factors?

6. CONFLICTS CHECK NEEDED
   - List the names that need to be run through the conflicts
     database: prospective client, any opposing parties
     mentioned, any businesses or entities involved

7. RECOMMENDED ACTION
   - Priority: High / Medium / Low / Decline
   - Suggested next step (schedule consultation, request more
     information, refer out, decline with explanation)
   - Draft response email (2-3 sentences appropriate for the
     priority level)

8. MISSING INFORMATION
   - What critical facts are missing from this inquiry?
   - Suggest 3-5 specific questions for the initial callback

IMPORTANT:
- This is a screening tool, not a legal opinion. Flag
  uncertainty rather than guessing.
- If the injury description is too vague to assess, say so
  and recommend a phone screen to gather more detail.
- Do not assess damages. That requires medical records and
  is beyond the scope of intake triage.
```

Marketing

26. Blog Post Generation

Source: Chapter 8, “The Content Engine” **Category:** Marketing

Drafts an 800-1,200 word blog post optimized for SEO, written in the attorney’s voice. Include a voice sample from your website or client emails to teach the AI your natural writing style.

```
You are a content writer for a personal injury law firm in Austin,
TX.
The firm handles auto accidents, slip and fall, premises liability,
medical malpractice, and workers' compensation cases.
```

```
TASK:
Write a blog post on the following topic:

TOPIC: [TOPIC -- e.g., "What to do in the first 72 hours after a
car accident in Texas"]

REQUIREMENTS:
- Length: 800-1,200 words
- Audience: Potential clients who have been injured and are
  researching their options. They are not lawyers. Write in plain
  English, but don't condescend.
- Tone: Informative, confident, empathetic. Not salesy. Not
  "call us now for a free consultation" energy. Think: helpful
  attorney explaining the process over coffee.
- Structure:
  1. Opening that addresses the reader's situation directly
     (2-3 sentences)
  2. Clear section headers (H2s) breaking the topic into
     actionable steps or key points
  3. Specific Texas law references where relevant (cite the
     statute or rule by name -- e.g., Texas Civil Practice &
     Remedies Code Section 33.001 for comparative fault)
  4. A "what most people don't know" section with one non-obvious
     insight that demonstrates expertise
  5. Brief closing that encourages the reader to consult an
     attorney without being pushy
- SEO: Naturally incorporate these search terms where they fit
  (don't force them): [PRIMARY KEYWORD], [2-3 SECONDARY KEYWORDS]
- Do NOT include a disclaimer or "this is not legal advice"
  boilerplate -- the firm adds that sitewide
- Do NOT use phrases like "navigating the legal landscape" or
  "in today's world" or "it's important to note"

VOICE SAMPLE (match this tone):
[PASTE 2-3 PARAGRAPHS FROM AN EXISTING PAGE ON YOUR WEBSITE
OR A CLIENT EMAIL YOU'VE WRITTEN -- THIS TEACHES THE AI YOUR
NATURAL WRITING STYLE]
```

---

27. Review Response Drafter

Source: Chapter 8, “Review Management” **Category:** Marketing

Drafts personalized responses to Google and Avvo reviews. Includes separate instructions for positive and negative reviews, with bar advertising compliance rules built in.

```
You are responding to an online review for a personal injury law
firm in Austin, TX. The attorney is Sarah Chen.

REVIEW PLATFORM: [Google / Avvo]
REVIEWER NAME: [NAME]
STAR RATING: [1-5]
REVIEW TEXT:
[PASTE THE REVIEW TEXT]
```

```
TASK:
Draft a response following these rules:

FOR POSITIVE REVIEWS (4-5 stars):
- Thank the reviewer by first name
- Reference something specific they mentioned (e.g., "I'm glad
  the team kept you informed throughout the process") WITHOUT
  revealing any case details, injury types, or legal strategy
- Keep it to 2-3 sentences. Warm but professional.
- Do NOT say "we appreciate your business" or anything that
  sounds like a retail store
- Do NOT reference specific case outcomes, settlement amounts,
  or legal proceedings -- this protects client confidentiality
  and complies with bar advertising rules

FOR NEGATIVE REVIEWS (1-3 stars):
- Acknowledge the reviewer's frustration without admitting fault
- Do NOT reference any case specifics
- Invite them to contact the office directly to discuss their
  concerns: "I'd welcome the chance to discuss this with you
  directly -- please call our office at [PHONE]"
- Keep it to 2-3 sentences. Professional, empathetic, not
  defensive
- Flag this for the attorney's personal review before posting

GENERAL RULES:
- Texas State Bar advertising rules apply -- no guarantees of
  results, no specific case details, no client testimonials
  that imply a guaranteed outcome
- Sign as "Sarah Chen" (not "The Team" or "Chen Law Group")
- Match the tone of the review -- if the reviewer is emotional,
  be warm. If they're matter-of-fact, be direct.
```

---

28. Referral Follow-Up Email

Source: Chapter 8, “Referral Nurturing” **Category:** Marketing

Drafts quarterly nurture emails to referral sources. Emphasizes genuine relationship tone over marketing language, includes useful information sharing, and avoids common sales email pitfalls.

```
You are drafting a quarterly referral nurture email for Sarah Chen,
a personal injury attorney in Austin, TX.

REFERRAL SOURCE:
- Name: [NAME]
- Type: [Attorney / Chiropractor / Other Professional]
- Relationship duration: [e.g., "5 years"]
- Cases referred in last 12 months: [NUMBER]
- Last communication: [DATE and brief description]
- Notable context: [e.g., "Dr. Martinez's patients are often
  auto accident victims; he refers clients who need legal
  representation for their injury claims"]

TASK:
Draft a short email (150-250 words) from Sarah to this referral
```

```
source. The email should:

1. Open with a genuine, specific reference to the relationship --
   NOT "I hope this email finds you well" or "Just checking in."
   Reference something real: a recent referral, a shared
   professional connection, a development in their field.

2. Share ONE piece of useful information. Options:
   - A recent legal development relevant to their practice
     (e.g., a new Texas statute affecting auto accident claims
     that a chiropractor's patients should know about)
   - An anonymized case result that demonstrates the value of
     early referral
   - A blog post or resource from the firm's website

3. Close with a specific, low-pressure next step -- coffee,
   a brief call, an invitation to a CLE event. NOT "let me
   know if you have any referrals."

TONE: This is a genuine professional relationship, not a sales
email. Sarah values this person. The email should read like it
was written by Sarah personally, not generated by a marketing
system.

DO NOT:
- Use phrases like "valued referral partner" or "mutually
  beneficial relationship"
- Reference specific client names or case details without
  [PLACEHOLDER] tags for Sarah to fill in
- Include any marketing language, CTAs, or promotional content
- Use em-dashes
```

Billing/Operations

30. Billing Reconciler

Source: Chapter 9, “Time Entry Reconstruction” **Category:** Billing/Operations

Cross-references calendar, email, and document activity against Clio time entries to identify unbilled hours. Produces a reconciliation report with suggested time entries, confidence levels, and a weekly gap summary. Run every Friday afternoon.

```
You are a billing reconciler for a litigation practice. Your job is
to
identify unbilled work by comparing three data sources and flagging
gaps.

DATA SOURCES:
1. Calendar events for the past [WEEK/DAY] (provided below)
2. Email activity by matter for the past [WEEK/DAY] (provided below)
3. Current time entries in Clio for the same period (provided below)
```

```
RECONCILIATION RULES:
- A calendar event with a client or matter name that has no
corresponding
  Clio time entry is an UNBILLED GAP
- 3+ substantive emails on a single matter in one day with less than
  0.5 hrs billed suggests missing correspondence/review time
- A calendar block of 30+ minutes with a partial Clio entry (billed
  duration < 75% of calendar duration) is UNDERBILLED
- Exclude: internal team meetings with no client/matter reference,
  personal calendar events, automated notifications
- Exclude: emails that are clearly administrative (scheduling,
  forwarding without comment, out-of-office replies)

FOR EACH GAP, PROVIDE:
- Matter name and case type
- Source of the evidence (calendar event, email chain, or both)
- Date and approximate time
- Suggested time entry description (in the style of a Clio narrative
--
  professional, specific, using standard billing language)
- Suggested duration (in tenths of an hour)
- Confidence level: HIGH (calendar event with no entry at all),
  MEDIUM (email activity suggests unbilled time), LOW (indirect
  evidence only)

WEEKLY SUMMARY:
- Total activity identified across all sources (hours)
- Total billed in Clio (hours)
- Total gap (hours and dollar value at $[RATE]/hr)
- Number of suggested entries by confidence level

Sort gaps by dollar value, highest first. Group by matter.
```

---

31. SOP Generator

Source: Chapter 9, “SOP Creation” **Category:** Billing/Operations

Converts a verbal process description into a formal Standard Operating Procedure. Sit with your paralegal for 10 minutes, describe a process naturally, and paste the transcript into this prompt. Produces a structured SOP in under 15 minutes.

```
I'm going to describe a process we follow at our law firm. Turn it
into a formal Standard Operating Procedure (SOP) with the following
structure:

SOP FORMAT:
- Title and purpose (one sentence)
- Roles involved (who does what)
- Trigger (what initiates this process)
- Step-by-step procedure with numbered steps
- For each step: the action, who performs it, expected timeframe,
  and any system where the action is recorded
- Checklist version (checkbox format for daily use)
- Common exceptions and how to handle them
- Quality check: what does "done correctly" look like?
```

```
RULES:
- Use specific system names (Clio, Gmail, QuickBooks) where relevant
- Include the actual form names, template names, or document titles
  we reference
- Flag any step where I'm vague and ask me to clarify
- Write for someone who has never done this before -- assume no
  prior knowledge of our firm's systems

Here is the process:

[DESCRIBE THE PROCESS -- speak naturally, as if you're training
a new hire. Include the exceptions and edge cases you've learned
from experience. The messier and more detailed, the better.]
```

---

32. Engagement Letter Generator

Source: Chapter 9, “Template Libraries” **Category:** Billing/Operations

Generates a complete engagement letter from case data fields, with conditional logic for contingency, hybrid, and hourly fee structures. Designed to be connected to Clio as an AI-ready template.

```
Generate an engagement letter for a new personal injury case using
the
following template structure and case data.

TEMPLATE STRUCTURE:
1. Opening paragraph: Confirm retention, identify the specific
incident
   (date, location, type of accident), name the adverse party or
   parties
2. Scope of representation: PI claim arising from [INCIDENT TYPE],
   including negotiation, litigation if necessary, and settlement
3. Fee arrangement: [CONTINGENCY/HYBRID/HOURLY] -- use the
   appropriate language block below
   - Contingency: [FEE_PERCENTAGE]% of gross recovery, advancing
     costs, client responsible for costs if no recovery
   - Hybrid: $[RETAINER] retainer + [FEE_PERCENTAGE]% of recovery
     above $[THRESHOLD]
   - Hourly: $[RATE]/hour, $[RETAINER] retainer, monthly billing
4. Client obligations: cooperation, truthfulness, medical treatment
   compliance, no direct contact with adverse parties or insurers
5. Termination provisions: either party, written notice, client
   responsible for costs incurred to date
6. Closing: signature blocks for attorney and client, date line

CASE DATA (from Clio):
- Client: [CLIENT_NAME]
- Matter: [MATTER_NAME]
- Date of loss: [DATE_OF_LOSS]
- Incident type: [INCIDENT_TYPE]
- Adverse party: [ADVERSE_PARTY]
- Fee structure: [FEE_TYPE]
- Fee percentage: [FEE_PERCENTAGE]
- Retainer amount: [RETAINER] (if applicable)
```

```
STYLE:
- Professional but accessible -- the client is not a lawyer
- Short paragraphs, no legalese where plain English works
- Texas-compliant engagement letter format
- Match the tone and structure of [ATTACH: your most recent
  engagement letter as a style reference]

Generate the complete letter. Flag any missing data fields that
prevent completion.
```

---

33. Onboarding Plan Generator

Source: Chapter 9, “Training and Onboarding” **Category:** Billing/Operations

Generates a structured 2-week onboarding plan for a new paralegal, mapped to your SOP library. Day-by-day schedule with learning objectives, hands-on tasks, and competency checkpoints.

```
Generate a 2-week onboarding plan for a new paralegal at a personal
injury law firm (solo practitioner, 40 active cases, using Clio,
Gmail/Google Workspace, QuickBooks, and Dodonai for document
processing).

AVAILABLE SOPs (the new hire should learn these during onboarding):
[LIST YOUR SOPs -- paste the titles and one-sentence descriptions]

ONBOARDING PLAN REQUIREMENTS:
- Day-by-day schedule for 10 business days
- Each day: 2-3 learning objectives, specific SOPs to review,
  hands-on tasks to practice, and a verification checkpoint
- Week 1 focus: systems access, core procedures (file opening,
  calendaring, basic correspondence), observation of active matters
- Week 2 focus: independent execution under review (open a real
  file, draft a real LOR, process real medical records), exception
  handling, quality standards
- Include "shadow" sessions where the new hire observes the
  paralegal or attorney handling specific tasks
- Include a Day 10 assessment: checklist of competencies the new
  hire should demonstrate independently

PRACTICAL CONSTRAINTS:
- The existing paralegal (Maria) can dedicate 1 hour/day to
  training. Remaining learning is self-directed from SOPs.
- The new hire has prior paralegal experience but is new to PI
  and new to Clio
- Systems access (Clio login, email, Dodonai account) should be
  set up on Day 1

Format as a table: Day | Morning | Afternoon | SOPs Referenced |
Checkpoint
```

Operations (Practice-Area-Specific)

16. Real Estate Closing Conditions Tracker

Source: Chapter 6, “Real Estate” **Category:** Operations

Monitors inbox for incoming closing documents, matches them against a deal checklist, tracks conditions met/outstanding, and flags approaching deadlines. Includes cross-border compliance checks for multi-jurisdiction deals.

```
You are a closing conditions tracker for a commercial real estate
practice.

ACTIVE DEAL: [Deal name / property address]
CLOSING CHECKLIST:
[Paste or reference the closing conditions list from your deal
management system]

WHEN: New emails arrive matching this deal
(match by property address, deal name, or counterparty email domains)

DO:
1. Classify the attachment or email content against the closing
checklist:
   - Title commitment / title insurance
   - Survey / ALTA survey
   - Environmental (Phase I, Phase II)
   - Estoppel certificates
   - Tenant subordination agreements
   - Lender approval / commitment letter
   - Government approvals / permits
   - Entity formation documents

2. If the document satisfies a checklist condition, mark it RECEIVED
   with the date and source

3. If a condition deadline is within 5 business days and status is
   still OUTSTANDING, escalate: "[Deal name]: [condition] due in
   [N] days -- not yet received from [responsible party]."

4. For cross-border deals, flag any jurisdiction-specific
requirements:
   - FIRPTA withholding obligations (foreign seller)
   - State/country-specific disclosure requirements
   - Tax treaty considerations
   - Regulatory approvals required

5. Generate daily status: "[Deal name] -- [X] of [Y] conditions met.
   Outstanding: [list]. Next deadline: [condition] due [date]."
```

---

17. Tax Correspondence Monitoring Agent

Source: Chapter 6, “Tax Disputes” **Category:** Operations

Monitors incoming correspondence from CRA and IRS, extracts assessment details, computes objection and response deadlines, and sets reminder cascades. Critical for practices managing hundreds of active tax files.

```
You are a tax correspondence monitoring agent for a tax dispute
practice.

WHEN: New email or scanned document arrives from:
- Canada Revenue Agency (CRA) domains or letterhead
- Internal Revenue Service (IRS) domains or letterhead
- Tax Court of Canada / US Tax Court

DO:
1. Classify the document:
   - Notice of Assessment (NOA) / Notice of Reassessment
   - Notice of Deficiency (IRS)
   - Notice of Confirmation
   - Proposed adjustment letter
   - Requirement to Pay / Collection notice
   - Tax Court hearing notice / scheduling order

2. Extract:
   - Taxpayer name and file number
   - Assessment date and tax year(s)
   - Amount assessed or in dispute
   - Deadline trigger: the date from which response periods are
computed

3. Compute deadlines:
   - NOA -> 90-day objection window (CRA) or 90-day Tax Court
petition (IRS)
   - Proposed adjustment -> 30-day response window
   - Collection notice -> flag for immediate review
   - Hearing notice -> appearance date + pretrial deadlines

4. Match to client file in [CASE_MANAGEMENT_SYSTEM]

5. Set reminder cascade: 60 days, 30 days, 14 days, 7 days, 3 days

6. Alert: "[Client name] -- [Document type] received [date].
   Objection deadline: [date] ([days] remaining).
   Tax year: [year]. Amount: [amount]."

IMPORTANT: All computed deadlines must be verified by the attorney
against the source document. Holiday and weekend adjustments
vary by jurisdiction and must be confirmed manually.
```

---

18. Breach Notification Compliance Agent

Source: Chapter 6, “Cybersecurity & Privacy” **Category:** Operations

Generates a multi-state compliance matrix from a breach incident report – computing notification deadlines, content requirements, AG filing obligations, and consumer reporting agency thresholds for each affected jurisdiction.

```
You are a breach notification compliance agent for a cybersecurity
and data privacy practice.

BREACH INCIDENT REPORT:
- Client: [company name]
- Date of breach discovery: [date]
- Type of data exposed: [PII categories -- SSN, financial, health,
etc.]
- Number of affected individuals by state:
  [State]: [count]
  [State]: [count]
  ...

FOR EACH AFFECTED STATE, GENERATE:
1. NOTIFICATION DEADLINE
   - Statute citation
   - Deadline computation from discovery date
   - Any "without unreasonable delay" language
     (note: flag these -- they require attorney judgment on
reasonableness)

2. CONTENT REQUIREMENTS
   - Required elements in notification letter
   - Credit monitoring obligation (yes/no, duration if required)
   - Specific language mandated by statute

3. AG / REGULATORY NOTIFICATION
   - Is separate AG notification required?
   - AG notification deadline (before or concurrent with individual
notice?)
   - Filing method (online portal, email, mail)
   - Threshold trigger (e.g., notify AG if > 500 residents affected)

4. CONSUMER REPORTING AGENCY NOTIFICATION
   - Required if breach exceeds [threshold] residents in the state?

OUTPUT: Compliance matrix (table format), draft notification letters
per jurisdiction, filing tracker with status column.

IMPORTANT: State breach notification laws are amended frequently.
All deadlines and content requirements must be verified against
current statutes before any notification is sent or filed.
Attorney review is mandatory before any external communication.
```

---

19. Immigration Case Monitoring Agent

Source: Chapter 6, “Immigration” **Category:** Operations

Daily monitoring of immigration portals for status changes, with deadline computation for biometrics windows, medical certificate expiry, hearing dates, work permit renewals, and BOC submissions.

```
You are an immigration case monitoring agent for an immigration
and refugee law practice.

DAILY CHECK (run at 7:00 AM):
```

```
FOR EACH ACTIVE CLIENT FILE:
1. Check IRCC portal / USCIS case status for updates
   - Flag any status change (new stage, request for docs, decision)

2. Compute deadline windows:
   - Biometrics: appointment scheduled? Completed? Window expiring?
   - Medical certificate: exam date + 12-month validity
     -> flag at 60 days before expiry if application still pending
   - Hearing date: days remaining, filing deadlines for supporting
materials
   - Work permit: current expiry date, renewal filing deadline
   - BOC (Basis of Claim): 15-day deadline from referral

3. If status changed:
   - Draft client update: "[Client name] -- your application status
     has changed to [new status]. Here's what this means and
     what happens next: [plain language explanation]."
   - Compute any new deadlines triggered by the status change
   - Update case management system

4. Daily summary:
   - Applications with status changes: [list]
   - Deadlines within 30 days: [list with days remaining]
   - Biometrics/medical expiring: [list]
   - No changes (but monitored): [count] files checked

IMPORTANT: Portal status information should be verified by the
practitioner. Deadline computations for hearing materials and
BOC submissions must be confirmed against the specific
tribunal's scheduling order or referral notice.
```

---

20. Estate Administration Lifecycle Agent

Source: Chapter 6, “Estates & Probate” **Category:** Operations

Generates a complete estate administration lifecycle timeline from date of death through final distribution. Tracks sequential deadlines for beneficiary notification, probate application, creditor notice, tax returns, and distribution across a 12-month period.

```
You are an estate administration lifecycle agent.

NEW ESTATE FILE OPENED:
- Deceased: [name]
- Date of death: [date]
- Jurisdiction: [province/state]
- Executor: [name, contact]
- Estimated estate value: [amount]
- Will located: [yes/no]

GENERATE LIFECYCLE TIMELINE:

1. IMMEDIATE (Days 1-14):
   - [ ] Obtain death certificate (original + [N] certified copies)
   - [ ] Locate and secure the will
```

```
   - [ ] Notify beneficiaries (deadline: [jurisdiction-specific,
typically 10-30 days])
   - [ ] Draft beneficiary notification letters
   - [ ] Secure estate assets (change locks, redirect mail, freeze
accounts)

2. PROBATE APPLICATION (Days 14-60):
   - [ ] Prepare estate inventory (assets + liabilities)
   - [ ] File probate application / Letters of Administration
   - [ ] Publish creditor notification (deadline: [jurisdiction-
specific])
   - [ ] Notify known creditors directly

3. ADMINISTRATION (Days 60-270):
   - [ ] Collect estate assets
   - [ ] Pay debts and administration expenses
   - [ ] File tax returns (terminal return deadline: [date])
   - [ ] Obtain tax clearance certificate
   - [ ] Prepare interim accounting if required

4. DISTRIBUTION (Days 270-365):
   - [ ] Prepare final estate accounting
   - [ ] Obtain beneficiary approval of accounting
   - [ ] File court accounting if required
   - [ ] Distribute estate assets per will / intestacy rules
   - [ ] Close estate accounts

FOR EACH ITEM: Track status (pending/in-progress/complete),
responsible party, deadline, and any dependencies.

IMPORTANT: Estate administration timelines and requirements vary
significantly by jurisdiction. All deadlines must be verified
against the applicable provincial/state rules. Tax return
deadlines should be confirmed with the estate's accountant.
```

---

21. Collaborative Law Session Follow-Up Agent

Source: Chapter 6, “Family & Collaborative Law” **Category:** Operations

Generates the complete post-session paper trail for collaborative family law proceedings: session summary, updated disclosure tracker, revised agreement sections, action items by party, distribution email drafts, and next session agenda.

```
You are a collaborative law session follow-up agent.

SESSION COMPLETED:
- Case: [case name]
- Date: [session date]
- Participants:
  - [Spouse 1 name] + [Attorney 1 name]
  - [Spouse 2 name] + [Attorney 2 name]
  - [Coach name, if applicable]
  - [Financial neutral name, if applicable]
- Session notes / transcript: [attached or referenced]

GENERATE:
```

```
1. SESSION SUMMARY (distribute to all participants):
   - Topics discussed (bulleted, factual, neutral tone)
   - Agreements reached (specific terms, not generalities)
   - Open items requiring further discussion
   - Next steps for each party

2. UPDATED DISCLOSURE TRACKER:
   - New financial information exchanged this session
   - Outstanding disclosure items by party
   - Deadlines for outstanding items

3. REVISED AGREEMENT SECTIONS:
   - Draft updated language for any terms agreed upon
   - Track changes from prior version
   - Flag any terms that were discussed but not finalized
     (mark as "UNDER DISCUSSION -- NOT AGREED")

4. ACTION ITEMS BY PARTY:
   - [Spouse 1]: [list]
   - [Spouse 2]: [list]
   - [Attorney 1]: [list]
   - [Attorney 2]: [list]
   - [Coach]: [list]
   - [Financial neutral]: [list]

5. DISTRIBUTION DRAFTS:
   - Draft email to each participant with their relevant sections
   - Each email includes: session summary, their action items,
     disclosure reminders specific to them

6. NEXT SESSION AGENDA (draft):
   - Open items carried forward
   - New items flagged during this session
   - Disclosure items due before next session

IMPORTANT: Session summaries in collaborative practice are shared
documents. Tone must be neutral and factual. Do not characterize
either party's position or emotional state. Agreements must be
clearly distinguished from items still under discussion.
Attorney must review all outputs before distribution.
```

---

23. Inbox Triage Classification Rules

Source: Chapter 7, “Agent 2: Inbox Triage” **Category:** Operations

Classification rules for an inbox triage agent. Defines how to categorize emails by examining sender domain, subject line patterns, and message content, then sort into Urgent/Action/FYI priority tiers.

```
### Classify each message

For each unread message in the inbox, classify by examining:
- Sender domain (court domains, known opposing counsel, client
emails)
- Subject line patterns ("ORDER", "NOTICE", "RE: [matter name]")
```

```
- Message content (filing references, case numbers, deadline
language)

Priority tiers:
- **Urgent:** Court orders, motions requiring response, client
  emergencies, messages from judges
- **Action:** Client questions, opposing counsel correspondence,
  scheduling requests, items requiring a reply within 48 hours
- **FYI:** Newsletters, vendor updates, automated notifications,
  marketing emails, CLE announcements

Apply labels: "Triage/Urgent", "Triage/Action", "Triage/FYI"
Archive all FYI items. Leave Urgent and Action in inbox.
```

Prompts by Chapter

For quick reference, here is the count of prompts originating from each chapter:

Chapter	Prompts
Ch 1: The 1% Problem	3
Ch 2: What AI-Native Looks Like	0
Ch 3: The Architecture	3
Ch 4: Dodonai	0
Ch 5: From Documents to Strategy	6
Ch 6: Seven Practice Areas	7
Ch 7: The Agent Playbook	4
Ch 8: Marketing, Intake, and Growth	4
Ch 9: Operations	4
Ch 10: The Safety Layer	0
Ch 11: Getting Started	2
Total	**35**

Appendix B: Dodonai Product Reference

This appendix is a quick-reference card for each Dodonai product. Keep it bookmarked for when you need to check what a product does, what it costs, or what formats it supports.

All products are available at **app.dodon.ai**. Dodonai is SOC 2 certified and HIPAA compliant. Zero data retention on all LLM processing. Your documents are never used to train AI models.

Deposition Summaries

What it does: Generates page-line, narrative, and issue-based deposition summaries from uploaded transcripts. A 100-page transcript produces a finished summary in about 3 minutes. Replaces outsourced deposition summary services that typically cost $300-$2,000 per transcript and take days to deliver.

Key features:

- **Three summary formats:** page-line (chronological with citations), narrative (flowing overview), and issue-based (grouped by topic)
- **Key topic highlighting** – specify topics at upload and the summary gives them extra focus where relevant
- **Interactive viewer** with side-by-side source linking
- **Downloadable** as PDF, Word, or TXT
- **Saved to your account** for retrieval anytime
- **Full-service option** – Dodonai's team produces summaries at your direction ($1/page, 24-hour turnaround)

Formats:

Direction	Supported
Input	PDF, TXT, DOC, DOCX, CSV, PPTX
Output	PDF, Word, TXT

Speed:

Transcript Length	Approximate Time
37 pages	~60 seconds
64 pages	~84 seconds

Transcript Length	Approximate Time
100 pages	~3 minutes

Pricing: 1 credit per page.

Tips:

1. Specify key topics when you upload. If you're summarizing for a PI case, enter the injury mechanism and key medical terms. The summary will flag those sections.
2. Use page-line format for motion practice and trial prep (it's the format most judges and opposing counsel expect). Use narrative for quick case assessment. Use issue-based for complex commercial disputes with multiple claims.
3. Run the summary before your deposition review session. Having a structured summary in hand before you annotate the transcript cuts your prep time significantly.

Medical Record Summaries & Chronologies

What it does: Processes medical records and produces both a narrative summary of the patient's healthcare history and a date-ordered chronology of clinical events. 500 pages process in under 15 minutes.

Key features:

- **Dual output** – narrative summary + chronological timeline from a single upload
- **Key event extraction** – diagnoses, surgeries, prescriptions, provider visits, lab results, imaging findings
- **Interactive viewer** with side-by-side source linking to original records
- **Sort and filter** chronology entries by provider or date
- **Custom topic highlighting** – enter specific diagnoses, providers, or conditions to emphasize
- **Index tables** – interactive, sortable tables for navigating large record sets
- **Visual graphic timeline** (advanced)
- **Reusable workflow templates** (advanced)
- **Full-service option** – upload records and let Dodonai's team produce the summary ($1/page, 24-hour turnaround)

Formats:

Direction	Supported
Input	PDF (including scanned documents), multi-hundred-page files
Output	PDF, Word, XLS, TXT

Speed:

Record Length	Approximate Time
200 pages	~5 minutes
500 pages	<15 minutes

Pricing: 1 credit per page.

Tips:

1. For PI cases, enter the injury date and key diagnoses as topics. The chronology will highlight gaps in treatment and pre-existing conditions, which are exactly what opposing counsel will target.
2. Download the XLS format when you need to share chronology data with experts. Physicians and life care planners prefer spreadsheets they can sort and annotate.
3. Use the sort-by-provider view to quickly identify which treaters have the most entries. That tells you who your key medical witnesses are before you start scheduling depositions.

Transcript Management

What it does: Centralizes your deposition transcript library with semantic search, AI-powered chat analysis, annotation tools, and automated Blue Book citations. Upload transcripts once, then search, highlight, and cite across your entire collection. Eliminates ctrl-F searching across separate PDF files and manual citation formatting.

Key features:

- **Semantic search** – search by meaning across all uploaded transcripts, not just keyword matching
- **AI Chat/Copilot** – ask questions about transcript content and get AI-assisted analysis with highlighted source passages
- **One-click Blue Book citations** – highlight any passage and get a properly formatted citation ready to paste into briefs, motions, or correspondence
- **Cross-transcript search** – search across all transcripts in a case simultaneously
- **Format conversion** – convert between PDF, TXT, and Word at any time
- **Annotation tools** – highlight, bookmark, and annotate directly in the viewer

Formats:

Direction	Supported
Input	PDF, TXT, DOC, DOCX, CSV, PPTX
Output	PDF, Word, TXT (format conversion)

Speed: Search results return in seconds, even across large transcript collections.

Pricing: Included with all plans. No additional credit cost for search, chat, or annotations.

Tips:

1. Upload all transcripts for a case as soon as you receive them. Cross-transcript search is most valuable when the full deposition record is in the system. You can find where Witness A contradicts Witness B without remembering which transcript to check.
2. Use the AI Copilot to prep for follow-up depositions. Ask it to identify every mention of a specific event or document across all depositions taken so far.
3. When drafting motions, highlight the passage you need and copy the Blue Book citation directly. No more manual page-line formatting.

OCR

What it does: Converts scanned PDFs, handwritten documents, tables, and images into editable, searchable text using AI-powered recognition. Significantly more accurate than legacy OCR on poor scans, mixed formatting, and handwriting. Unlike traditional OCR that processes characters individually, Dodonai's AI OCR understands document context to produce cleaner, more reliable output.

Key features:

- **Handwriting recognition** – transcribes handwritten notes and annotations with higher accuracy than legacy OCR
- **Table detection and formatting** – identifies tables within documents and preserves their structure
- **Image interpretation** – provides meaningful descriptions of images embedded in documents
- **Context-aware extraction** – understands document context, not just individual characters
- **Multi-document PDF indexing** – automatically indexes and labels individual documents within multi-page PDFs
- **Direct integration** with Dodonai's summarization, chronology, and eDiscovery tools

Formats:

Direction	Supported
Input	PDF, TIFF (scanned and digital)
Output	Word, TXT, PDF (searchable)

Speed: Processes hundreds of pages per hour. Processing time varies based on scan quality and document complexity.

Pricing: Free with all plans. No credit cost.

Tips:

1. Run OCR first on any scanned document production before using eDiscovery search. Dodonai's eDiscovery will auto-OCR documents, but pre-processing lets you verify the text extraction quality on critical documents.
2. For medical records received as scanned faxes (common with older provider offices), OCR will handle the degraded quality far better than legacy tools. Check the output on a few pages before running the full set through summarization.
3. OCR output feeds directly into all other Dodonai products. Once a scanned document is OCR-processed, it's fully searchable in eDiscovery and ready for summarization without re-uploading.

---

eDiscovery

What it does: Searches across document sets using semantic (meaning-based) search instead of keyword matching. Describe what you're looking for in plain language, get ranked results with source citations, and click through to verify in the original document. Where legacy eDiscovery platforms like Relativity and Everlaw charge $50,000+/year, Dodonai starts at $30/month.

Key features:

- **Semantic concept search** – finds relevant documents even when they use different terminology than your query
- **Ranked results** with relevance scoring and highlighted matching passages
- **Source verification** – every result links directly to the passage in the original document
- **Audit trail** – defensible review methodology you can document for the court
- **Automatic OCR** – scanned and image-based documents are processed automatically before indexing
- **Zero manual tagging** – upload documents and start searching immediately, no coding or classification required
- **Cross-format search** – search across PDFs, Word documents, emails, text files, and scanned documents in a single query

Formats:

Direction	Supported
Input	PDF, Word, TXT, email files, ZIP archives, scanned documents (auto-OCR)
Output	Ranked search results with source links; export search results

Speed: Search results return in seconds across document sets of any size.

Pricing: Included with all plans. No additional credit cost for searching.

Tips:

1. Write your search queries in natural language, not keywords. Instead of searching for "breach" AND "contract" AND "2024," describe what you're looking for: "communications about the defendant's failure to deliver under the supply agreement." Semantic search handles the rest.
2. Use eDiscovery early in case assessment. Upload the initial document production and run broad concept searches to identify hot documents before committing to a full manual review.
3. For privilege reviews, search for "attorney-client communications" or "legal advice regarding [topic]" to quickly flag potentially privileged documents for attorney review.

Extract & Draft Agents (XD Agents)

What it does: Automates data extraction and report generation using reusable templates you define. Specify the fields and format you need, point the agent at source documents, and it extracts the data and assembles the report. Build once, reuse indefinitely. Especially valuable for firms that process the same document types repeatedly: intake forms, insurance declarations, demand packages, settlement statements.

Key features:

- **Custom extraction templates** – define the data fields, sections, structure, and formatting you need
- **Multi-source extraction** – pull data from multiple source documents into a single report
- **Reusable templates** – build a template once, run it against any number of source documents
- **AI-powered understanding** – extracts based on meaning, not field position, so it handles documents with varying formats and layouts
- **Human-in-the-loop editing** – adjust extracted data directly in the platform before exporting
- **Export** finalized reports in your preferred format
- **Integration** with existing document management workflows

Formats:

Direction	Supported
Input	PDF, Word, TXT, spreadsheets, scanned documents (auto-OCR)
Output	Customizable report format (PDF, Word, TXT)

Speed: Extraction runs in minutes. Speed depends on the number of source documents and template complexity.

Pricing: 2 credits per page.

Tips:

1. Start with your most repetitive document task. If you extract the same data from intake forms, insurance declarations, or demand packages every week, that's the template to build first. The ROI is immediate once you stop doing it manually.
2. Use multi-source extraction for demand package assembly. Point the agent at the medical chronology, billing records, and lost wages documentation, and let it populate your demand template from all three.
3. After the first extraction, review the output carefully and refine your template. Tightening the field definitions in your template improves accuracy on every subsequent run.

Pricing Tiers

All Dodonai products are included in every plan. The difference is credit volume.

Plan	Monthly Price	Annual Price	Monthly Credits	Annual Credits
Lite	$30/mo	$300/yr ($25/mo)	200	4,800
Basic	$60/mo	$600/yr ($50/mo)	1,000	24,000
Pro	$100/mo	$1,000/yr ($84/mo)	5,000	120,000
Enterprise	Custom	Custom	Unlimited	Unlimited

How credits work:

- 1 credit = 1 page (~400 words)
- Deposition Summaries and Medical Records: **1 credit/page**
- XD Agents: **2 credits/page**
- Transcript Management, eDiscovery, and OCR: **Included (no credit cost)**
- Annual credits roll over month to month
- No per-seat fees – your whole team uses one account
- As low as **$0.02/page** on Pro annual

Included with all plans:

- AI deposition summaries

- Medical record chronologies
- AI-powered OCR
- Extract & Draft Agents
- Custom rule engine
- eDiscovery tools
- Transcript management
- Chat and email support

Pro and Enterprise add:

- Dedicated onboarding
- Phone support
- Custom integrations (Enterprise)
- SLA guarantees (Enterprise)

Choosing a plan: A solo practitioner handling 10-15 PI cases processes roughly 200-500 pages of depositions and medical records per month. Lite covers light usage; Basic covers most solo and small-firm workloads; Pro handles high-volume litigation teams. If you're processing more than 5,000 pages per month or need custom integrations, contact Dodonai for Enterprise pricing.

For the latest pricing and feature details, visit dodon.ai/pricing or contact hello@dodon.ai.

Appendix C: Tools & Integrations

This appendix covers the tools most commonly connected to AI agents in a law firm environment. For each, you'll find what data flows through it, how agents connect, and practical notes from real implementations.

You don't need all of these on day one. Most firms start with Clio and Gmail, add calendar and Dodonai outputs, then layer in financial tools and Slack as their agent system matures.

---

1. Clio

What it is: Cloud-based practice management. Cases, contacts, billing, calendars, documents, tasks, and time entries in one system. Clio is the most widely used practice management software among solo and small firms (Clio Legal Trends 2026).

What data flows through it:

Data Type	Examples
Cases/Matters	Case name, status, practice area, responsible attorney, open date, SOL date
Contacts	Clients, opposing counsel, adjusters, expert witnesses, judges
Calendar entries	Court dates, deadlines, mediations, depositions, SOL expirations
Tasks	Follow-ups, filing deadlines, document requests, review assignments
Time entries	Billable hours, activity descriptions, billing rates
Billing	Invoices, payments, trust balances, AR aging
Documents	Filed documents, correspondence, pleadings (via Clio Manage)

How agents connect: Clio provides a REST API with OAuth 2.0 authentication. Your AI agent reads and writes data through standard HTTP requests. The API covers matters, contacts, activities, tasks, calendar entries, bills, and documents.

What agents can do:

- **Read** case status, deadlines, contact information, billing history, time entries, and task lists
- **Create** tasks, time entries, calendar events, and notes
- **Update** matter status, task completion, and contact information
- **Search** across matters, contacts, and activities by keyword or date range

Practical notes:

- Clio's API requires a registered application with OAuth credentials. You'll need to create an app in the Clio Developer portal and authorize it for your firm's account.
- Rate limits are generous for single-firm use (typically 100 requests/minute). A morning health-check agent that scans 40 active matters won't come close.
- Clio's webhook support lets agents react to changes in real time. When a matter status changes or a deadline approaches, the webhook fires and your agent can take action (send an alert, draft an email, update a task list).
- **Start here:** A read-only agent that pulls this week's deadlines and upcoming court dates every morning, formats them as a daily briefing, and delivers it to Slack or email. No writes, no risk, immediate value.
- Clio Manage (document management) has more limited API access than the core Clio features. Document upload and retrieval work, but some advanced document management features require the web interface.

API documentation: developer.clio.com

---

2. Gmail / Outlook (Email)

What it is: Your firm's email system. Gmail (Google Workspace) or Outlook (Microsoft 365) are where client communications, opposing counsel correspondence, court notices, and vendor updates land.

What data flows through it:

Data Type	Examples
Incoming email	Client questions, opposing counsel motions, court notices, adjuster updates
Outgoing email	Client updates, demand letters, discovery responses, scheduling
Drafts	AI-prepared responses awaiting attorney review
Labels/Folders	Triage categories, matter-specific organization
Attachments	Filed documents, medical records, expert reports

How agents connect:

- **Gmail:** Google's Gmail API with OAuth 2.0. Supports reading messages, creating drafts, applying labels, archiving, and searching. The gmail.modify scope lets agents label, archive, and create drafts without send or delete permission.
- **Outlook:** Microsoft Graph API with OAuth 2.0. Similar capabilities: read, draft, categorize, move, search.

What agents can do:

- **Triage incoming email** – classify messages by priority and category (client, opposing counsel, court, vendor, newsletter), apply labels, archive low-priority items
- **Detect outreach replies** – monitor for responses to previously sent emails and update tracking logs
- **Create drafts** – prepare response drafts for attorney review (never auto-send)
- **Search** – find specific correspondence by sender, date, subject, or content
- **Extract data** – pull court dates from notices, settlement figures from demand letters, deadlines from motions

Practical notes:

- **Draft-only is the right default.** AI agents should create drafts for attorney review, not send emails autonomously. This is both an ethical safeguard and a practical one. Attorneys must review outgoing communications before they leave the firm.
- Gmail's label system maps well to agent triage. Create a label taxonomy (Urgent, Action Required, FYI, by matter) and let the agent classify incoming mail against it.
- Outlook's categories and folders serve the same purpose. The Microsoft Graph API uses slightly different terminology but the same concepts.
- For high-volume firms, an inbox triage agent running twice daily (morning and afternoon) catches time-sensitive items you'd otherwise miss until evening.
- **Attachment handling** requires separate API calls to download and process attachments. For court filings received as email attachments, agents can extract the filing, OCR it if needed, and update your case management system.

---

3. Google Calendar / Outlook Calendar

What it is: Your firm's scheduling system. Court dates, depositions, mediations, client meetings, SOL deadlines, and internal reviews.

What data flows through it:

Data Type	Examples
Events	Court hearings, depositions, mediations, client meetings, IMEs
Deadlines	SOL expirations, discovery cutoffs, motion response deadlines, filing deadlines
Attendees	Client names, opposing counsel, judges, expert witnesses, adjusters

Data Type	Examples
Locations	Courtrooms, conference rooms, virtual meeting links

How agents connect:

- **Google Calendar:** Google Calendar API with OAuth 2.0. Read events, create events, update events, check free/busy.
- **Outlook Calendar:** Microsoft Graph API (same OAuth flow as email). Same capabilities.

What agents can do:

- **Morning briefing** – pull today's and tomorrow's meetings, enrich attendee information from Clio and web search, deliver a prep document before each meeting
- **Deadline monitoring** – scan for approaching deadlines (SOL, discovery cutoffs, motion responses) and send alerts at configurable intervals (30 days, 14 days, 7 days, 48 hours)
- **Meeting prep** – for each external meeting, pull the attendee's case history, recent correspondence, and open tasks, then compile a one-page briefing
- **Scheduling** – find available time slots across attorney calendars for depositions, mediations, or client meetings
- **Post-meeting follow-up** – after a meeting ends, create follow-up tasks in Clio and draft a confirmation email to the client

Practical notes:

- Calendar agents are among the highest-ROI integrations. A deadline monitoring agent that alerts you to approaching SOL expirations and filing deadlines is essentially a malpractice insurance supplement.
- Sync your Clio calendar to Google/Outlook if they aren't already connected. Most agents read from one calendar system; having Clio deadlines visible in that system ensures nothing falls through.
- For meeting prep, the agent needs access to both the calendar (who is attending) and your case management system (what's the context). This is a two-tool integration, but the combination is where the value lives.
- **Court date monitoring is passive, not active.** Your agent can alert you to upcoming deadlines, but calculating legal deadlines (counting backward from hearing dates, accounting for service rules, holidays, and local rules) requires verified legal logic. Always confirm deadline calculations against your jurisdiction's rules.

---

4. QuickBooks

What it is: Accounting software. Invoicing, expense tracking, profit/loss, trust accounting (with add-ons), and financial reporting.

What data flows through it:

Data Type	Examples
Invoices	Client billing, amount, status (paid/outstanding/overdue)
Payments	Payment receipts, payment methods, dates
Expenses	Operating costs, vendor payments, marketing spend
Accounts	Trust accounts, operating accounts, IOLTA balances
Reports	P&L, AR aging, cash flow, profitability by matter

How agents connect: QuickBooks Online REST API with OAuth 2.0. Read and write access to invoices, payments, customers, expenses, and reports.

What agents can do:

- **AR monitoring** – flag invoices overdue by 30, 60, 90+ days and draft collection follow-ups
- **Profitability analysis** – cross-reference Clio time entries with QuickBooks billing to calculate effective hourly rate and profitability by matter, client, or practice area
- **Expense tracking** – categorize expenses, flag anomalies, generate monthly spending summaries
- **Cash flow forecasting** – analyze payment patterns to project cash flow for the next 30-60-90 days
- **Financial reporting** – generate weekly or monthly summaries for the managing partner without manual spreadsheet work

Practical notes:

- QuickBooks' API is mature and well-documented. If your firm already uses QuickBooks Online, the API connection is straightforward.
- Trust accounting has specific compliance requirements (state bar rules on IOLTA handling). An agent can monitor trust balances and flag issues, but trust account transactions should always have attorney approval before execution.
- The highest-value QuickBooks agent for most firms is an AR aging monitor that runs weekly, identifies overdue invoices, and drafts polite follow-up emails. Firms that automate AR follow-up consistently collect faster.
- **Cross-reference with Clio** for the real insight: time entries in Clio + invoices in QuickBooks = profitability by matter. An agent that runs this analysis monthly tells you which cases are profitable and which are bleeding margin.

---

5. Dodonai

What it is: AI document processing platform. Generates structured outputs (summaries, chronologies, search results, extracted data) from uploaded documents.

What data flows through it:

Data Type	Examples
Document uploads	Deposition transcripts, medical records, discovery productions, contracts
Processed outputs	Page-line summaries, narrative summaries, medical chronologies, search results
Extracted data	XD Agent outputs (custom-templated data extraction from source documents)
Search results	eDiscovery semantic search hits with source citations

How agents connect: Dodonai outputs are structured documents (PDF, Word, TXT, XLS) that serve as input to AI agents. Upload a medical chronology or deposition summary to your AI assistant's context, and it becomes a data source for further analysis, draft preparation, or case assessment.

What agents can do with Dodonai outputs:

- **Demand package drafting** – feed the medical chronology and deposition summaries to your AI assistant as context for drafting the damages narrative, treatment timeline, and liability analysis sections
- **Case health assessment** – combine medical chronology (treatment gaps, provider count) with deposition summary (key admissions, liability facts) to generate a case strength evaluation
- **Expert witness prep** – extract key findings from the medical chronology and deposition testimony to build a concise briefing document for expert review
- **Cross-referencing** – compare what a deponent testified to against what the medical records show, flagging inconsistencies

Practical notes:

- Dodonai's outputs are designed to be consumed by both humans and AI. The structured format (chronological entries, page-line citations, labeled sections) makes them reliable context for AI analysis.
- For best results, feed the AI assistant the specific Dodonai output that's relevant to the task, not the raw source documents. A 500-page medical record is unwieldy; the 15-page chronology Dodonai produces from it is exactly what the AI needs.
- XD Agent templates are where automation compounds. Once you build a template for extracting data from (for example) insurance declarations,

every subsequent declaration runs through the same template with no additional setup.

- **Think of Dodonai as the document processing layer and your AI assistant as the reasoning layer.** Dodonai turns raw documents into structured data. Your AI assistant turns structured data into analysis, drafts, and decisions.

---

6. Slack

What it is: Team messaging platform. Channels, direct messages, and integrations. Increasingly used by law firms for internal communication and agent output delivery.

What data flows through it:

Data Type	Examples
Agent alerts	Deadline warnings, churn detection, new filing notifications
Digests	Morning briefings, weekly health reports, inbox triage summaries
Action items	Tasks requiring attorney decision, draft approvals, client follow-ups
Team communication	Case discussions, status updates, document sharing

How agents connect: Slack's Incoming Webhooks allow agents to post messages to specific channels or direct messages with a single HTTP POST request. For more complex interactions, the Slack API (Bot tokens, OAuth) supports reading messages, reacting, threading, and interactive buttons.

What agents can do:

- **Deliver digests** – post morning briefings, weekly health reports, and inbox triage summaries to a dedicated channel or DM
- **Send alerts** – notify attorneys of approaching deadlines, new court filings, client cancellations, or urgent emails
- **Request decisions** – post action items that require attorney input (approve a draft, review a deadline, respond to a client)
- **Thread updates** – post follow-up information in threads to keep conversations organized

Practical notes:

- Incoming Webhooks are the simplest starting point. Create a webhook URL for a channel, and any agent can post to it with a single API call. No OAuth flow, no bot setup.
- **Channel design matters.** Don't dump all agent output into one channel. Use dedicated channels: #deadlines, #inbox-digest, #billing-alerts, #case-

updates. This lets attorneys subscribe to what matters and mute what doesn't.
- For solo practitioners, direct messages work better than channels. The agent posts your daily briefing as a DM, and it's waiting for you when you open Slack in the morning.
- Slack is the delivery layer, not the processing layer. Agents do their work (scanning Clio, triaging email, running reports) and post the results to Slack. Keep the messages concise and actionable. If the full report is 2 pages, post a 3-line summary in Slack with a link to the full document.
- **Microsoft Teams** serves the same role for firms on Microsoft 365. The integration pattern is similar (webhooks for simple posting, Bot Framework for interactive features), though the API is more complex than Slack's.

---

7. Court Filing Systems

What it is: Electronic filing portals. CM/ECF (federal courts), state e-filing systems (varies by jurisdiction), and local court portals. These are where filings are submitted, docketed, and served.

What data flows through it:

Data Type	Examples
Filed documents	Motions, briefs, orders, judgments, notices
Docket entries	Filing dates, hearing schedules, case assignments
Service notifications	Electronic service receipts, notices of electronic filing (NEFs)
Deadlines	Response dates triggered by filings, scheduling orders

How agents connect: This is where you need to be honest about limitations.

- **CM/ECF (PACER):** Federal courts use CM/ECF for filing and PACER for retrieval. PACER has an API (PACER Case Locator) that supports basic case search, but full docket retrieval and document download remain limited. Third-party services like CourtListener, RECAP, and Docketbird provide structured access to PACER data.
- **State e-filing portals:** Most state court e-filing systems do not provide APIs. Access is through web portals with manual login. Some jurisdictions have modernized (e.g., Tyler Technologies' Odyssey platform powers many state courts), but API access is typically restricted to authorized e-filing service providers.
- **Email monitoring:** The most reliable automated connection to court filing systems is monitoring the notification emails they send. CM/ECF sends NEFs (Notices of Electronic Filing) by email. State courts send similar

notifications. An email triage agent can detect these, extract the filing type and deadline, and create corresponding calendar entries and tasks.

What agents can realistically do:

- **Monitor NEF emails** – detect new filings from email notifications, extract filing type and relevant dates
- **Calculate response deadlines** – from a detected filing, calculate the response deadline (with appropriate caveats about jurisdiction-specific rules)
- **Download filed documents** – via PACER for federal cases (cost: $0.10/page), or from email attachments where courts attach the filing
- **Update case management** – create tasks in Clio for required responses, add hearing dates to the calendar
- **Docket monitoring** via third-party services (CourtListener, Docketbird) that provide structured data feeds

What agents cannot reliably do:

- **File documents** – e-filing requires authenticated login through the court's portal. No court system allows filing via a third-party API without being an authorized e-filing service provider.
- **Access sealed or restricted documents** – these require authenticated court portal access.
- **Replace a docket monitoring service** – for comprehensive docket tracking, commercial services (Lex Machina, CourtListener, PacerPro) provide more complete and reliable data than DIY scraping.

Practical notes:

- Email monitoring is the pragmatic starting point. Your agent watches for NEFs and court notices in your inbox, extracts the relevant information, and updates your calendar and task list. This alone prevents missed deadlines.
- **Deadline calculation is high-stakes.** An agent can calculate "motion filed on [date], response due in 14 days," but computing legal deadlines correctly requires accounting for service method, weekends, holidays, local rules, and scheduling orders. Always verify agent-calculated deadlines against your jurisdiction's rules. This is a "trust but verify" workflow, not a "set and forget" automation.
- Third-party docket monitoring services are worth the cost for active litigators. They provide structured data, push notifications, and historical docket analytics that are difficult to replicate with DIY solutions.
- The court filing ecosystem is the least API-friendly part of legal technology. Plan for manual steps in this workflow and focus your automation on what happens after a filing is detected (deadline tracking, task creation, document processing) rather than on the filing system itself.

Getting Started: Your First API Connection

If you're setting up tool integrations for the first time, start with **Clio**. Here's why:

1. **Clio has the richest data.** Cases, contacts, deadlines, billing, tasks, documents – it's the hub of your practice. An agent connected to Clio can answer more questions about your firm than any other single data source.
2. **The API is well-documented and well-supported.** Clio's developer portal has clear documentation, a sandbox environment for testing, and a responsive support team. You won't be debugging undocumented behavior.
3. **Read-only is valuable from day one.** You don't need to write data back to Clio to get value. A read-only agent that pulls your weekly deadlines, active case counts, and billing summaries provides immediate ROI with zero risk to your data.
4. **It connects to everything else.** Once Clio data is available to your agents, you can cross-reference it with email (who's writing about which case), calendar (what's coming up on which matter), and financials (which cases are profitable). Clio is the connective tissue.

Step-by-step for Clio:

1. **Create a developer account** at developer.clio.com.
2. **Register an application** – choose "private app" for internal use. You'll get a client ID and client secret.
3. **Authorize the app** – follow Clio's OAuth flow to connect your firm's Clio account to your application. This generates an access token.
4. **Start with a read** – make your first API call: GET /api/v4/matters.json?status=open&fields=id,display_number,description,status,practice_area. This returns your open matters.
5. **Build from there** – add calendar entries (/api/v4/calendar_entries.json), contacts, and tasks as you need them.

If you're not technical: This is where Dodonai Consulting comes in. The integration setup – API credentials, authentication flows, data mapping, agent configuration – is exactly what Dodonai's consulting team handles. You describe what you want the agent to do; they build and maintain the connections.

Recommended integration order:

Priority	Tool	First Agent
1	Clio	Weekly case health briefing

Priority	Tool	First Agent
2	Gmail/Outlook	Inbox triage and priority classification
3	Google/Outlook Calendar	Deadline monitoring and meeting prep
4	Dodonai	Document output analysis and draft prep
5	Slack	Agent output delivery and alerts
6	QuickBooks	AR monitoring and profitability reporting
7	Court filing systems	NEF monitoring and deadline tracking

Start with one. Get it working. Then add the next. The compounding value comes when agents can cross-reference data across tools – but each tool is useful on its own before you connect the others.

This appendix covers the most common integrations for small and mid-size litigation firms. For specialized tools (document management systems, practice-specific databases, state bar portals), contact Dodonai Consulting at hello@dodon.ai.

Appendix D: Sample Agent Outputs

This appendix shows what AI agents actually produce – realistic sample outputs that demonstrate the quality readers can expect. Each output is sourced from a prompt or agent described elsewhere in the book. For the full prompts, see Appendix A.

Chapter 2 includes three sample outputs (meeting intel briefing, case health digest, inbox triage summary). Chapter 7 includes excerpts from the Deadline Monitor, Inbox Triage, Client Update Drafter, Billing Reconciler, and Case Health Dashboard. The outputs below are different – they cover strategic document analysis, operational workflows, and marketing content that appear throughout the book but weren't shown in full.

All examples use Sarah Chen's practice at Chen Law Group in Austin, TX (40 active matters, PI/med mal/workers' comp).

---

1. Damages Narrative from Medical Chronology

Source: Chapter 5, Workflow 1 **Agent/Workflow:** /demand-narrative skill – medical chronology to demand letter section **Prompt used:** Two-prompt chain (damages narrative + medical specials). See Appendix A, Workflow 1.

> **INJURY SUMMARY**
>
> On March 15, 2025, Maria Gonzalez sustained cervical disc herniation at C5-C6, lumbar disc protrusion at L4-L5 with left-sided radiculopathy, and left shoulder impingement as a result of a rear-end collision on IH-35 in Austin, Texas. MRI findings at Austin Diagnostic Imaging on April 3, 2025 confirmed the cervical herniation and lumbar protrusion, with neuroforaminal narrowing at L4-L5 impinging the traversing nerve root. Over 14 months of treatment across four providers, Ms. Gonzalez progressed from conservative care through epidural steroid injections, with her treating orthopedist recommending surgical consultation for an anterior cervical discectomy and fusion. At her most recent appointment on May 8, 2026, Dr. Ravi Mehta assigned a whole-person permanent impairment rating of 12% under the AMA Guides, 6th Edition.
>
> **TREATMENT TIMELINE**

Ms. Gonzalez presented to St. David's Medical Center ER on March 15, 2025, reporting acute neck pain, lower back pain with radiating symptoms into her left leg, and left shoulder pain. ER records document cervical and lumbar muscle spasm, positive straight-leg raise on the left, and decreased left grip strength (Chronology: 03/15/2025). She was prescribed a Medrol dose pack and Flexeril.

Chiropractic care began March 22, 2025 with Dr. Lisa Okonkwo at Capital Chiropractic (3x/week, 48 visits over four months). After plateauing, Dr. Okonkwo referred to Austin Spine Associates (Chronology: 03/22/2025 - 07/18/2025).

Dr. Mehta's evaluation on August 1, 2025 confirmed persistent C5-C6 radicular symptoms. Following MRI confirmation, he initiated three cervical epidural steroid injections at St. David's Surgical Center on September 12, October 10, and November 7, 2025. Ms. Gonzalez reported temporary relief of 40-60% after each injection, with symptoms recurring within two to three weeks.

Concurrent physical therapy at Capital PT (16 sessions, September - December 2025) focused on cervical stabilization and lumbar core strengthening.

At follow-up on January 15, 2026, Dr. Mehta documented that injections provided only temporary relief and recommended surgical consultation for ACDF at C5-C6. Ms. Gonzalez has deferred the surgical decision pending resolution of this claim.

IMPACT ON DAILY LIFE

Functional capacity evaluation (07/14/2025) documented: unable to lift more than 15 pounds, cannot sit longer than 30 minutes without repositioning, difficulty with overhead tasks. As an administrative assistant, these limitations directly affect job performance. Ms. Gonzalez reports she can no longer carry her two-year-old daughter for extended periods.

PROGNOSIS

Dr. Mehta's May 8, 2026 report assigns a 12% whole-person permanent impairment rating. Ms. Gonzalez will require either surgical intervention or ongoing pain management. If surgery is elected, estimated 6-8 month recovery with expected improvement but some degree of permanent cervical limitation.

GAPS AND FLAGS

- **45-day treatment gap** between final PT session (12/19/2025) and January follow-up with Dr. Mehta (01/15/2026). Defense will cite this to argue symptom resolution. Address proactively – gap coincides with the holiday period and FCE documented ongoing symptoms throughout.
- No pre-existing conditions documented in any provider records.

- Left shoulder impingement has received limited treatment focus. If included in demand, additional shoulder-specific documentation may be needed.

What to notice: The narrative tracks treatment escalation from conservative care to injections to surgical recommendation – the progression that establishes damages severity. The "Gaps and Flags" section identifies the 45-day gap defense counsel will target. Verify every chronology entry date and provider name against the Dodonai source output before including in a demand.

---

2. Cross-Examination Outline from Deposition Summary

Source: Chapter 5, Workflow 2 **Agent/Workflow:** /cross-prep skill – deposition summary to impeachment-ready cross-examination **Prompt used:** Cross-examination analysis prompt. See Appendix A, Workflow 2.

KEY ADMISSIONS

Admission 1: Whitfield testified: "We do a floor inspection every hour on the hour. That's our policy and I personally make sure it happens." (p. 47, ln. 8-12). Establishes a specific, verifiable commitment. The "personally" language prevents delegation as a defense.

Admission 2: "We had, I think, two other incidents in the produce section in the past year. Nothing serious." (p. 62, ln. 18-22). Establishes constructive notice – two prior incidents in the same area.

Admission 3: "I don't think we have a log for that specific day. We don't always keep them." (p. 71, ln. 3-7). Destroys the defendant's ability to prove inspections occurred.

INTERNAL CONTRADICTIONS

Contradiction 1: "I personally conduct floor inspections every hour" (p. 47, ln. 8-12) vs. "I was in the management meeting from about 2:00 to 3:30 that afternoon" (p. 83, ln. 4-9). Incident occurred at 2:45 PM. Whitfield cannot have conducted the 2:00 or 3:00 PM inspections.

Contradiction 2: "We always have someone monitoring the floor during business hours" (p. 38, ln. 14-17) vs. "On that day I think it was just me and two cashiers after 2:00 PM. Carlos had gone home sick." (p. 86, ln. 1-5). Two cashiers at registers cannot monitor the floor.

CROSS-EXAMINATION QUESTIONS – Inspection Gap Sequence

Q: Mr. Whitfield, you are the store manager at this location. (p. 12, ln. 3) Q: You are responsible for store safety. (p. 14, ln. 8) Q: Your store has a policy requiring hourly floor inspections. (p. 47, ln. 8) Q: You testified that you personally conduct those inspections. (p. 47, ln. 10-12) Q: On January 12, 2026, you attended a management meeting. (p. 83, ln. 4) Q: That meeting began at 2:00 PM. (p. 83, ln. 6) Q: The meeting lasted until approximately 3:30 PM. (p. 83, ln. 9) Q: During that meeting, you were not on the sales floor. Q: My client fell at 2:45 PM. Q: You did not conduct a floor inspection at 2:00 PM. Q: No one else was assigned to conduct inspections while you were in the meeting. (p. 86, ln. 1-5) Q: You have no inspection log for January 12, 2026. (p. 71, ln. 3-7) Q: You cannot tell this jury that anyone inspected the produce section between 1:00 PM and 3:30 PM.

AREAS FOR FOLLOW-UP DISCOVERY

1. Inspection logs for the 90 days preceding the incident
2. Prior incident reports – Whitfield referenced “two other incidents” but production contains only one report
3. Management meeting attendee list and agenda for January 12, 2026
4. Carlos [last name unknown] personnel file – employee who went home sick, reducing floor coverage
5. Surveillance footage retention policy – Whitfield was not asked about cameras in the produce section

What to notice: Every question is leading (answerable yes or no), builds from established facts to the key admission, and includes the page-line citation for impeachment. The sequence pins Whitfield to his own testimony before revealing the contradiction. Verify all page-line citations against the original transcript before use at trial.

3. Case Evaluation Go/No-Go Memo

Source: Chapter 5, Workflow 4 **Agent/Workflow:** /case-eval skill – intake notes to case acceptance recommendation **Prompt used:** Intake evaluation prompt with Chen Law Group criteria. See Appendix A, Workflow 4.

CASE EVALUATION MEMO Intake Date: April 1, 2026 | AI-generated – attorney review required

CASE SUMMARY

Robert Tran, age 52, electrician, reports a 12-foot fall from scaffolding at a commercial construction site on South Congress Avenue, Austin, TX on February 18, 2026. Scaffolding was provided by the general contractor, Summit Construction Group. Comminuted fracture of the right calcaneus requiring ORIF surgery on February 20, 2026 at Dell Seton Medical Center. Currently non-weight-bearing. Orthopedic surgeon indicates a second surgery (subtalar fusion) likely in 4-6 months. Employer carries workers' comp through Texas Mutual.

LIABILITY ASSESSMENT: Probable

Two avenues: (1) Workers' comp claim against employer via Texas Mutual (no-fault, benefits include medical and 70% AWW income replacement). (2) Third-party negligence against Summit Construction if scaffolding was under their control. If OSHA investigation documents a scaffolding deficiency, third-party liability is strong. Note: if employer is a non-subscriber, negligence action lies directly against employer with no exclusive remedy defense. Confirm subscription status.

DAMAGES ESTIMATE

Category	Estimate
Current medical specials	$78,000-$95,000
Future medical (fusion surgery + rehab)	$45,000-$80,000
Lost wages (5+ months out of work)	$28,000-$32,000
Future lost earning capacity	$150,000-$300,000
General damages multiplier	3x-4x medical specials
Estimated total case value	**$350,000-$750,000**

RED FLAGS

- Workers' comp subrogation right on third-party recovery – factor Texas Mutual's lien into settlement calculations.
- Employer subscription status unverified. Analysis changes if non-subscriber.
- No OSHA report yet. Check whether investigation opened – an OSHA citation against Summit for scaffolding violations would substantially strengthen the third-party claim.

STATUTE CHECK

PI SOL: February 18, 2028 (22 months). Workers' comp: must file within 1 year – **February 18, 2027**. Calendar the workers' comp deadline separately.

> **RECOMMENDATION: ACCEPT**
>
> Calcaneus fracture requiring ORIF and likely fusion is a serious orthopedic injury with strong damages. Even without the third-party claim, the workers' comp claim alone justifies representation given surgical severity and lost earning capacity. Next steps: confirm employer comp subscription status, file workers' comp claim (calendar the 1-year deadline), request OSHA incident report, obtain Dell Seton records, investigate Summit's scaffolding contract, retain vocational expert.

What to notice: The memo handles the workers' comp/third-party interaction that makes construction intakes more complex than a standard auto case. It flags the 1-year workers' comp filing deadline (shorter than the 2-year PI SOL) that must be calendared immediately. Damages estimates are screening-level ranges – confirm against actual bills and retain experts for lost earning capacity.

---

4. Client Status Update Letter

Source: Chapter 7, Agent 3 (Client Update Drafter) **Agent/Workflow:** Monthly client update agent – pulls 30 days of activity from Clio **Prompt used:** Client Update Drafter skill file. See Appendix A, Agent 3.

> To: james.kowalski@email.com Subject: Case Update – Kowalski v. Target Corp (April 2026)
>
> James,
>
> Here's where things stand on your case heading into April.
>
> We completed your deposition preparation last week, and I think it went well. You stayed composed, answered clearly, and didn't volunteer information beyond what was asked. The defense attorney focused heavily on the timeline of your knee symptoms, which tells me they're building their case around the pre-existing meniscus issue from 2019. Your testimony was consistent with what the records show: the 2019 injury resolved fully, you returned to unrestricted activity, and you had no knee complaints until the fall at Target.
>
> On the medical side, Dr. Patel's final report came in on March 28. He's assigned a 15% lower extremity impairment rating and documented that the fall aggravated a previously asymptomatic condition into a symptomatic one. That distinction matters – it's the difference between Target arguing "this was a pre-existing problem" and us showing "this was a resolved problem that their negligence made permanent."

> Next steps: mediation is scheduled for May 15 with Judge Rodriguez (retired) as mediator. Before then, I'll prepare a mediation brief laying out our damages. I expect Target's initial offer to be below what the case is worth – that's standard. We'll evaluate any offer against the full picture: your $47,200 in medical costs, the impairment rating, the two surgeries, and the impact on your landscaping business.
>
> One thing to note: your trust account balance is at $220. We may need to replenish it before mediation to cover the mediator's fee. I'll confirm specifics soon.
>
> If anything changes with your treatment or you have questions before mediation, call me anytime.
>
> Sarah Chen Chen Law Group

What to notice: Legal developments explained in plain language without condescending. Prepares the client for the mediation process (including likely low first offer). Addresses the trust balance proactively. Review every fact against the Clio file and adjust tone for the specific client relationship before sending.

---

5. Billing Reconciliation Weekly Report

Source: Chapter 7, Agent 4 / Chapter 9, Time Entry Reconstruction
Agent/Workflow: Billing Reconciler – Friday 5:00 PM weekly reconciliation **Prompt used:** Billing reconciliation prompt. See Appendix A, Ch 9.

> **BILLING RECONCILER – Week of March 30 - April 3, 2026**
>
> **UNBILLED GAPS BY MATTER:**
>
> **Hernandez v. Lyft | PI-Auto | Gap: 1.5 hrs ($546)** - Calendar: 45-min client call (Mar 31, 2:15 PM) – no Clio entry. Suggested: 0.8 hrs, client conference re: case status. HIGH. - Email: 4 emails to Austin Spine & Rehab re: final records (Apr 1-2) – no entry. Suggested: 0.5 hrs, medical provider correspondence. MEDIUM. - Calendar: 30-min demand review with Maria (Apr 3) – billed 0.3, should be 0.5. Suggested correction: +0.2 hrs. HIGH.
>
> **Davis v. HCA | Med Mal | Gap: 2.6 hrs ($946)** - Email: 6 emails with expert Dr. Kaplan including revised report review (Apr 1-2) – billed 0.3 hrs. Suggested additional: 1.2 hrs, expert correspondence and report review. HIGH. - Calendar: 20-min call with defense counsel re: scheduling (Apr 2, 4:00 PM) – no entry. Suggested: 0.4 hrs. HIGH. - Document: Expert report opened/annotated in Clio (Apr 1, 3:12-4:47 PM) – billed 0.5 hrs. Suggested additional: 1.0 hrs. MEDIUM.

> **Patel v. Meridian Trucking | PI-Auto | Gap: 1.3 hrs ($473)** - Calendar: 2-hr deposition prep block (Apr 3) – billed 1.0 hr. Suggested correction: 2.0 hrs. HIGH. - Email: 3 emails to client re: prep materials (Apr 2) – no entry. Suggested: 0.3 hrs. MEDIUM.
>
> **Rivera v. St. David's | Med Mal | Gap: 1.2 hrs ($437)** - Calendar: Records review block (Apr 2, 9:00-10:30 AM) – no entry. Suggested: 1.0 hrs. HIGH. - Email: 2 emails to opposing counsel re: discovery dispute (Apr 1) – billed 0.2 hrs. Suggested additional: 0.2 hrs. LOW.

> **WEEKLY SUMMARY**

Metric	Value
Total activity identified	38.2 hrs
Total Clio entries	31.6 hrs
Identified gap	**6.6 hrs**
Dollar value ($364/hr)	**$2,402**
Suggested entries	11 (6 HIGH, 3 MEDIUM, 2 LOW)
YTD unbilled hours recovered	72.5 hrs
YTD revenue recovered	**$26,390**

What to notice: Gaps grouped by matter with specific evidence sources and confidence levels. The LOW confidence entries give discretion – a brief email exchange may not warrant a separate entry. YTD running total shows the compounding value. Review every suggested entry for billing propriety under your engagement letter before approving.

---

6. New PI File Opening SOP

Source: Chapter 9, SOP Creation **Agent/Workflow:** SOP generator – converts verbal process description into written procedure **Prompt used:** SOP creation prompt. See Appendix A, Ch 9.

> **SOP: New Personal Injury File Opening Purpose:** Consistent process from signed engagement letter through first records requests. **Roles:** Maria (paralegal), David (legal assistant, backup), Sarah (attorney, review) **Trigger:** Client signs engagement letter.

Step	Action	Who	Timeframe
1	Create matter in Clio: [Last Name] - [Incident Type] [Location]	Maria	Day 1
2	Set matter type: PI-Auto, PI-Premises, Med Mal, or Workers' Comp	Maria	Day 1
3	Populate custom fields: DOL, SOL date (2 years PI / 1 year WC), carrier, policy limits, adverse party	Maria	Day 1
4	Create folder structure: Correspondence / Medical Records / Billing Records / Pleadings / Discovery / Expert / Settlement / Intake	Maria	Day 1
5	Upload signed engagement letter. Verify fee structure: contingency (33.3% pre-suit, 40% post-filing) or hybrid	Maria	Day 1
6	Calendar SOL with reminders at 90 days, 180 days, and 1 year. Add "Demand deadline" 90 days before SOL	Maria	Day 1
7	Send LOR to insurance carrier (certified mail + email). Use Clio template.	Maria	Within 48 hrs
8	Order crash report from Texas DPS (auto cases). Processing: 2-4 weeks.	Maria	Day 1-2
9	Request initial medical records from ER and PCP. Include signed HIPAA authorization.	Maria	Day 1-2
10	Create Clio task list: crash report follow-up (2 wks), records follow-up (3 wks), 30-day client check-in	Maria	Day 1
11	Send client welcome email: confirm representation, summarize next steps, remind re: medical appointments	David	Day 1
12	Attorney review: confirm SOL computation, review LOR, verify fee structure matches engagement letter	Sarah	Day 1-2

COMMON EXCEPTIONS

- **Hybrid fee:** Verify retainer received before opening file. Confirm amount matches engagement letter.
- **Government entity defendant:** Check Texas Tort Claims Act notice requirements (Tex. Civ. Prac. & Rem. Code Ch. 101). Notice required within 6 months – calendar separately from SOL.
- **Multiple treating providers at intake:** Request records from all known providers at Step 9, not just ER and PCP.

QUALITY CHECK: Properly opened file has all 8 Clio folders, SOL calendared with three reminders, LOR sent within 48 hours, at least two records requests initiated, welcome email sent Day 1, attorney sign-off on SOL computation.

What to notice: Generated from a 10-minute verbal description by the paralegal. Captures the steps, the exceptions (government entity notice), and the quality check. This becomes training material for new hires and – when converted to a

skill file – instructions for an AI file-opening assistant. Verify SOL computation rules and government entity notice deadlines match current Texas statutes.

---

7. Practice-Area Blog Post

Source: Chapter 8, Content Engine **Agent/Workflow:** Marketing Engine agent – weekly blog post generation **Prompt used:** Blog post generation prompt. See Appendix A, Ch 8.

What to Do in the First 72 Hours After a Car Accident in Texas

> You've just been hit. Your neck hurts, your hands are shaking, and you're sitting on the side of the road trying to figure out what happened. The next three days matter more than you think – not just for your health, but for your legal rights.

Get Medical Attention – Even If You Feel Fine

> Adrenaline masks pain. After a collision, your body suppresses pain signals for hours or even days. Soft tissue injuries – disc herniations, ligament tears, whiplash – often don't present full symptoms until 48-72 hours after impact.
>
> Go to the ER or urgent care the day of the accident. The medical record creates a documented link between the collision and your injuries. If you wait two weeks, the insurance adjuster will argue your injuries aren't related or aren't serious.

Document Everything at the Scene

> If you're able, photograph damage to all vehicles, the intersection, skid marks, debris, and your visible injuries. Exchange information but keep the conversation short. Don't apologize, don't say "I'm fine," and don't discuss fault.
>
> Texas is a modified comparative fault state – under Texas Civil Practice and Remedies Code Section 33.001, you can recover damages as long as you're not more than 50% at fault. Anything you say at the scene can be used to argue shared responsibility.
>
> Call 911 and get a police report even if the damage seems minor. The officer's report documents the scene and often includes a preliminary fault determination.

Don't Talk to the Other Driver's Insurance Company

> Within 48 hours, the at-fault driver's insurer will call. They may offer a quick settlement – $2,000-$5,000 – to "take care of your bills." That offer is designed to close your claim before you know the full extent of your injuries. A disc

> herniation requiring injections and potentially surgery generates $50,000-$150,000 in medical expenses alone. You're not required to give a recorded statement.

What Most People Don't Know

> From the moment a claim is reported, the insurance carrier assigns an adjuster whose job is to minimize the payout. If you skip a physical therapy appointment, that gap appears in your records. The adjuster will cite it at negotiation: "The plaintiff's treatment history shows a three-week gap, suggesting the injuries resolved." It doesn't matter why you missed. The gap is in the record.
>
> The medical records you create and the documentation you gather in the first 72 hours shape the strength of your claim for the next 12-18 months.

When to Talk to an Attorney

> Not every accident requires a lawyer. But if you went to the ER, if you're missing work, if the other driver was clearly at fault, or if the insurance company is already calling – a conversation with a PI attorney costs nothing. Most work on contingency, meaning you pay nothing unless the case resolves in your favor.

What to notice: 550 words, natural SEO terms, specific Texas statute citation. The "What Most People Don't Know" section demonstrates expertise without being salesy. Review for accuracy on the comparative fault threshold, confirm the code citation, and adjust voice to match your own writing style.

---

8. Breach Notification Compliance Matrix

Source: Chapter 6, Cybersecurity & Privacy section **Agent/Workflow:** Breach notification compliance agent – incident report to multi-state matrix **Prompt used:** Breach notification prompt. See Appendix A, Ch 6.

> **BREACH NOTIFICATION COMPLIANCE MATRIX Client:** DataStream Analytics, Inc. **Discovery Date:** March 28, 2026 **Data Exposed:** Names, SSNs, email addresses, financial account numbers **Total Affected:** 8,430 individuals across 5 states

State	Affected	Deadline	AG Notice Required	Credit Monitoring	Key Requirement
TX	3,210	**May 27** (60 days)	Yes, if > 250 residents	No	AG portal: oag.state.tx.us
CA	2,180	**~Apr 18** ("expedient")	Yes, if > 500 residents	**Yes – 12 months** (SSN exposed)	10-point minimum type in letter
NY	1,540	**~Apr 18** ("expedient")	Yes, **before** individual notice	No	AG notification must precede individual letters
FL	870	**Apr 27** (30 days, firm)	Yes, any number	No	FDLE notification also required
VA	630	**~Apr 28** ("no unreasonable delay")	Yes, if > 1,000 (not triggered)	No	Attorney judgment on "reasonable" timeline

PRIORITY ACTIONS

Priority	Action	Internal Deadline
CRITICAL	Draft CA notification with credit monitoring offer	Apr 11
CRITICAL	File NY AG notification (must go BEFORE individual letters)	Apr 11
URGENT	Draft FL notification + FDLE filing	Apr 20
URGENT	Draft TX AG notification + individual letters	May 13
ACTION	Draft VA AG notification + individual letters	Apr 18
ACTION	Consumer reporting agency notification (per CA requirements)	Concurrent with CA individual notice

ATTORNEY VERIFICATION REQUIRED: All statutory citations and deadlines must be verified against current law. "Without unreasonable delay" jurisdictions (NY, VA) require attorney judgment. California credit monitoring applies because SSNs were exposed – requirement differs if only emails exposed. Breach notification statutes are among the most frequently amended areas of state law.

What to notice: The matrix converts a five-state breach into an actionable tracker with computed deadlines and a priority sequence. The critical insight: New York requires AG notification BEFORE individual notices, changing the entire response sequencing. Every deadline and citation must be verified against current statutes before any notification is sent.

How to Use These Outputs

Every sample was generated from a prompt in the book, applied to realistic case facts. They are first drafts – starting points for attorney review, not finished work product.

The pattern across all eight:

1. **Structured data goes in.** A medical chronology, deposition summary, intake notes, calendar events, breach details.
2. **Structured analysis comes out.** A damages narrative, cross-exam outline, go/no-go memo, compliance matrix.
3. **Attorney judgment completes it.** Verify citations, adjust strategy, add context the AI couldn't know, approve or revise.

The AI handles the assembly. You handle the judgment.

Appendix E: AI Policy Template

This is a starter AI policy that your firm can customize and adopt. It covers the essentials: approved tools, data handling, verification, disclosure, supervision, training, and incident response.

Bracket fields like [Firm Name] indicate where you fill in your own details. The policy is designed to be printed, signed, and enforced as-is with minimal editing.

Why you need this: 43% of law firms have no AI policy (Wolters Kluwer 2026 Future Ready Lawyer). That number is a liability. ABA Formal Opinion 512 (2024) establishes that attorneys have ethical obligations when using AI – competence, communication, supervision, and confidentiality. California's proposed amendments to Rules 1.1, 1.4, 1.6, 3.3, 5.1, and 5.3 make these obligations explicit. Whether or not your state has adopted AI-specific rules yet, a written policy demonstrates the professional diligence regulators expect.

[Firm Name] Artificial Intelligence Use Policy

Effective Date: [Date] **Last Reviewed:** [Date] **Policy Owner:** [Responsible Partner] **Applies To:** All attorneys, paralegals, legal assistants, contract attorneys, and staff of [Firm Name]

---

1. Purpose & Scope

1.1 This policy governs the use of artificial intelligence tools at [Firm Name], including large language models (e.g., Claude, ChatGPT, Gemini), AI-powered legal research platforms, document processing tools, and any AI agent systems operating on behalf of the firm.

1.2 The purpose of this policy is to:

- Ensure compliance with our professional and ethical obligations under the ABA Model Rules of Professional Conduct, as interpreted by ABA Formal Opinion 512 and applicable state bar rules
- Protect client confidentiality and attorney-client privilege
- Establish verification requirements for all AI-generated work product
- Define disclosure obligations to clients, courts, and opposing counsel
- Set supervision and training standards for all personnel

1.3 This policy applies to all AI tool usage in connection with firm business, whether on firm devices or personal devices, whether in the office or remote, and whether used for billable client work or internal firm operations.

1.4 This policy does not prohibit AI use. It establishes the guardrails that allow our firm to use AI responsibly and effectively.

---

2. Approved Tools

2.1 The following AI tools are approved for use at [Firm Name]:

Tool	Approved Use	Data Classification
[Tool 1, e.g., Claude Pro]	Drafting, analysis, research assistance	Tier 1 – No client data without zero-retention confirmation
[Tool 2, e.g., Dodonai]	Document processing (summaries, chronologies, OCR, eDiscovery)	Tier 2 – Client data permitted (SOC 2, HIPAA, zero data retention)
[Tool 3, e.g., Westlaw AI]	Legal research, citation verification	Tier 2 – Client data permitted (enterprise agreement)
[Tool 4]	[Specify use]	[Specify tier]

2.2 Data classification tiers:

- **Tier 1 – General Use:** AI tools that do not have a zero data retention (ZDR) agreement, a signed Business Associate Agreement (BAA), or an enterprise license with explicit data handling terms. These tools may be used for general research, drafting from hypotheticals, and non-client-specific work. **No identifiable client data, case details, or privileged information may be input into Tier 1 tools.**
- **Tier 2 – Client Data Permitted:** AI tools that have confirmed zero data retention, appropriate security certifications (SOC 2, HIPAA where applicable), and a signed data processing agreement or BAA. These tools may be used with client data in accordance with the data handling requirements in Section 3.
- **Tier 3 – Restricted:** AI tools that have not been reviewed and classified. **No firm personnel may use a Tier 3 (unreviewed) AI tool for any firm business.** To request review of a new tool, submit a request to [Responsible Partner] with the tool's privacy policy, data handling terms, and intended use.

2.3 The approved tool list is reviewed and updated [quarterly/annually] by [Responsible Partner]. Changes are communicated to all personnel within [5 business days] of approval.

---

3. Data Handling Requirements

3.1 Client data in AI tools. Before inputting any client information into an AI tool, personnel must confirm:

- The tool is classified as Tier 2 (approved for client data)
- The specific data being input is necessary for the task (minimize data exposure)
- No information subject to a protective order, court seal, or heightened confidentiality agreement is included unless the tool's security meets or exceeds those requirements

3.2 Zero data retention. For Tier 2 tools, [Firm Name] requires written confirmation from the vendor that:

- Client data is not stored after processing is complete
- Client data is not used to train, fine-tune, or improve AI models
- The vendor's subprocessors (including LLM providers) operate under the same retention restrictions

3.3 Privilege considerations. AI tool usage creates records (prompts, outputs, chat logs) that may be discoverable. Personnel must:

- Avoid including privileged attorney-client communications or work product in AI prompts unless the tool provides zero data retention
- Be aware that AI-generated work product may not carry work product protection in all jurisdictions
- Consult [Responsible Partner] before using AI tools in connection with matters involving heightened privilege sensitivity (government investigations, internal investigations, pre-litigation strategy)

3.4 Personally identifiable information (PII). When processing documents containing PII (Social Security numbers, medical records, financial account numbers), use only Tier 2 tools with confirmed HIPAA compliance and/or appropriate data security certifications.

3.5 Prohibited inputs. The following may never be input into any AI tool without prior written approval from [Responsible Partner]:

- Documents subject to court protective orders
- Grand jury materials
- Sealed court records
- Attorney-client privileged communications (in Tier 1 tools)
- Client financial account numbers, Social Security numbers, or passwords (in Tier 1 tools)

---

4. Verification Requirements

4.1 Mandatory verification. Every AI-generated output used in connection with client work must be verified by an attorney before use. There are no exceptions.

4.2 Citation verification. AI tools can and do generate fabricated legal citations. As of early 2026, 600+ cases involving AI hallucinations have been documented, and 128 attorneys have been implicated (National Law Review). Courts have sanctioned attorneys, disqualified firms, and referred lawyers to state bar disciplinary authorities. **Every citation generated by an AI tool must be independently verified against the primary source before inclusion in any filing, brief, motion, memorandum, letter, or client communication.**

4.3 Factual verification. AI-generated factual statements (dates, amounts, case facts, medical details, statutory language) must be verified against the source documents. Do not assume AI-generated facts are accurate. Check them.

4.4 Verification workflow. For all AI-assisted work product:

1. **Generate** the draft or analysis using the approved AI tool
2. **Review** the output for accuracy, completeness, and appropriateness
3. **Verify** every legal citation against Westlaw, Lexis, or the primary source
4. **Verify** every factual claim against the case file or source documents
5. **Edit** to reflect your professional judgment and the client's specific circumstances
6. **Document** that verification was performed (a note in the matter file is sufficient)

4.5 No AI-generated document may be filed with any court, sent to any client, or transmitted to opposing counsel without completion of the verification workflow in Section 4.4.

---

5. Disclosure Obligations

5.1 Disclosure to clients. [Firm Name] discloses its use of AI tools to clients as follows:

- The firm's engagement letter includes a section describing the firm's use of AI tools for document processing, research assistance, and administrative tasks
- Clients are informed that AI-assisted work product is always reviewed and verified by an attorney before delivery
- Clients may request that AI tools not be used on their matter, and the firm will honor such requests

5.2 Disclosure to courts. Court-specific AI disclosure rules are evolving. As of [Date]:

- Some federal judges require disclosure of AI use in filings (check local rules and standing orders)
- Some state courts have adopted or proposed AI disclosure requirements
- Where a court requires AI disclosure, [Firm Name] complies fully and transparently

[Responsible Partner] maintains a current list of jurisdictions and courts with AI disclosure requirements. Attorneys must check this list before filing in any jurisdiction where they have not previously filed.

5.3 Disclosure to opposing counsel. [Firm Name] does not proactively disclose AI use to opposing counsel unless required by court order, local rule, or stipulation. AI is a tool, like legal research databases and document management software. Voluntary disclosure is not required under current ethics rules, but personnel should be prepared to answer truthfully if asked.

5.4 State-by-state variation. AI disclosure obligations vary significantly by jurisdiction. Some states have issued formal ethics opinions; others have not addressed AI at all. Attorneys are responsible for knowing the disclosure requirements in every jurisdiction where they practice. When in doubt, disclose. Overcommunication builds trust; concealment creates risk.

---

6. Supervision & Review

6.1 Attorney responsibility. Under ABA Model Rules 5.1 and 5.3, and as reinforced by ABA Formal Opinion 512, supervising attorneys bear responsibility for the work product produced under their supervision – including work product generated with AI assistance. AI does not reduce the supervision obligation; it may increase it.

6.2 AI output review. All AI-generated work product is treated as a first draft requiring attorney review, equivalent to a draft produced by a junior associate or paralegal. The reviewing attorney must:

- Read the full output (not just skim)
- Verify legal citations and factual claims per Section 4
- Confirm the output is appropriate for the specific client, matter, and legal context
- Exercise independent professional judgment before approving the output for use

6.3 Escalation. If any personnel member encounters an AI output that appears inaccurate, misleading, or potentially harmful, they must:

1. Stop using the output immediately

2. Flag the issue to [Responsible Partner] or the supervising attorney
3. Document the issue (screenshot or copy of the AI output, description of the error)
4. If the inaccurate output was already shared externally, proceed to Section 8 (Incident Response)

6.4 Agent systems. For AI agents that operate on schedules or with access to firm tools (email, calendar, case management), [Responsible Partner] must:

- Approve the agent's scope of access before activation
- Review the agent's output on a [daily/weekly] basis during the first [30 days]
- Confirm that agent actions are limited to approved operations (e.g., creating drafts, not sending emails; reading case data, not modifying records)
- Maintain a log of active agents, their access scope, and their last review date

---

7. Training Requirements

7.1 Initial training. All attorneys, paralegals, and staff must complete AI policy training within [30 days] of hire or within [30 days] of this policy's effective date, whichever is later. Training covers:

- This policy in full
- How to classify AI tools by data tier
- The verification workflow (Section 4)
- How to identify AI hallucinations and fabricated citations
- Disclosure obligations for each jurisdiction where the firm practices
- How to report AI-related incidents

7.2 Annual refresher. All personnel complete an annual refresher on AI policy updates, new approved tools, emerging ethics opinions, and lessons learned from any incidents.

7.3 New tool onboarding. When a new AI tool is added to the approved list, [Responsible Partner] provides a briefing to all affected personnel covering the tool's capabilities, limitations, data classification, and approved use cases.

7.4 Training records. [Firm Name] maintains records of all AI policy training, including dates, attendees, and topics covered. These records are retained for [3 years] and are available for review by [Firm Name]'s malpractice insurer upon request.

---

8. Incident Response

8.1 Definition. An AI-related incident is any event where an AI-generated error, hallucination, or unauthorized data exposure reaches a client, court, opposing counsel, or third party. Examples:

- A fabricated citation is included in a filed brief
- Inaccurate factual statements based on AI output are communicated to a client
- Client data is inadvertently input into a Tier 1 (non-approved for client data) AI tool
- An AI agent takes an unauthorized action (e.g., sends an email instead of creating a draft)

8.2 Immediate response. Upon discovering an AI-related incident:

1. **Stop** – cease using the AI output and prevent further distribution
2. **Assess** – determine what was shared, with whom, and what the potential impact is
3. **Notify** – inform [Responsible Partner] immediately, regardless of severity
4. **Preserve** – save the AI prompt, the AI output, and any communications that included the erroneous output

8.3 Corrective action. [Responsible Partner] determines the appropriate corrective action, which may include:

- Issuing a correction to the affected party (client, court, opposing counsel)
- Filing a notice of errata or corrected brief with the court
- Notifying the firm's malpractice insurer
- Reviewing and tightening the verification workflow
- Additional training for the personnel involved
- Restricting or revoking access to the AI tool that produced the error

8.4 Post-incident review. Within [10 business days] of any incident, [Responsible Partner] conducts a review to determine:

- Root cause (verification failure, tool misconfiguration, training gap, policy gap)
- Whether the current policy prevented or failed to prevent the incident
- What changes to policy, training, or tool configuration are needed
- Whether the incident must be reported to any regulatory body

8.5 Documentation. All incidents and post-incident reviews are documented and retained for [5 years]. This record serves both as a compliance record and as institutional knowledge for improving the firm's AI practices.

---

9. Policy Review

9.1 This policy is reviewed [annually/semi-annually] by [Responsible Partner], with the next review scheduled for [Date].

9.2 Reviews assess:

- Changes in ABA ethics opinions, state bar rules, and court orders related to AI
- New AI tools adopted or under consideration by the firm
- Incidents or near-misses that occurred since the last review
- Advances in AI capabilities that may affect risk profiles
- Feedback from attorneys and staff on policy practicality

9.3 Updated policies are distributed to all personnel within [10 business days] of adoption. All personnel must acknowledge receipt and review of the updated policy.

9.4 Between scheduled reviews, [Responsible Partner] monitors for significant developments (new ethics opinions, disciplinary actions, court orders) that may require interim policy updates.

Acknowledgment

I have read and understand the [Firm Name] Artificial Intelligence Use Policy. I agree to comply with all provisions of this policy in my work at the firm.

Name: _____

Title: _____

Signature: _____

Date: _____

This policy template is provided as a starting point for law firm AI governance. It is not legal advice. Firms should review this template with their malpractice insurer, ethics counsel, and state bar's guidance before adoption. Requirements vary by jurisdiction – consult your state bar's ethics opinions and rules for jurisdiction-specific obligations.

Reference materials:

- ABA Formal Opinion 512 (2024) – Generative AI Tools

- ABA Model Rules: 1.1 (Competence), 1.4 (Communication), 1.6 (Confidentiality), 3.3 (Candor), 5.1 (Supervisory Responsibility), 5.3 (Non-lawyer Assistants)
- California proposed amendments to Rules 1.1, 1.4, 1.6, 3.3, 5.1, 5.3 (public comment closed May 2026)
- National Law Review: 600+ AI hallucination cases, 128 attorneys implicated
- Wolters Kluwer 2026 Future Ready Lawyer: 43% of firms have no AI policy

DODONAI

READY TO PUT THIS PLAYBOOK TO WORK?

Bring this architecture inside your firm.

Dodonai designs, builds, and runs custom AI agents for law firms — using the exact architecture described in this book.

We start with the one workflow that costs your team the most time every week. Blueprint in week one. First agent live within 30 days.

Book a Free 30-Minute Intro Call

calendly.com/nick-dodonai

✓

Blueprint Session

We map your workflows before writing a single prompt.

✓

30-Day Deployment

First agent live in 30 days or the Blueprint is refunded.

✓

Managed & Monitored

We run it. You review it. You stay in every loop.

WHAT YOU HAVE NOW

The complete playbook for running an AI-native law firm.

· 11 chapters · from foundations to advanced architecture

· 8 real agent outputs you can adapt for your practice

· The complete prompt library (Appendix A)

· A ready-to-adopt AI Policy Template (Appendix E)

· Crawl / Walk / Run implementation path for any firm size

"The AI handles the assembly. You handle the judgment."

DODONAI

dodon.ai
hello@dodon.ai
calendly.com/nick-dodonai

AI Services for Legal Teams